Fundamentals of Statistics

A. Reeves Johnson

Contents

Preface viii

Notation x

I The Kiss Experiment and the Statistical Model 1

1 The Kiss Experiment and What Statistics Is For 2

1.1 Probability and statistics are the same picture read in opposite directions 2

1.2 The kiss experiment 4

1.3 Modeling assumptions 4

1.4 Population versus sample 5

1.5 The trinity of statistical inference 7

1.6 What this book is and is not 8

1.7 What comes next 9

1.8 Summary 9

1.9 Exercises 9

2 Statistical Models and Identifiability 11

2.1 The rationale behind statistical modeling 11

2.2 Statistical models 12

2.2.1 Examples of parametric models 14

2.3 From the sample space to the data and back 16

2.4 Well-specified models 17

2.5 Parametric, nonparametric and semiparametric models 18

2.6 Identifiability 20

2.6.1 Example: a non-identifiable parameter 21

2.6.2 More examples 21

2.7 Gaussian mixture models . . . 23
2.7.1 Possible Gaussian mixture models . . . 24
2.8 What comes next . . . 26
2.9 Summary . . . 26
2.10 Exercises . . . 26

II Probability for Statistics **28**

3 The Law of Large Numbers **29**
3.1 Why limits of random variables need new tools . . . 30
3.2 A first look: the sample mean on two scales . . . 31
3.3 Convergence in probability . . . 34
3.3.1 Example: the minimum of uniforms . . . 36
3.3.2 A workhorse: shrinking variance . . . 37
3.4 The Weak Law of Large Numbers (WLLN) . . . 38
3.5 What comes next . . . 42
3.6 Summary . . . 42
3.7 Exercises . . . 43

4 The Central Limit Theorem **44**
4.1 Convergence in distribution . . . 44
4.1.1 Why only the continuity points? . . . 45
4.1.2 Two equivalent definitions . . . 45
4.1.3 Example 1: i.i.d. random variables . . . 47
4.1.4 Example 2: the maximum of uniforms . . . 47
4.1.5 Limits that are constants . . . 49
4.2 The Central Limit Theorem . . . 50
4.2.1 The essence of the CLT: a 'stochastic' bound . . . 52
4.3 Hoeffding's inequality: help for small sample sizes . . . 53
4.4 What the limit theorems buy you . . . 54
4.5 What comes next . . . 56
4.6 Summary . . . 56
4.7 Exercises . . . 57

5 Working with Limits **58**

5.1 Arithmetic and convergence . . . 58
5.2 Slutsky's theorem . . . 59
5.3 Continuous mapping theorem (CMT) . . . 60
5.4 The delta method . . . 61
5.5 Where the limit theorems are used in the rest of the book . . . 65
5.6 What comes next . . . 66
5.7 Summary . . . 67
5.8 Exercises . . . 67

6 The Gaussian Distribution and Quantiles **68**

6.1 Why the Gaussian is everywhere . . . 68
6.2 The density and its CDF . . . 69
6.3 Properties of Gaussian distributions . . . 70
6.4 Computing probabilities, forwards and backwards . . . 72
6.5 Quantiles . . . 73
6.6 What comes next . . . 76
6.7 Summary . . . 76
6.8 Exercises . . . 77

III Foundations of Inference **78**

7 Estimators: Bias, Variance and Risk **79**

7.1 Statistics and estimators . . . 79
7.2 A statistic reduces the data . . . 80
7.2.1 The sufficiency principle . . . 80
7.3 Consistency . . . 82
7.4 Asymptotic normality . . . 83
7.5 Bias . . . 85
7.5.1 Jensen's inequality . . . 85
7.6 Variance . . . 87
7.7 Quadratic risk . . . 87
7.8 Strengths and weaknesses of estimators . . . 91
7.9 What comes next . . . 91

7.10 Summary . . . 91
7.11 Exercises . . . 92

8 Confidence Intervals **93**
8.1 What a confidence interval is . . . 93
8.2 The recipe, on an exact case . . . 94
8.3 The kiss experiment: an interval that isn't one . . . 96
8.4 Three ways to get unstuck . . . 97
8.5 What 'asymptotic level' buys at a finite n . . . 100
8.6 The subway: when the scale depends on the parameter in a new way . . . 101
8.7 What makes an interval wide . . . 103
8.8 What the 95% means, and three things it doesn't . . . 104
8.9 What comes next . . . 105
8.10 Summary . . . 105
8.11 Exercises . . . 106

9 Hypothesis Tests **107**
9.1 A yes/no question about a parameter . . . 107
9.2 Null and alternative hypotheses . . . 108
9.3 What a test is . . . 110
9.4 Two kinds of error, level and power . . . 111
9.5 Building a test from the Central Limit Theorem . . . 114
9.6 The kiss experiment, exactly and by the CLT . . . 116
9.7 p-values . . . 118
9.8 Tests and confidence intervals are one construction . . . 120
9.9 What comes next . . . 122
9.10 Summary . . . 122
9.11 Exercises . . . 122

IV Methods of Estimation **124**

10 Distances Between Distributions **125**
10.1 What 'close to the truth' should mean . . . 125
10.2 Total variation . . . 127
10.3 What total variation says about tests . . . 130

10.4 Why total variation can't be estimated by an average . . . 132
10.5 The Kullback–Leibler divergence . . . 132
10.5.1 Independent observations add . . . 136
10.5.2 Near the truth, KL is a parabola . . . 136
10.5.3 KL controls total variation . . . 137
10.6 Estimating the KL divergence: the road to maximum likelihood . . . 138
10.7 When the model is wrong: the closest member of the family . . . 140
10.8 What comes next . . . 143
10.9 Summary . . . 143
10.10 Exercises . . . 144

11 The Likelihood and Maximum Likelihood 145
11.1 From the KL divergence to the likelihood . . . 145
11.2 The likelihood of a discrete model . . . 146
11.3 The likelihood of a continuous model . . . 148
11.4 The maximum likelihood estimator . . . 150
11.5 Finding a maximum . . . 151
11.6 Maximum likelihood estimators of the standard families . . . 153
11.7 When the maximum likelihood estimator fails to exist or isn't unique . . . 157
11.8 No formula: mixtures and the EM algorithm . . . 158
11.9 What comes next . . . 160
11.10 Summary . . . 160
11.11 Exercises . . . 161

12 Properties of the MLE 162
12.1 Three questions about an estimator . . . 163
12.2 Consistency . . . 163
12.3 The score and the Fisher information . . . 165
12.3.1 Fisher information of the standard families . . . 168
12.4 Asymptotic normality . . . 170
12.5 The Cramér–Rao bound . . . 173
12.6 What the theorems need, and the models they do not cover . . . 174
12.7 What comes next . . . 175
12.8 Summary . . . 175
12.9 Exercises . . . 176

13 Method of Moments and M-Estimation 177

V Parametric Hypothesis Tests 178

14 The t-Test 179

15 Wald and Likelihood Ratio Tests 180

VI Nonparametric Hypothesis Tests 181

16 Goodness of Fit for Discrete Distributions 182

17 Kolmogorov-Smirnov, Lilliefors and QQ Plots 183

VII Bayesian Statistics 184

18 Priors, Posteriors and Conjugacy 185

19 Jeffreys Priors and Bayesian Intervals 186

VIII Linear Regression 187

20 Univariate Linear Regression 188

21 Multivariate Linear Regression and Inference 189

IX Generalized Linear Models 190

22 Exponential Families 191

23 Link Functions and Fitting 192

X Principal Component Analysis 193

24 Principal Component Analysis 194

Appendices **196**

A Mathematical Toolkit **196**

A.1 Limits of sequences of numbers . . . 196
A.2 Supremum and infimum . . . 197
A.3 Sets, maps and one-to-one . . . 197
A.4 Indicator functions . . . 198
A.5 Random variables as functions, and 'almost surely' . . . 200
A.6 Densities, CDFs and quantiles . . . 201
A.7 Three inequalities . . . 202
A.8 Convex functions . . . 203
A.9 Tangent lines, the mean value theorem and Taylor's expansion . . . 203
A.10 Exponentials and logarithms . . . 204
A.11 Orders of magnitude: how fast is 'fast' . . . 205
A.12 Functions of several variables: gradient and Hessian . . . 205
A.13 To be added with later chapters . . . 207

B The Modes of Convergence in Full **208**

B.1 Intuition for the three main modes . . . 208
B.2 Convergence almost surely . . . 209
B.2.1 The Strong Law of Large Numbers (SLLN) . . . 210
B.2.2 Convergence almost surely versus convergence in probability . . . 210
B.3 Convergence in mean . . . 212
B.4 The hierarchy of convergence types . . . 212
B.5 Summary . . . 214
B.6 Exercise . . . 214

C Supplementary Proofs **215**

C.1 Pinsker's inequality . . . 215
C.2 Each step of the EM algorithm raises the likelihood . . . 218
C.3 The uniform law of large numbers . . . 219

Solutions to the Exercises **221**

Preface

This is a book about the mathematics underneath the statistical methods you already use. It covers what a graduate introduction to statistical inference covers: the limit theorems that justify large-sample reasoning, statistical models, estimators and how to judge them, confidence intervals, hypothesis tests, maximum likelihood, the chi-squared family, Bayesian inference, linear regression, generalized linear models and principal component analysis. It assumes you have met most of these before as recipes, and it asks the question an introductory course has no room for: why is any of this allowed?

Who it is for. A reader with a solid undergraduate statistics course behind them. You know the sample mean, the normal distribution, confidence intervals and t-tests, and you have had a semester of calculus. Every step of every derivation is written out; nothing is left as an exercise that the text then relies on.

The mathematics is all here: the ϵ-N definition of a limit, the proofs, the inequalities. At the same time, the limit theorems of Part II are used constantly in the rest of the book, and nobody carries their exact statements around in their head. So every later use of the Law of Large Numbers, the Central Limit Theorem, Slutsky's theorem or the delta method restates in one line what that theorem gives you, at the point where it is used, marked *Recall*. You can read any chapter after Chapter 6 without turning back to Part II.

How it is organized. One chapter per idea, ten to eighteen pages, in parts ordered by what has to be understood before what. Chapter 1 introduces the example that runs through the whole book and the three questions statistics asks about it. Chapter 2 defines the object every later chapter lives inside. Part II is the probability: where an average settles, how it fluctuates on the way, and what you may do to a limit and still know where you are. Part III answers the three questions with guarantees, and the rest of the book is methods.

Every chapter opens with a box, 'In this chapter', that says what is coming and how it connects to the chapter before, and closes with a paragraph on what comes next, a summary and exercises. Five kinds of box run through the text:

- *Intuition*: the picture or the argument in words behind a result, usually before the result is proved.
- *From undergraduate statistics*: the version of the idea you already met, and what the general version adds. There are many of these: the statistics you already know is the natural place to start from.
- *Pitfall*: a mistake that is easy to make.
- *Worked example*: the computation carried through with numbers.
- *Does this make sense?*: how to tell that an answer is absurd without redoing the work. Extreme cases, bounds, a second route to the same number.

Every definition is followed by a short block that says the same thing in plain words and then says what it means for the running example. Every proof is preceded by one line of strategy, marked either *core* (follow it; there are about ten such proofs) or *for completeness* (the statement is what you need; read the proof when you want to). Solutions to all exercises are at the back of the book, and each exercise says on which page.

The mathematics you need. A solid introductory statistics course and one semester of calculus, provided you are willing to read a definition slowly. Where a proof leans on a mathematical idea that an introductory course leaves implicit (the ϵ-N definition of a limit, the mean value theorem, Markov's inequality, what 'injective' means), the text marks the step with a gray pointer such as '(toolkit A.9)' to the Mathematical Toolkit, Appendix A, which explains that idea in a page with a picture and says where else the book uses it. The toolkit is where to go when a step feels unjustified. It is not a prerequisite to read first.

The minimal path through Part II. Chapters 3 to 5 are about limits of random variables, and they exist to deliver three facts: averages settle on their expectations, the error of an average is Gaussian on the $1/\sqrt{n}$ scale, and there are rules for transforming limits and plugging estimates into them. Read those three chapters for those three facts and take Chapter 3's opening box as the working translation of the word 'asymptotic' wherever it appears. The remaining modes of convergence are in Appendix B, for orientation only.

Two conventions that bite. This book writes the Gaussian distribution as $\text{Gauss}(\mu, \sigma^2)$ with the *variance* as the second parameter, so $\text{Gauss}(85, 4)$ has standard deviation 2. Much of the software you will meet takes the standard deviation instead. And it writes q_α for the quantile with probability α *above* it, so $q_{0.05} = 1.645$ and $q_{0.025} = 1.96$, while most tables and software index a quantile by the probability below. Both conventions are standard in mathematical statistics and both are traps when you move between sources. Chapter 6 explains them with pictures; every formula in this book follows them.

Figures. Every figure is generated by a short script that ships with the book's source, so any of them can be redrawn with other numbers. Nothing in a figure is schematic: when a curve is drawn, it is the real curve, and when a probability is shaded, the shaded area is that probability.

Where the material comes from. The syllabus is that of MIT's 18.6501x, *Fundamentals of Statistics*, which I worked through and took notes on, and this book grew out of those notes. You do not need the course to read this book, and after this paragraph it is not mentioned again.

Notation

The symbols that recur, in the order a reader meets them. These are the standard conventions of mathematical statistics; the two that differ from most software are marked.

Symbol	Meaning
$X_1, \dots, X_n$	the sample: n random variables, i.i.d. (independent, identically distributed) unless said otherwise
$x_1, \dots, x_n$	the realized values, numbers; likewise $\bar{X}_n$ is a random variable and $\bar{x}_n$ a number (Chapter 3)
$X_1, \dots, X_n \overset{iid}{\sim} \text{Bern}(p)$	the X_i are i.i.d. with the named distribution; Bern, Bin, Pois, Expo, Unif, Gauss
$\text{Gauss}(\mu, \sigma^2)$	Gaussian with mean μ and *variance* σ^2 (most software takes the standard deviation)
φ, Φ	density and CDF of $\text{Gauss}(0,1)$
q_α	the quantile with probability α *above* it: $\mathbb{P}(Z > q_\alpha) = \alpha$, so $q_{0.05} = 1.645$, $q_{0.025} = 1.96$ (software indexes by the probability below)
$\mathbb{P}$, $\mathbb{E}$, $\mathbb{V}\text{ar}$	probability, expectation, variance; $\mathbb{P}_\theta$, $\mathbb{E}_\theta$ when they are computed under the distribution with parameter θ
$\mathbb{1}(A)$	the indicator: 1 if A holds, 0 otherwise; $\mathbb{E}[\mathbb{1}(A)] = \mathbb{P}(A)$
$\bar{R}_n$, $\hat{p}$	the kiss experiment: $R_i = 1$ if couple i turns right, $\hat{p} = \bar{R}_n$; $n = 124$, 80 right turns, $\bar{R}_n = 0.645$
$\bar{T}_n$, $\hat{\lambda}$	the subway inter-arrival times, $T_i \overset{iid}{\sim} \text{Expo}(\lambda)$, $\hat{\lambda} = 1/\bar{T}_n$; $n = 64$, $\bar{T}_n = 6.23$
$X_n \xrightarrow[n\to\infty]{\mathbb{P}} X$	convergence in probability (Definition 3.1)
$X_n \xrightarrow[n\to\infty]{\text{(d)}} X$	convergence in distribution (Definition 4.1); the limit may be written as a distribution, $X_n \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$
$X_n \xrightarrow[n\to\infty]{\text{a.s.}} X$, $X_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} X$	convergence almost surely (Appendix B); both arrows at once, which is how the two laws of large numbers are written together
$(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$	a statistical model: sample space E and a family of candidate distributions indexed by $\theta \in \Theta$ (Definition 2.1)
θ, θ^*, $\hat{\theta}_n$	a parameter, the true parameter, an estimator of it from n observations
Bias, $R(\hat{\theta}_n)$	bias $\mathbb{E}_\theta[\hat{\theta}_n] - \theta$ and quadratic risk $\mathbb{E}_\theta[(\hat{\theta}_n - \theta)^2]$ (Chapter 7)
$\mathcal{I} \ni \theta$	the random interval $\mathcal{I}$ contains θ; same as $\theta \in \mathcal{I}$, written to put the random object first (Chapter 8)
α	a small probability: a test's level, or $1-$ (an interval's level)
r.v., CDF, PDF, PMF	random variable; cumulative distribution function $F(x) = \mathbb{P}(X \le x)$; density; probability mass function
CLT, LLN, CMT	Central Limit Theorem, Law of Large Numbers, continuous mapping theorem
(toolkit A.k)	a pointer to entry k of the Mathematical Toolkit, Appendix A, which explains the piece of mathematics being used

Part I

The Kiss Experiment and the Statistical Model

The book opens with its running example and with the object every later chapter lives in: the kiss experiment (Chapter 1) and the statistical model (Chapter 2), which turns the three assumptions of that experiment into a family of candidate distributions with a parameter to learn. The model comes first so that the limit theorems of Part II arrive with a job to do.

Chapter 1

The Kiss Experiment and What Statistics Is For

In this chapter

- Probability and statistics as the same picture read in opposite directions: from a known truth to the data it will produce, and from observed data back to the truths that could have produced it.
- The book's running example, the kiss experiment, and the three modeling assumptions that turn 80 out of 124 into a mathematical problem.
- Why the sample proportion is a random variable, what population it is a sample from, and the three questions about it that the next chapters answer.
- The trinity of statistical inference: estimation, confidence intervals and hypothesis testing, each stated for the kiss experiment and each pointed to the chapter that delivers it.

1.1 Probability and statistics are the same picture read in opposite directions

Probability studies randomness. Sometimes the physical process is completely known: dice, cards, roulette, fair coins. Roll one die. Alice gets \$1 if the number of dots is at most 3; Bob gets \$2 if the number of dots is at most 2. Who would you rather be? You don't need any data to answer: Alice expects $1 \times \frac{3}{6} = \$0.50$ per roll, Bob expects $2 \times \frac{2}{6} \approx \$0.67$, so Bob. Roll two dice and win \$100 if you called the sum correctly: you call 7 because six of the thirty-six equally likely outcomes sum to 7 and no other total has as many. In both cases the mechanism is known, and probability deduces what the data will look like.

Statistics is the reverse problem. The mechanism is unknown, you observe data it produced, and you want to say something about the mechanism. The two directions fit into a single diagram, the central dogma of probability and statistics.



A drug trial shows both directions at work. *Probability:* previous studies showed the drug to be 80% effective. Then in a study of 100 patients you can anticipate that on average 80 will be cured, and that at least 65 will be cured with probability 99.99%. *Statistics:* you observe that 78 of 100 patients were cured. You will be able to conclude, by the end of Chapter 8, that you are 95% confident the drug is effective on between 69.9% and 86.1% of patients. (That sentence is the standard shorthand for a confidence interval; Section 8.8 says exactly what it does and does not claim.) Figure 1.1 draws both halves.

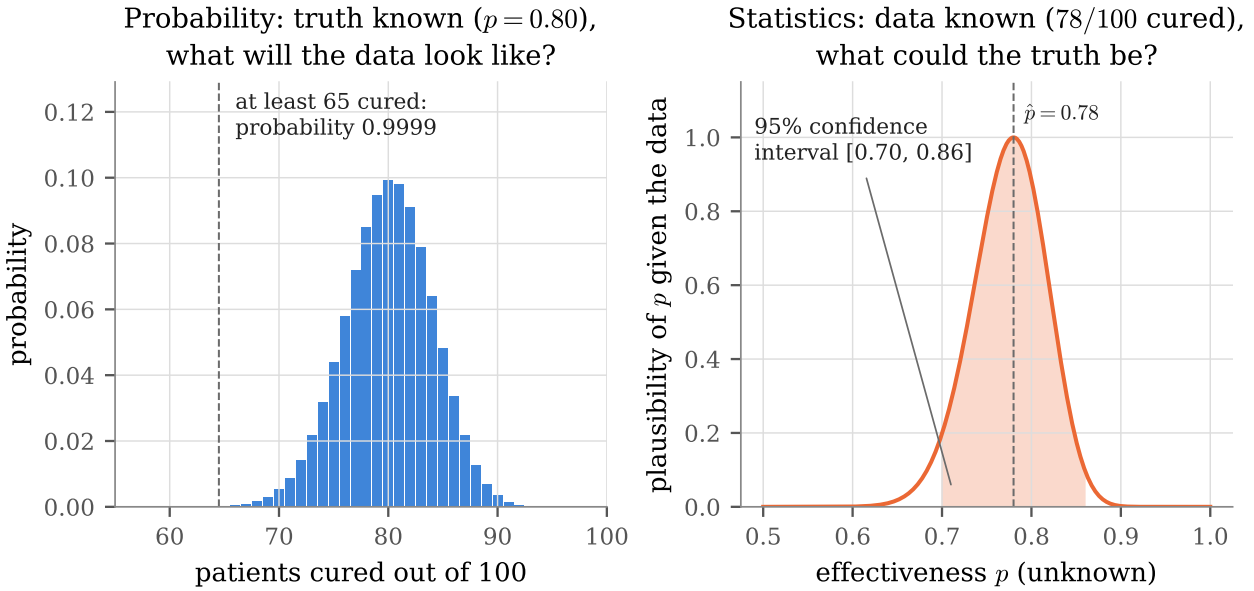


Figure 1.1: The drug example in both directions. Left, probability: the effectiveness $p = 0.80$ is known, so the number cured out of 100 has a known distribution (Binomial), and the probability of at least 65 cures is a computable 0.9999. Right, statistics: 78 of 100 were cured and p is unknown; the curve shows how plausible each value of p makes the data (it is the likelihood of Chapter 11, scaled to peak at 1), and the shaded band is the 95% confidence interval $[0.70, 0.86]$ of Chapter 8.

Notice what the right-hand panel is *not*. It isn't a probability distribution for p: in this book p is a fixed unknown number, not a random variable, and the curve is a plausibility ranking of candidate truths, not a density. The one place where p does get a distribution is Chapter 18, and the difference is the whole point of that chapter.

From undergraduate statistics: descriptive and inferential statistics

An introductory course divides statistics into descriptive (summarize the data you have) and inferential (say something about the population you don't have). This book is entirely about the second, and it adds the step the introductory course leaves implicit: a mathematical model of how the data were produced, precise enough that the inferential statements can be proved rather than asserted. The model is the thing on the left of the diagram, and every method in the book is a way of traveling the orange arrow with a guarantee attached.

1.2 The kiss experiment

Newborn babies, lying on their backs, turn their heads to the right more often than to the left. A 2003 paper in *Nature* asked whether that preference survives into adulthood, and looked for it in an unexpected place: which way couples turn their heads when they kiss.[1] The author watched kissing couples in public places (airports, train stations, beaches and parks) in the United States, Germany and Turkey, and recorded for each couple which way their heads turned.

Let p denote the proportion of couples who turn their heads to the right when kissing. The statistical experiment: observe n kissing couples and record each outcome, 1 for right and 0 for left. Estimate p with the proportion $\hat{p}$ of couples who turned right. Of the $n = 124$ couples observed, 80 turned to the right, so

$$\hat{p} = \frac{80}{124} = 64.5\%.$$

Two things about the number 124 because the whole book turns on them. It is small, and it is small for a reason: each observation needs an observer in a public place at the right moment, close enough to see which way the heads turn, and a kiss is over in seconds. And it is the only sample there is: the study is published, and the question is what these 124 observations can be made to say. Everything that follows is about how much a sample of this size can be made to say.

Now some random intuition. 64.5% is much larger than 50%, so there seems to be a preference for turning right. But suppose the data had been right, right, left, that is $n = 3$: that's 66.7% to the right, an even bigger number. Better evidence? Obviously not; intuitively you need a large enough sample size to make a call. How large? Another way to put the problem: for $n = 124$, what is the minimum number of couples turning right that you would need to see to be convinced that $p > 50\%$? 63? 72? 75? 80? Mathematical modeling is needed to understand the accuracy of this procedure, and the number turns out to be 72 (Section 1.5).

1.3 Modeling assumptions

Formally, the procedure consists of the following. For $i = 1, \ldots, n$, define $R_i = 1$ if the i-th couple turns to the right and $R_i = 0$ otherwise. The estimator of p is the sample average

$$\hat{p} = \bar{R}_n = \frac{1}{n} \sum_{i=1}^{n} R_i.$$

What is the accuracy of this estimator? To answer, you propose a statistical model that describes, or approximates well, the experiment. You think of the R_i's as random variables so that $\hat{p}$ is also a random variable, and then you study its fluctuations.

Coming up with a model consists of making assumptions on the observations R_i, $i = 1, \ldots, n$, in order to draw statistical conclusions. Here are the three assumptions this model makes:

1. Each R_i is a random variable.
2. Each of the r.v.s R_i is Bernoulli with parameter p.
3. $R_1, \ldots, R_n$ are mutually independent.

[1] O. Güntürkün, 'Adult persistence of head-turning asymmetry', *Nature* 421 (2003), 711.

Let's discuss these assumptions, since each one is a decision rather than a fact.

1. Randomness is a way of modeling lack of information. With perfect information about the conditions of the kiss (the couple's heights, which side each was standing on, which way each habitually turns, what is going on in their minds), physics and psychology might let you predict the outcome. You don't have that information, so you treat the outcome as random.
2. Given (1), the R_i's are necessarily Bernoulli r.v.s since $R_i \in \{0, 1\}$; the only question is their parameters. They could have a different parameter for each couple, $R_i \sim \text{Bern}(p_i)$, but you don't have enough information in the data to estimate 124 separate p_i's from 124 observations. So you simply assume the observations come from the same process: $p_i = p$ for all i.
3. Independence is reasonable because the couples were observed in many different places, in three countries, at different times. It would fail if, say, several of the 124 observations were the same couple seen more than once, or if couples watched each other and copied.

A collection of r.v.s $X_1, \ldots, X_n$ satisfying (2) and (3), each with the same distribution and all mutually independent, is called **independent and identically distributed**, abbreviated i.i.d., and the book writes $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$. Nearly every model in this book starts from an i.i.d. sample.

Intuition: complicated process = simple process + random noise

That equation is the one-line philosophy of statistical modeling. Some processes are random in fact (a biased coin, measurement error, which customer walks in next). Others are deterministic but far too complex to describe (which way a particular couple turns). Statistical modeling treats the second kind like the first: pick a plausible simple process (every couple turns right with the same probability p) and a noise distribution (Bernoulli), and let the noise absorb everything you chose not to model. Good modeling consists in choosing a simple process and a noise distribution that are plausible for the data at hand, and the trade is always the same: more assumptions buy fewer parameters to estimate, at the price of being wrong if the assumptions are. Chapter 2 makes this precise with the definition of a statistical model.

1.4 Population versus sample

In the model, what is $\hat{p} = \bar{R}_n$ a sample *from*? Here is a thought experiment that answers it. Assume there is a total population of 5,000 couples who kiss in airports, and assume for the sake of argument that $p = 35\%$, or that $p = 65\%$. What do samples of size 124 look like in each case? Figure 1.2 answers by drawing the exact distribution of $\bar{R}_n$ under the i.i.d. model for each assumed p, together with the distribution for samples of size 3. (For a population of 5,000 the draws are without replacement, which changes the standard deviation of the picture by about one percent; the bridge box below says why the book works with the i.i.d. model instead.)

The figure is the whole subject in miniature. Under each candidate truth, probability tells you what the data would look like. Statistics runs that backwards: the observed data are typical under some candidate truths and absurd under others, and the first group is your answer. The rest of the book is about doing this with guarantees, and about doing it when the candidate truths form a continuum rather than two values.

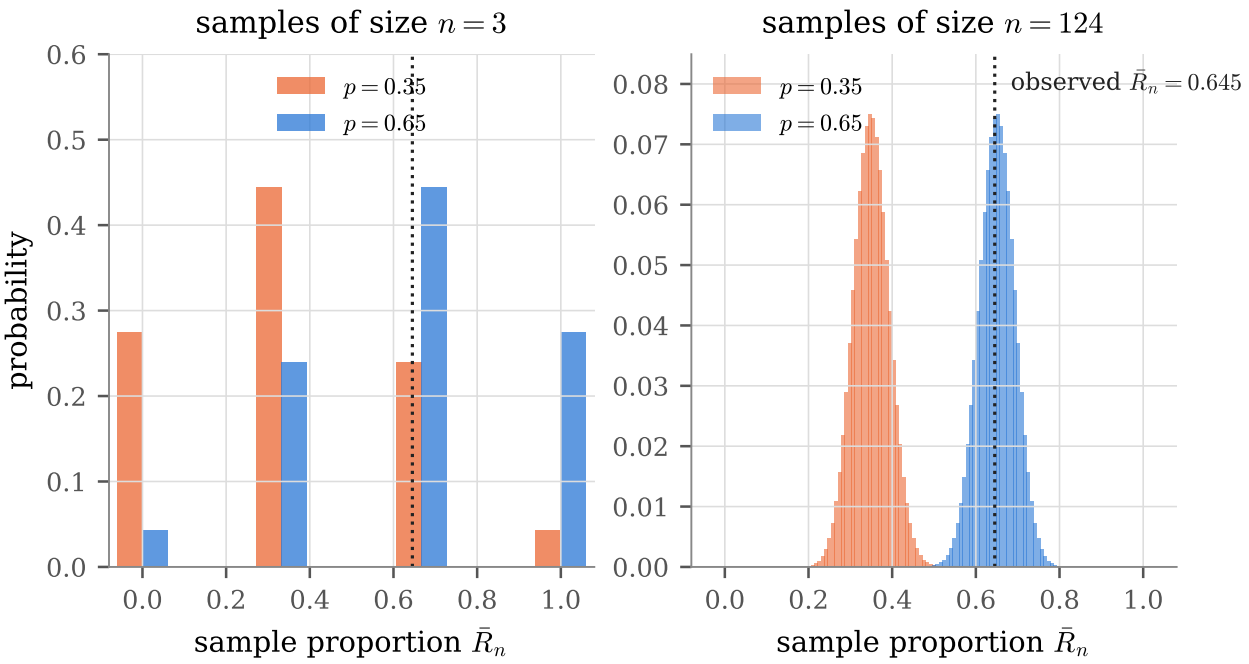


Figure 1.2: What samples look like under two hypothetical truths. The bars are the exact (Binomial) distribution of the sample proportion $\bar{R}_n$ when the true proportion is $p = 0.35$ (orange) and when it is $p = 0.65$ (blue). With $n = 3$ (left) the two distributions overlap heavily: two or more right turns out of 3 has probability 0.28 under $p = 0.35$ and 0.72 under $p = 0.65$, so three couples settle nothing. With $n = 124$ (right) the distributions have separated, and the observed 0.645 (dotted line) is typical under $p = 0.65$ and essentially impossible under $p = 0.35$ (probability about 2×10^{-11} of a proportion that large).

From undergraduate statistics: population, sample, parameter, statistic

An introductory course describes a population (all kissing couples), a parameter of it (p, the proportion who turn right), a sample (124 couples) and a statistic computed from the sample ($\hat{p} = 0.645$) that estimates the parameter. This book keeps all four words and changes what sits behind the first one. The population is no longer a finite list of 5,000 couples; it is a probability distribution, Bern(p), and a sample is n independent draws from it. The change is what makes the mathematics go: a finite population would make the draws slightly dependent (drawing without replacement) and would make the answer depend on the population size, while the distribution model gives clean results that hold whatever the size of the population and whether or not you can even list it.

Why does the model need probability? Because you need to understand probabilistic aspects of the distribution of the random variable

$$\hat{p} = \bar{R}_n = \frac{1}{n} \sum_{i=1}^{n} R_i.$$

Specifically, you need to be able to answer questions such as:

- Is the expected value of $\hat{p}$ close to the unknown p?
- Does $\hat{p}$ take values close to p with high probability?

- Is the variance of $\hat{p}$ large? That is, does $\hat{p}$ fluctuate a lot?

Most of the probability tools needed are about averages of independent random variables, which is why Part II is about exactly that, once Chapter 2 has said what a model is. For the kiss experiment the three answers are: $\mathbb{E}[\hat{p}] = p$ exactly; $\hat{p}$ is close to p with high probability once n is large (the Law of Large Numbers, Chapter 3); and $\mathbb{V}\mathrm{ar}(\hat{p}) = p(1-p)/n$, which at $n = 124$ (and evaluated at $p = 0.645$, since p itself is unknown) makes the typical fluctuation of $\hat{p}$ about 0.043 (the Central Limit Theorem, Chapter 4, says the fluctuations are Gaussian).

Pitfall: is 64.5% an estimator?

No: 64.5% is an estimate, and an estimate is not an estimator. An **estimator** is a rule for turning data into a guess, $\hat{p} = \bar{R}_n$, and because the data are random the estimator is a random variable; it has an expectation, a variance and a distribution, and those are what Figure 1.2 draws. An **estimate** is the number the rule produced on the one sample you have, $\bar{r}_n = 0.645$. It is one draw from the estimator's distribution: before the study, $\bar{R}_n$ could have come out anywhere from 0 to 1 with probabilities set by p, and the study produced one of those values. Once observed it is a fixed number; the distribution belongs to the rule that produced it, not to the number. Every statement about accuracy in this book is a statement about the estimator; the estimate is what you report. Chapter 3 returns to this distinction under the name 'theoretical sample mean', and Section 2.3 walks the whole chain from the sample space to the number. Outside that discussion and this box, the book follows the usual practice and writes the upper-case symbol with its realized value, $\bar{R}_n = 0.645$, trusting the context to say which of the two is meant.

1.5 The trinity of statistical inference

Everything statistical inference does falls under three headings, together called the trinity of statistical inference. Here is each one as a sentence about the kiss experiment, with the chapter that makes the sentence precise.

1. **Estimation.** '$\hat{p} = \bar{R}_n$ is an estimator for the proportion p of couples who turn their heads to the right.' What makes it a good one? It is unbiased, its variance shrinks like $1/n$, and it is consistent. Chapter 7 defines these properties and compares estimators by them; Chapter 11 shows that $\bar{R}_n$ is also the maximum likelihood estimator.

2. **Confidence intervals.** '$[0.56, 0.73]$ is a 95% confidence interval for p.' A range of truths under which the observed data are not surprising, with a guarantee about how often the procedure captures the truth. Chapter 8 builds this interval three ways and explains what the 95% does and does not mean.

3. **Hypothesis testing.** 'We find statistical evidence that more couples turn their heads to the right when kissing.' A yes/no decision about p with a controlled chance of being wrong. Chapter 9 sets up the decision and Chapter 15 supplies the general tests. Here is the answer to the question of Section 1.2: if the truth were $p = 1/2$, that is, if couples had no preference for either side, then the number turning right among 124 couples would reach 72 or more with probability 0.044 and 75 or more with probability 0.012. So 63 right turns would be unremarkable (probability 0.46 of at least that many), 72 would be mild evidence, and 80,

which has probability 0.0008 of being reached under $p = 1/2$, is strong evidence. So 72 is the threshold at the conventional 5% level for this one-sided question; the two-sided version of Chapter 9, which also guards against $p < 1/2$, asks for 74. (Both numbers come from the exact Binomial distribution of the count; Chapter 9 also derives the Gaussian-approximation versions, 72 and 73, and says why they can differ.)

All three are statements about the same random variable $\bar{R}_n$, built from the same three modeling assumptions, and Figure 1.3 draws them on one axis. Change an assumption, for instance drop independence, and all three change.

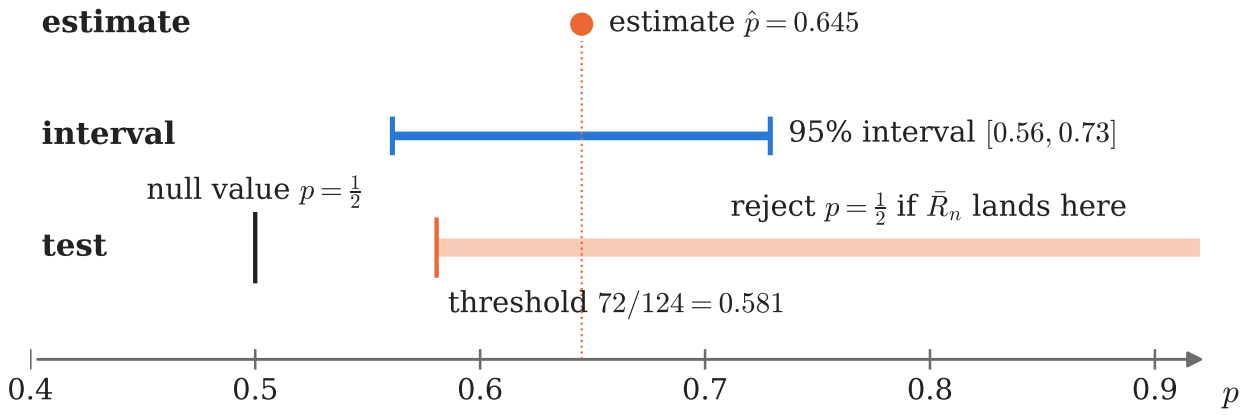


Figure 1.3: The trinity on one axis. Every statement is about where p sits on the line. The estimate is a point, $\hat{p} = 0.645$. The 95% confidence interval is a bracket around it, $[0.56, 0.73]$: the truths under which the data are not surprising. The test compares the data to one particular truth, $p = \frac{1}{2}$: if the observed proportion lands at or beyond the threshold $72/124 = 0.581$ (shaded), the data count as evidence against $p = \frac{1}{2}$ and for $p > \frac{1}{2}$. The threshold is set so that this happens only 5% of the time when $p = \frac{1}{2}$ really is the truth. Chapter 8 and Chapter 9 build the bracket and the threshold; the point is where both start.

1.6 What this book is and is not

This book has two aims: to introduce the mathematical theory behind statistical methods, and to supply theoretical guarantees for the methods you already use. By the end you should be able to take a real-life situation like the kiss experiment, state it as a statistical problem in mathematical terms, choose methods appropriate to it, and say what those methods do and do not license you to conclude. Four things recur throughout: the derivation of a standard method rather than its recipe, the justification of a quantitative statement from a stated set of modeling assumptions, the mathematics that arises in statistics and is worth seeing in its own right, and a toolbox of arguments that transfers to models the book never mentions.

Three things are outside its scope. Applied modeling, meaning which data to collect and how to clean what you collected, is a separate subject, and the kiss experiment illustrates the point: Section 1.3 takes the 124 observations as given and says nothing about where and when to watch, or how to decide which way a pair of heads turned. Implementation is not treated at all, and no formula here is tied to a particular piece of software. And the aim is not a list of named methods:

the general machinery of Chapter 15 delivers whole families of tests at once, including most of the named ones, which is the reason for building it. Where the undergraduate version of an idea makes the general version easier to see, that version appears in a box labeled 'From undergraduate statistics'.

1.7 What comes next

The three questions of Section 1.4 are all about the random variable $\hat{p} = \bar{R}_n$: its expectation, where it concentrates, how much it fluctuates. The first has a one-line answer, $\mathbb{E}[\hat{p}] = p$. Before any of them is answered, Chapter 2 turns the three assumptions of Section 1.3 into an object, the statistical model, so that everything later can be stated for any model and not only this one. The other two questions are about what happens as n grows, and 'as n grows' is a limit, so Part II is about limits of random variables: Chapter 3 says where $\bar{R}_n$ ends up, Chapter 4 says what its fluctuations look like on the way, and Chapter 5 says what you may do to those limits (transform them, plug estimates into them) and still know where you are. Chapter 6 then collects what you need about the Gaussian distribution, which is where all of the limits land. If you are comfortable with limits of random variables already, Chapter 3 opens with a box that says which parts of Part II are the minimal path.

1.8 Summary

- Probability goes from a known mechanism to the data it produces; statistics goes from observed data back to the mechanisms that could have produced them. Same diagram, opposite arrows.
- The kiss experiment: $n = 124$ couples, 80 turned right, $\hat{p} = \bar{R}_n = 0.645$. It is the running example of the book.
- Three modeling assumptions turn the data into mathematics: each R_i is random, each is $\text{Bern}(p)$ with the same p, and they are independent. Together: $R_1, \ldots, R_n \overset{iid}{\sim} \text{Bern}(p)$. Each assumption is a choice with a justification.
- The population is a probability distribution, not a list; a sample is n independent draws. The estimator $\hat{p}$ is a random variable; the estimate 0.645 is one realization of it.
- The trinity of inference: estimation (is $\hat{p}$ a good guess?), confidence intervals (which truths are compatible with the data?), hypothesis testing (is $p > 1/2$?). Every inferential method in this book is one of the three.

1.9 Exercises

Exercise 1.1. Two dice are rolled and you win \$100 if you called the sum. Which total should you call, and what is your expected winning? Why is this a probability question rather than a statistics question? *(Solution on page 221.)*

Exercise 1.2. In the drug example, verify the probability statement: if the drug cures each patient independently with probability 0.8, show that the number cured out of 100 has mean 80 and standard deviation 4, and explain why 'at least 65 cured' has probability very close to 1. *(Solution on page 221.)*

Exercise 1.3. Suppose the kiss data had been $n = 3$ couples, two turning right. Under the hypothetical truths $p = 0.35$ and $p = 0.65$, compute the probability of observing at least two right turns. What does the comparison tell you? *(Solution on page 221.)*

Exercise 1.4. Which of the three modeling assumptions would be violated in each case, and in which direction would it push $\hat{p}$? (a) The observer recorded each couple twice by mistake. (b) Half the couples were observed at a single wedding reception where the photographer told everyone to turn right. (c) The observer only recorded couples whose faces were visible from where the observer stood. *(Solution on page 221.)*

Chapter 2

Statistical Models and Identifiability

In this chapter

- What a statistical model is: a sample space and a family of candidate distributions indexed by a parameter. The whole of the rest of the book happens inside this object.
- The chain from the abstract sample space to the numbers in the data file, walked once for the kiss experiment: outcomes, random variables as functions on them, and observations as those functions evaluated at the outcome that happened.
- The trade that modeling always makes: assumptions buy fewer parameters to learn, at the price of being wrong if the assumptions are. The siblings example, drawn both ways.
- Well-specified models, and parametric versus nonparametric ones.
- Identifiability: when the distribution of the data pins down the parameter, and when it doesn't, with the standard example drawn in the parameter space so you can see what fails.
- Gaussian mixture models, as a generative story and as an example of a model whose parameter is identifiable only up to relabeling.

Chapter 1 made three assumptions about the kiss experiment and called the result a model. This chapter says what a model is in general, so that everything from Part III on can be stated for any model and not only the Bernoulli one.

2.1 The rationale behind statistical modeling

Chapter 1 was about one estimator, $\hat{p} = \bar{R}_n$, under one set of assumptions. To say what estimators, intervals and tests are in general, you first have to fix the object they all live in: a model, which is the three assumptions of Section 1.3 written as a mathematical object. The rationale behind statistical modeling can be outlined as follows:

- Let $X_1, \ldots, X_n$ be n independent copies of X. The goal of statistics is to learn the distribution of X.

  NB: It's crucial to understand this notation, this way of specifying problems, and, most importantly, the relationship between the n i.i.d. r.v.s $X_1, \ldots, X_n$ and the r.v. X. Each of the n i.i.d. r.v.s $X_1, \ldots, X_n$ has the same distribution as the r.v. X. The relationship between X and $X_1, \ldots, X_n$ can be understood as follows:

1. X is the underlying r.v. of interest, with an unknown distribution that you want to learn.
2. $X_1, \ldots, X_n$ are n r.v.s that form your sample of observations drawn from the same distribution as X. This is all-important: When you see something like 'Suppose $X_1, \ldots, X_n$ are n independent copies of X' you should *automatically interpret these n variables as your sample data.*

By observing the values of $X_1, \ldots, X_n$, you can make inferences about the distribution of X.

- If $X \in \{0, 1\}$, then it's Bernoulli and you only have to learn the parameter p. That was the kiss experiment.
- It can be more complicated. Take a (partial) dataset of the number of siblings, including self, collected from college students:

$$2, 3, 2, 4, 1, 3, 1, 1, 1, 1, 1, 2, 2, 3, 2, 2, 2, 3, 2, 1, 3, 1, 2, 3, \ldots$$

  If $X \in \{1, 2, 3, 4, 5, 6, \geq 7\}$, each with a different probability $\mathbb{P}(X = x)$, then you could make no assumption and try to learn the whole PMF: that's 7 parameters to learn (six, really, since they sum to one).
- Or you could assume that $X - 1 \sim \text{Pois}(\lambda)$. That's only 1 parameter, λ, to learn. Figure 2.1 shows both on the siblings data.
- **NB:** It's important to keep in mind that parameters define distributions. What is meant is that, suppose you're considering the family of Gaussian distributions. Once you specify μ and σ^2, you've fixed the distribution. Learning the parameter and learning the distribution are the same task. (The Gaussian family appears in several examples of this chapter, written $\text{Gauss}(\mu, \sigma^2)$ with the variance second, and with Φ for the CDF of $\text{Gauss}(0, 1)$; Chapter 6 is its full treatment, and the notation page has what is needed here.)

Intuition: assumptions are currency

Every assumption you add to a model buys you something and costs you something. '$X - 1$ is Poisson' reduces seven unknowns to one, which means 24 observations go a long way instead of being spread over seven cells, and it lets you say things about families of 8 siblings that you never observed. What it costs is the possibility of being wrong: if the real PMF isn't Poisson-shaped, no amount of data will make the model right, and the estimate of λ will be a precise estimate of the wrong thing. The phrase for this is that a model should be plausible, and the mathematics that follows assumes it is. Checking the assumption is the job of the goodness-of-fit tests in Chapter 16.

2.2 Statistical models

Definition 2.1 (Statistical model)**.** Let the observed outcome of a statistical experiment be a sample $X_1, \ldots, X_n$ of n i.i.d. r.v.s in some sample space E (usually $E \subseteq \mathbb{R}$; the requirement that E be measurable is a technicality, see (toolkit A.5)) and denote by $\mathbb{P}$ their common distribution. A **statistical model** associated to that statistical experiment is a pair:

$$(E, (\mathbb{P}_\theta)_{\theta \in \Theta}),$$

where:

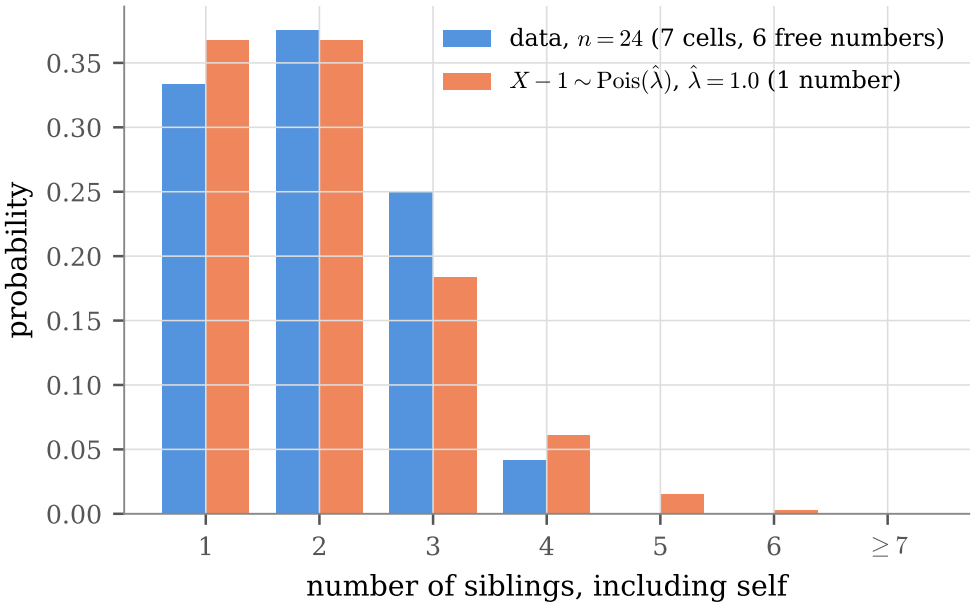


Figure 2.1: The siblings data ($n = 24$ values) as an empirical PMF, next to the one-parameter model $X - 1 \sim \text{Pois}(\hat{\lambda})$ with $\hat{\lambda}$ set to the sample mean minus one. Seven numbers on the left, six of them free since they sum to one; one on the right. The model fits well here; where it didn't, that would be the assumption's fault, not the data's.

- E is the **sample space** for X; that is, the set containing all possible outcomes of X. The sample space is therefore also where your observations, $X_1, \ldots, X_n$, live.
- $(\mathbb{P}_\theta)_{\theta \in \Theta}$ is a family of probability distributions on E, the candidate distributions for X.
- Θ is the **parameter space**; that is, a set consisting of the possible values of θ.

> **In words.** A model is a list of candidate explanations for the data: where an observation can land (E), and every distribution you are willing to consider as the one that produced it, with a label θ on each.
>
> **For the kiss data.** $E = \{0, 1\}$, the candidates are the Bern(p) distributions, one for each $p \in (0, 1)$, and the label is p itself. 'Learn the model' means 'find out which label is right'.

Note: A **probability measure** or **probability distribution** takes an event $A \subseteq E$ as input and returns $\mathbb{P}_\theta(A)$ (toolkit A.5). It is assumed to satisfy the usual axioms of probability. Most importantly, a probability measure is not a PMF or PDF. A probability measure assigns probabilities to events in the sample space, while a PMF or PDF assigns probabilities (or densities) to specific values of a discrete or continuous r.v. The two determine each other, so in practice you write down the family by writing down its PMFs or PDFs, as in the examples below, but the object in the definition is the distribution.

Figure 2.2 is the picture to keep in mind: a model is a family of distributions, and the parameter is the label that picks one out.

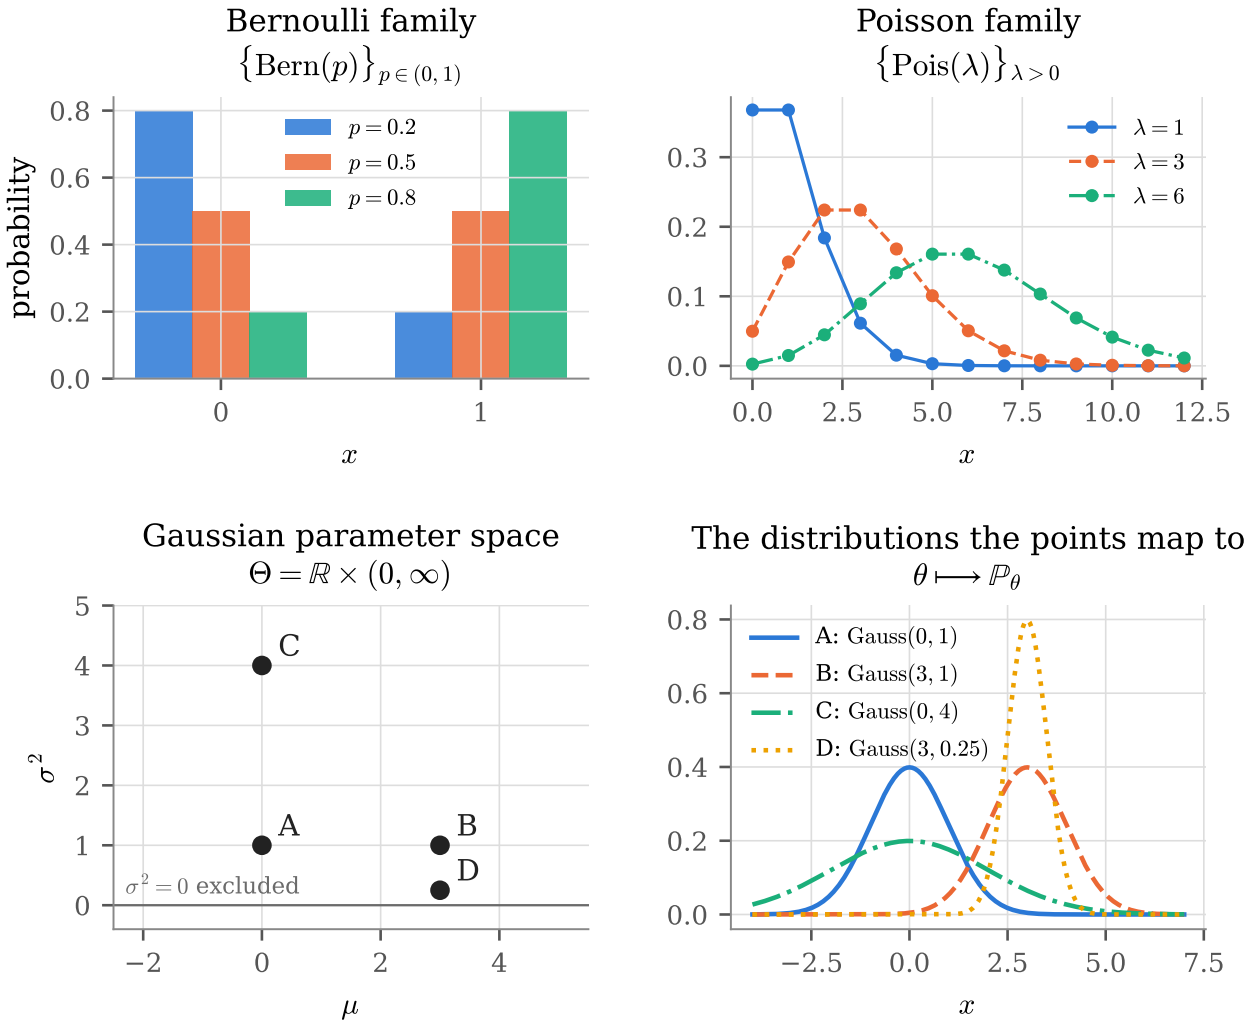


Figure 2.2: A statistical model as a family of candidate distributions. Top: two one-parameter families, Bernoulli and Poisson; sliding the parameter moves through the family. Bottom: the Gaussian family has a two-dimensional parameter space (left), and each point of it is one distribution (right). The arrow $\theta \mapsto \mathbb{P}_\theta$ from the left panel to the right one is what the definition is about, and Section 2.6 asks whether it ever sends two points to the same distribution.

2.2.1 Examples of parametric models

For parametric models (defined in Section 2.5: the parameter is a list of finitely many numbers), the parameters live in some space that's a subset of $\mathbb{R}^d$ and the parameters can be represented by a vector. Here are some canonical examples of parametric models:

1. For n Bernoulli trials: $\big(\{0,1\}, (\mathrm{Bern}(p))_{p\in(0,1)}\big)$. The vectors of parameters thus consist of singletons (p).

2. If $X_1, \ldots, X_n \overset{iid}{\sim} \mathrm{Pois}(\lambda)$ for some unknown $\lambda > 0$: $(\mathbb{N}, (\mathrm{Pois}(\lambda))_{\lambda>0})$. The vectors of parameters thus consist of singletons (λ).

3. If $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$ for some unknown $\mu \in \mathbb{R}$ and $\sigma^2 > 0$:

$$\left(\mathbb{R}, (\text{Gauss}(\mu, \sigma^2))_{(\mu,\sigma^2)\in\mathbb{R}\times(0,\infty)}\right).$$

   The vector of parameters has length 2 and is given by (μ, σ^2).

4. If $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}_d(\mu, \mathbb{I}_d)$ for some unknown $\mu \in \mathbb{R}^d$: $\left(\mathbb{R}^d, (\text{Gauss}_d(\mu, \mathbb{I}_d))_{\mu\in\mathbb{R}^d}\right)$. The vector of parameters is μ itself, of length d.

5. If $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, a])$ for some unknown $a > 0$: $(\mathbb{R}_+, (\text{Unif}([0, a]))_{a>0})$. Not $([0, a], \ldots)$: see the pitfall below.

Sometimes there's no simple notation like $\text{Bern}(p)$ for the family and you have to be more explicit. The linear regression model: $(X_1, Y_1), \ldots, (X_n, Y_n) \in \mathbb{R}^d \times \mathbb{R}$ are i.i.d. with $Y_i = \beta^\top X_i + \epsilon_i$, $\epsilon_i \overset{iid}{\sim} \text{Gauss}(0, 1)$ independent of $X_i \sim \text{Gauss}_d(0, \mathbb{I}_d)$, for an unknown $\beta \in \mathbb{R}^d$. Here $E = \mathbb{R}^d \times \mathbb{R}$, $\Theta = \mathbb{R}^d$, and $\mathbb{P}_\beta$ is the joint distribution of the pair (X, Y) that this recipe produces. Chapter 21 is about this model in a slightly different form: the X_i treated as fixed numbers rather than Gaussian draws, and the error variance σ^2 unknown, which makes $\Theta = \mathbb{R}^d \times (0, \infty)$ and makes the Y_i independent but no longer identically distributed, since each has its own mean $\beta^\top x_i$. Treating the predictors as fixed is what lets the regression chapter condition on them and still have a model in the sense of Definition 2.1.

Pitfall: the sample space must not depend on the parameter

Which of these is a statistical model for n Bernoulli trials: $\left(\{1\}, (\text{Bern}(p))_{p\in(0,1)}\right)$ or $\left(\{0, 1\}, (\text{Bern}(p))_{p\in(0.2,0.4)}\right)$? Only the second. The first has a sample space that leaves out an outcome the distributions put mass on; a Bernoulli r.v. takes the value 0, so $\{1\}$ isn't where the observations live. The second is a model, just one that assumes p is between 0.2 and 0.4 (which may or may not be well-specified, the next section's question). Similarly, for $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, a])$ with a unknown, the model is $(\mathbb{R}_+, (\text{Unif}([0, a]))_{a>0})$ and not $([0, a], \ldots)$: the sample space is fixed before you know a, so it has to be a set that contains every observation for every candidate a. Any such set will do ($\mathbb{R}$ works too); the convention in this book is the smallest natural one. A sample space that mentions the unknown parameter is not a sample space.

From undergraduate statistics: the model was always there

In an introductory course the model is a sentence: 'assume the data are normally distributed', 'assume each trial succeeds with the same probability'. This book makes the sentence into an object with three parts, and the point of doing so is that every later statement can name exactly what it depends on. 'The sample mean is unbiased' is a statement about $\mathbb{E}_\theta[\bar{X}_n]$ for every $\theta \in \Theta$; 'the t-test has level 5%' is a statement that holds for every member of the Gaussian family; 'the estimator is consistent' means $\hat{\theta}_n \to \theta$ under $\mathbb{P}_\theta$ for every θ. When a method fails in practice, it's almost always because the data came from outside the family, and you can only say that if the family is written down.

2.3 From the sample space to the data and back

The definition says the observations live in a sample space E, and a random variable is, by definition, a function on a sample space S (toolkit A.5); Chapter 3 returns to that definition when it needs sequences of random variables. These are two uses of the same phrase for two different sets, and the gap between them is where the single most common confusion in this subject lives: what is the relationship between the abstract space on which probability is defined and the concrete numbers in the data file? This section walks the chain once, slowly, for the kiss experiment, with Figure 2.3 as the map.

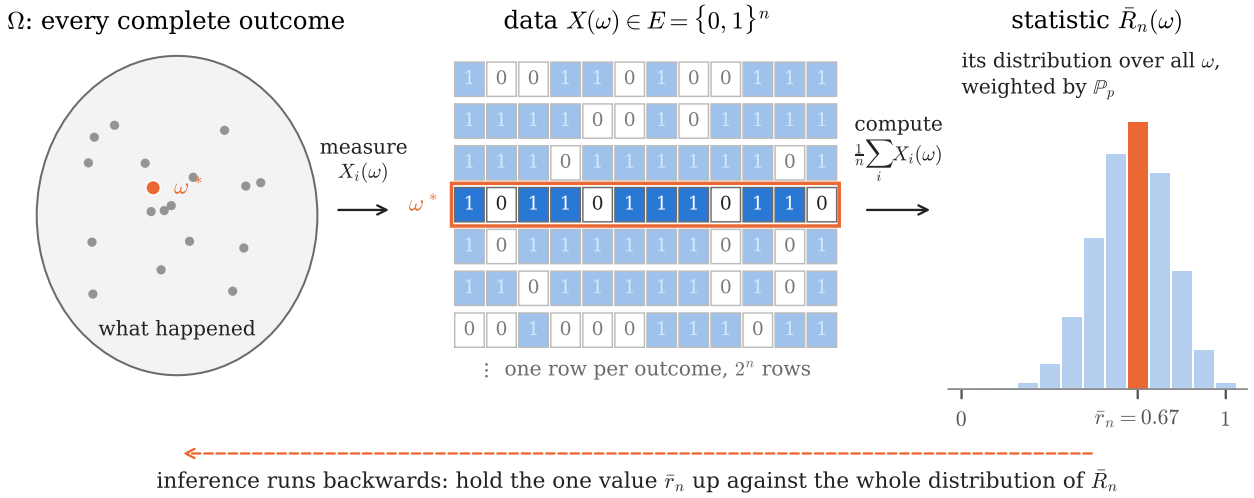


Figure 2.3: Three layers between the sample space and a number, for a kiss experiment with $n = 12$ couples. Left: the sample space Ω, one point per complete possible outcome of the study; the one that happened is marked. Middle: the data are a function on Ω, one row of 0s and 1s per outcome, and the marked row is the data set you have. Right: the statistic $\bar{R}_n$ is another function on Ω; over all outcomes weighted by their probabilities it has a distribution (the bars), and the marked outcome gives the one value you observe, $\bar{r}_n = 8/12$. Nothing travels from right to left: inference holds the one value up against the distribution of the rule.

Layer one: the sample space. An outcome ω of the study is everything that could happen, in full detail: which 124 couples were observed, where and in which order, and which way each one turned. The set of all such outcomes is the sample space, written Ω here (S in the toolkit and in Chapter 3). It is as abstract as the subject gets, and you never see an element of it. What you see is what your measuring instruments report about the one ω that occurred.

Layer two: random variables are the measuring instruments. A random variable is a function on Ω (toolkit A.5). The instrument for couple i is R_i, with $R_i(\omega) = 1$ if, in outcome ω, couple i turned right and 0 otherwise. The data vector $(R_1, \ldots, R_n)$ is a function on Ω too, with values in $E^n = \{0,1\}^n$; one row of Figure 2.3 per outcome. And $\bar{R}_n = \frac{1}{n}\sum_i R_i$ is yet another function on Ω. That is all 'the estimator is a random variable' means: a rule that could be applied to any outcome, written down before you know which outcome you will get.

Layer three: the observations are those functions evaluated at the outcome that happened. Nature picked one ω^*. Your data are $r_1 = R_1(\omega^*), \ldots, r_n = R_n(\omega^*)$, and your estimate is $\bar{r}_n = \bar{R}_n(\omega^*) = 0.645$. Going from the sample space to the observations is 'plug in the ω that

occurred'. A realization is one evaluation of a function; a distribution is how often each value would come up across all of Ω, weighted by probability. The realized number 0.645 is one draw from the distribution of $\bar{R}_n$; after it is observed it is a fixed number, and the distribution belongs to the rule, not to the number (the pitfall in Section 1.2).

Going back is not an inverse. You never recover ω^* from the data, and you do not need to. What inference does instead is ask, for each candidate θ: across all the outcomes that could have occurred under $\mathbb{P}_\theta$, how often would $\bar{R}_n$ have landed near the value it took? The one number is held up against the whole distribution of the rule, and the candidates under which it would have been typical are the plausible ones. That is the dashed arrow in Figure 2.3, and it is what every confidence interval and test in this book does.

Intuition: the trick that makes Ω concrete

Since Ω only ever enters through the random variables defined on it, you may as well take Ω to be the set of data sets you could have gotten. For the kiss experiment that is $\Omega = \{0,1\}^{124}$, a finite set with 2^{124} elements, and under the model $\mathbb{P}_p(\omega) = p^{\#\text{ones}}(1-p)^{\#\text{zeros}}$. Every row of Figure 2.3 is a point of this Ω, and $R_i(\omega)$ is just 'the i-th entry of the row'. This is why Definition 2.1 mentions E and never Ω: for a fixed n the observable sample space E^n *is* the sample space, and a model is a family of probabilities on it. The abstract Ω of probability theory is needed only when one space has to carry infinitely many random variables at once, $R_1, R_2, R_3, \ldots$ without end, which is exactly what the limit theorems of Part II need in order to say 'as $n \to \infty$'. For any fixed n, think $\Omega = E^n$ and nothing is lost.

From undergraduate statistics: the sampling distribution, defined rather than simulated

An introductory course introduces the sampling distribution of the sample mean by rerunning the study: draw many samples, compute $\bar{x}$ for each, draw the histogram. That histogram is the right panel of Figure 2.3, approximated by simulation; this chapter defines it instead, without rerunning anything, by putting a probability on Ω and asking how $\bar{R}_n(\omega)$ is distributed under it. The two pictures agree: each simulated sample is one ω, the intro course's 'many samples' are many points of Ω, and the limit of the histogram as the number of samples grows is the distribution of the random variable. Figure 3.5 in Chapter 3 draws the same picture as rows and columns: a row is an ω, the row average is $\bar{R}_n(\omega)$, and the column of row averages is the sampling distribution. Your data set is the one row nature handed you.

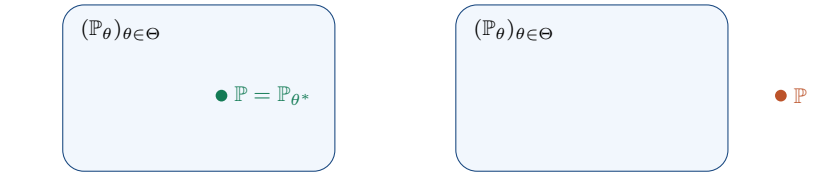
2.4 Well-specified models

It is almost always assumed that a statistical model is well-specified.

Definition 2.2 (Well-specified)**.** A model $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ is **well-specified** if there exists some $\theta \in \Theta$ such that $\mathbb{P} = \mathbb{P}_\theta$, where $\mathbb{P}$ is the true distribution that generates the observed data $X_1, \ldots, X_n$. In other words, $\mathbb{P} \in (\mathbb{P}_\theta)_{\theta\in\Theta}$. If $\mathbb{P} \notin (\mathbb{P}_\theta)_{\theta\in\Theta}$, then the model is **misspecified**. That particular θ such that $\mathbb{P} = \mathbb{P}_\theta$ is the **true parameter**, often written θ^*, and is unknown.

> **In words.** The true distribution is on your list of candidates. If it isn't, no label is right, and 'the true parameter' has no meaning.

> **For the kiss data.** If couples really do turn right with a single probability p^*, independently, the model is well-specified and $p^* \in (0, 1)$ is the true parameter. If, say, the probability differed by country, no single Bern(p) reproduces the data-generating process and the model is misspecified.



well-specified: the truth is in the family misspecified: no θ reproduces the truth

Here's an example to illustrate the meaning of well-specified models:

- Suppose you're modeling the outcome of a coin toss using a Bernoulli distribution with parameter p, where p represents the probability of getting heads. Suppose further that the *true* probability of getting heads is $p^* = 0.6$.
- Now, if your model considers all possible values of $p \in (0, 1)$, then it's a well-specified model because it includes the true probability distribution, the Bern(0.6) distribution, as one of the possible distributions in the family.
- On the other hand, if your model only considers values of p between, say, 0 and 0.5, then it's not well-specified because it doesn't include the true probability distribution, Bern(0.6), in the family of distributions.

The above example addresses misspecification when the true parameter value is not contained in your model's parameter space. But a model misspecification could also be due to applying the 'wrong' distribution: for example, by modeling the toss of a coin using a Gaussian distribution, or the siblings data of Figure 2.1 with a shifted Poisson when the true PMF has a second bump at 5.

One aim of a statistical experiment is to *estimate* the true parameter θ^*. Another is to check θ^*'s properties when they have a specific meaning determined by the context of the study. As an example of this second aim you may want to know if $\theta^* > 2$ or if $\theta^* \neq 1/2$, which is the question of Chapter 9.

2.5 Parametric, nonparametric and semiparametric models

It will often be assumed that the parameter space Θ is a subset of $\mathbb{R}^d$ for some $d \geq 1$. In this case, when $\Theta \subseteq \mathbb{R}^d$ for $d \geq 1$, the model is called **parametric**. Hence, Θ contains vectors of parameters with length d; a finite-dimensional space.

On the other hand, when Θ is infinite-dimensional, the model is called **nonparametric**. The elements of Θ are then functions (a density, a CDF) rather than vectors of numbers.

Thus, the distinction between parametric and nonparametric models is a distinction between the *length* of the vectors in the parameter space Θ. It is not a distinction of the *number* of vectors living in the space. For example, there is an uncountably infinite number of vectors living in the parameter space of a Bernoulli model where $\Theta = (0, 1)$, but all the vectors are singletons: any number between

0 and 1. What makes a model nonparametric is that no finite list of numbers can specify a member of the family; you'd need the value of a function at infinitely many points.

Nonparametric models typically include assumptions on the functions that describe the r.v. For a continuous r.v. there are two: the PDF and the CDF. Here are two examples of nonparametric models:

1. If $X_1, \dots, X_n \in \mathbb{R}$ are i.i.d. r.v.s with unknown unimodal[1] PDF f:
$$\left(\mathbb{R}, (\mathbb{P}_f)_{\{f : f \text{ is a unimodal PDF}\}}\right).$$
2. If $X_1, \dots, X_n \in [0,1]$ are i.i.d. r.v.s with unknown invertible CDF F:
$$\left([0,1], (\mathbb{P}_F)_{\{F : F \text{ is an invertible CDF}\}}\right).$$

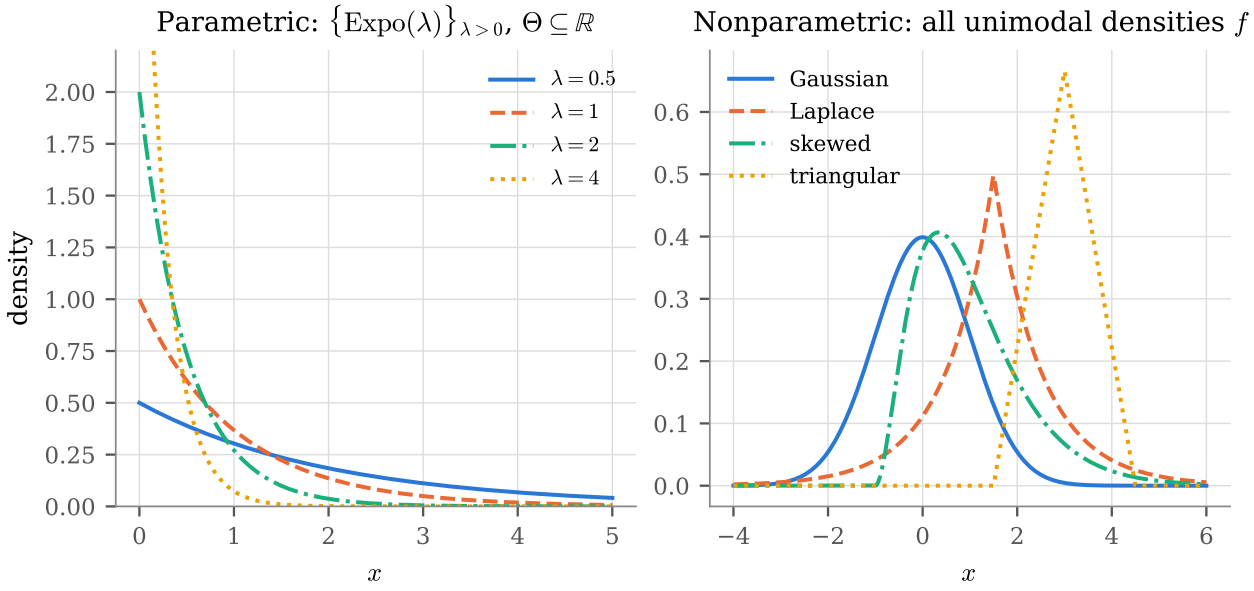


Figure 2.4: Left: a parametric family. One number, λ, moves you along a single curve through the space of densities; the whole family is this one-dimensional path. Right: four members of the nonparametric family 'all unimodal densities'. No finite set of numbers indexes them; the family is the whole space of such shapes, and these are four points in it.

A third kind: if $\Theta = \Theta_1 \times \Theta_2$ where Θ_1 is finite-dimensional and Θ_2 is infinite-dimensional, the model is **semiparametric**. In these models you only care to estimate the finite-dimensional parameter, and the infinite-dimensional one is called a **nuisance parameter**. The standard example is the Cox proportional hazards model from survival analysis, where the conditional CDF of a survival time Y given covariates $X = x$ has the form
$$F(t) = 1 - \exp\left(-\int_0^t h(u)e^{\beta^\top x}\,\mathrm{d}u\right),$$
with h an unknown non-negative nuisance function and $\beta \in \mathbb{R}^d$ the parameter of interest. This book does not develop such models; the example is here so that you have seen what a semiparametric parameter space looks like, with $\Theta_1 = \mathbb{R}^d$ for β and Θ_2 the space of non-negative functions h.

[1] A unimodal PDF increases on $(-\infty, a)$ and decreases on (a, ∞) for some a.

2.6 Identifiability

The first thing to ask of a model is whether it's well-specified. You can't verify it, since the true $\mathbb{P}$ is exactly what you don't know, but you can ask whether it's plausible, and the mathematics that follows assumes it. If it isn't, θ^* doesn't exist and 'estimating the true parameter' has no meaning; what the estimators of Part IV estimate in that case is the subject of Chapter 10. Presuming your model is well-specified, you also need to ascertain whether the parameter of the model is *identifiable*. Identifiability is a property of the parameterization: of the map from Θ to the family, not of the family itself, and not of any estimator or estimation procedure.

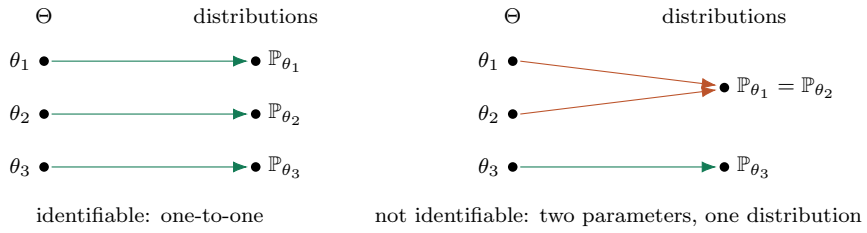
Definition 2.3 (Identifiability)**.** The parameter $\theta \in \Theta$ is **identifiable** if the map $\theta \mapsto \mathbb{P}_\theta$ is injective (one-to-one, (toolkit A.3)). That is:

$$\theta_1 \neq \theta_2 \implies \mathbb{P}_{\theta_1} \neq \mathbb{P}_{\theta_2},$$

or equivalently, taking the contrapositive:

$$\mathbb{P}_{\theta_1} = \mathbb{P}_{\theta_2} \implies \theta_1 = \theta_2.$$

> **In words.** Different labels, different distributions. Then knowing the distribution (which is what infinitely many data would tell you) pins down the label.
>
> **For the kiss data.** $p_1 \neq p_2$ gives $\text{Bern}(p_1) \neq \text{Bern}(p_2)$, since the two put different probabilities on 1; p is identifiable. The example below shows a model where two different labels give the same distribution.



This definition means that distinct values of θ should correspond to distinct probability distributions. Intuitively, if the parameter θ is identifiable you can go from $\mathbb{P}_\theta$ to θ. But if θ is *not* identifiable you can have multiple θ's that give the same $\mathbb{P}$. This means that even if you had an infinite number of observations, so that you knew $\mathbb{P}$ exactly, you'd have no way of deciding which of those θ's produced it.

Alternatively, think of it this way: if you are handed $\mathbb{P}_\theta$, can you actually recover θ? If not, if θ is not identifiable, then there are multiple θ that yield the distribution you were handed. So if you're going to use the parameter to learn about a distribution, then you must make sure the parameter is identifiable. Data can only ever tell you about $\mathbb{P}$; identifiability is what lets you translate that into a statement about θ.

The formal way of convincing someone that the parameter is *not* identifiable is to exhibit two different parameter values that give the same distribution, as the next example shows.

2.6.1 Example: a non-identifiable parameter

Example 2.4. Let $X_i = \mathbb{1}(Y_i \geq 0)$ for i.i.d. r.v.s $Y_1, \ldots, Y_n \sim \text{Gauss}(\mu, \sigma^2)$, where μ and σ^2 are unknown and the Y_i are *not observed*; you only see whether each one is non-negative. Then μ and σ^2 are *not* identifiable from the common distribution of the X_i's.

To see why, first note that, by definition, the X_i are Bernoulli r.v.s since the indicator function $\mathbb{1}(Y_i \geq 0)$ simply tells you if Y_i is non-negative or not and assigns 1 or 0 accordingly. So you want to determine whether you can find μ and σ^2 just from observing the Bernoulli r.v.s X_i.

Now, as the X_i's are Bernoulli, what is the relevant parameter? It's the probability that $Y \geq 0$, so that $X_i \sim \text{Bern}(\mathbb{P}(Y \geq 0))$, where Y represents any one of the n i.i.d. r.v.s Y_i. Because Y is a non-standard Gaussian r.v., standardize it to get:

$$\mathbb{P}(Y \geq 0) = \mathbb{P}\left(\frac{Y-\mu}{\sigma} \geq -\frac{\mu}{\sigma}\right) = 1 - \Phi\left(-\frac{\mu}{\sigma}\right) = \Phi\left(\frac{\mu}{\sigma}\right).$$

So the parameter of this Bernoulli is $p = \Phi(\mu/\sigma)$. As Φ is invertible, you can solve for μ/σ to get:

$$\frac{\mu}{\sigma} = \Phi^{-1}(p).$$

So you can learn p (it's a Bernoulli proportion), apply Φ^{-1}, and this gives you μ/σ. Hence the ratio μ/σ is identifiable.

However, μ and σ^2 by themselves are *not* identifiable. Why? One way to see it is to count dimensions. The parameter of a Bernoulli r.v. is one number between 0 and 1. The parameter of the Gaussian is two numbers, μ and σ^2. A continuous map from a two-dimensional space onto a one-dimensional one can't be one-to-one: it has to send whole curves of (μ, σ^2) pairs to the same p.

But the formal way mentioned above requires a counterexample. So, let $(\mu_1, \sigma_1^2) = (1, 4)$ and $(\mu_2, \sigma_2^2) = (2, 16)$. Then:

$$\frac{\mu_1}{\sqrt{\sigma_1^2}} = \frac{1}{\sqrt{4}} = \frac{1}{2}, \quad \text{and} \quad \frac{\mu_2}{\sqrt{\sigma_2^2}} = \frac{2}{\sqrt{16}} = \frac{1}{2}.$$

Hence different parameter vectors, $(1, 4)$ and $(2, 16)$, both yield the same Bernoulli parameter $p = \Phi(1/2) = 0.691$, so $\mathbb{P}_{(1,4)} = \mathbb{P}_{(2,16)}$ while $(1, 4) \neq (2, 16)$. Figure 2.5 shows the two Gaussians, which look nothing alike, putting the same mass above zero, and shows in the parameter space what the counterexample is a sample of: a whole ray of parameters producing one distribution.

2.6.2 More examples

Three variations are worth doing in your head because the pattern is always the same: write down the distribution of what you actually observe as a function of θ, and ask whether that function is one-to-one.

Example 2.5 (Censored Exponential)**.** Let $X \sim \text{Expo}(\lambda)$ with $\lambda > 0$ unknown, and suppose you observe only $Y = \mathbb{1}(X > 5)$, a censored version of X: you can't see X, only whether it exceeded 5. Then $Y \sim \text{Bern}(f(\lambda))$ with $f(\lambda) = \mathbb{P}(X > 5) = \int_5^\infty \lambda e^{-\lambda x}\, \mathrm{d}x = e^{-5\lambda}$. The model for Y is $\big(\{0, 1\}, (\text{Bern}(e^{-5\lambda}))_{\lambda > 0}\big)$, and λ *is* identifiable because $\lambda \mapsto e^{-5\lambda}$ is strictly decreasing, hence

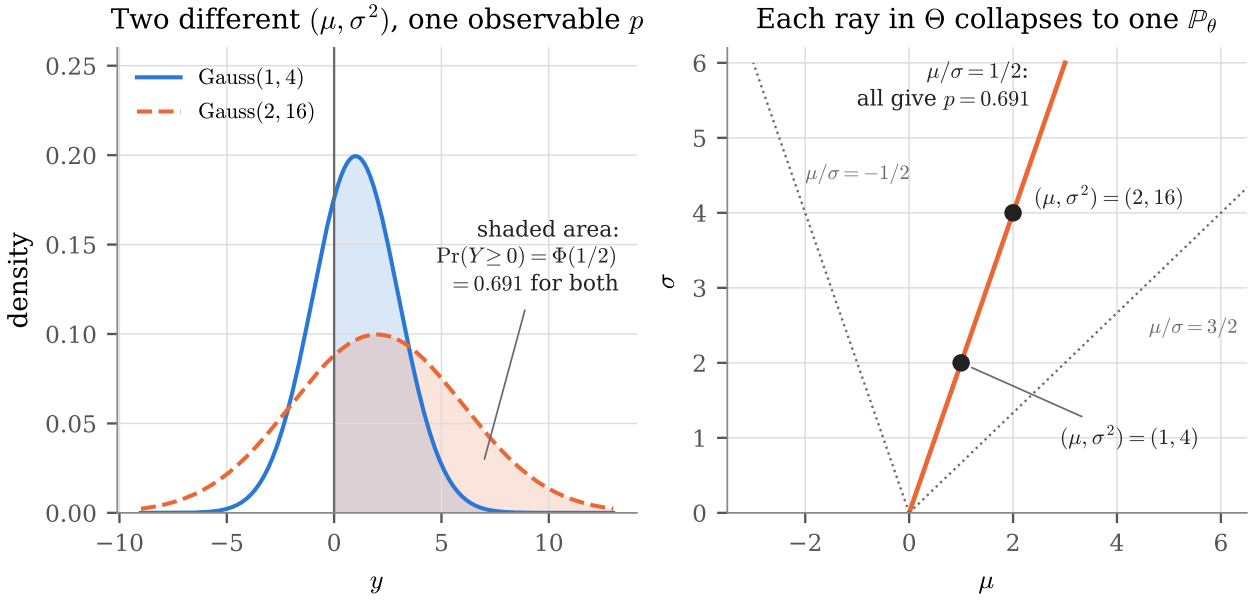


Figure 2.5: Example 2.4. Left: Gauss$(1, 4)$ and Gauss$(2, 16)$ are different distributions, but the only thing you observe is whether $Y \geq 0$, and both put probability $\Phi(1/2) = 0.691$ there. Right: in the (μ, σ) plane every point on the ray $\mu/\sigma = 1/2$ produces that same Bernoulli, and every other ray produces some other one. The observable distribution identifies the ray (that is, μ/σ) and nothing more; the counterexample in the text is two points on the orange ray.

injective: from $p = \mathbb{P}(Y = 1)$ you recover $\lambda = -\ln(p)/5$. Censoring loses information (you'll estimate λ less precisely than from the X's), but it doesn't lose identifiability. Compare Example 2.4, where the trouble was not the censoring but the two-dimensional parameter behind a one-dimensional observable.

Example 2.6 (Squaring a uniform)**.** Let $X_i = Y_i^2$ where $Y_1, \ldots, Y_n \overset{iid}{\sim} \text{Unif}([0, a])$ are unobserved. Is a identifiable? Yes: X_i takes values in $[0, a^2]$, and $\mathbb{P}(X_i \leq x) = \mathbb{P}(Y_i \leq \sqrt{x}) = \sqrt{x}/a$ for $0 \leq x \leq a^2$. Different a give different CDFs (for instance, different supports), so the map $a \mapsto \mathbb{P}_a$ is injective.

Example 2.7 (Thresholding a uniform at half its range)**.** Let $X_i = \mathbb{1}(Y_i \geq a/2)$ where $Y_1, \ldots, Y_n \overset{iid}{\sim} \text{Unif}([0, a])$ are unobserved. Is a identifiable? No: $\mathbb{P}(Y_i \geq a/2) = 1/2$ for *every* $a > 0$, so $X_i \sim \text{Bern}(1/2)$ whatever a is. The map $a \mapsto \mathbb{P}_a$ is constant, the worst possible failure of injectivity: the data carry no information about a at all.

Figure 2.6 draws the three maps side by side; identifiability is the shape of the curve.

From undergraduate statistics: non-identifiability you have already met: collinearity

In a regression course, if two predictors are perfectly collinear (say $x_2 = 2x_1$), the software drops one of them or leaves its coefficient blank. That is non-identifiability. The model $Y = \beta_1 x_1 + \beta_2 x_2 + \epsilon$ with $x_2 = 2x_1$ produces exactly the same distribution of Y for every (β_1, β_2) with the same value of $\beta_1 + 2\beta_2$: a whole line of parameter vectors, one distribution, exactly like the ray in Figure 2.5. The data can tell you $\beta_1 + 2\beta_2$ and nothing more, and no sample size fixes it. Dropping a predictor is the software's way of choosing one point on the line. Chapter 21 returns to this when it needs the design matrix to have full rank.

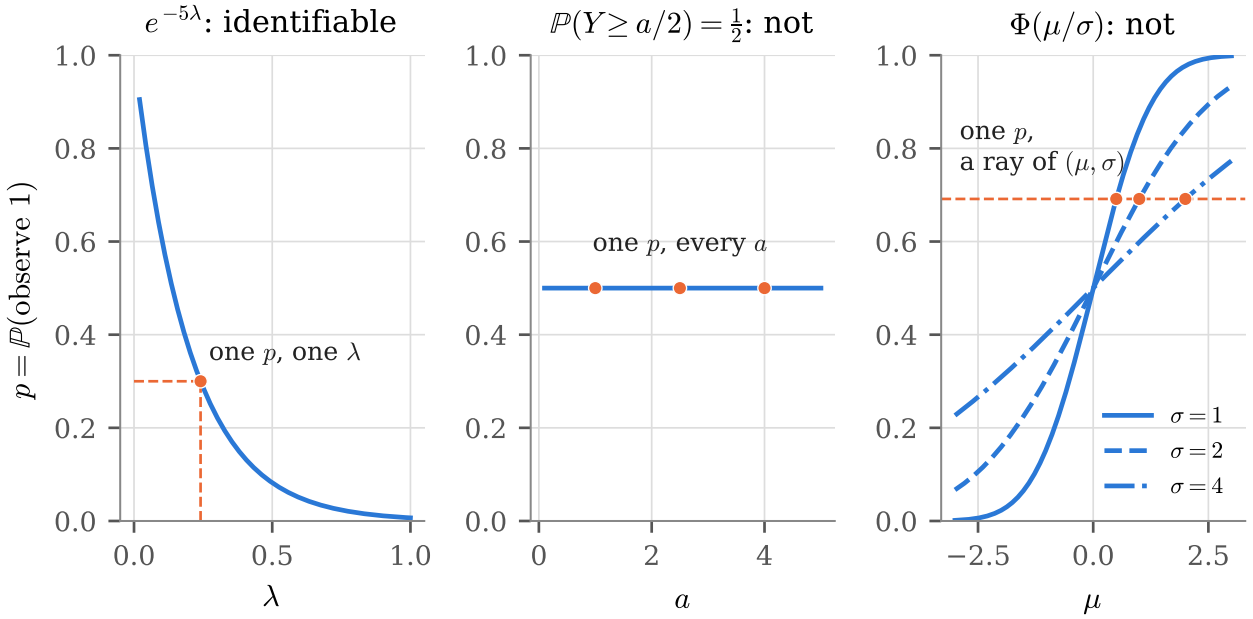


Figure 2.6: Identifiability as the shape of the map from the parameter to the distribution of what you observe. In each example the observation is a Bernoulli whose parameter p is a function of θ. Left: the censored Exponential, $p = e^{-5\lambda}$, strictly decreasing, so each p comes from exactly one λ. Middle: the uniform thresholded at half its range, $p = \frac{1}{2}$ for every a, so the observation says nothing about a. Right: the Gaussian sign example, $p = \Phi(\mu/\sigma)$, drawn against μ for three values of σ: one value of p is reached by a whole ray of (μ, σ) pairs, the ray of Figure 2.5. Identifiable means the curve never takes the same value twice.

2.7 Gaussian mixture models

A mixture of Gaussians is best described in terms of the generative model. You first generate a discrete variable that determines which of the component Gaussians to use, and then generate an observation from that chosen Gaussian density.

More formally, suppose you have two Gaussian r.v.s, $X^{(1)}$ and $X^{(2)}$, such that $X^{(1)} \sim \text{Gauss}(\mu_1, \sigma_1^2)$ and $X^{(2)} \sim \text{Gauss}(\mu_2, \sigma_2^2)$. Their PDFs are:

$$f_1(x) = \frac{1}{\sigma_1\sqrt{2\pi}} \exp\left(-\frac{(x-\mu_1)^2}{2\sigma_1^2}\right),$$
$$f_2(x) = \frac{1}{\sigma_2\sqrt{2\pi}} \exp\left(-\frac{(x-\mu_2)^2}{2\sigma_2^2}\right).$$

Let $\pi \in (0, 1)$ be the *mixture weight*. A **mixture** of the two Gaussians has the PDF:

$$\begin{aligned} f(x) &= \pi f_1(x) + (1-\pi) f_2(x) \\ &= \pi \left(\frac{1}{\sigma_1\sqrt{2\pi}}\right) \exp\left(-\frac{(x-\mu_1)^2}{2\sigma_1^2}\right) + (1-\pi)\left(\frac{1}{\sigma_2\sqrt{2\pi}}\right) \exp\left(-\frac{(x-\mu_2)^2}{2\sigma_2^2}\right). \end{aligned}$$

(The symbol π does double duty here, as the weight and as the constant 3.14159... inside the normalizing constants; both uses are standard and the context always tells them apart.)

To better understand a mixture of Gaussian distributions, think of it as a latent variable model.[2] Proceed in 3 steps:

1. **Latent component selection:** Sample $Z \sim \text{Bern}(\pi)$, so $Z = 1$ with probability π and $Z = 0$ with probability $1 - \pi$.
2. **Independent sampling:** Sample $X^{(1)} \sim \text{Gauss}(\mu_1, \sigma_1^2)$ and $X^{(2)} \sim \text{Gauss}(\mu_2, \sigma_2^2)$, independent of Z and of each other. Both values are conceptually sampled regardless of Z's value.
3. **Mixture realization:** Define the r.v. X by:

$$X = ZX^{(1)} + (1 - Z)X^{(2)}.$$

   Only one of the two is used: $X = X^{(1)}$ if $Z = 1$ and $X = X^{(2)}$ if $Z = 0$.

Question: How can this be a *mixture* if each sample comes from exactly *one* Gaussian?

Answer: The answer lies in distinguishing between:

- **Sample generation**: Each individual observation is generated from only one component, determined by Z.
- **Distribution properties**: The overall probability distribution of X is a weighted mixture of both Gaussians because you don't get to see Z.

Mathematically, for any possible value x, the law of total probability over Z gives the density

$$f(x) = \mathbb{P}(Z = 1)f_1(x) + \mathbb{P}(Z = 0)f_2(x) = \pi f_1(x) + (1 - \pi)f_2(x),$$

so both components contribute to the overall distribution, weighted by their probabilities. From the population perspective, if you generate n samples from this model, approximately $n\pi$ of them come from $\text{Gauss}(\mu_1, \sigma_1^2)$ and approximately $n(1 - \pi)$ from $\text{Gauss}(\mu_2, \sigma_2^2)$. Figure 2.7 shows both perspectives.

An alternative specification of the same r.v. is $X = (X^{(1)})^Z (X^{(2)})^{1-Z}$, which picks the same component as before. The latent variable perspective is especially useful for inference algorithms like EM (expectation-maximization), where you need to estimate both the component parameters $(\mu_1, \sigma_1^2, \mu_2, \sigma_2^2)$ and the mixture weight π without ever observing which component produced which observation.

2.7.1 Possible Gaussian mixture models

There are several scenarios for a model with mixtures of Gaussians, depending on the constraints imposed on its 5 parameters. Here are 3 possible examples. It should be noted that, for clarity but with an abuse of notation, the PDFs of the two Gaussian r.v.s are not explicitly written below; each distribution is written generically as $\text{Gauss}(\mu_1, \sigma_1^2)$.

1. Model where all 5 parameters are unknown:

$$\Big(\mathbb{R}, (\pi \cdot \text{Gauss}(\mu_1, \sigma_1^2) + (1 - \pi) \cdot \text{Gauss}(\mu_2, \sigma_2^2))_{\pi \in (0,1),\, \mu_1, \mu_2 \in \mathbb{R},\, \sigma_1^2, \sigma_2^2 \in (0,\infty)}\Big).$$

[2] Recall that a latent variable is an unobserved variable that can only be inferred indirectly from observable variables.

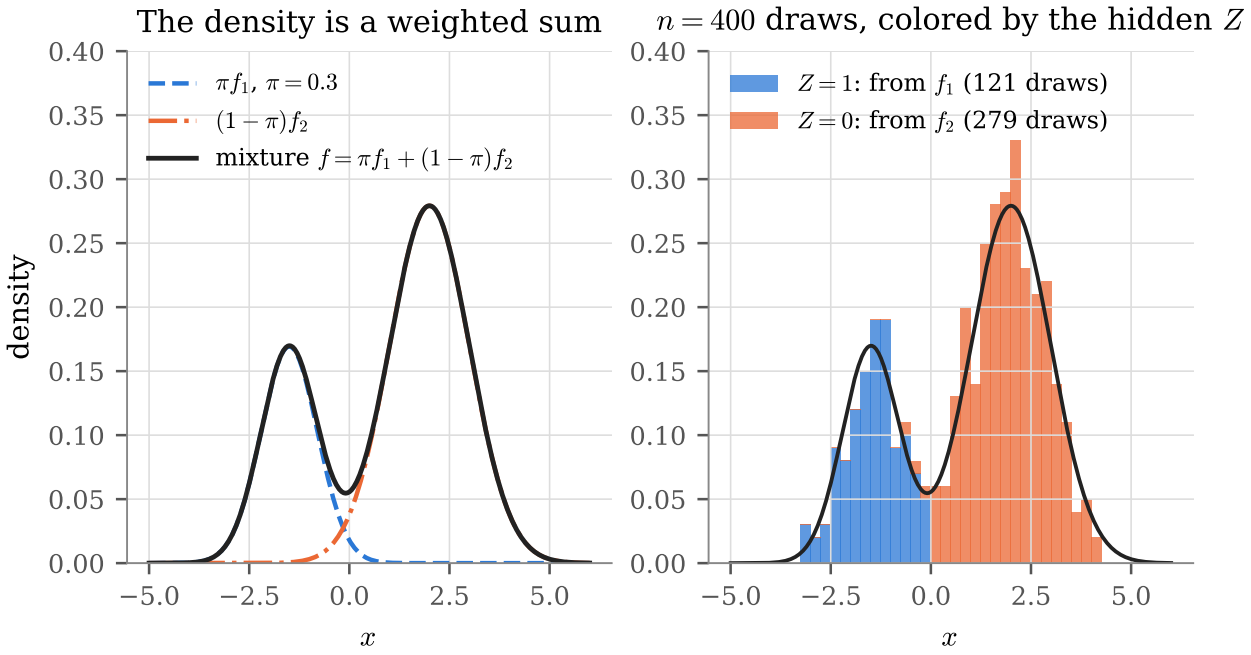


Figure 2.7: A two-component Gaussian mixture with $\pi = 0.3$, Gauss$(-1.5, 0.5)$ and Gauss$(2, 1)$. Left: the mixture density is the sum of the two weighted components. Right: 400 draws from the generative model, histogrammed and colored by the latent Z that produced each one. Each draw came from exactly one component; the mixture density (black) is what you see when you cannot tell which.

2. Model with known variances, and for simplicity suppose $\sigma_1^2 = \sigma_2^2 = 1$:

$$\big(\mathbb{R}, (\pi \cdot \text{Gauss}(\mu_1, 1) + (1-\pi) \cdot \text{Gauss}(\mu_2, 1))_{\pi \in (0,1),\, \mu_1, \mu_2 \in \mathbb{R}}\big)\,.$$

3. Model with only unknown means, and for simplicity suppose $\sigma_1^2 = \sigma_2^2 = 1$ and $\pi = 1/2$:

$$\Big(\mathbb{R}, \big(\tfrac{1}{2} \cdot \text{Gauss}(\mu_1, 1) + \tfrac{1}{2} \cdot \text{Gauss}(\mu_2, 1)\big)_{\mu_1, \mu_2 \in \mathbb{R}}\Big)\,.$$

Pitfall: mixtures are identifiable only up to relabeling

Take the third model and the two parameter vectors $(\mu_1, \mu_2) = (-1, 2)$ and $(\mu_1, \mu_2) = (2, -1)$. They're different points of $\Theta = \mathbb{R}^2$, and they give exactly the same density: $\frac{1}{2} f_{-1} + \frac{1}{2} f_2 = \frac{1}{2} f_2 + \frac{1}{2} f_{-1}$. So by Definition 2.3 the parameter is not identifiable; the same happens in the first two models with $(\pi, \mu_1, \sigma_1^2, \mu_2, \sigma_2^2) \leftrightarrow (1-\pi, \mu_2, \sigma_2^2, \mu_1, \sigma_1^2)$. This failure is called *label switching*, and it's harmless in a way Example 2.4 is not: the two parameter vectors describe the same two components with their names swapped, so the data do pin down the components, just not which one is 'first'. The fix is a convention such as $\mu_1 < \mu_2$, which restricts Θ to half the plane and removes the swap; that nothing else collides, so that the map is then injective, is a theorem about finite Gaussian mixtures that is not proved here. The lesson is that non-identifiability is sometimes a defect of the labeling rather than of the information in the data, and the first thing to ask when you find it is which of the two you're looking at.

2.8 What comes next

A model fixes the family of candidate distributions. Three things are done inside it, and they are Part III: guess θ (Chapter 7), attach a range to the guess (Chapter 8), and decide a yes/no question about θ (Chapter 9). All three are stated for a general model $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ and computed, as always, for the kiss model. All three also rest on what an average of n draws does as n grows, and that is Part II, which comes first: where the average settles (Chapter 3), how it fluctuates on the way (Chapter 4), what may be done to those limits (Chapter 5), and the Gaussian distribution the limits land on (Chapter 6).

2.9 Summary

- A statistical model is $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$: a sample space and a family of candidate distributions for X, indexed by a parameter. Parameters define distributions; learning one is learning the other. The sample space must not depend on θ.
- Assumptions buy fewer parameters: the siblings PMF has six, the shifted Poisson one. The price is misspecification if the assumption is wrong.
- Three layers: the sample space Ω (every complete outcome), random variables (functions on Ω, the measuring instruments and the rules applied to their readings), and the observations (those functions evaluated at the one outcome that happened). An estimate is one draw from the estimator's distribution. Inference does not invert the chain; it holds the one value up against the distribution of the rule. For fixed n, take $\Omega = E^n$.
- Well-specified: the true $\mathbb{P}$ is some $\mathbb{P}_{\theta^*}$ in the family. Everything from Part III on assumes this; Chapter 10 says what estimators estimate when it fails.
- Parametric: $\Theta \subseteq \mathbb{R}^d$. Nonparametric: Θ is a space of functions. Semiparametric: a finite-dimensional parameter of interest plus an infinite-dimensional nuisance parameter.
- Identifiable: $\theta \mapsto \mathbb{P}_\theta$ is injective, so knowing the distribution of the data would tell you θ. To show a parameter is not identifiable, exhibit two parameter values with the same distribution. Observing only a function of the data can destroy identifiability (two numbers behind one observable proportion) or leave it intact (one number behind one proportion, injectively); what matters is whether the map from the parameter to the observable distribution is one-to-one.
- A Gaussian mixture is a latent-variable model: pick a component with $Z \sim \text{Bern}(\pi)$, then draw from it. Its density is $\pi f_1 + (1-\pi) f_2$. Its parameter is identifiable only up to swapping the component labels.

2.10 Exercises

Exercise 2.1. Which of the following is a statistical model? (1) $\big(\{1\}, (\text{Bern}(p))_{p\in(0,1)}\big)$; (2) $\big(\{0,1\}, (\text{Bern}(p))_{p\in(0.2,0.4)}\big)$; (3) both; (4) neither. *(Solution on page 221.)*

Exercise 2.2. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0,a])$ for some unknown $a > 0$. Which is the associated statistical model: (1) $([0,a], (\text{Unif}([0,a]))_{a>0})$; (2) $(\mathbb{R}_+, (\text{Unif}([0,a]))_{a>0})$; (3) $(\mathbb{R}, (\text{Unif}([0,a]))_{a>0})$; (4) none? *(Solution on page 221.)*

Exercise 2.3. Which of these models are parametric? (A) $E = \{0,1,2,\ldots\}$ and $\{\mathbb{P}_\theta\}$ is the set of all distributions on E. (B) $E = \{0,1\}$, $(\text{Bern}(\theta))_{\theta\in[0,1]}$. (C) $E = \mathbb{R}$, $(\text{Gauss}(0,\sigma^2))_{\sigma^2>0}$. (D)

$E = \{1, 2, 3, 4\}$ and $\mathbb{P}_{(p_1,\ldots,p_4)}$ puts probability p_i on i, with $p_i \geq 0$ and $\sum p_i = 1$. (E) $E = \mathbb{R}$, $(\text{Gauss}(\mu, \sigma^2))_{\mu\in\mathbb{R},\sigma^2>0}$. (F) $E = [0, \infty)$, $(\text{Unif}([0, \theta]))_{\theta>0}$. (G) $E = [0, 1]$ and $\{\mathbb{P}_f\}$ is the set of all distributions with a continuous density f on $[0, 1]$. *(Solution on page 222.)*

Exercise 2.4. You observe $X_i = \mathbb{1}(Y_i \geq c)$ for $Y_i \overset{iid}{\sim} \text{Gauss}(\mu, 1)$ with $\sigma^2 = 1$ known and c a known constant. Is μ identifiable? What if instead σ^2 is unknown and $\mu = 0$ is known? *(Solution on page 222.)*

Exercise 2.5. In the two-component mixture with all five parameters unknown, is the parameter $(\pi, \mu_1, \sigma_1^2, \mu_2, \sigma_2^2)$ identifiable on the restricted parameter space where $\mu_1 < \mu_2$? What about the point $\pi = 1$, if it were allowed? *(Solution on page 222.)*

Part II

Probability for Statistics

The three questions of Chapter 1 are about the random variable $\bar{R}_n$ as n grows, and 'as n grows' is a limit. This part has the three limit facts every later chapter uses, the Law of Large Numbers (Chapter 3), the Central Limit Theorem (Chapter 4) and the rules for working with limits (Chapter 5), and then the Gaussian distribution and the quantile notation that the rest of the book is written in (Chapter 6). If limits are not your difficulty, Chapter 3's opening box says what can be skipped.

Chapter 3

The Law of Large Numbers

In this chapter

- Why limits of random variables need different tools than limits of numbers.
- The sample mean on two scales: collapsing on μ (this chapter) and, magnified by $\sqrt{n}$, settling into a Gaussian (the next).
- Convergence in probability, the mode of convergence that says an estimator gets close to what it estimates.
- The Weak Law of Large Numbers, with Chebyshev's inequality as the proof.

The running example is the kiss experiment of Chapter 1: $n = 124$ couples, 80 of whom turned right, so $\hat{p} = \bar{R}_n = 0.645$.

Chapter 2 fixed the object inference happens in; this part supplies the three facts about averages that Part III stands on.

The next three chapters. The rest of the book rests on three limit facts about averages: the Law of Large Numbers (averages settle on their expectations, this chapter), the Central Limit Theorem (the error of an average is Gaussian on the $1/\sqrt{n}$ scale, Chapter 4), and the rules for transforming limits and plugging in estimates (continuous mapping, Slutsky, the delta method, Chapter 5). Those three chapters are the minimal path. Two more modes of convergence, almost sure and in mean, complete the standard hierarchy; that material is in Appendix B because nothing later is proved with it; where a later chapter mentions almost sure convergence, the in-probability version is the one in use. Read it once for orientation and don't spend time mastering it.

Intuition: what 'asymptotic' means

An asymptotic result is a license to replace an unknown sampling distribution by its limiting distribution, usually a Gaussian, with an error the theorem does not tell you. That is all 'asymptotic' ever means in this book. The definitions below are about infinite sequences, while every application is about one fixed sample size, usually $n = 124$; the theorems say $n \to \infty$ and they get applied at $n = 64$ without further comment. The gap is real, and Section 4.4 in the next chapter says what the alternatives are and why statistical practice accepts the trade. When a later chapter says 'by the CLT' or 'asymptotically normal', read it as: for the n at hand, treat this as Gaussian and carry on.

3.1 Why limits of random variables need new tools

Chapter 1 asked whether $\hat{p}$ is close to p once n is large. That is a question about a limit, and limits of random variables need their own definition. In calculus you say that a sequence of real numbers x_n converges to a limit x if, for any $\epsilon > 0$, $|x_n - x| < \epsilon$ for all large n (toolkit A.1). Convergence is more subtle in probability. Returning to calculus, suppose that $x_n = x$ for all n. Then, trivially, $\lim_{n\to\infty} x_n = x$.

Consider a probabilistic version of this example. Suppose that $X_1, X_2, \ldots$ is a sequence of i.i.d. r.v.s such that $X_i \overset{iid}{\sim} \text{Gauss}(0, 1)$. Since these all have the same distribution, you might think X_n 'converges' to $X \sim \text{Gauss}(0, 1)$. But this can't be right: $X_n - X \sim \text{Gauss}(0, 2)$ for every n, so $\mathbb{P}(|X_n - X| < \epsilon)$ is the same small number for every n and never improves. (In fact $\mathbb{P}(X_n = X) = 0$ for all n, since two independent continuous r.v.s are equal with probability zero.[1]) This suggests limiting processes involving probabilities require different tools than those limiting processes from calculus.

The central ideas in this chapter and the next are the two most famous results in all of probability theory: the Law(s) of Large Numbers and the Central Limit Theorem. An understanding of why these results are valid requires an understanding of limiting processes that involve r.v.s. These limits are conceptually more general than the limits of undergraduate calculus since, as limits of r.v.s, these limits are not limits of numbers but limits of *functions*.

So what is meant by 'convergence of random variables'? And what is a sequence of r.v.s? Recall that, by definition, an r.v. is a mapping from a sample space, S, to the real number line (toolkit A.5). Suppose, to fix ideas, that the outcomes can be listed (a continuous r.v. needs a richer sample space, but the picture is the same), so you can write:

$$S = \{s_1, \ldots, s_k\}.$$

This means an r.v. X defined on S can be written as:

$$X(s_i) = x_i, \text{ for } i = 1, \ldots, k.$$

When you have a sequence of r.v.s it's useful to keep in mind that there's always an underlying sample space. In particular, each X_n is a function from S to $\mathbb{R}$:

$$X_n(s_i) = x_{n,i}, \text{ for } i = 1, \ldots, k.$$

So a sequence of r.v.s is a sequence of functions $X_n : S \to \mathbb{R}$. Pick an outcome s_i and you get an ordinary sequence of numbers $x_{1,i}, x_{2,i}, x_{3,i}, \ldots$, one such sequence for every outcome. A *realization* of the sequence is what you see when the experiment is run once and the outcome happens to be s_i.

This chapter and the next examine the two modes of convergence this book works with; Appendix B has the full set. Why are there different modes of convergence? Remember that to say a sequence

[1] Why is this true? Let Y_1 and Y_2 be continuous r.v.s with joint density f_{Y_1,Y_2}. Then

$$\mathbb{P}(Y_1 = Y_2) = \iint_D f_{Y_1,Y_2}(y_1, y_2)\, \mathrm{d}y_1\, \mathrm{d}y_2, \qquad D = \{(y_1, y_2) \in \mathbb{R}^2 : y_1 = y_2\}.$$

The set D is a line living in two dimensions and hence has no area: you're integrating a bounded function over a region of zero area, so the integral is 0. This is the two-dimensional version of the fact that the 'area' under a density at a single point is zero.

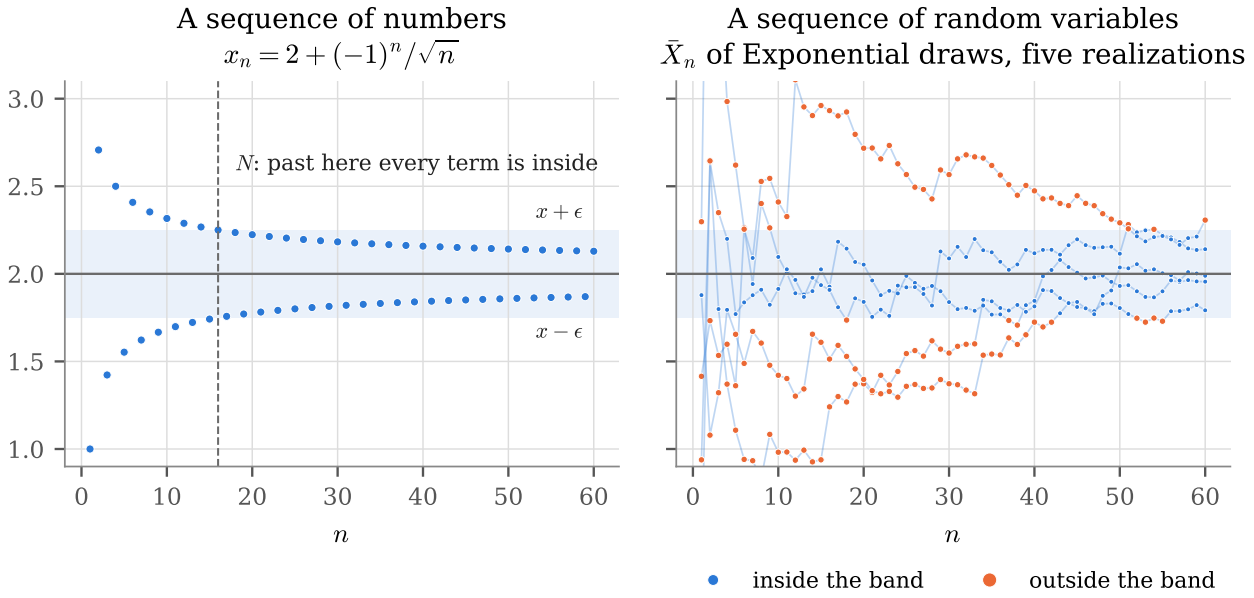


Figure 3.1: Left: a sequence of numbers converging to 2. For the band of half-width ϵ there is an index N past which every term is inside. Right: five realizations of the running sample mean $\bar{X}_n$ of Exponential draws with mean 2, one path per outcome of the experiment. Each path eventually settles inside the band, but at a different index, and no N works for every path: some realization will leave the band after any N you name. The modes of convergence are different ways of saying 'the orange points become rare.'

$(X_i)_{i=1}^{\infty}$ converges to X means that the X_i's are getting 'closer and closer' to X. Different modes of convergence relate to different ways of defining what 'closer' means. Figure 3.1 shows the difficulty: for a sequence of numbers there is a single index N past which every term is within ϵ of the limit, while for a sequence of r.v.s the index that works depends on the outcome of the experiment and may not exist for some outcomes.

Two of the modes, convergence almost surely and convergence in probability, compare the *values* of X_n and X outcome by outcome, so X_n and X must live on the same sample space. The third, convergence in distribution, compares the CDFs of X_n and X and never looks at values at all; X_n and X can be defined on entirely different probability spaces. But as r.v.s have CDFs, the difference between sequences of r.v.s and sequences of CDFs isn't as significant as you might expect.

So it's important to emphasize here that the 'things' converging are r.v.s, or their CDFs, and that they converge to an r.v., or a CDF.[2]

3.2 A first look: the sample mean on two scales

To motivate the notion of convergence of r.v.s, consider the following example. Suppose that $X_1, \ldots, X_n$ are i.i.d. r.v.s with mean μ and variance σ^2 and define the (theoretical, non-realized)

[2]Just because a sequence of r.v.s converges to an r.v. doesn't mean the limiting r.v. has an interesting distribution or what you normally think of as a 'distribution'. In fact, one mode of convergence explained below (convergence in probability), is almost entirely concerned with the convergence of sequences of r.v.s to a constant or 'degenerate' r.v.

sample mean:

$$\bar{X}_n = \frac{1}{n}\sum_{i=1}^{n} X_i.$$

First, consider the behavior of the distribution of $\bar{X}_n$ when n is large. It seems reasonable that $\bar{X}_n$ will be close to μ if n is sufficiently large; that is, the r.v. $\bar{X}_n - \mu$ should have a distribution that, for large n, is concentrated around 0 or, more precisely:

$$\mathbb{P}(|\bar{X}_n - \mu| \leq \epsilon) \approx 1,$$

for any small positive number ϵ. This type of limiting process, and more specifically, that as n increases without limit the probability equals 1, is an example of *convergence in probability*, and also illustrates a convergence process to a constant (here, μ), as discussed in the previous footnote.

Moreover because $\mathbb{V}\text{ar}(\bar{X}_n) = \sigma^2/n$,[3] it follows that $\mathbb{V}\text{ar}(\bar{X}_n) \to 0$ as $n \to \infty$. Because the variance of $\bar{X}_n$ eventually vanishes as n increases, the limit tells you nothing about the *shape* of the distribution of $\bar{X}_n$: every density in the left panel of Figure 3.2 is on its way to becoming a spike at μ.

It's possible though to look at the difference between $\bar{X}_n$ and μ on a 'magnified' scale. To that end, multiply the difference $\bar{X}_n - \mu$ by $\sqrt{n}$, which makes the mean and variance of the resulting r.v. constant. Specifically, define the r.v. Y_n by:

$$Y_n = \sqrt{n}(\bar{X}_n - \mu),$$

and note that:

$$\mathbb{E}[Y_n] = \mathbb{E}[\sqrt{n}(\bar{X}_n - \mu)] = \sqrt{n}(\mathbb{E}[\bar{X}_n] - \mu) = \sqrt{n}(\mu - \mu) = 0,$$

and

$$\mathbb{V}\text{ar}(Y_n) = \mathbb{V}\text{ar}(\sqrt{n}(\bar{X}_n - \mu)) = n\,\mathbb{V}\text{ar}(\bar{X}_n - \mu) = n\,\mathbb{V}\text{ar}(\bar{X}_n) = \sigma^2.$$

Intuition: what's special about $\sqrt{n}$?

Why does multiplying by $\sqrt{n}$, and not by n or by $\log n$, yield a constant mean and variance?

- **The mean:** Multiplying $\bar{X}_n - \mu$ by anything doesn't impact its expected value because $\bar{X}_n$ is an unbiased estimator of μ, so the expected value of the difference is 0 to begin with.
- **The variance:** As n increases, $\mathbb{V}\text{ar}(\bar{X}_n) = \sigma^2/n$ decreases. When you multiply an r.v. by $\sqrt{n}$ the variance is scaled by a factor of n, which exactly cancels the division by n. Multiplying by n instead would scale the variance by n^2 and the result would blow up; multiplying by $\log n$ would not be enough and the result would still collapse to a spike.

The standard deviation of $\bar{X}_n$ is $\sigma/\sqrt{n}$, so $\sqrt{n}$ is the magnification that holds the spread fixed.

You can now consider the behavior of the CDF of Y_n as n increases. If this sequence of CDFs has a limit, then you can use the limiting CDF to approximate the CDF of Y_n (and hence of $\bar{X}_n$).

[3] That $\mathbb{V}\text{ar}(\bar{X}_n) = \sigma^2/n$ follows from the properties of the variance and the independence of the X_i. Specifically:

$$\mathbb{V}\text{ar}(\bar{X}_n) = \frac{1}{n^2}\,\mathbb{V}\text{ar}(X_1 + \cdots + X_n) = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}.$$

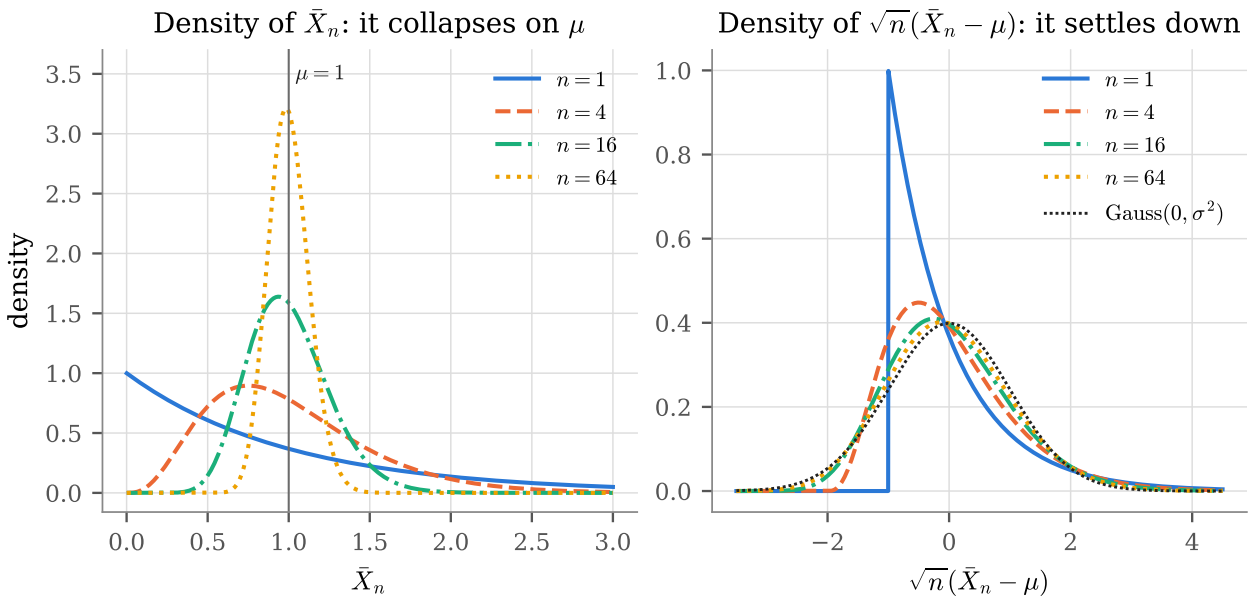


Figure 3.2: The same sample mean of Expo(1) draws on two scales. Left: the density of $\bar{X}_n$ collapses onto $\mu = 1$ as n grows (this is the Law of Large Numbers). Right: after the $\sqrt{n}$ magnification the density of $\sqrt{n}(\bar{X}_n - \mu)$ keeps a fixed spread and its shape settles down to a Gaussian (this is the Central Limit Theorem). At $n = 1$ the magnified density is the skewed exponential itself, shifted; by $n = 64$ it is hard to tell from Gauss$(0, \sigma^2)$.

For example, suppose you have that:

$$\mathbb{P}(Y_n \leq y) = \mathbb{P}(\sqrt{n}(\bar{X}_n - \mu) \leq y) \approx F_0(y).$$

Then, to say something about the CDF of $\bar{X}_n$, subtract μ from $\bar{X}_n$ and multiply this difference by $\sqrt{n}$, to get:

$$\mathbb{P}(\bar{X}_n \leq x) = \mathbb{P}(\sqrt{n}(\bar{X}_n - \mu) \leq \sqrt{n}(x - \mu)) \approx F_0(\sqrt{n}(x - \mu)),$$

provided n is sufficiently large to make the approximation valid. This type of limiting process is an example of *convergence in distribution*, which is the subject of Chapter 4.

What if instead of transforming $\bar{X}_n$ into Y_n, you 'standardized' it by dividing Y_n by σ. That is, define the r.v. Z_n by:

$$Z_n = \sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right).$$

Now, calculate $\mathbb{E}[Z_n]$ and $\mathbb{V}\text{ar}(Z_n)$:

$$\mathbb{E}[Z_n] = \mathbb{E}\left[\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right)\right] = \frac{\sqrt{n}}{\sigma}(\mathbb{E}[\bar{X}_n] - \mu) = \frac{\sqrt{n}}{\sigma}(\mu - \mu) = 0.$$

Thus, $\mathbb{E}[Y_n] = \mathbb{E}[Z_n] = 0$. But observe that $\mathbb{V}\text{ar}(Z_n) \neq \mathbb{V}\text{ar}(Y_n) = \sigma^2$:

$$\begin{aligned}
\mathbb{V}\text{ar}(Z_n) &= \mathbb{V}\text{ar}\left(\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right)\right) \\
&= \frac{n}{\sigma^2}\,\mathbb{V}\text{ar}\left(\bar{X}_n\right) \\
&= \frac{n}{\sigma^2}\,\mathbb{V}\text{ar}\left(\frac{1}{n}\sum_{i=1}^{n} X_i\right) \\
&= \frac{1}{n\sigma^2}\,\mathbb{V}\text{ar}\left(\sum_{i=1}^{n} X_i\right) \\
&= \frac{n\sigma^2}{n\sigma^2} \\
&= 1.
\end{aligned}$$

Thus, despite the extra scaling factor, $\sqrt{n}$, the r.v. Z_n has a 'standard' look to it: Its mean is 0 and its variance is 1. It may not come as a surprise that Z_n converges in distribution to a standard Gaussian distribution.

From undergraduate statistics: the standard error

An introductory course states, usually without proof, that the sampling distribution of the sample mean has mean μ, has standard deviation $\sigma/\sqrt{n}$ (the *standard error* of the mean), and is approximately normal when n is large. The two panels of Figure 3.2 are the two halves of that statement. The standard error shrinking to zero is what drives the Law of Large Numbers. The shape becoming normal is the Central Limit Theorem. Dividing $\bar{X}_n - \mu$ by the standard error $\sigma/\sqrt{n}$ is exactly the standardization that produces Z_n:

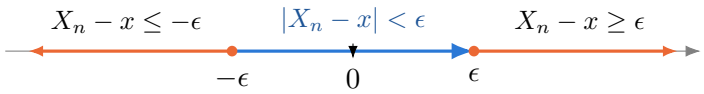
$$Z_n = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} = \sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right).$$

This chapter proves the first half; Chapter 4 makes 'approximately normal when n is large' into a theorem and says what 'large' buys you.

3.3 Convergence in probability

Convergence in probability considers the distance between a sequence of r.v.s X_n and some number x being at least as large as some arbitrary $\epsilon > 0$. That is, in formal terms you're concerned with whether the probability of $|X_n - x| \geq \epsilon$ converges to 0.

A first step in such problems, and limit problems generally, is to grasp the meaning of $|X_n - x| \geq \epsilon$. An absolute value can be thought of as the distance from 0. So $|X_n - x|$ is the distance of the number $X_n - x$ from 0, and the event $|X_n - x| \geq \epsilon$ says that this distance is at least ϵ, which happens either when $X_n - x \geq \epsilon$ or when $X_n - x \leq -\epsilon$. A visualization along the real line is instructive.



If a sequence of r.v.s converges in probability to a number, say, 0, then you say the sequence converges to the zero *random variable* (not the number 0). Now for the definition.

Definition 3.1 (Convergence in probability)**.** A sequence of r.v.s, $(X_n)_{n\geq 1}$, **converges in probability** to an r.v. X if, for every $\epsilon > 0$:

$$\lim_{n\to\infty} \mathbb{P}(|X_n - X| \geq \epsilon) = 0, \text{ or equivalently } \lim_{n\to\infty} \mathbb{P}(|X_n - X| < \epsilon) = 1.$$

If $(X_n)_{n\geq 1}$ converges in probability to some r.v. X, then this is denoted by $X_n \xrightarrow[n\to\infty]{\mathbb{P}} X$.

> **In words.** Read the definition from the inside out. Fix a tolerance ϵ. For each n, $\mathbb{P}(|X_n - X| \geq \epsilon)$ is a number, the probability mass that X_n puts outside the ϵ-band around X. The definition says this sequence of *numbers* goes to 0 in the ordinary calculus sense (toolkit A.1), and that it must do so for every tolerance, however small. Nothing is said about any single outcome; everything is said about probabilities.
>
> **For the kiss data.** $\bar{R}_n \xrightarrow[n\to\infty]{\mathbb{P}} p$ means: pick any tolerance, say 0.01; the probability that the sample proportion turning right misses p by more than 0.01 can be made as small as you like by observing enough couples. It does not say how many couples; Section 3.4 and Chapter 4 do.

Figure 3.3 shows the picture for the sample mean, where X is the constant μ.

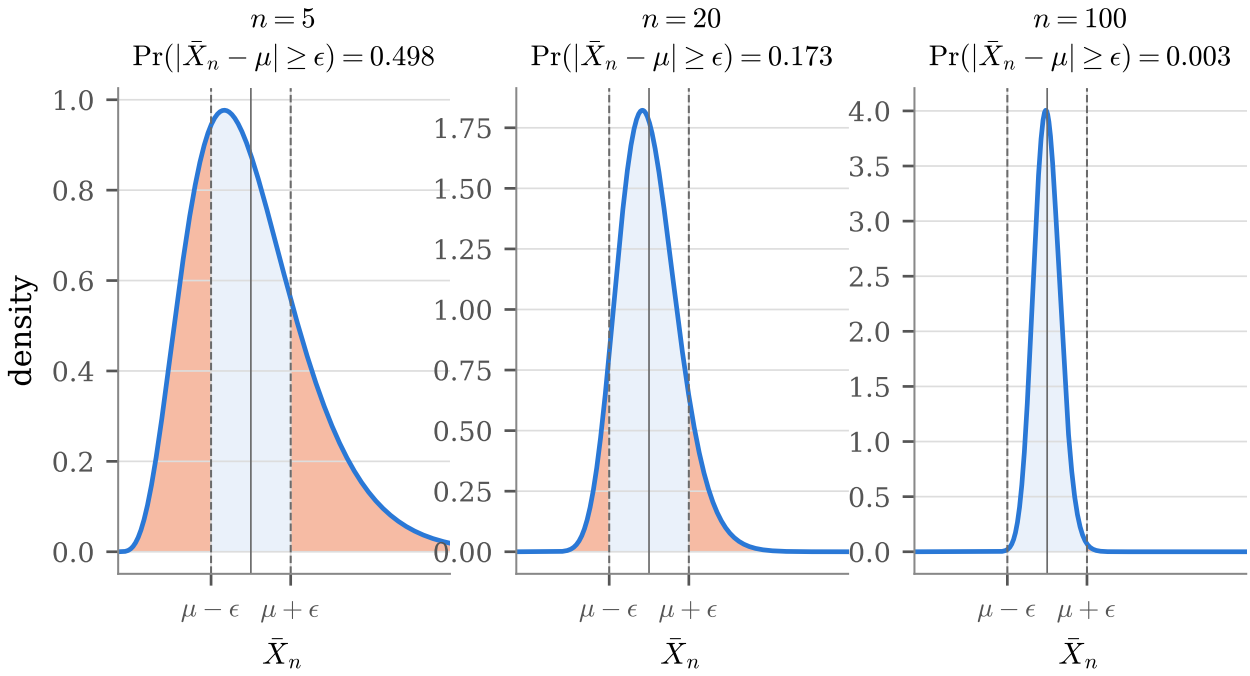


Figure 3.3: Convergence in probability as a picture. The density of $\bar{X}_n$ for Expo(1) draws at three sample sizes, with the band $[\mu - \epsilon, \mu + \epsilon]$ for $\epsilon = 0.3$. The shaded orange mass is $\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon)$; it goes from about one half at $n = 5$ to essentially zero at $n = 100$. Shrink ϵ and the same thing happens, only it takes a larger n.

Though convergence in probability is defined as a sequence of r.v.s converging to an r.v., in statistical applications the limit is almost always a constant: the sample mean converges to μ, an estimator

converges to the parameter it estimates. There's a reason the interesting limits are constants. If the r.v.s in the sequence are independent and the sequence converges in probability, then the limiting r.v. must be a constant.[4] The sample mean $\bar{X}_n$ isn't a sequence of independent r.v.s (each $\bar{X}_n$ shares X_1 with all the others), so this fact doesn't apply to it directly, but the conclusion is the same: the randomness averages out.

3.3.1 Example: the minimum of uniforms

Example 3.2. Let $X_1, X_2, \ldots \overset{iid}{\sim} \text{Unif}([0,1])$ and define the minimum of the first n by $X_{(1)} = \min_{i=1,\ldots,n} X_i$. Then:

$$X_{(1)} \xrightarrow[n\to\infty]{\mathbb{P}} 0.$$

Proof strategy (core). Write the event $|X_{(1)} - 0| \geq \epsilon$ as an event about the whole sample, use independence to turn its probability into a product of n equal factors, and let $n \to \infty$. This is the template for every direct convergence proof that follows.

Proof.

- Let $\epsilon > 0$. Write the claim using the definition of convergence in probability: you need $\mathbb{P}(|X_{(1)} - 0| \geq \epsilon) \to 0$.
- Now, you'll simplify the expression in such a way as to eventually take the limit of the expression as $n \to \infty$.
- To that end, observe that:

$$\begin{aligned}
\mathbb{P}(|X_{(1)} - 0| \geq \epsilon) &= \mathbb{P}(|X_{(1)}| \geq \epsilon) \\
&= \mathbb{P}(X_{(1)} \geq \epsilon) \\
&= \mathbb{P}(X_1 \geq \epsilon, X_2 \geq \epsilon, \ldots, X_n \geq \epsilon) \\
&= \mathbb{P}(X_1 \geq \epsilon)\mathbb{P}(X_2 \geq \epsilon) \cdots \mathbb{P}(X_n \geq \epsilon) \\
&= \mathbb{P}(X_1 \geq \epsilon)^n.
\end{aligned}$$

  The movement from the second to the third equality holds because, as $X_{(1)}$ is the minimum of the X_i's, the minimum is at least ϵ exactly when every X_i is at least ϵ: the single event is equivalent to the joint event. The fourth equality results since the X_i's are independent, and the fifth since they're identically distributed.

- You know that $X_1 \sim \text{Unif}([0,1])$. Therefore, $f_{X_1}(x) = 1$ for all values in X_1's support.

[4] A sketch, not a proof, and not needed later: the limit of a sequence can't depend on any finite number of its terms (throw away the first million terms and the limit is unchanged). If the terms are independent, the limit is therefore independent of every one of them, hence independent of everything that's random in the experiment, hence independent of itself. An r.v. independent of itself satisfies $\mathbb{V}\text{ar}(X) = \mathbb{C}\text{ov}(X, X) = 0$, so it's a constant. The precise statement is Kolmogorov's zero-one law.

- This implies that $\mathbb{P}(X_1 \geq \epsilon) = \int_{\epsilon}^{1} 1 \, \mathrm{d}x = 1 - \epsilon$, and further that:

$$\mathbb{P}(X_1 \geq \epsilon)^n = (1 - \epsilon)^n.$$

  Also, note that $1 - \epsilon$ is some number between 0 and 1 (for $\epsilon \geq 1$ the probability is already 0, so there's nothing to prove).

- Now you're ready to take the limit of this probability and confirm the statement:

$$\lim_{n \to \infty} \mathbb{P}(|X_{(1)} - 0| \geq \epsilon) = \lim_{n \to \infty} (1 - \epsilon)^n = 0,$$

  and hence $X_{(1)} \xrightarrow[n \to \infty]{\mathbb{P}} 0$, as desired.

□

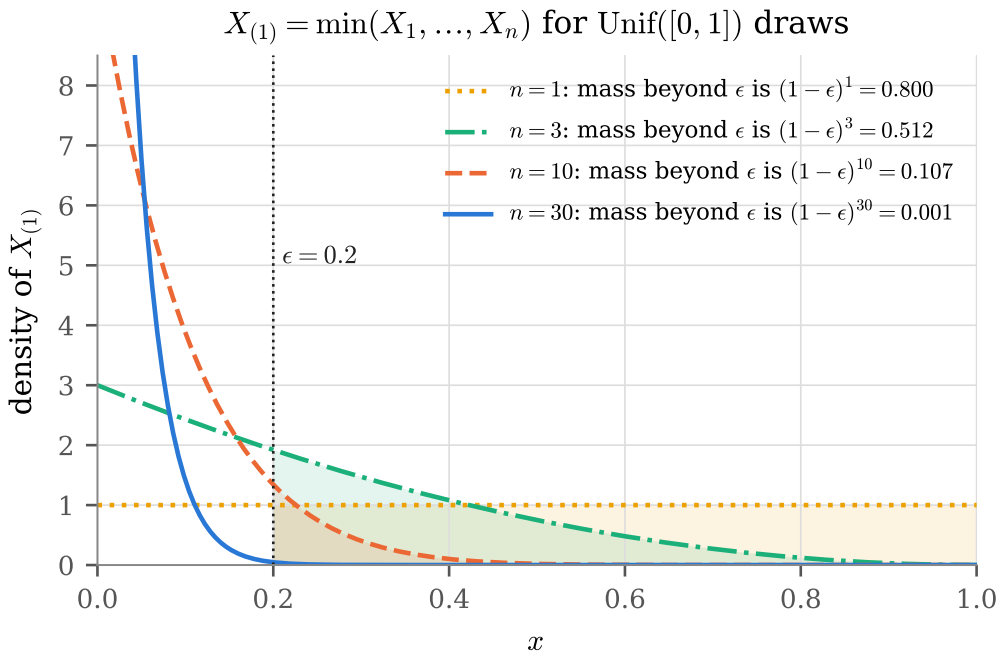


Figure 3.4: Example 3.2 as a picture. The density of $X_{(1)} = \min(X_1, \ldots, X_n)$ is $n(1-x)^{n-1}$ on $[0, 1]$. As n grows the mass piles up against 0, and the mass beyond ϵ, which is the $(1 - \epsilon)^n$ of the proof, goes to zero. The proof computes the shaded tail; the picture shows why it has to shrink. (The $n = 10$ and $n = 30$ curves are cut off at the top; their peaks are at 10 and 30.)

The proof is the template for every direct proof of convergence in probability: write the event $|X_n - X| \geq \epsilon$, rewrite it as an event about the original sample, use independence to turn a joint probability into a product, and take the limit.

3.3.2 A workhorse: shrinking variance

Most of the time you won't compute $\mathbb{P}(|X_n - x| \geq \epsilon)$ directly. You'll use the following criterion, which only needs a mean and a variance.

Proposition 3.3 (Shrinking-variance criterion)**.** *Let $(T_n)_{n\geq 1}$ be a sequence of r.v.s with $\mathbb{E}[T_n] \to \mu$ and $\mathbb{V}ar(T_n) \to 0$. Then $T_n \xrightarrow[n\to\infty]{\mathbb{P}} \mu$.*

Proof strategy (core). Chebyshev's inequality bounds the probability of a large deviation by a variance; a variance going to 0 drags the probability with it. The triangle inequality handles the mean not being exactly μ (toolkit A.7).

Proof. The argument is Chebyshev's inequality. Recall its two ingredients. *Markov's inequality:* if $Y \geq 0$ then $\mathbb{P}(Y \geq a) \leq \mathbb{E}[Y]/a$ for every $a > 0$. *Chebyshev's inequality:* applying Markov's inequality to $Y = (T - \mathbb{E}[T])^2$ with $a = \epsilon^2$ gives, for any r.v. T with finite variance,

$$\mathbb{P}(|T - \mathbb{E}[T]| \geq \epsilon) \leq \frac{\mathbb{V}\text{ar}(T)}{\epsilon^2}.$$

Now fix $\epsilon > 0$. For n large enough that $|\mathbb{E}[T_n] - \mu| < \epsilon/2$, the event $|T_n - \mu| \geq \epsilon$ forces $|T_n - \mathbb{E}[T_n]| \geq \epsilon/2$ (triangle inequality), so

$$\mathbb{P}(|T_n - \mu| \geq \epsilon) \leq \mathbb{P}(|T_n - \mathbb{E}[T_n]| \geq \epsilon/2) \leq \frac{4\,\mathbb{V}\text{ar}(T_n)}{\epsilon^2} \to 0. \qquad \square$$

Example 3.4. Let $X_n \sim \text{Bern}(1/n)$, so X_n is 1 with probability $1/n$ and 0 otherwise. For $0 < \epsilon < 1$, $\mathbb{P}(|X_n - 0| \geq \epsilon) = \mathbb{P}(X_n = 1) = 1/n \to 0$, so $X_n \xrightarrow[n\to\infty]{\mathbb{P}} 0$. Alternatively, $\mathbb{E}[X_n] = 1/n \to 0$ and $\mathbb{V}\text{ar}(X_n) = \frac{1}{n}(1 - \frac{1}{n}) \to 0$, and Proposition 3.3 gives the same conclusion.

3.4 The Weak Law of Large Numbers (WLLN)

Frequently, statisticians are concerned with situations in which the limiting r.v. is a constant and the r.v.s in the sequence are sample means. The most famous result along these lines is the Weak Law of Large Numbers (WLLN), which is stated here but explained and proved below.

Theorem 3.5 (Weak Law of Large Numbers)**.** *Let $X_1, X_2, \ldots$ be i.i.d. r.v.s with $\mathbb{E}[X_i] = \mu$ and $\mathbb{V}ar(X_i) = \sigma^2 < \infty$, and let $\bar{X}_n$ be the sample mean. Then, for every $\epsilon > 0$:*

$$\lim_{n\to\infty} \mathbb{P}(|\bar{X}_n - \mu| < \epsilon) = 1.$$

That is, $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} \mu$.

Let's motivate this all-important result. Suppose you have some probability distribution with a certain mean, μ, and variance, σ^2, which is assumed finite. Then you draw n independent r.v.s, $X_1, \ldots, X_n$, from the distribution so that the X_i's are independent and identically distributed (i.i.d. r.v.s). At this point, make sure you understand that each X_i is an exact copy of the parent r.v. you are 'drawing' these r.v.s from. Also, make sure to keep in mind that you're 'drawing' r.v.s and not values of r.v.s (i.e., numbers). You can think of this as carrying out a long experiment during which these r.v.s are drawn. Once you've drawn all of these r.v.s, you calculate their average that yields the **(theoretical) sample mean**, $\bar{X}_n$:

$$\bar{X}_n = \frac{X_1 + \cdots + X_n}{n}.$$

Question: Why 'theoretical' sample mean?

Answer: In order to avoid confusion with the descriptive statistic that also goes by 'sample mean', $\bar{X}_n$ is the theoretical sample mean because *it's a random variable*, not a number. The sample here, $X_1, \ldots, X_n$, is a sample of n r.v.s, each of which has the same distribution. If the values of the X_i's are realized, then the descriptive sample mean, the realized value, is easily calculated as:

$$\bar{x}_n = \frac{x_1 + \cdots + x_n}{n},$$

where the x_i's are the realized values of the X_i's. This explanation is given in the hopes that when you see $\bar{X}_n$ your mind doesn't go to the number $\bar{x}_n$. In the kiss experiment, $\bar{R}_n$ is the r.v. 'proportion turning right in a sample of n couples' and $\bar{r}_n = 80/124 = 0.645$ is the number the one sample happened to produce.

Notice that $\bar{X}_n$ is itself an r.v. because it's a function of r.v.s. The sample mean isn't the true mean, μ, of the parent distribution; that is, $\bar{X}_n \neq \mu$, although as demonstrated below $\mathbb{E}[\bar{X}_n] = \mu$. But, $\mu = \mathbb{E}[X_i]$ for all i and is a number (not an r.v.). This is how it should be since the distribution of every X_i is an exact copy of X's distribution. As such, the mean of X must equal the mean of the X_i's. The sample mean is the simplest and most natural way for trying to estimate the true mean, and the Weak Law of Large Numbers (WLLN) offers support to this notion.

Before explaining in more detail the WLLN consider some properties of the sample mean, $\bar{X}_n$. By linearity, its expectation is:

$$\mathbb{E}[\bar{X}_n] = \mathbb{E}\left[\frac{X_1 + \cdots + X_n}{n}\right] = \frac{n\mu}{n} = \mu.$$

The expectation $\mathbb{E}[\bar{X}_n]$ involves two different kinds of averaging, and it's worth separating them. The sample mean, $\bar{X}_n$, averages over the n draws of one experiment. The expectation, $\mathbb{E}[\bar{X}_n]$, averages over *all* possible outcomes of that experiment, weighted by their probabilities, which is a theoretical average because you never observe all possible outcomes. The statement $\mathbb{E}[\bar{X}_n] = \mu$ says that the second kind of average of the first kind of average lands on the true mean. Figure 3.5 draws the two directions.

Now, consider the variance of $\bar{X}_n$:

$$\mathbb{V}\mathrm{ar}(\bar{X}_n) = \frac{\mathbb{V}\mathrm{ar}(X_1 + \cdots + X_n)}{n^2} = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}.$$

This follows since the X_i's are independent so that the variance of the sum is the sum of the individual variances.

Intuition: why the variance of the sample mean is scaled by $1/n$

Averaging n independent draws cancels noise. Think of each X_i as $\mu + U_i$ where the U_i are independent, mean-zero errors. Then $\bar{X}_n = \mu + \bar{U}_n$, and the errors in $\bar{U}_n$ have no reason to line up: some are positive, some negative, and summing n of them produces a total whose typical size grows like $\sqrt{n}$, not like n (variances add, so standard deviations add in quadrature). Dividing by n then leaves a typical error of size $\sigma\sqrt{n}/n = \sigma/\sqrt{n}$. That's why $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \sigma^2/n$ and why doubling the precision costs four times the data. Figure 3.6 is that claim drawn: the sum of n cancelling errors reaches about $\sqrt{n}$, not n.

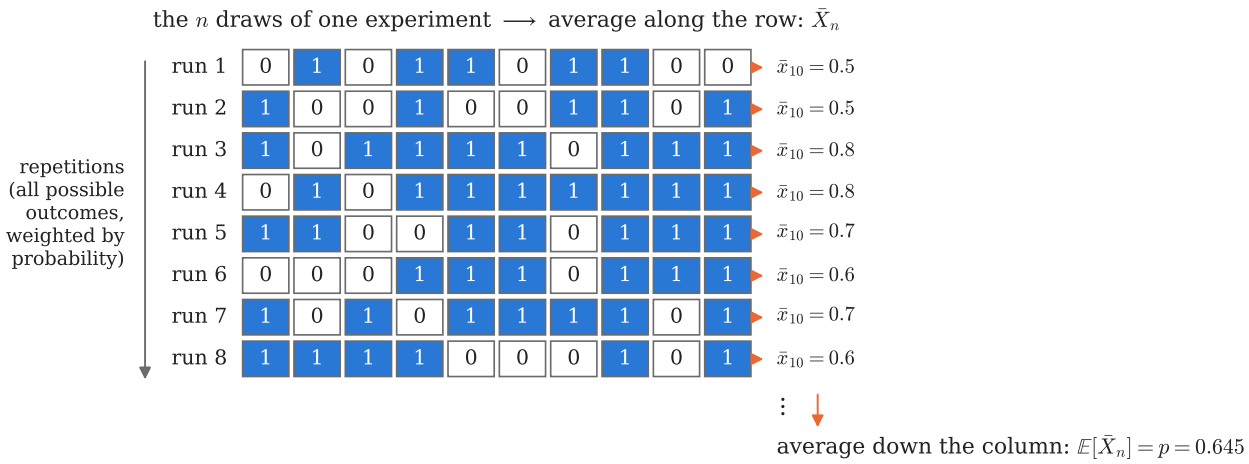


Figure 3.5: Two kinds of averaging. Each row is one run of the experiment, here $n = 10$ Bernoulli draws with $p = 0.645$; averaging along the row gives that run's sample mean, a different number each time. The expectation averages down the column, over every possible run weighted by its probability (only eight runs are shown; the expectation uses all 2^{10} of them). The first average is a random variable; the second is the number p.

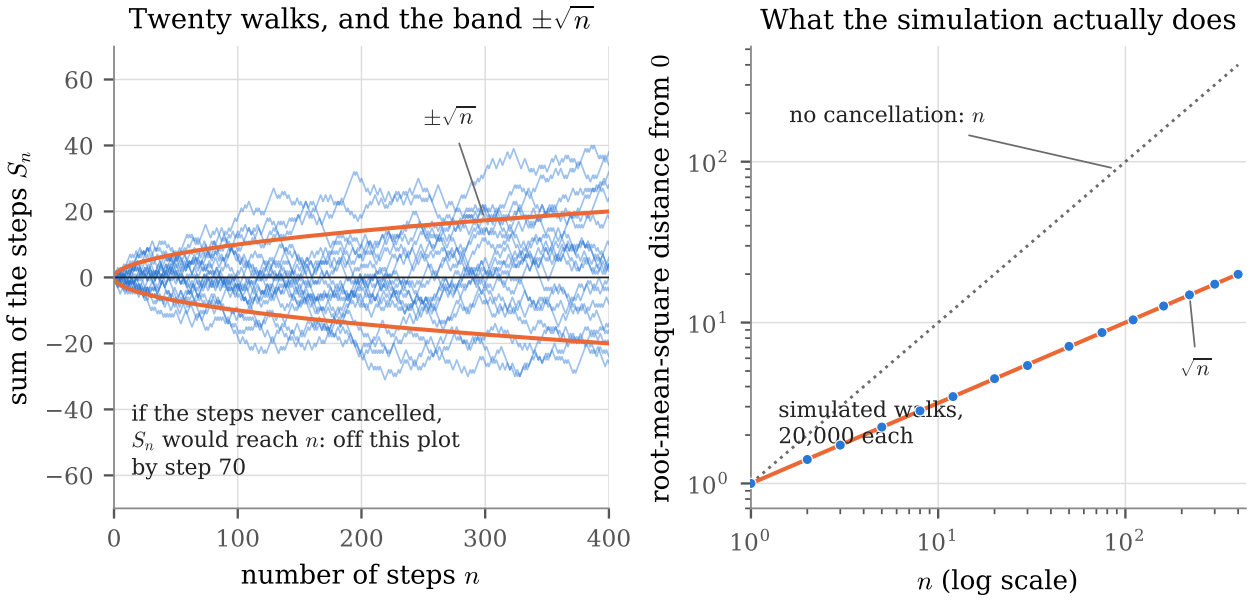


Figure 3.6: Where the $\sqrt{n}$ in every formula in this book comes from. Left: twenty running sums of steps of $+1$ and -1, each step equally likely, with the band $\pm\sqrt{n}$ drawn over them. If the steps never cancelled, the sum after n steps would be n, which leaves the top of this figure by step 70. Because they cancel, the sum stays inside a band that widens like $\sqrt{n}$. Right: the root-mean-square distance of the sum from zero, measured over 20,000 simulated runs at each n, against $\sqrt{n}$ and against n on logarithmic axes. The simulated points sit on $\sqrt{n}$, which is exact here: $\mathbb{E}[S_n^2] = n$, since each step has variance 1 and the cross terms have expectation zero.

Does this make sense? The cost of precision

Any time you produce a standard error, ask what it would take to halve it. The answer is always four times the data, and it follows from the $\sqrt{n}$: to replace $\sigma/\sqrt{n}$ by $\sigma/(2\sqrt{n})$ you need $\sqrt{n}$ to double, so n must quadruple. If a calculation of yours implies that twice the data halves the error, you have divided by n somewhere that should have been $\sqrt{n}$. For the kiss experiment, the standard error at $n = 124$ is 0.043; getting it to 0.02 takes not 250 couples but about 570.

Now, Chebyshev's inequality says that the probability that the distance of some r.v. from its mean is larger than some number, say ϵ, is bounded above by the variance of the r.v. divided by ϵ^2. In the present context, the variance is σ^2/n, and so the upper bound is $\dfrac{\sigma^2}{n\epsilon^2}$. Formally, this is expressed as:

$$\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) \leq \frac{\sigma^2}{n\epsilon^2}.$$

The WLLN states that if ϵ is a fixed number and you take the limit of this probability as $n \to \infty$, then the limit converges to 0. That is, the WLLN states that for any $\epsilon > 0$, as $n \to \infty$ you have:

$$\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) = \mathbb{P}\left(\left|\frac{X_1 + \cdots + X_n}{n} - \mu\right| \geq \epsilon\right) \to 0.$$

So the probability of falling far from the mean diminishes to zero as you draw more and more samples.

To make the interpretation of the WLLN more concrete, suppose again that you're conducting one long experiment, and during that experiment you draw n i.i.d. r.v.s, $X_1, \ldots, X_n$, from the same parent distribution. One way of thinking about the X_i is that each one equals the true mean, μ, plus some measurement noise, U_i, where $\mathbb{E}[U_i] = 0$ for all i and the U_i are independent. You can then express these noisy measurements as $X_i = \mu + U_i$ for $i = 1, \ldots, n$.

So you've got a collection of noisy measurements and then you compute their average. What the Weak Law of Large Numbers says is that the sample mean, $\bar{X}_n$, is unlikely to be far off from the true mean, μ, of the parent distribution. And by far off is meant at least ϵ distance away. The sample mean is, thus, a good estimator of the true mean.

Here's the proof of the WLLN, which is easily demonstrated with Chebyshev's inequality.

Proof strategy (core). Two lines: Chebyshev's inequality with $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \sigma^2/n$, then $n \to \infty$. This is the proof to know by heart.

Proof.

- Let ϵ be some positive real number. Then by Chebyshev's inequality:

$$\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) \leq \frac{\mathbb{V}\mathrm{ar}(\bar{X}_n)}{\epsilon^2} = \frac{\sigma^2}{n\epsilon^2}.$$

- As $n \to \infty$, the right-hand side goes to 0, which implies $\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon)$ goes to 0, too.

□

Worked example: the kiss experiment and Chebyshev

In the kiss experiment $R_i \overset{iid}{\sim} \text{Bern}(p)$, so $\mu = p$ and $\sigma^2 = p(1-p)$. How likely is it that the sample proportion misses p by more than 0.084 (about 10 couples out of 124)? Chebyshev's inequality with $n = 124$, $\epsilon = 0.084$ and the worst case $p(1-p) \leq 1/4$ gives

$$\mathbb{P}(|\bar{R}_n - p| \geq 0.084) \leq \frac{p(1-p)}{n\epsilon^2} \leq \frac{0.25}{124 \times 0.084^2} \approx 0.29,$$

and with $p(1-p)$ evaluated at 0.645 the bound is 0.26. The bound is true but loose: the CLT (Chapter 4) says the probability is really about 0.05. Chebyshev's inequality is the right tool for proving that something converges, and the wrong tool for saying how fast. Figure 4.5 in the next chapter shows all the bounds on the same axes.

From undergraduate statistics: consistency

An estimator that converges in probability to the parameter it estimates is called *consistent*. The WLLN says the sample mean is a consistent estimator of the population mean, and the sample proportion is a consistent estimator of the population proportion. In an introductory course you were told 'with a bigger sample the estimate gets closer to the truth'; the WLLN is that sentence with every word defined. Note what it does not say: it doesn't say $\bar{X}_n$ equals μ for any finite n, and it doesn't say how big n must be for a given accuracy. The CLT and Hoeffding's inequality, both in Chapter 4, handle the second question.

3.5 What comes next

The Law of Large Numbers says where $\bar{X}_n$ ends up: at μ, with the variance σ^2/n shrinking to nothing. That's the left panel of Figure 3.2, and it leaves the obvious question open. If every density in that panel is on its way to becoming a spike, what does $\bar{X}_n$ look like *along the way*? The right panel already gave the answer in a picture: magnify the error by $\sqrt{n}$ and its shape settles down to a Gaussian, whatever the parent distribution. Making that picture into a theorem needs the second mode of convergence, and it's the subject of the next chapter.

3.6 Summary

- A sequence of r.v.s is a sequence of functions on the sample space; one outcome fixes one ordinary sequence of numbers. No single index N works for every outcome, so convergence has to be defined through probabilities.
- 'Asymptotic' means: for the n at hand, replace the unknown sampling distribution by its limit, usually a Gaussian, and carry on; the theorem does not say how large n must be.
- $\bar{X}_n$ is a random variable, the theoretical sample mean; $\bar{x}_n$ is the number one sample produced. Every theorem is about the first. Its expectation is a second average, over all possible samples, of the first.
- **Convergence in probability:** $X_n \xrightarrow[n\to\infty]{\mathbb{P}} X$ if $\mathbb{P}(|X_n - X| \geq \epsilon) \to 0$ for every $\epsilon > 0$. The mass outside any band around the limit disappears.

- **Shrinking-variance criterion:** $\mathbb{E}[T_n] \to \mu$ and $\mathbb{V}\mathrm{ar}(T_n) \to 0$ imply $T_n \xrightarrow[n\to\infty]{\mathbb{P}} \mu$. This is Chebyshev's inequality, and it's how consistency is usually proved.
- **Weak Law of Large Numbers:** $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} \mu$ whenever the variance is finite because $\mathbb{E}[\bar{X}_n] = \mu$ and $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \sigma^2/n$. Estimators built from averages are consistent.
- Chebyshev's bound $\sigma^2/(n\epsilon^2)$ is a guarantee at every n but a loose one: 0.26 in the kiss experiment where the truth is 0.05.

3.7 Exercises

Exercise 3.1. Let $X_n \sim \mathrm{Bern}(1/n)$. (a) For $0 < \epsilon < 1$, what is $\mathbb{P}(|X_n| > \epsilon)$? (b) Does (X_n) converge in probability, and to what? (c) Does it converge in distribution? (d) What is $\lim_n \mathbb{E}[\cos(X_n)]$? *(Solution on page 222.)*

Exercise 3.2. Let $X_1, \ldots, X_n \overset{iid}{\sim} \mathrm{Expo}(\lambda)$, so $\mathbb{E}[X_1] = 1/\lambda$ and $\mathbb{V}\mathrm{ar}(X_1) = 1/\lambda^2$. Use the shrinking-variance criterion to show that $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} 1/\lambda$, and find the smallest n for which Chebyshev's inequality guarantees $\mathbb{P}(|\bar{X}_n - 1/\lambda| \geq 0.1/\lambda) \leq 0.05$. *(Solution on page 222.)*

Exercise 3.3. In the kiss model with $p = 0.645$, how many couples would Chebyshev's inequality require before $\mathbb{P}(|\bar{R}_n - p| \geq 0.02) \leq 0.05$? Is the answer sensitive to p? *(Solution on page 222.)*

Chapter 4

The Central Limit Theorem

In this chapter

- Convergence in distribution, the mode of convergence that compares CDFs rather than values, and why its definition only looks at continuity points.
- The Central Limit Theorem in its three equivalent forms, and what it says at the sample sizes you actually have.
- Hoeffding's inequality: a finite-sample guarantee for bounded data when n is too small to trust the CLT.
- What the limit theorems buy you, and which of their conclusions could be reached without limits.

Chapter 3 ended with $\bar{X}_n$ collapsing onto μ and the question of what it looks like on the way there. This chapter answers it.

4.1 Convergence in distribution

Convergence in distribution is weaker than convergence in probability (and, in the full hierarchy of Appendix B, the weakest mode there is). All it says is that the CDF of X_n converges to the CDF of X as n goes to infinity, *but doesn't require any dependence between the X_n's and X*. That is, the r.v.s X_n and X can be defined on entirely different probability spaces. The CDF of these r.v.s is used instead of a PDF or PMF because all r.v.s, whether they be discrete, continuous or mixed, have CDFs.

Question: Why does not requiring any dependence between the X_n's and X make convergence in distribution weak? You might think that, by not introducing a dependence assumption, you're actually making the statement stronger.

Answer: By not imposing any dependence structure between X_n and X, convergence in distribution allows for a wider range of scenarios where the CDFs of X_n converge to the CDF of X. It doesn't guarantee that the actual values of X_n get closer to X as n increases, which is a requirement for stronger forms of convergence like convergence in probability or convergence almost surely. Convergence in distribution only concerns the limiting behavior of the distributions, not the actual values of the r.v.s. A statement that holds in more scenarios is a weaker statement because it tells you less about any one of them.

Definition 4.1 (Convergence in distribution)**.** A sequence $(X_n)_{n\geq 1}$ of r.v.s **converges in distribution** to an r.v. X, denoted $X_n \xrightarrow[n\to\infty]{(\mathrm{d})} X$, if:

$$\lim_{n\to\infty} F_{X_n}(x) = F_X(x),$$

for all x at which F_X is continuous.

> **In words.** For each fixed x, $F_{X_n}(x) = \mathbb{P}(X_n \leq x)$ is a number, and the definition asks that this sequence of numbers converge to $\mathbb{P}(X \leq x)$, at every x where the limit's CDF has no jump (toolkit A.6). It compares two distributions, not two random variables: nothing is said about X_n and X being close on any outcome, and they need not even live on the same sample space.
>
> **For the kiss data.** The CLT will say that the standardized error $\sqrt{n}(\bar{R}_n - p)/\sqrt{p(1-p)}$ converges in distribution to Gauss$(0,1)$: for every x, the probability that the standardized error is below x approaches $\Phi(x)$. With $n = 124$ that is the license to compute $\mathbb{P}(|\bar{R}_n - p| \leq 0.084) \approx \Phi(1.95) - \Phi(-1.95) = 0.95$.

4.1.1 Why only the continuity points?

The qualifier 'at which F_X is continuous' looks like a technicality. It isn't; without it the most basic examples fail.

Example 4.2. Let $T_n \sim \mathrm{Unif}\left(\left[5 - \frac{1}{2n}, 5 + \frac{1}{2n}\right]\right)$, a uniform on an interval of width $1/n$ centered at 5. Every sensible notion of convergence should say $T_n \to 5$, and indeed $T_n \xrightarrow[n\to\infty]{\mathbb{P}} 5$ since $\mathbb{P}(|T_n - 5| > \delta) = 0$ as soon as $\frac{1}{2n} \leq \delta$. The CDF F_n of T_n is 0 to the left of the interval, rises linearly to 1 across it, and in particular $F_n(5) = \frac{1}{2}$ for *every* n. The CDF of the constant 5 is the step function $F(t) = \mathbb{1}(t \geq 5)$, with $F(5) = 1$. So $F_n(t) \to F(t)$ for every $t \neq 5$ but not at $t = 5$, the one point where F jumps. With the continuity-point qualifier, $T_n \xrightarrow[n\to\infty]{(\mathrm{d})} 5$; without it, T_n wouldn't converge in distribution to anything. Figure 4.1 shows the picture.

The general principle: a jump in F_X at x means X has positive probability of equaling x exactly, and a sequence approaching from the left or the right can straddle that jump forever. Away from the jumps there's nothing to straddle.

4.1.2 Two equivalent definitions

Alternatively, convergence in distribution can be defined using the expectation of a function, provided the function is continuous and bounded. This is the form that generalizes: on spaces where there is no CDF to write down, it becomes the definition of convergence in distribution.

Proposition 4.3 (Equivalent characterizations)**.** *For real-valued r.v.s the following are equivalent:*

1. $F_{X_n}(x) \to F_X(x)$ *for all* x *at which* F_X *is continuous;*
2. $\mathbb{E}[f(X_n)] \to \mathbb{E}[f(X)]$ *for every continuous and bounded function* $f : \mathbb{R} \to \mathbb{R}$.

If moreover X has a density (for instance X is Gaussian), then $X_n \xrightarrow[n\to\infty]{(d)} X$ implies convergence of interval probabilities:

$$\mathbb{P}(a \leq X_n \leq b) \to \mathbb{P}(a \leq X \leq b), \qquad \text{for all } a < b.$$

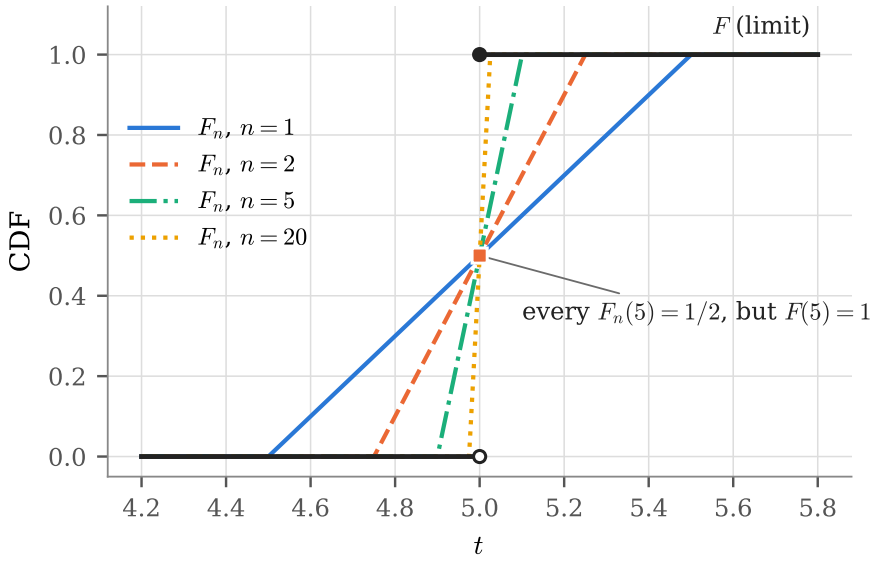


Figure 4.1: The CDFs of $T_n \sim \text{Unif}([5 - \frac{1}{2n}, 5 + \frac{1}{2n}])$ for several n and the CDF of the limit, the constant 5. The F_n converge to the step function everywhere except at the jump, where all of them pass through $\frac{1}{2}$. The definition of convergence in distribution excludes exactly that point.

Note that in (2) the things converging are numbers, $\mathbb{E}[f(X_n)]$, so the convergence is ordinary convergence of numbers. The second form is the useful one in proofs because it doesn't single out any particular x; the first form is the useful one in computations because CDFs are what you can write down. The last statement is why the CLT is enough to compute probabilities: $\mathbb{P}(a \le Z_n \le b) \approx \Phi(b) - \Phi(a)$, which is what every asymptotic confidence interval and test rests on.

One implication between the two modes is used on nearly every page from here on, so here it is as a statement with a proof.

Proposition 4.4 (In probability implies in distribution)**.** *If* $X_n \xrightarrow[n\to\infty]{\mathbb{P}} X$*, then* $X_n \xrightarrow[n\to\infty]{(d)} X$*.*

Proof strategy (for completeness; skip on a first reading). Sandwich $F_{X_n}(x)$ between $F_X(x-\epsilon)$ and $F_X(x+\epsilon)$, up to the probability that X_n and X differ by more than ϵ, which goes to 0; then let $\epsilon \to 0$ and use continuity of F_X at x. The statement is what matters; the proof is a standard ϵ argument.

Proof. Fix a continuity point x of F_X and an $\epsilon > 0$. If $X_n \le x$ then either $X \le x+\epsilon$ or $|X_n - X| > \epsilon$, so

$$F_{X_n}(x) \le F_X(x+\epsilon) + \mathbb{P}(|X_n - X| > \epsilon).$$

In the same way, if $X \le x - \epsilon$ then either $X_n \le x$ or $|X_n - X| > \epsilon$, so

$$F_X(x-\epsilon) - \mathbb{P}(|X_n - X| > \epsilon) \le F_{X_n}(x).$$

Let $n \to \infty$: the probability terms vanish by convergence in probability, and $F_{X_n}(x)$ is squeezed, in the limit, between $F_X(x-\epsilon)$ and $F_X(x+\epsilon)$. Now let $\epsilon \to 0$: both bounds tend to $F_X(x)$ because F_X is continuous at x. □

The continuity-point qualifier is what the last line uses: at a jump of F_X the two bounds would not meet, and Example 4.2 shows the conclusion can indeed fail there.

Example 4.5. In Example 3.4, $X_n \sim \text{Bern}(1/n)$ converges in probability, hence in distribution, to 0. The function $f(x) = \cos x$ is continuous and bounded, so $\mathbb{E}[\cos(X_n)] \to \cos(0) = 1$. You can check this directly: $\mathbb{E}[\cos X_n] = (1 - \frac{1}{n})\cos 0 + \frac{1}{n}\cos 1 \to 1$.

4.1.3 Example 1: i.i.d. random variables

If $(X_n)_{n\geq 1}$ is a sequence of i.i.d. r.v.s, then convergence in distribution isn't very interesting because they converge to the CDF of any one r.v. in the sequence. That is, if $X_1, X_2, \ldots$ is a sequence of i.i.d. r.v.s with CDF $F_X(x)$, then $X_n \xrightarrow[n\to\infty]{\text{(d)}} X$. This is because:

$$F_{X_n}(x) = F_X(x), \text{ for all } x \text{ implies that } \lim_{n\to\infty} F_{X_n}(x) = F_X(x), \text{ for all } x.$$

This is the example from Section 3.1: the sequence of standard Gaussians converges in distribution to a standard Gaussian even though $\mathbb{P}(X_n = X) = 0$. Convergence in distribution is the mode that makes sense of that example.

4.1.4 Example 2: the maximum of uniforms

Example 4.6. Let $X_1, X_2, \ldots \overset{iid}{\sim} \text{Unif}([0,1])$ and define the maximum of the first n by $X_{(n)} = \max_{i=1,\ldots,n} X_i$. Then:

$$n(1 - X_{(n)}) \xrightarrow[n\to\infty]{\text{(d)}} \text{Expo}(1).$$

Before the proof, a word on why the maximum is being magnified by n rather than $\sqrt{n}$. The maximum of n uniforms piles up just below 1; by the same argument as in Example 3.2 it converges in probability to 1, so the unmagnified sequence has a degenerate limit. The gap $1 - X_{(n)}$ is of order $1/n$ (with n points spread over $[0,1]$ the largest one sits about $1/n$ from the top), so n is the magnification that holds its size fixed. The CLT's $\sqrt{n}$ is for sums; a maximum is not a sum and gets a different rate and a different (non-Gaussian) limit.

Proof strategy (core). Write the CDF of $n(1-X_{(n)})$ as a probability about the whole sample, use independence to get a power $(1 - x/n)^n$, and recognize the limit e^{-x} (toolkit A.10). Same template as Example 3.2, with a CDF in place of a tail probability.

Proof.

- Define the CDF of the r.v. $n(1 - X_{(n)})$ at a point $x > 0$:

$$F_{n(1-X_{(n)})}(x) = \mathbb{P}(n(1 - X_{(n)}) \leq x).$$

- Now, simplify the CDF in such a way as to eventually take the limit as $n \to \infty$.

- To that end, observe that:

$$\begin{aligned}\mathbb{P}(n(1-X_{(n)}) \leq x) &= \mathbb{P}(n - nX_{(n)} \leq x) \\ &= \mathbb{P}\left(X_{(n)} \geq 1 - \frac{x}{n}\right) \\ &= 1 - \mathbb{P}\left(X_{(n)} < 1 - \frac{x}{n}\right) \\ &= 1 - \mathbb{P}\left(X_1 < 1 - \frac{x}{n}, X_2 < 1 - \frac{x}{n}, \ldots, X_n < 1 - \frac{x}{n}\right) \\ &= 1 - \mathbb{P}\left(X_1 < 1 - \frac{x}{n}\right)\mathbb{P}\left(X_2 < 1 - \frac{x}{n}\right)\cdots\mathbb{P}\left(X_n < 1 - \frac{x}{n}\right) \\ &= 1 - \mathbb{P}\left(X_1 < 1 - \frac{x}{n}\right)^n.\end{aligned}$$

  The movement from the third to the fourth equality holds because, as $X_{(n)}$ is the maximum of the X_i's, the maximum is below a level exactly when every X_i is below it: the single event is equivalent to the joint event. The fifth equality results since the X_i's are independent.

- Now, recall that the X_i's, and thus X_1, are uniformly distributed on $(0, 1)$. This means its CDF is linear, emanating from the origin rising to 1 with a slope of 1; in other words, plug any value of the support into the CDF and you get the same value back. Hence, for n large enough that $1 - x/n \in (0, 1)$:

$$\mathbb{P}\left(X_1 < 1 - \frac{x}{n}\right) = 1 - \frac{x}{n}.$$

- Returning to the last right-hand side from above, you have that:

$$1 - \mathbb{P}\left(X_1 < 1 - \frac{x}{n}\right)^n = 1 - \left(1 - \frac{x}{n}\right)^n,$$

  which by taking the limit as $n \to \infty$ gives:

$$\lim_{n\to\infty}\left[1 - \left(1 - \frac{x}{n}\right)^n\right] = 1 - e^{-x},$$

  using the calculus fact $(1 - x/n)^n \to e^{-x}$.

- To summarize, then, you have:

$$\lim_{n\to\infty} \mathbb{P}(n(1 - X_{(n)}) \leq x) = 1 - e^{-x} \quad \text{for every } x > 0,$$

  and $1 - e^{-x}$ is the CDF of the Exponential with $\lambda = 1$ (for $x \leq 0$ both sides are 0). This proves the result.

□

Pitfall: convergence in distribution says nothing about closeness

Let $Z \sim \text{Gauss}(0, 1)$ and set $X_n = -Z$ for every n. By the symmetry of the Gaussian, $X_n \sim \text{Gauss}(0, 1)$ for every n, so $X_n \xrightarrow[n\to\infty]{\text{(d)}} Z$. But $|X_n - Z| = 2|Z|$, which is not small and

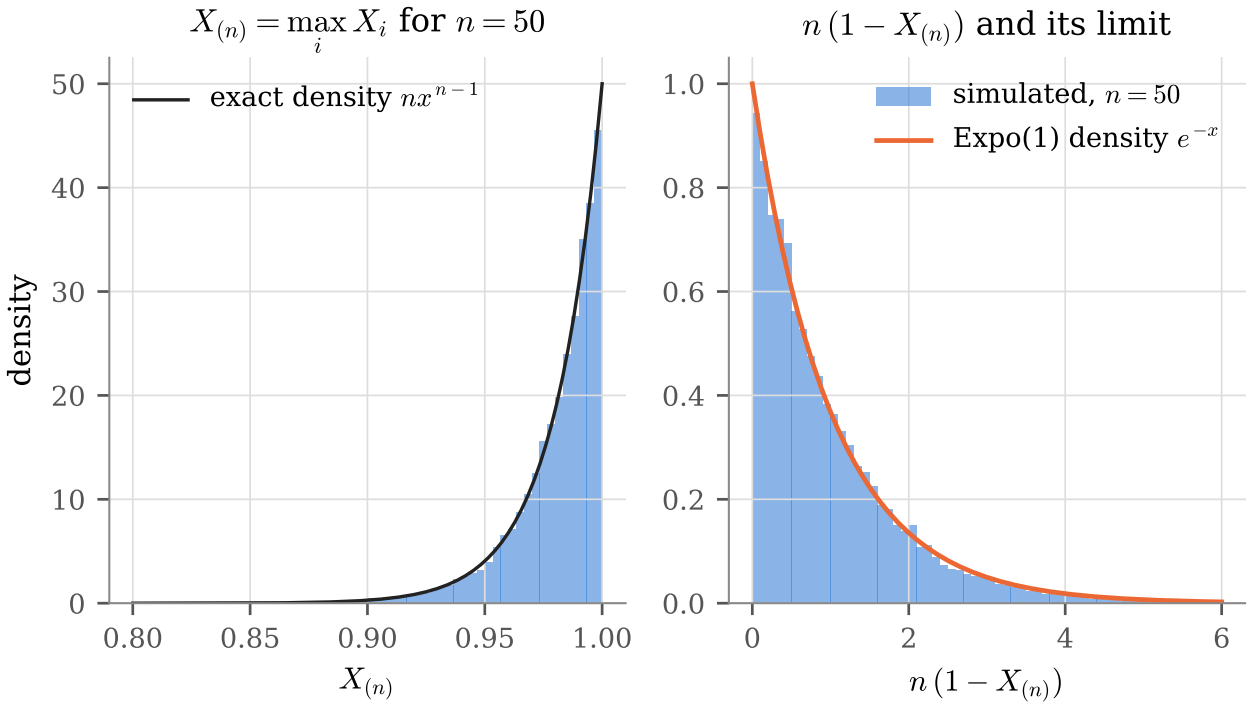


Figure 4.2: Example 4.6 by simulation, 20,000 replications. Left: the maximum of $n = 50$ uniforms crowds the top of the interval; its exact density nx^{n-1} is drawn over the histogram. Right: the magnified gap $n(1 - X_{(n)})$ is already very close to Expo(1) at $n = 50$.

never becomes small. So X_n does not converge in probability to Z, even though it converges in distribution to it. Convergence in distribution is a statement about two CDFs, not about two columns of numbers getting closer. The same example shows why you can't add limits in distribution: $X_n + Z = 0$ for every n, while the sum of two independent Gauss$(0, 1)$ r.v.s is Gauss$(0, 2)$. Knowing the limiting distributions of X_n and Y_n separately doesn't determine the limiting distribution of $X_n + Y_n$; you need to know how they're related. Chapter 5 says what you can do instead.

4.1.5 Limits that are constants

One case of convergence in distribution is used constantly later, so here it is as a statement with a proof.

Proposition 4.7. *If* $X_n \xrightarrow[n\to\infty]{(d)} c$ *for a constant* c*, then* $X_n \xrightarrow[n\to\infty]{\mathbb{P}} c$*.*

Proof strategy (for completeness; skip on a first reading). The CDF of a constant is a step; convergence of the CDFs at two continuity points, one on each side of c, squeezes the mass outside the ϵ-band to 0.

Proof. The CDF of the constant c is the step $F(x) = \mathbb{1}(x \geq c)$, continuous everywhere except at c. Fix $\epsilon > 0$. The points $c - \epsilon$ and $c + \epsilon/2$ are continuity points, so $F_{X_n}(c - \epsilon) \to F(c - \epsilon) = 0$ and $F_{X_n}(c + \epsilon/2) \to F(c + \epsilon/2) = 1$. Therefore

$$\mathbb{P}(|X_n - c| \geq \epsilon) = \mathbb{P}(X_n \leq c - \epsilon) + \mathbb{P}(X_n \geq c + \epsilon) \leq F_{X_n}(c - \epsilon) + (1 - F_{X_n}(c + \epsilon/2)) \to 0 + 0. \quad \square$$

For a constant limit, then, 'in probability' and 'in distribution' are the same thing, which is why the Law of Large Numbers could equally be written with the arrow for convergence in distribution, and why Slutsky's theorem in Chapter 5 needs only convergence in probability for the sequence with the constant limit.

4.2 The Central Limit Theorem

As in the previous section, let $X_1, X_2, \ldots$ be i.i.d. r.v.s with mean μ and variance σ^2. The Law of Large Numbers says that as $n \to \infty$, $\bar{X}_n$ converges to the constant μ. But what is $\bar{X}_n$'s distribution along the way to becoming the constant μ? This is addressed by the Central Limit Theorem (CLT), an all-important theorem in statistics.

The CLT states that for large n, the distribution of $\bar{X}_n$ *after standardization* approaches a standard Gaussian distribution. Recall that if you standardize an r.v., the resulting r.v. is measured in units of standard deviation. In the present case, the convergence process examines the typical deviations of $\bar{X}_n$ around μ. Now the statement of the CLT.

Theorem 4.8 (Central Limit Theorem)**.** *Let $X_1, \ldots, X_n$ be i.i.d. r.v.s such that $\mathbb{E}[X_i] = \mu$ and $\mathbb{V}ar(X_i) = \sigma^2 < \infty$ for $i = 1, \ldots, n$. Then:*

$$\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right) \xrightarrow[n\to\infty]{(d)} Gauss(0,1).$$

In words, the CDF of the left-hand side converges to Φ, the CDF of the standard Gaussian distribution.

An equivalent statement of the CLT in which $\bar{X}_n$ remains mean-centered by μ but omits the standardization by σ is given as follows:

$$\sqrt{n}\left(\bar{X}_n - \mu\right) \xrightarrow[n\to\infty]{(\mathrm{d})} \mathrm{Gauss}(0, \sigma^2).$$

Question 1: Why does this second statement hold?

Answer 1: It holds because of the way the Gaussian family behaves under scaling. Let $X \sim \mathrm{Gauss}(0,1)$, so its PDF is:

$$f_X(x) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{x^2}{2}\right).$$

Now let $Y = \sigma X$ for some $\sigma > 0$. Then $\mathbb{E}[Y] = \mathbb{E}[\sigma X] = \sigma\,\mathbb{E}[X] = 0$. Hence the mean remains 0. But note what happens to the variance of X after multiplying it by σ:

$$\begin{aligned}
\mathbb{V}\mathrm{ar}(Y) = \mathbb{E}[(Y - \mathbb{E}[Y])^2] &= \mathbb{E}[(\sigma X - 0)^2] \\
&= \sigma^2\,\mathbb{E}[X^2] \\
&= \sigma^2\,\mathbb{V}\mathrm{ar}(X) \\
&= \sigma^2,
\end{aligned}$$

and Y is still Gaussian (a linear function of a Gaussian is Gaussian), so $\sigma X \sim \mathrm{Gauss}(0, \sigma^2)$. Multiplying both sides of the first statement by the constant σ therefore gives the second; the step

'multiply both sides of a limit in distribution by a constant' is justified by the continuous mapping theorem of Chapter 5.

Question 2: In the statement of the CLT the r.v. $\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right)$ converges to the standard Gaussian distribution. The term in parentheses standardizes $\bar{X}_n$; that is, the mean, μ, is subtracted from $\bar{X}_n$, and the result divided by the standard deviation, σ. But why multiply this standardized term by $\sqrt{n}$?

Answer 2: Because σ is the standard deviation of a single X_i, not of $\bar{X}_n$. The standard deviation of the sample mean equals $\sigma/\sqrt{n}$.[1] Standardizing $\bar{X}_n$ properly means dividing by *its* standard deviation, and division by a fraction is multiplication by its reciprocal:

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} = \sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right).$$

So the $\sqrt{n}$ isn't an extra factor; it's the part of the standardization that's easy to lose track of.

A third and equivalent statement of the CLT uses a standardized sum. Define the standardized sum of $X_1 + \cdots + X_n$ by:

$$Z_n = \frac{n\bar{X}_n - n\mu}{\sqrt{n}\sigma},$$

where $n\bar{X}_n$ is just $X_1 + \cdots + X_n$. As Z_n is standardized it follows that $\mathbb{E}(Z_n) = 0$ and $\mathbb{V}\mathrm{ar}(Z_n) = 1$. Further, let Z be a standard Gaussian r.v. Then the CLT states that:

$$\lim_{n\to\infty} \mathbb{P}(Z_n \leq z) = \mathbb{P}(Z \leq z), \text{ for every } z,$$

where $\mathbb{P}(Z \leq z)$ is the standard Gaussian CDF $\Phi(z)$. The reason why the CLT states that the CDF of Z_n (and not its PDF) converges to the standard Gaussian CDF is because the result applies to *all* r.v.s with finite variance, and all r.v.s have CDFs. The X_i can be discrete, as in the kiss experiment, in which case Z_n has no PDF at all.

Question 3: Why do you standardize the sums of the X_i's for the approximation?

Answer 3: Suppose the $X_1 + \cdots + X_n$ aren't standardized. Then $\mathbb{E}(X_1 + \cdots + X_n) = n\mu$ and $\mathbb{V}\mathrm{ar}(X_1 + \cdots + X_n) = n\sigma^2$. Both are finite for every n, but neither stays bounded as $n \to \infty$: the unstandardized sum drifts off (its mean is $n\mu$) and spreads out (its standard deviation is $\sigma\sqrt{n}$), so no fixed distribution can describe it and its CDF has no limit that is a CDF. Standardizing pins the sum down, with $\mathbb{E}(Z_n) = 0$ and $\mathbb{V}\mathrm{ar}(Z_n) = 1$ for all n, so that a limiting distribution has a chance to exist.

From undergraduate statistics: the '$n \geq 30$' rule

Introductory courses offer a rule of thumb: the normal approximation to the sample mean is fine once $n \geq 30$. The CLT contains no such number. It's a limit statement, and how fast the limit is reached depends on the parent distribution: above all on how asymmetric and heavy-tailed it

[1] To see why, let $X_1, \ldots, X_n$ be i.i.d. r.v.s. By independence, you have that:

$$\mathbb{V}\mathrm{ar}(\bar{X}_n) = \mathbb{V}\mathrm{ar}\left(\frac{X_1}{n} + \cdots + \frac{X_n}{n}\right) = \frac{\sigma^2}{n^2} + \cdots + \frac{\sigma^2}{n^2} = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}.$$

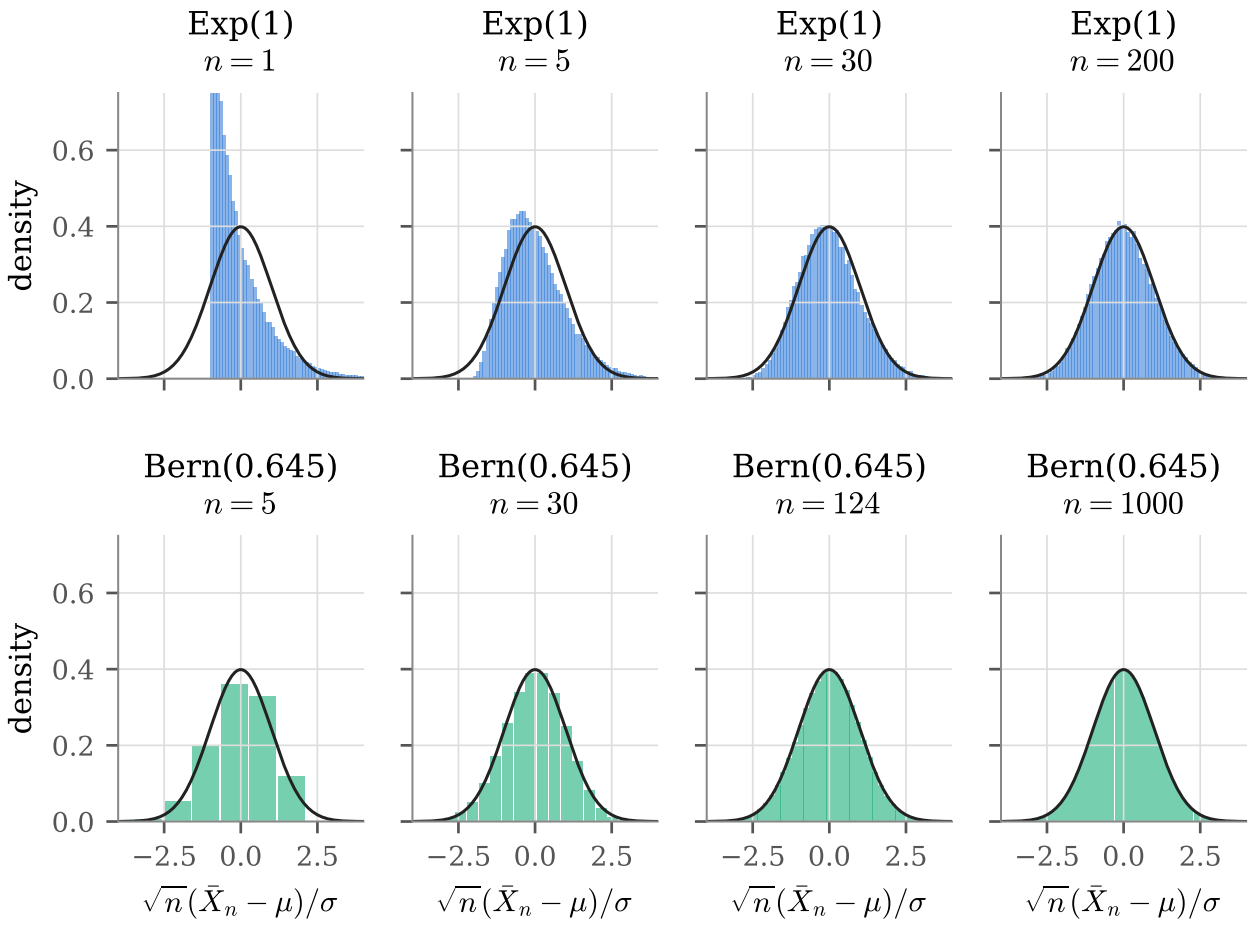


Figure 4.3: The CLT in action. Each panel shows the distribution of the standardized sample mean $\sqrt{n}(\bar{X}_n - \mu)/\sigma$ with the Gauss$(0, 1)$ density drawn over it. Top row: Expo(1) draws (histograms from 50,000 simulated samples). The parent is strongly skewed and the skew is still visible at $n = 30$. Bottom row: the kiss experiment, Bern(0.645) draws, drawn from the exact Binomial distribution. Here the standardized mean lives on a lattice, visible as bars; by $n = 124$ the Gaussian is already an excellent description.

is. Figure 4.3 shows a skewed Exponential parent where $n = 30$ is visibly not enough, and a Bernoulli parent with $p = 0.645$ where it's plenty. For a symmetric parent, $n = 5$ can be enough; for a heavily skewed one, $n = 200$ may not be. The number 30 is a convention, not a theorem. What is a theorem is the *rate*: the Berry–Esseen bound says the gap between the CDF of Z_n and Φ is at most a constant times $\mathbb{E}|X - \mu|^3/(\sigma^3\sqrt{n})$, so the error shrinks like $1/\sqrt{n}$, with a constant that is large for heavy-tailed parents. A finer expansion (the Edgeworth expansion) shows that the leading term of the error is proportional to the parent's skewness, which is why symmetric parents reach the Gaussian faster in practice.

4.2.1 The essence of the CLT: a 'stochastic' bound

Due to the so-called 'empirical rule', values between -3 and 3 in the support of a standard Gaussian account for nearly 100% of its probability mass (99.7% of it). By the CLT, this implies that, for n

large enough for the approximation to apply, the following inequality holds with probability about 99.7%:

$$|\bar{X}_n - \mu| \leq \frac{3\sigma}{\sqrt{n}},$$

since, for every n, the two events are the same:

$$\begin{aligned}\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right) \in [-3, 3] &\Leftrightarrow -3 \leq \sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right) \leq 3 \\ &\Leftrightarrow \left|\sqrt{n}\left(\frac{\bar{X}_n - \mu}{\sigma}\right)\right| \leq 3 \\ &\Leftrightarrow \frac{\sqrt{n}}{\sigma}|\bar{X}_n - \mu| \leq 3 \\ &\Leftrightarrow |\bar{X}_n - \mu| \leq \frac{3\sigma}{\sqrt{n}}.\end{aligned}$$

The essence of the CLT is that, despite $\bar{X}_n - \mu$ being an r.v. with an unbounded support (when the X_i are unbounded), it's *stochastically* bounded in the sense that, with high probability, it's between $-3\sigma/\sqrt{n}$ and $3\sigma/\sqrt{n}$. Replace 3 by 1.96 and 'overwhelming' by 95% and you have the recipe for every confidence interval in Chapter 8. Figure 4.4 draws the bound as a funnel around the running average, and it is also the chapter's two theorems in one picture.

4.3 Hoeffding's inequality: help for small sample sizes

What if your sample size, n, is not large enough to apply the CLT? Hoeffding's inequality is a helpful tool when n is small that effectively delivers the same essential message as the CLT: $\bar{X}_n$ is close to μ within some stochastic bound. The statement isn't as strong as the CLT and requires additional restrictions on the r.v.s. Namely, the r.v.s must be bounded almost surely, that is, with probability one (toolkit A.5), which rules out any r.v. with an arbitrarily large or small support (e.g., Gaussian and Exponential r.v.s). But despite the stronger restrictions, its power lies in providing a bound on the distance between $\bar{X}_n$ and μ that is a theorem rather than an approximation, for every n, however small. Here's the statement.

Theorem 4.9 (Hoeffding's inequality, 1963)**.** *Let n be some positive integer and $X, X_1, \ldots, X_n$ be i.i.d. r.v.s such that $\mu = \mathbb{E}[X]$ and X is almost surely bounded within $[a, b]$ where $a < b$. Then:*

$$\mathbb{P}\left(|\bar{X}_n - \mu| \geq \epsilon\right) \leq 2\exp\left(-\frac{2n\epsilon^2}{(b-a)^2}\right), \quad \textit{for all } \epsilon > 0.$$

Compare this with Chebyshev's bound $\sigma^2/(n\epsilon^2)$. Both go to zero as n grows, but Hoeffding's bound goes to zero *exponentially* fast in n while Chebyshev's goes like $1/n$. The price is the boundedness assumption, which Chebyshev doesn't need. Note also what Hoeffding's inequality is not: it's a bound, not an approximation. The CLT says $\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) \approx 2(1 - \Phi(\epsilon\sqrt{n}/\sigma))$ for large n without saying how large; Hoeffding says $\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) \leq 2e^{-2n\epsilon^2/(b-a)^2}$ for every $n \geq 1$, full stop.

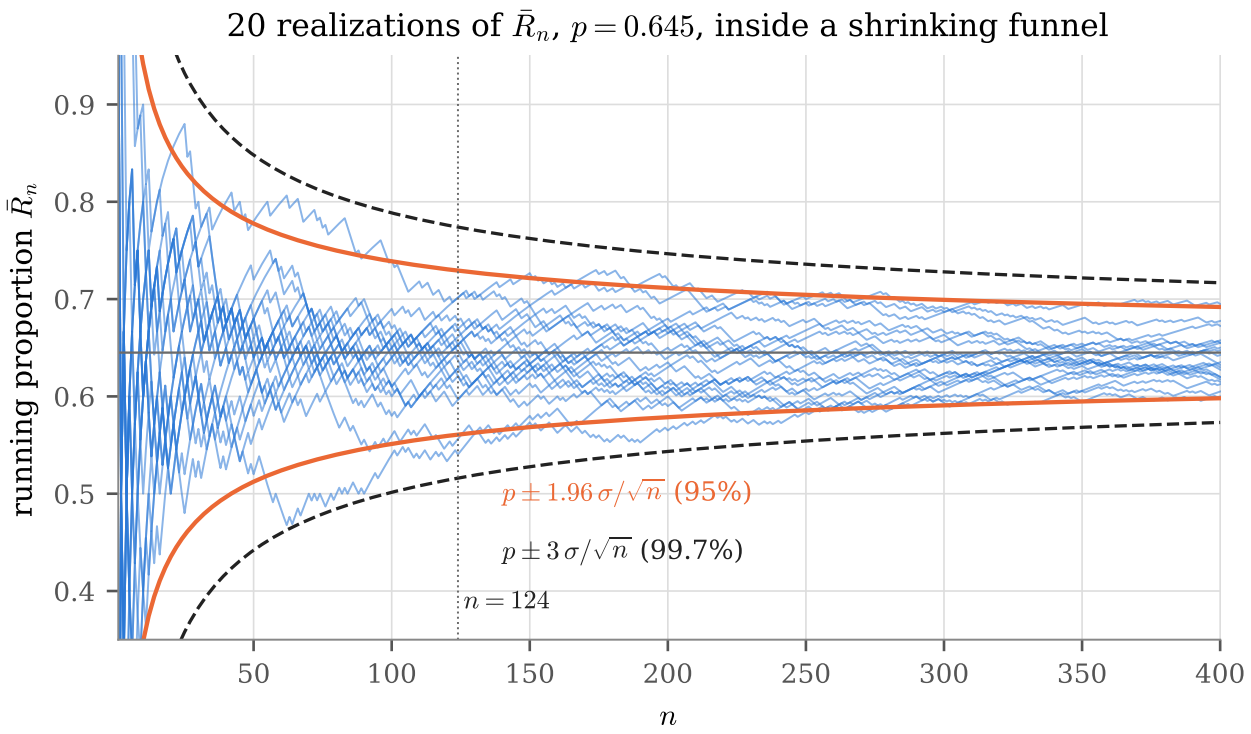


Figure 4.4: The Law of Large Numbers and the CLT in one picture. Twenty realizations of the running proportion $\bar{R}_n$ in the kiss model ($p = 0.645$), with the curves $p \pm 1.96\,\sigma/\sqrt{n}$ and $p \pm 3\,\sigma/\sqrt{n}$, $\sigma = \sqrt{p(1-p)}$. The funnel narrows like $1/\sqrt{n}$: that is the Law of Large Numbers. At any fixed n past the first few dozen, about 5% of the paths are outside the inner funnel and almost none outside the outer one (in this simulation, 6% of the points with $n \geq 30$ are outside the inner funnel): that is the stochastic bound. The inner funnel at $n = 124$ is the width of the confidence interval of Chapter 8.

Worked example: the kiss experiment, three ways

The R_i are Bernoulli, so they're bounded in $[0, 1]$ and $(b-a)^2 = 1$. With $n = 124$ and $\epsilon = 0.084$:

- Chebyshev (with $p(1-p)$ evaluated at 0.645): $\mathbb{P}(|\bar{R}_n - p| \geq 0.084) \leq 0.26$; with the worst case $p(1-p) \leq 1/4$ instead, 0.29.
- Hoeffding: $\mathbb{P}(|\bar{R}_n - p| \geq 0.084) \leq 2\exp(-2 \times 124 \times 0.084^2) = 2e^{-1.75} \approx 0.35$.
- CLT: $\mathbb{P}(|\bar{R}_n - p| \geq 0.084) \approx 2\left(1 - \Phi\left(\frac{0.084\sqrt{124}}{\sqrt{0.645 \times 0.355}}\right)\right) = 2(1 - \Phi(1.95)) \approx 0.05$.

At this n the exact Binomial probability is 0.05, so the CLT approximation is right on and both bounds are loose, with Chebyshev slightly tighter than Hoeffding. Figure 4.5 shows that Hoeffding overtakes Chebyshev at about $n = 165$ and is far ahead by $n = 400$, as the exponential must eventually beat $1/n$. The two bounds are guarantees; the CLT number is an approximation that happens, here, to be excellent.

4.4 What the limit theorems buy you

A fair question at this point: couldn't the results of this chapter and the last be proved without limits? More of them than you would guess, and it's worth seeing which because it shows what the

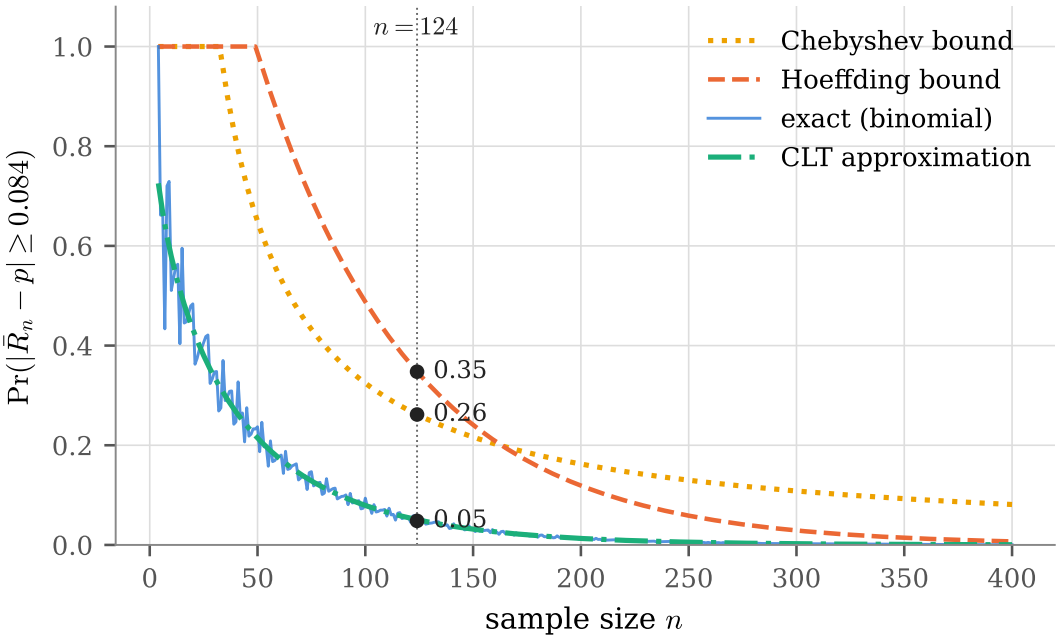


Figure 4.5: $\mathbb{P}(|\bar{R}_n - p| \geq 0.084)$ in the kiss model with $p = 0.645$, as a function of n: the exact Binomial probability, the CLT approximation (on top of it for n beyond a few dozen), and the Chebyshev and Hoeffding upper bounds. The jaggedness of the exact curve is real: it comes from the lattice of possible values of $\bar{R}_n$ moving relative to the fixed ϵ. The dots mark $n = 124$.

limit theorems are actually for.

Results that have a finite-sample version.

- The Law of Large Numbers is Chebyshev's inequality with n sent to infinity. The finite-n content is the inequality itself, $\mathbb{P}(|\bar{X}_n - \mu| \geq \epsilon) \leq \sigma^2/(n\epsilon^2)$; the limit only tidies the statement. Hoeffding's inequality (Section 4.3) is the same idea with a better bound; the 'conservative' confidence interval of Chapter 8 uses the same trick of replacing the unknown variance by its worst case, $p(1-p) \leq 1/4$. Both are guarantees that hold at $n = 2$ as well as at $n = 10^6$.
- When the parent distribution is known, an exact sampling distribution replaces the CLT. For Bernoulli data the sum $n\bar{R}_n$ is Binomial, so the kiss experiment's proportion has exact tests and an exact interval (the Clopper–Pearson interval, which for these data is $[0.55, 0.73]$). For Exponential data the sum is Gamma, which gives an exact interval for the subway rate λ (an exercise in Chapter 8). For Gaussian data the standardized mean has a t distribution and the ratio of variances an F distribution, which is why the t-test of Chapter 14 and Gaussian linear regression in Chapter 21 need no limit theorem at all.
- The CLT itself has a finite-sample version: the Berry–Esseen theorem mentioned in Section 4.2 bounds the gap between the CDF of the standardized mean and Φ at every n. The limit statement is that theorem with the error term dropped.

What has no finite-sample substitute. What you can't do without a limit theorem is get a statement that is both universal and sharp. Chebyshev and Hoeffding hold for every parent too, but as bounds that may sit far from the truth; an approximation that is accurate for every parent at once

is the whole content of the CLT. It's the only reason the same recipe, estimate $\pm 1.96\times$ standard error, works for proportions, means, rates, regression coefficients and maximum likelihood estimators of every model in this book. Without it you must either know the parent exactly, and then compute its exact sampling distribution, which is possible for the three families above and intractable for most others, or accept an inequality, which is honest but often far from the truth: Hoeffding says 0.35 where the exact answer is 0.05 in the kiss experiment (Figure 4.5). The asymptotic normality of a general maximum likelihood estimator (Chapter 12) and Wilks' theorem for the likelihood ratio test (Chapter 15) have no finite-sample versions; their exact distributions depend on the model and are usually unobtainable.

So the asymptotic route is a trade: one formula for every model, in exchange for not knowing how large n must be. The price is paid in the places where the approximation is poor, which the Berry–Esseen bound says are the heavy-tailed and asymmetric parents, and small n.

4.5 What comes next

You now have two limits for the sample mean: it collapses onto μ (Chapter 3) and, magnified by $\sqrt{n}$, its error is Gaussian (this chapter). Almost nothing in statistics is literally a sample mean. The estimator of the subway rate is $1/\bar{T}_n$, the standard error you can compute has $\hat{p}$ where the CLT has p, and the odds are a ratio of two things. The next chapter is the toolkit for carrying the two limits through such operations: Slutsky's theorem for plugging in, the continuous mapping theorem for transforming, and the delta method for the CLT of a function of an average.

4.6 Summary

- **Convergence in distribution:** $X_n \xrightarrow[n\to\infty]{\text{(d)}} X$ if $F_{X_n}(x) \to F_X(x)$ at every continuity point of F_X; equivalently $\mathbb{E}[f(X_n)] \to \mathbb{E}[f(X)]$ for every bounded continuous f. It compares distributions, not values: $X_n = -Z$ converges in distribution to Z without ever being close to it.
- Convergence in probability implies convergence in distribution (Proposition 4.4). The converse can fail, as the $X_n = -Z$ example shows, unless the limit is a constant, in which case it holds (Proposition 4.7).
- **Central Limit Theorem:** $\sqrt{n}(\bar{X}_n - \mu)/\sigma \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$, whatever the parent distribution with finite variance. With overwhelming probability $|\bar{X}_n - \mu| \le 3\sigma/\sqrt{n}$: a funnel that narrows like $1/\sqrt{n}$. How fast the Gaussian arrives depends on the parent, above all on its asymmetry; the '$n \ge 30$' rule is a convention.
- Averages are not the only things with limits, and $\sqrt{n}$ and the Gaussian are not the only rate and limit: the maximum of n uniforms has an error of order $1/n$ and an Exponential limit (Example 4.6).
- **Hoeffding's inequality:** for data in $[a,b]$, $\mathbb{P}(|\bar{X}_n - \mu| \ge \epsilon) \le 2e^{-2n\epsilon^2/(b-a)^2}$ for every n. A guarantee, not an approximation, and exponentially better than Chebyshev for large n.
- The limit theorems' content is universality: one formula for every parent. The price is not knowing how large n must be. When the parent is known, exact distributions (Binomial, Gamma, t, F) need no limits at all.

4.7 Exercises

Exercise 4.1. Let $T_n \sim \text{Unif}([5 - \frac{1}{2n}, 5 + \frac{1}{2n}])$. (a) For a fixed $0 < \delta < 1$, find N such that $\mathbb{P}(|T_n - 5| > \delta) = 0$ whenever $n \geq N$. (b) What is the limit of (T_n)? (c) Let F_n be the CDF of T_n and F the CDF of the limit. For which t does $F_n(t) \to F(t)$? *(Solution on page 222.)*

Exercise 4.2. Let $Y_1, Y_2, \ldots \overset{iid}{\sim} \text{Unif}([0, 1])$ and $M_n = \max(Y_1, \ldots, Y_n)$. (a) For $0 < \delta < 1$, find $\mathbb{P}(|M_n - 1| > \delta)$. (b) Does M_n converge in probability to a constant? (c) Find the CDF of M_n on $[0, 1]$. *(Solution on page 222.)*

Exercise 4.3. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, b])$, so $\sigma^2 = b^2/12$, with n too small to trust the CLT. Find the Chebyshev and Hoeffding upper bounds on $\mathbb{P}(|\bar{X}_n - \mathbb{E}[X]| \geq c\,\sigma/\sqrt{n})$ for $c = 2$ and $c = 6$. *(Solution on page 223.)*

Exercise 4.4. In the kiss experiment, use the CLT to find the smallest n for which $\mathbb{P}(|\bar{R}_n - p| \geq 0.02) \approx 0.05$ when $p = 0.645$, and compare with the Chebyshev answer of Section 3.7. *(Solution on page 223.)*

Chapter 5

Working with Limits: Slutsky, the Continuous Mapping Theorem and the Delta Method

In this chapter

- Which arithmetic survives a limit: sums, products and quotients for convergence in probability, and the one case (Slutsky's theorem) where it works for convergence in distribution.
- The continuous mapping theorem: a continuous function of a convergent sequence converges.
- The delta method: the Central Limit Theorem for a function of an average, with the derivative as the scale factor.
- A map from every asymptotic result in the rest of the book back to these tools.

You now have two limits in hand: $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} \mu$ from Chapter 3 and $\sqrt{n}(\bar{X}_n - \mu) \xrightarrow[n\to\infty]{\text{(d)}}$ Gauss$(0, \sigma^2)$ from Chapter 4. Almost nothing in statistics is literally an average, so the question is what you may do to these limits and still know where you end up.

5.1 Arithmetic and convergence

For convergence in probability and convergence almost surely, the arithmetic properties of limits of real-valued sequences have analogous arithmetic properties. (Almost sure convergence is defined in Appendix B; on the minimal path, read the rules below for convergence in probability, the only one of the two the rest of the book uses.)

Let $(X_n)_{n\geq 1}$ and $(Y_n)_{n\geq 1}$ be sequences of r.v.s and suppose that:

$$X_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} X \text{ and } Y_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} Y.$$

Then:

1. **Additive property:** $X_n + Y_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} X + Y$.

2. **Multiplicative property:** $X_n Y_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} XY$.

3. **Division property:** If $Y \neq 0$ almost surely, then:

$$\frac{X_n}{Y_n} \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} \frac{X}{Y}.$$

For almost sure convergence these are just the calculus rules applied outcome by outcome: on the probability-one set where both sequences of numbers converge, their sums, products and quotients converge to the sums, products and quotients of the limits. For convergence in probability the rules hold too, though the proof is a little longer.

NB: These rules **don't** apply to convergence in distribution. The pitfall box in Section 4.1 gave the counterexample: $X_n = -Z$ and $Y_n = Z$ both converge in distribution to Gauss$(0,1)$, but $X_n + Y_n = 0$, while two independent sequences with the same marginal limits would have sum converging to Gauss$(0,2)$. The marginal limits don't determine the limit of the sum because they say nothing about how X_n and Y_n are related. This is where Slutsky's theorem comes in: it's the special case where one of the two limits is a constant, so there's no relationship to worry about.

5.2 Slutsky's theorem

For convergence in distribution, Slutsky's theorem is the main tool, and the kiss experiment shows why you need one. The CLT says

$$\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{p(1-p)}} \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1),$$

but the denominator involves the unknown p, so you cannot compute the left side from data. You'd like to replace $\sqrt{p(1-p)}$ by $\sqrt{\hat{p}(1-\hat{p})}$ with $\hat{p} = \bar{R}_n$. Write the thing you want as a product:

$$\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{\hat{p}(1-\hat{p})}} = \underbrace{\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{p(1-p)}}}_{\xrightarrow[n\to\infty]{\text{(d)}}\text{Gauss(0,1) by the CLT}} \cdot \underbrace{\sqrt{\frac{p(1-p)}{\hat{p}(1-\hat{p})}}}_{\xrightarrow[n\to\infty]{\mathbb{P}}\text{1 by the LLN and CMT}}.$$

The first factor converges in distribution and the second in probability, to the constant 1. What is the limit of the product? The pitfall of Section 4.1 says that in general the two marginal limits don't determine it. Slutsky's theorem says that when one of the limits is a constant, they do.

Theorem 5.1 (Slutsky)**.** *Let $(X_n)_{n\geq 1}$ and $(Y_n)_{n\geq 1}$ be sequences of r.v.s such that:*

$$X_n \xrightarrow[n\to\infty]{(d)} X \text{ and } Y_n \xrightarrow[n\to\infty]{\mathbb{P}} y,$$

where X is an r.v. and y is a given real number; that is, the limit of Y_n is deterministic. Then:

1. ***Additive property:*** $X_n + Y_n \xrightarrow[n\to\infty]{(d)} X + y$.

2. ***Multiplicative property:*** $X_n Y_n \xrightarrow[n\to\infty]{(d)} yX$.

3. ***Division property:*** *If $y \neq 0$, then:*

$$\frac{X_n}{Y_n} \xrightarrow[n\to\infty]{(d)} \frac{X}{y}.$$

Slutsky's theorem is useful precisely because it allows you to combine multiple r.v.s with known asymptotics, and retain this knowledge; that is, you know what the resultant distribution will converge to as $n \to \infty$. Its typical use: the CLT hands you a sequence converging in distribution to a Gaussian, the Law of Large Numbers hands you a sequence converging in probability to a constant, and Slutsky lets you multiply or divide the one by the other. Note that, by Proposition 4.7, the hypothesis $Y_n \xrightarrow[n\to\infty]{\mathbb{P}} y$ could equally be written $Y_n \xrightarrow[n\to\infty]{(\mathrm{d})} y$; the theorem needs the limit to be a constant, not a particular mode.

Worked example: the plug-in standardization in the kiss experiment, finished

Back to the product above. The second factor converges in probability to 1 because $\hat{p} \xrightarrow[n\to\infty]{\mathbb{P}} p$ (Law of Large Numbers) and $x \mapsto \sqrt{p(1-p)/(x(1-x))}$ is continuous at p (continuous mapping theorem, next). Slutsky's multiplicative property then gives

$$\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{\hat{p}(1-\hat{p})}} \xrightarrow[n\to\infty]{(\mathrm{d})} 1 \cdot \mathrm{Gauss}(0,1) = \mathrm{Gauss}(0,1).$$

This is the plug-in route to the confidence interval for p, and it's the justification for every $\hat{p} \pm 1.96\sqrt{\hat{p}(1-\hat{p})/n}$ you've ever computed.

From undergraduate statistics: why you're allowed to use s in place of σ

Every introductory course replaces the unknown σ by the sample standard deviation s in the standard error and carries on as if nothing happened (sometimes switching from z to t as an apology). Slutsky's theorem is the reason this works for large n: $s \xrightarrow[n\to\infty]{\mathbb{P}} \sigma$ by the Law of Large Numbers applied to the squares, so

$$\frac{\bar{X}_n - \mu}{s/\sqrt{n}} = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \cdot \frac{\sigma}{s} \xrightarrow[n\to\infty]{(\mathrm{d})} \mathrm{Gauss}(0,1) \cdot 1.$$

The t distribution is the exact finite-sample answer when the data are Gaussian (Chapter 14); Slutsky is the large-sample answer for any parent.

5.3 Continuous mapping theorem (CMT)

The Continuous Mapping Theorem (CMT) states that if there's some random variable X for which $X_n \xrightarrow[n\to\infty]{(\mathrm{d})} X$, then $f(X_n) \xrightarrow[n\to\infty]{(\mathrm{d})} f(X)$, too, provided f is a continuous function (a function whose graph has no jumps; a small change in the input gives a small change in the output). In short, CMT states that a continuous function preserves the asymptotic limits of an r.v.

But the CMT is actually more general than explained above. Specifically, it applies to all three modes of convergence, almost surely (Appendix B) and in probability as well as in distribution. Here's its more general statement.

Theorem 5.2 (Continuous mapping theorem)**.** *If f is a continuous function and $(X_n)_{n\geq 1}$ is a sequence of r.v.s, then:*

$$X_n \xrightarrow[n\to\infty]{a.s./\mathbb{P}/(d)} X \quad \Rightarrow \quad f(X_n) \xrightarrow[n\to\infty]{a.s./\mathbb{P}/(d)} f(X).$$

For almost sure convergence the theorem is the calculus fact that a continuous function of a convergent sequence of numbers converges, applied outcome by outcome. For convergence in distribution it follows from the $\mathbb{E}[g(X_n)]$ characterization of Proposition 4.3, since $g \circ f$ is continuous and bounded whenever g is. Only continuity at the points where the limit lives is needed: if $X_n \xrightarrow[n\to\infty]{\mathbb{P}} p$ with $p > 0$, then $1/X_n \xrightarrow[n\to\infty]{\mathbb{P}} 1/p$ even though $f(x) = 1/x$ isn't continuous at 0.

Example 5.3. By the Law of Large Numbers, $\bar{R}_n \xrightarrow[n\to\infty]{\mathbb{P}} p$. Therefore $f(\bar{R}_n) \xrightarrow[n\to\infty]{\mathbb{P}} f(p)$ for any continuous f: the sample odds $\bar{R}_n/(1-\bar{R}_n)$ converge to the true odds $p/(1-p)$, the sample log-odds converge to the true log-odds, and so on. Consistency survives continuous transformations.

Pitfall: what the CMT does not give you

By the CLT, $\sqrt{n}(\bar{R}_n - p) \xrightarrow[n\to\infty]{(\mathrm{d})} Y$ with $Y \sim \mathrm{Gauss}(0, p(1-p))$, so by the CMT

$$f\left(\sqrt{n}(\bar{R}_n - p)\right) \xrightarrow[n\to\infty]{(\mathrm{d})} f(Y).$$

That is a true statement, and almost never the one you want. What you want is the limit of $\sqrt{n}\left(f(\bar{R}_n) - f(p)\right)$: the magnified error of the transformed *estimator*, not a transformation of the magnified error. The CMT is silent on this. The delta method is the tool, and Figure 5.1 shows the two objects side by side.

5.4 The delta method

The delta method is what you need when you're taking the limit of something like $\sqrt{n}\left(g(\bar{X}_n) - g(\mu)\right)$: you know the CLT for $\bar{X}_n$, and you want the CLT for a function of it.

Theorem 5.4 (Delta method)**.** *Let $(Z_n)_{n\geq 1}$ be a sequence of r.v.s such that:*

$$\sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{(d)} Gauss(0, \sigma^2),$$

for some $\theta \in \mathbb{R}$ and $\sigma^2 > 0$. A sequence like $(Z_n)_{n\geq 1}$ that satisfies this condition is said to be asymptotically normal around θ, and σ^2 is its asymptotic variance. Let $g : \mathbb{R} \to \mathbb{R}$ be continuously differentiable at the point θ. Then:

1. *$g(Z_n) \xrightarrow[n\to\infty]{\mathbb{P}} g(\theta)$;*
2. *$(g(Z_n))_{n\geq 1}$ is also asymptotically normal, around $g(\theta)$;*
3. *more precisely,*

$$\sqrt{n}\left(g(Z_n) - g(\theta)\right) \xrightarrow[n\to\infty]{(d)} Gauss\left(0,\, g'(\theta)^2\sigma^2\right).$$

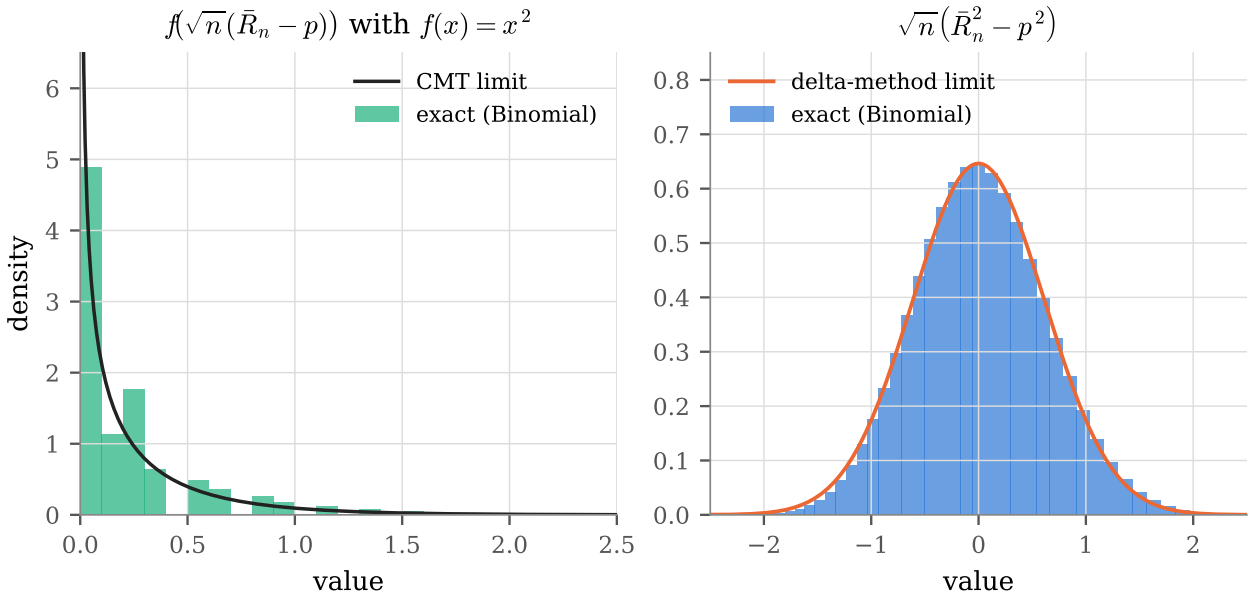


Figure 5.1: Two different random variables, for $f(x) = x^2$ in the kiss model ($p = 0.645$, $n = 124$; p^2 is the probability that two independent couples both turn right). Left: f applied to the magnified error, $(\sqrt{n}(\bar{R}_n - p))^2$, which by the CMT converges to the square of a Gaussian, $\sigma^2\chi_1^2$ with $\sigma^2 = p(1-p)$, a chi-squared shape (the curve). True, and useless for saying how accurate $\bar{R}_n^2$ is. Right: the magnified error of the transformed estimator, $\sqrt{n}(\bar{R}_n^2 - p^2)$, which the delta method says is Gauss$(0, (2p)^2 p(1-p))$, standard deviation 0.62 (the curve). The bars in both panels are the exact Binomial distribution of $n\bar{R}_n$, binned.

Intuition: a differentiable function is locally linear

Figure 5.2 is the whole idea. For large n, Z_n sits in a narrow Gaussian blob of width $\sigma/\sqrt{n}$ around θ. On a blob that narrow, the curve g is indistinguishable from its tangent line at θ, which has slope $g'(\theta)$. A linear function maps a Gaussian to a Gaussian, shifting the center to $g(\theta)$ and multiplying the width by the slope. So $g(Z_n)$ sits in a Gaussian blob of width $|g'(\theta)|\,\sigma/\sqrt{n}$ around $g(\theta)$, which is statement (3). Where g is steep, errors in Z_n are amplified; where it's flat, they're damped.

Proof strategy (core). The mean value theorem (toolkit A.9) writes $g(Z_n) - g(\theta)$ as a slope times $Z_n - \theta$. The slope converges in probability to $g'(\theta)$ (CMT), and $\sqrt{n}(Z_n - \theta)$ converges in distribution (the assumption); Slutsky multiplies the two.

Proof. By the mean value theorem there's a point $\tilde{\theta}_n$ between Z_n and θ such that

$$g(Z_n) - g(\theta) = g'(\tilde{\theta}_n)(Z_n - \theta).$$

Multiply through by $\sqrt{n}$:

$$\sqrt{n}\,(g(Z_n) - g(\theta)) = g'(\tilde{\theta}_n) \cdot \sqrt{n}(Z_n - \theta).$$

Now take the two factors in turn. First, $Z_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$: by Slutsky, $Z_n - \theta = \frac{1}{\sqrt{n}} \cdot \sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{(\mathrm{d})} 0 \cdot \text{Gauss}(0, \sigma^2) = 0$, and convergence in distribution to a constant is convergence in probability

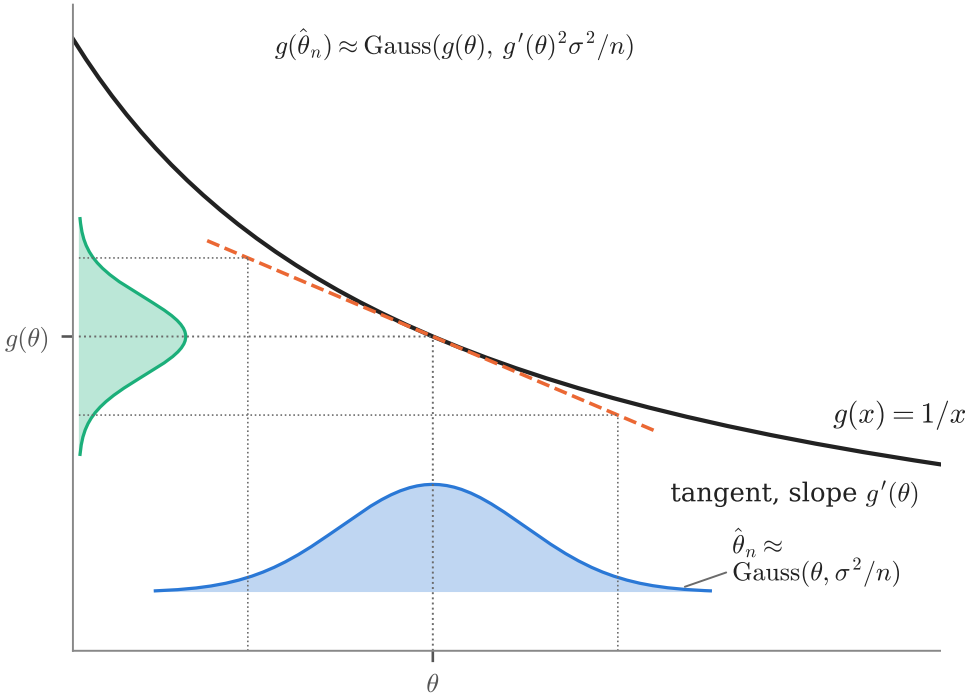


Figure 5.2: The delta method for $g(x) = 1/x$, with $\theta = 6.23$ and $\sigma/\sqrt{n} = 0.78$ (the subway example below). The Gaussian around θ on the horizontal axis is pushed through g. Because the blob is narrow, g acts like its tangent line, and the image on the vertical axis is a Gaussian around $g(\theta)$ whose width is the input width times $|g'(\theta)|$. The dotted lines map $\theta \pm 2\sigma/\sqrt{n}$ across.

(Proposition 4.7). This proves (1) via the CMT, and since $\tilde{\theta}_n$ is squeezed between Z_n and θ, also $\tilde{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$, hence $g'(\tilde{\theta}_n) \xrightarrow[n\to\infty]{\mathbb{P}} g'(\theta)$ by the CMT and the continuity of g' at θ. Second, $\sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{(\mathrm{d})} \text{Gauss}(0, \sigma^2)$ by assumption. Slutsky's multiplicative property with $y = g'(\theta)$ gives

$$g'(\tilde{\theta}_n) \cdot \sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{(\mathrm{d})} g'(\theta) \cdot \text{Gauss}(0, \sigma^2) = \text{Gauss}(0, g'(\theta)^2\sigma^2),$$

which is (3), and (2) is (3) read as a sentence. □

The proof is three lines because Slutsky and the CMT did the work in advance. Note that nothing requires g to be invertible or monotone; $g(x) = \sqrt{2}\sin x$ is a perfectly good choice. Note also that the asymptotic variance $g'(\theta)^2\sigma^2$ generally depends on the unknown θ, even when σ^2 doesn't; the plug-in trick of the previous section handles that.

Worked example: inter-arrival times of the T

You observe the times between arrivals of the subway at Kendall, $T_1, \ldots, T_n$, and model them as i.i.d. Expo(λ), so $\mathbb{E}[T_1] = 1/\lambda$ and $\mathbb{V}\text{ar}(T_1) = 1/\lambda^2$. The natural estimator of λ is $\hat{\lambda} = 1/\bar{T}_n$.

- *Consistency.* $\bar{T}_n \xrightarrow[n\to\infty]{\mathbb{P}} 1/\lambda$ by the Law of Large Numbers, so $\hat{\lambda} = 1/\bar{T}_n \xrightarrow[n\to\infty]{\mathbb{P}} \lambda$ by the CMT ($f(x) = 1/x$ is continuous at $1/\lambda > 0$).
- *The CLT for $\bar{T}_n$.* $\sqrt{n}\left(\bar{T}_n - \frac{1}{\lambda}\right) \xrightarrow[n\to\infty]{(\mathrm{d})} \text{Gauss}(0, \lambda^{-2})$.

- *The delta method for* $\hat{\lambda}$. Take $g(x) = 1/x$, so $g'(x) = -1/x^2$ and $g'(1/\lambda)^2 = \lambda^4$. Then

$$\sqrt{n}\left(\hat{\lambda} - \lambda\right) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}\left(0,\, \lambda^4 \cdot \lambda^{-2}\right) = \text{Gauss}(0, \lambda^2).$$

With $n = 64$ and $\bar{T}_n = 6.23$ minutes: $\hat{\lambda} = 0.1605$ per minute, and the standard error of $\hat{\lambda}$ (the standard deviation $\lambda/\sqrt{n}$ of its approximating Gaussian, with $\hat{\lambda}$ plugged in for λ) is about 0.020. So λ is pinned down to roughly 0.16 ± 0.04, which becomes the plug-in confidence interval $[0.12, 0.20]$ of Chapter 8.

Worked example: odds in the kiss experiment

The odds of turning right are $\omega = p/(1-p)$, estimated by $\hat{\omega} = \bar{R}_n/(1-\bar{R}_n)$. With $g(p) = p/(1-p)$ you get $g'(p) = 1/(1-p)^2$, and the CLT gives asymptotic variance $p(1-p)$ for $\bar{R}_n$, so

$$\sqrt{n}\,(\hat{\omega} - \omega) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}\left(0,\, \frac{p(1-p)}{(1-p)^4}\right) = \text{Gauss}\left(0,\, \frac{p}{(1-p)^3}\right).$$

At $p = 0.645$ and $n = 124$: $\hat{\omega} = 1.82$ and the standard error is $\sqrt{0.645/0.355^3/124} \approx 0.34$. The derivative $1/(1-p)^2 \approx 7.9$ is large, so the odds are estimated far less precisely, in absolute terms, than the proportion they're computed from: a steep g amplifies the error. Figure 5.3 puts this next to a flat g, where the same input error nearly disappears.

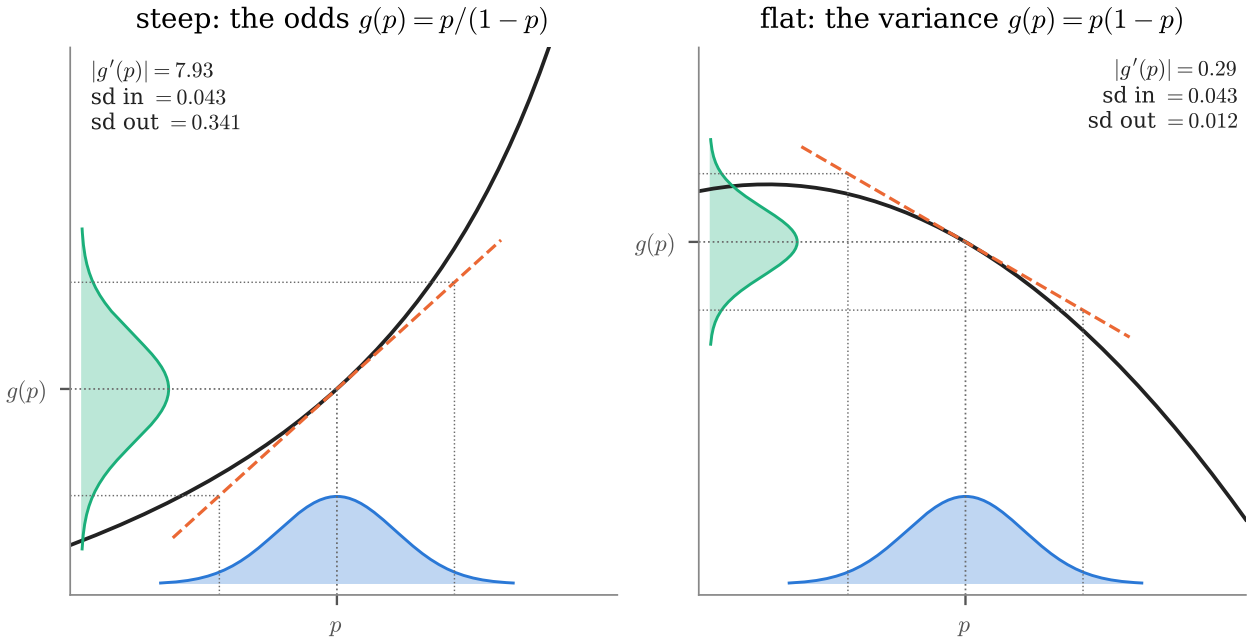


Figure 5.3: The same input, $\bar{R}_n \approx \text{Gauss}(p, p(1-p)/n)$ with standard deviation 0.043, pushed through two functions. Left: the odds $g(p) = p/(1-p)$ have slope 7.9 at $p = 0.645$, and the output standard deviation is $7.9 \times 0.043 = 0.34$. Right: the variance $g(p) = p(1-p)$ has slope -0.29 there, and the output standard deviation is $0.29 \times 0.043 = 0.012$. Same data, same precision in p, very different precision in $g(p)$: the delta method's $|g'(\theta)|$ is the whole difference.

From undergraduate statistics: propagation of error

If you took a laboratory course you met the rule 'if $y = g(x)$ and x has uncertainty δx, then y has uncertainty $\delta y \approx |g'(x)|\,\delta x$'. The delta method is that rule for standard errors, with a theorem attached:

$$\operatorname{se}\left(g(\hat{\theta}_n)\right) \approx |g'(\theta)|\,\operatorname{se}(\hat{\theta}_n), \qquad \operatorname{se}(\hat{\theta}_n) = \frac{\sigma}{\sqrt{n}}.$$

It also says what the lab rule leaves out: the transformed estimate isn't just 'uncertain by δy', it's approximately Gaussian, so you can attach a probability to the error bar.

Remark. When $g'(\theta) = 0$ the theorem still holds but says only that $\sqrt{n}(g(Z_n) - g(\theta)) \xrightarrow[n\to\infty]{\text{(d)}} 0$, that is, the error of $g(Z_n)$ is of smaller order than $1/\sqrt{n}$. To see the actual limit you magnify by n instead of $\sqrt{n}$ and get a second-order expansion: $n(g(Z_n) - g(\theta)) \xrightarrow[n\to\infty]{\text{(d)}} \frac{\sigma^2 g''(\theta)}{2}\chi^2_1$. No result later in this book needs this case.

Remark (Vector versions). Everything in this chapter has a version for vectors, with a covariance matrix in place of σ^2 and a matrix of partial derivatives in place of $g'(\theta)$. If $\sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}_d(0, \Sigma)$ and $g : \mathbb{R}^d \to \mathbb{R}^k$ is continuously differentiable at θ, with $\nabla g(\theta)$ the $d \times k$ matrix of partial derivatives, then

$$\sqrt{n}\,(g(Z_n) - g(\theta)) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}_k\left(0,\, \nabla g(\theta)^\top \Sigma\, \nabla g(\theta)\right),$$

and the multivariate CLT says that the vector of sample means of i.i.d. vectors with covariance Σ satisfies the hypothesis with Σ itself. The proofs are the ones above with matrices in place of numbers. The method of moments (Chapter 13) and the chi-squared goodness-of-fit test (Chapter 16) use these, and the book states them again there in the form each needs.

5.5 Where the limit theorems are used in the rest of the book

Every asymptotic statement in the later chapters is one of the facts of Chapters 3 and 4 or this chapter in disguise, with two exceptions the table names (the two theorems about the whole empirical CDF behind the Kolmogorov–Smirnov test). The table says which. 'LLN' means the Law of Large Numbers (Theorem 3.5), 'CLT' the Central Limit Theorem (Theorem 4.8), 'CMT' the continuous mapping theorem (Theorem 5.2), 'Slutsky' Theorem 5.1, and 'delta' the delta method (Theorem 5.4).

Result	Chapter	Rests on
Consistency of $\bar{X}_n$, $\hat{p}$, and any estimator built from an average	Chapter 7	LLN, then CMT for functions of averages
Asymptotic confidence intervals; the plug-in interval	Chapter 8	CLT; Slutsky for replacing p by $\hat{p}$ in the standard error
Confidence interval for $1/\bar{T}_n$ and other transformed estimators	Chapter 8	delta
Level and p-value of a test built on a standardized average	Chapter 9	CLT, Slutsky
Consistency of the maximum likelihood estimator	Chapter 12	LLN applied to the average log-likelihood, plus the KL divergence argument
Asymptotic normality of the MLE, with variance $1/I(\theta)$	Chapter 12	CLT applied to the average score, then a Taylor expansion exactly as in the delta method proof, with Fisher information as the slope
Method of moments: consistency and asymptotic variance	Chapter 13	LLN, then the (multivariate) delta method
Wald test	Chapter 15	asymptotic normality of the MLE, CMT (squaring gives a χ^2)
Likelihood ratio test (Wilks' theorem)	Chapter 15	a second-order Taylor expansion, the same device as the delta method, plus Slutsky
t-test	Chapter 14	nothing here: exact for Gaussian data (Cochran's theorem)
Chi-squared goodness-of-fit test	Chapter 16	multivariate CLT for the vector of counts, then CMT
Kolmogorov–Smirnov and Lilliefors tests	Chapter 17	two named theorems stated there without proof: Glivenko–Cantelli (a uniform LLN) and Donsker (a CLT for the whole empirical CDF)
Bayesian posteriors concentrate on the truth; Bernstein–von Mises	Chapter 18	LLN and CLT for the likelihood; stated, not proved
Inference for least squares with non-Gaussian errors	Chapter 21	CLT (exact t and F when errors are Gaussian)
Asymptotic normality of GLM estimates	Chapter 23	the MLE result again
Sample covariance converges to the true covariance	Chapter 24	LLN entrywise, then CMT for the eigenvectors

5.6 What comes next

Every limit in these three chapters landed on the same distribution. Chapter 6 collects what the rest of the book needs about it: the density, the three properties that let you standardize anything Gaussian, how to turn a Gaussian probability into a table lookup, and this book's convention for quantiles, which is the notation every confidence interval and test from Chapter 8 on is written in. If you know the Gaussian well, the quantile section is the one part not to skip.

5.7 Summary

- For convergence in probability (and almost surely), limits add, multiply and divide like limits of numbers. For convergence in distribution they do not: the marginal limits of X_n and Y_n don't determine the limit of $X_n + Y_n$.
- **Slutsky's theorem:** if $X_n \xrightarrow[n\to\infty]{\text{(d)}} X$ and $Y_n \xrightarrow[n\to\infty]{\mathbb{P}} y$ with y a constant, then $X_n + Y_n \xrightarrow[n\to\infty]{\text{(d)}} X + y$, $X_n Y_n \xrightarrow[n\to\infty]{\text{(d)}} yX$ and $X_n / Y_n \xrightarrow[n\to\infty]{\text{(d)}} X/y$. This justifies every plug-in: replacing σ by s, or p by $\hat{p}$, in a standard error.
- **Continuous mapping theorem:** $X_n \to X$ in any of the three modes implies $f(X_n) \to f(X)$ in the same mode, for f continuous at the limit. Consistency survives transformations. What it does not give you is the limit of $\sqrt{n}(f(\bar{X}_n) - f(\mu))$; that is the delta method's job.
- **Delta method:** if $\sqrt{n}(Z_n - \theta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, \sigma^2)$ and g is continuously differentiable at θ, then $\sqrt{n}(g(Z_n) - g(\theta)) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, g'(\theta)^2\sigma^2)$. A narrow Gaussian blob pushed through a smooth curve comes out Gaussian, with its width multiplied by the slope. This is propagation of error with a theorem attached.

5.8 Exercises

Exercise 5.1. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$. Use the delta method to find the asymptotic distribution of $\sqrt{n}\left(\bar{X}_n(1 - \bar{X}_n) - p(1-p)\right)$. For which p does the delta method fail to give a Gaussian limit, and why? *(Solution on page 223.)*

Exercise 5.2. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Expo}(\lambda)$. Show that $\sqrt{n}(\log \bar{X}_n + \log \lambda) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1)$, whatever λ is. Why might a statistician prefer to build a confidence interval for $\log(1/\lambda)$ rather than for $1/\lambda$? *(Solution on page 223.)*

Exercise 5.3. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$ with $0 < p < 1$. Show that

$$\sqrt{n}\,\frac{\bar{X}_n - p}{\sqrt{\bar{X}_n(1 - \bar{X}_n)}} \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1),$$

naming the theorem used at each step. *(Solution on page 223.)*

Chapter 6

The Gaussian Distribution and Quantiles

In this chapter

- The Gaussian density: what each piece of the formula does, why the CDF has no closed form, and why the tails are so thin that the distribution lives in a finite interval for practical purposes.
- The three properties used constantly from here on: closure under affine transformations, standardization, and symmetry. Plus the one that usually goes unstated: sums of independent Gaussians are Gaussian.
- How to turn a Gaussian probability into a table lookup, forwards and backwards.
- Quantiles, in this book's convention q_α (the number with mass α to its *right*), the four values worth memorizing, and the one-sided versus two-sided distinction that every confidence interval and test depends on.

This chapter is a review. If you know it, read the quantile section for the notation and move on.

6.1 Why the Gaussian is everywhere

Chapters 3 to 5 kept producing Gaussian limits; this chapter collects what you need about the Gaussian to use them, and fixes the quantile notation of the rest of the book. Because of the Central Limit Theorem (Chapter 4), the Gaussian (also called the normal) distribution is ubiquitous in statistics: every average of enough independent, identically distributed draws with finite variance is approximately Gaussian, whatever the draws come from, so every estimator built from an average inherits a Gaussian error distribution for large n. That is where the Gaussian enters statistical inference, and it stays: confidence intervals, test statistics, the asymptotic distribution of the maximum likelihood estimator, the error term of linear regression. The distribution is named after Carl Friedrich Gauss (1777–1855), who arrived at it in the context of the method of least squares, which is the subject of Chapter 20.

6.2 The density and its CDF

Definition 6.1 (Gaussian distribution)**.** A continuous r.v. X has the **Gaussian distribution** with parameters $\mu \in \mathbb{R}$ and $\sigma^2 > 0$ if its PDF is given by:

$$f_{\mu,\sigma^2}(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad \text{for } x \in (-\infty, \infty).$$

You write this as $X \sim \text{Gauss}(\mu, \sigma^2)$, and then $\mathbb{E}[X] = \mu$ and $\mathbb{V}\text{ar}(X) = \sigma^2$. The case $\mu = 0$, $\sigma^2 = 1$ is the **standard Gaussian**, whose density and CDF get their own letters:

$$\varphi(z) = \frac{1}{\sqrt{2\pi}} e^{-z^2/2}, \qquad \Phi(z) = \int_{-\infty}^{z} \varphi(t)\,\mathrm{d}t.$$

> **In words.** A Gaussian variable is one whose probabilities are areas under the bell-shaped curve above (toolkit A.6); μ says where the bell is centered and σ how wide it is. Two numbers fix the whole distribution.
>
> **For the kiss data.** The CLT says $\bar{R}_n$ behaves, for large n, like a Gaussian with $\mu = p$ and $\sigma^2 = p(1-p)/n$: centered at the truth, with width 0.043 at $n = 124$ and $p = 0.645$.

Take the formula apart. The constant $\dfrac{1}{\sigma\sqrt{2\pi}}$ is a *normalizing constant*: it normalizes the total area under the PDF to 1, and plays no other role. Everything about the shape is in the exponent. Consider the quadratic there:

$$-\frac{(x-\mu)^2}{2\sigma^2} = -\left(\frac{x^2 - 2\mu x + \mu^2}{2\sigma^2}\right).$$

It is zero at $x = \mu$ and negative everywhere else, so the density peaks at μ and falls off symmetrically on both sides. If σ is small, that is, if the standard deviation is small, then the quadratic is more negative at the same distance from μ, so the exponential falls off faster and the density is narrower. A small standard deviation and a narrow density are the same thing seen twice. Figure 6.1 shows the effect of each parameter.

Two consequences of the exponent are worth stating on their own.

- **The tails are thin.** The density decays like $e^{-x^2/(2\sigma^2)}$, faster than any exponential $e^{-c|x|}$, so for practical purposes a Gaussian lives in a finite interval: 99.7% of the mass is within 3σ of μ, and the probability of being more than 5σ away is about one in two million. This is what licenses a statement like 'with overwhelming probability $|\bar{X}_n - \mu| \le 3\sigma/\sqrt{n}$' (Section 4.2).
- **The CDF has no closed form.**

  $$\mathbb{P}(X \le x) = F_{\mu,\sigma^2}(x) = \frac{1}{\sigma\sqrt{2\pi}} \int_{-\infty}^{x} \exp\left(-\frac{(t-\mu)^2}{2\sigma^2}\right) \mathrm{d}t$$

  must be left as an integral: it's mathematically impossible to write the antiderivative of $e^{-t^2/2}$ in terms of elementary functions. Still, F_{μ,σ^2} is a perfectly well-defined function: input some value x and it returns the accumulated area under the PDF from $-\infty$ up to x. You get its values from tables or from software, and Section 6.4 says how.

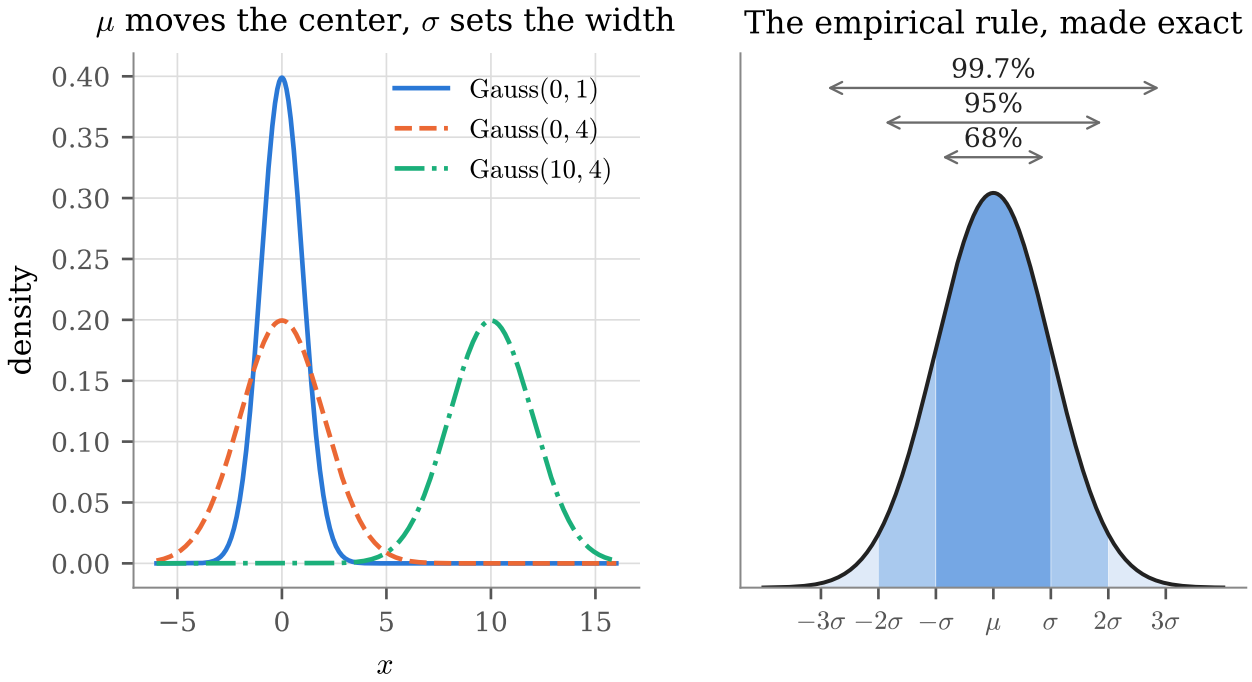


Figure 6.1: Left: μ moves the density without changing its shape; σ stretches it without moving its center, and the height drops to keep the area at 1. Right: the standard Gaussian with the regions within one, two and three standard deviations of the mean. Their probabilities are 0.6827, 0.9545 and 0.9973, which is the '68–95–99.7 rule' of an introductory course stated exactly.

6.3 Properties of Gaussian distributions

1. **Affine transformation.** Perhaps the most useful property of the Gaussian family is that it is invariant under affine transformations. Given the r.v. $X \sim \text{Gauss}(\mu, \sigma^2)$, then

$$aX + b \sim \text{Gauss}(a\mu + b, a^2\sigma^2) \qquad \text{for all } a \neq 0,\ b \in \mathbb{R}.$$

Shifting and rescaling a Gaussian produces another Gaussian; the mean and variance change in the way means and variances always do under $aX + b$, and the shape stays Gaussian, which is the content of the statement.

Proof strategy (for completeness; skip on a first reading). Write the CDF of $aX + b$ as a CDF of X at a transformed point, differentiate (chain rule) to get the density, and recognize the Gaussian formula with new parameters (toolkit A.6).

Proof. Take $a > 0$ (the case $a < 0$ is the same with the inequality flipped). The CDF of $Y = aX + b$ is

$$F_Y(y) = \mathbb{P}(aX + b \leq y) = \mathbb{P}\left(X \leq \frac{y-b}{a}\right) = F_X\left(\frac{y-b}{a}\right).$$

Differentiate with respect to y (chain rule) to get the density:

$$f_Y(y) = \frac{1}{a} f_X\left(\frac{y-b}{a}\right) = \frac{1}{a\sigma\sqrt{2\pi}} \exp\left(-\frac{\left(\frac{y-b}{a} - \mu\right)^2}{2\sigma^2}\right)$$
$$= \frac{1}{(a\sigma)\sqrt{2\pi}} \exp\left(-\frac{(y-(a\mu+b))^2}{2(a\sigma)^2}\right),$$

which is the Gauss$(a\mu + b, a^2\sigma^2)$ density. □

2. **Standardization.** If $X \sim \text{Gauss}(\mu, \sigma^2)$, then:

$$Z = \frac{X-\mu}{\sigma} \sim \text{Gauss}(0, 1).$$

This is the affine property with $a = 1/\sigma$ and $b = -\mu/\sigma$, and it is the single most used fact about the Gaussian because it reduces every Gaussian probability to a probability about the one fixed distribution that is tabulated. That is because of standardization, you have that:

$$\mathbb{P}(u \le X \le v) = \mathbb{P}\left(\frac{u-\mu}{\sigma} \le Z \le \frac{v-\mu}{\sigma}\right) = \Phi\left(\frac{v-\mu}{\sigma}\right) - \Phi\left(\frac{u-\mu}{\sigma}\right),$$

where Z on the right-hand side is as above, $Z = \frac{X-\mu}{\sigma}$.

Question: Why does this last result follow?

Answer: Consider $\mathbb{P}(u \le X \le v)$ where $X \sim \text{Gauss}(\mu, \sigma^2)$. First, mean-center the variables by subtracting the mean of X in the inequalities and then divide each term by the standard deviation.

$$\mathbb{P}(u \le X \le v) = \mathbb{P}(u - \mu \le X - \mu \le v - \mu)$$
$$= \mathbb{P}\left(\frac{u-\mu}{\sigma} \le \frac{X-\mu}{\sigma} \le \frac{v-\mu}{\sigma}\right)$$
$$= \mathbb{P}\left(\frac{u-\mu}{\sigma} \le Z \le \frac{v-\mu}{\sigma}\right).$$

Observe that the inequalities are preserved because adding a number and multiplying by a positive number doesn't change an inequality's direction.

3. **Symmetry.** If $X \sim \text{Gauss}(0, \sigma^2)$, then $-X \sim \text{Gauss}(0, \sigma^2)$ (the affine property with $a = -1$, $b = 0$). Consequently, for $x > 0$:

$$\mathbb{P}(|X| > x) = \mathbb{P}(X > x \text{ or } X < -x)$$
$$= \mathbb{P}(X > x) + \mathbb{P}(X < -x)$$
$$= \mathbb{P}(X > x) + \mathbb{P}(-X > x)$$
$$= 2\mathbb{P}(X > x).$$

Intuitively, when you negate a Gaussian r.v. you mirror its distribution about the vertical axis. This mirroring doesn't change the shape or the spread of the distribution; it only changes the sign of the values. As a result, the negated Gaussian random variable $-X$ maintains the same mean (here, 0) and the same variance σ^2 as the original r.v. X. For the standard Gaussian the symmetry reads $\Phi(-z) = 1 - \Phi(z)$, which is why tables only print positive z.

4. **Sums of independent Gaussians are Gaussian.** If $X_1 \sim \text{Gauss}(\mu_1, \sigma_1^2)$ and $X_2 \sim \text{Gauss}(\mu_2, \sigma_2^2)$ are independent, then

$$X_1 + X_2 \sim \text{Gauss}(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2).$$

The mean and variance are what they would be for any independent pair; the point is again that the shape stays Gaussian. (The proof is omitted here; it follows from the convolution formula, or in one line from moment generating functions.) Combined with the affine property, it gives the fact behind every exact small-sample result in this book: if $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$ then

$$\bar{X}_n \sim \text{Gauss}\left(\mu, \frac{\sigma^2}{n}\right) \qquad \text{exactly, for every } n,$$

with no appeal to the Central Limit Theorem. For Gaussian data the CLT's conclusion holds at $n = 2$.

From undergraduate statistics: z-scores

Standardization is the z-score of an introductory course: 'how many standard deviations above the mean is this observation?' What the introductory course sometimes leaves unsaid is that the z-score of a Gaussian observation is itself a standard Gaussian random variable, which is the reason a single table serves every μ and σ. Here the same move is applied not to an observation but to an estimator: $\sqrt{n}(\bar{X}_n - \mu)/\sigma$ is the z-score of $\bar{X}_n$, and the CLT says it is approximately standard Gaussian. Every 'z-interval' and 'z-test' you learned is that sentence plus a quantile from Section 6.5.

6.4 Computing probabilities, forwards and backwards

Since the CDF has no formula, every Gaussian probability is computed by standardizing and then looking up Φ, in a table or in software. There are two directions. *Forwards:* given a value, find a probability; standardize to a z-score and read $\Phi(z)$. *Backwards:* given a probability, find the value; read the table in reverse to get $z = \Phi^{-1}(\text{probability})$ and unstandardize, $x = \mu + \sigma z$. Figure 6.2 shows both on the graph of Φ.

Worked example: exam scores

Assume that the score distribution for a final exam is approximately $X \sim \text{Gauss}(85, 4)$, that is, mean 85 and standard deviation $\sigma = 2$.

- $\mathbb{P}(X > 90)$. Standardize: $z = (90 - 85)/2 = 2.5$. Then $\mathbb{P}(X > 90) = \mathbb{P}(Z > 2.5) = 1 - \Phi(2.5) = 1 - 0.9938 = 0.0062$. About six students in a thousand score above 90.
- $\mathbb{P}(80 < X < 90)$. The endpoints standardize to -2.5 and 2.5, so $\mathbb{P}(80 < X < 90) = \Phi(2.5) - \Phi(-2.5) = 2\Phi(2.5) - 1 = 0.9876$, using the symmetry $\Phi(-z) = 1 - \Phi(z)$.
- The score x such that $\mathbb{P}(X < x) = 0.90$ (the 90th percentile). For this you read the table backwards: the z with $\Phi(z) = 0.90$ is $z = 1.2816$, so $x = 85 + 2 \times 1.2816 = 87.6$.

Pitfall: variance or standard deviation?

This book writes $\text{Gauss}(\mu, \sigma^2)$: the second parameter is the *variance*. So $\text{Gauss}(85, 4)$ has standard deviation 2, and $\text{Gauss}(10, 4)$ in Figure 6.1 is twice as wide as $\text{Gauss}(0, 1)$, not four

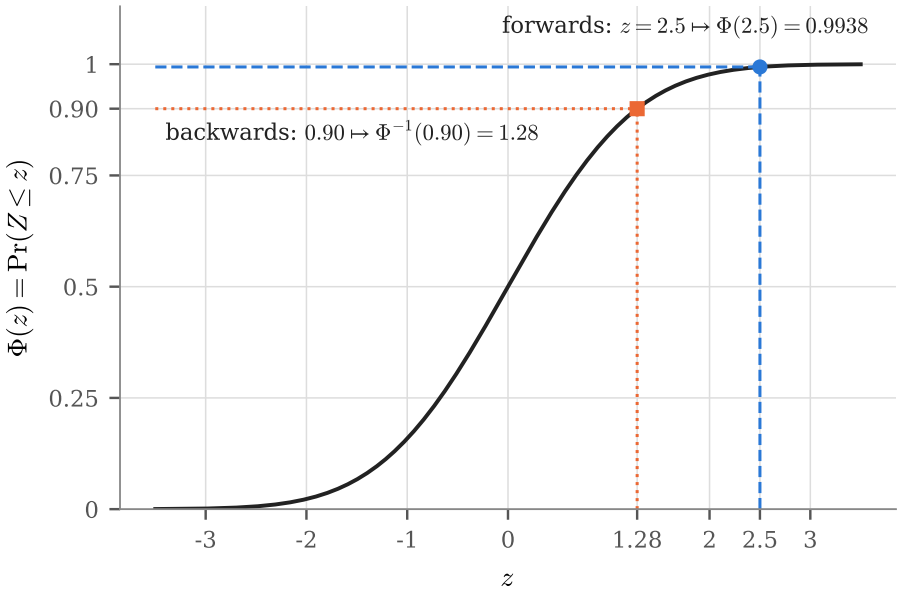


Figure 6.2: The standard Gaussian CDF read in both directions, with the exam example below. Forwards: a z-score goes in on the horizontal axis and a probability comes out, $\Phi(2.5) = 0.9938$. Backwards: a probability goes in on the vertical axis and a z-score comes out, $\Phi^{-1}(0.90) = 1.28$. Tables print the forward direction; the backward direction is found by scanning the table for the probability, or by asking software for the quantile directly.

times. Most software and many textbooks take the standard deviation instead, and the mix-up fails silently: for the exam scores, $\mathbb{P}(X \leq 90)$ is 0.994 with standard deviation 2, but a 4 entered where the 2 was meant answers a different question and gives 0.894, with no warning that anything went wrong. Check which convention is in force every time, and when you write Gauss(μ, σ^2) yourself, write the square.

6.5 Quantiles

To grasp the idea of a quantile it's instructive to consider an example involving a closely related idea, a percentile. Suppose you score a 72 on an exam. Then your percentile is the fraction of students who score *below* a 72. In terms of probability, percentiles are computed using the CDF of exam scores: Letting X denote the score of a random student on the exam, and to make things easy supposing X is a continuous r.v., then your percentile is $\mathbb{P}(X \leq 72) = F_X(72)$.

In general, a student with score x has percentile $F_X(x)$. Going the other way, if you start with a percentile, say 0.95, then $F^{-1}(0.95)$ is the score on the exam that has that percentile. A percentile is also called a quantile, which is why F^{-1} is called the **quantile function**.[1] The function F converts scores to quantiles, and the function F^{-1} converts quantiles to scores: the forward and backward readings of Figure 6.2. And finally, note that a quantile (of a specific order) is a number in the support of an r.v., not a probability. Now for the formal definition, in the convention this book uses throughout.

[1] More precisely, a percentile is a quantile expressed as a percentage.

Definition 6.2 (Quantile of order $1-\alpha$)**.** Let $\alpha \in (0,1)$. The **quantile of order** $1-\alpha$ of an r.v. X is the number q_α such that

$$\mathbb{P}(X \le q_\alpha) = 1-\alpha.$$

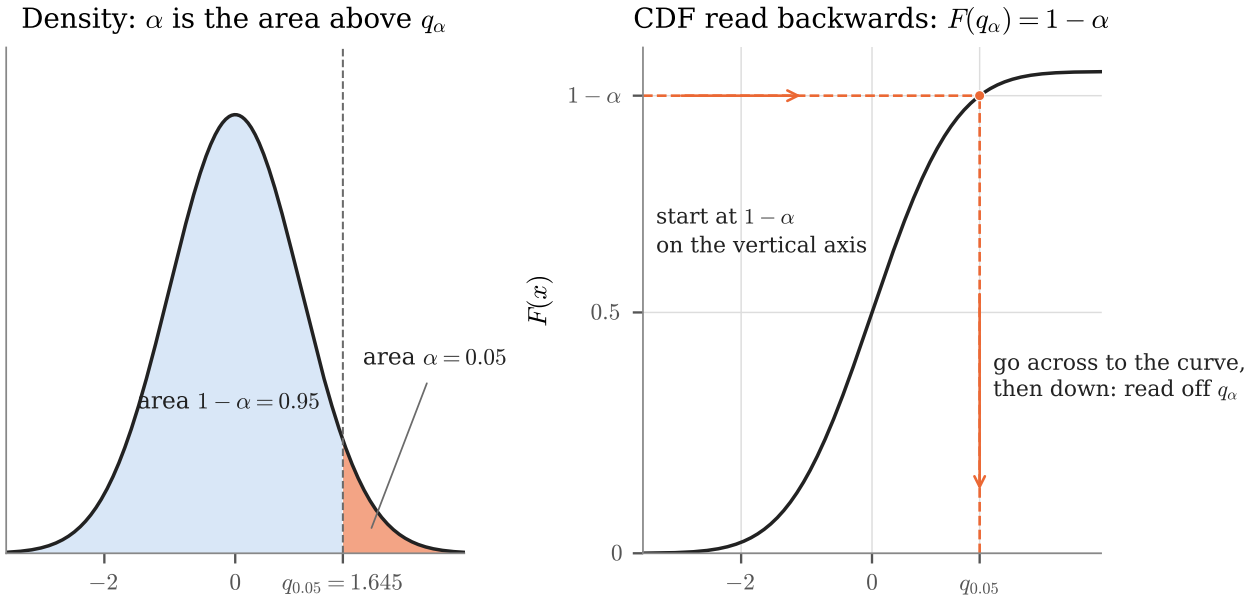


Figure 6.3: The definition as a picture, for $\alpha = 0.05$. Left: q_α is the cutoff on the horizontal axis with area α to its right and $1-\alpha$ to its left; the subscript is the area of the upper tail. Right: to find it, read the CDF backwards. Start at height $1-\alpha$ on the vertical axis, go across to the curve, then drop down to the horizontal axis: the number you land on is q_α. For the standard Gaussian and $\alpha = 0.05$ that number is 1.645.

> **In words.** q_α is the cutoff with probability α above it and $1-\alpha$ below: read the CDF backwards at $1-\alpha$, as in Figure 6.3 (toolkit A.6). The subscript names the tail it cuts off, so read $q_{0.05}$ as 'the cutoff with 5% above it'.
>
> **For the kiss data.** The 1.96 in $\bar{R}_n \pm 1.96 \times 0.043$ is $q_{0.025}$: it leaves 2.5% in each tail, 5% in both, so the interval catches the standardized error 95% of the time.

How to read the notation. The symbol q_α carries two numbers and it is easy to attach the subscript to the wrong one. The subscript α is a *probability*, the area of the upper tail in Figure 6.3; the value q_α is a *number on the horizontal axis*, a cutoff in the same units as X. So $q_{0.05} = 1.645$ says: the probability that a standard Gaussian exceeds 1.645 is 0.05. The phrase 'quantile of order $1-\alpha$' in the definition refers to the area *below* the cutoff, 0.95 for $q_{0.05}$, which is how the inverse CDF is indexed; the subscript refers to the area *above*. Both describe the same point.

Remark: It might seem weird to define quantiles in terms of their complements; that is, the quantile is of order $1-\alpha$ and not of order α. The reason is what the quantile is for. In the formulas that follow, α is usually a small tail probability (a 5% level, a 2.5% tail), and q_α is the cutoff that leaves exactly that much probability in the upper tail. Indexing the cutoff by the tail it cuts off, rather than by the mass below it, is what makes $q_{0.05}$ mean 'the 5% critical value' in every formula that follows. In an intro course the same number arrived as $z^* = 1.645$ or $z_{0.05}$ from the back of the book; $q_{0.05}$ is that z^* with the tail it cuts off written into the name.

Properties of quantiles. In the following let F be the CDF of some r.v. X, assumed continuous and strictly increasing so that F^{-1} exists.

- By the definitions of CDFs and quantiles, it follows that $F(q_\alpha) = 1 - \alpha$.
- How do you find q_α, the quantile of order $1 - \alpha$? Since $F(q_\alpha) = 1 - \alpha$, you just take the inverse of F, meaning $F^{-1}(1 - \alpha) = q_\alpha$.
- Note also that $\mathbb{P}(X > q_\alpha) = \alpha$. This follows from the definition of quantiles by considering the complement of $\mathbb{P}(X \leq q_\alpha)$, and it's the form you'll use most: q_α is the number exceeded with probability α.
- Let $X = Z \sim \text{Gauss}(0, 1)$. Then α equals the probability that $|Z|$ is greater than $q_{\alpha/2}$. That is, $\mathbb{P}(|Z| > q_{\alpha/2}) = \alpha$. This is the symmetry property: $\mathbb{P}(|Z| > q_{\alpha/2}) = 2\mathbb{P}(Z > q_{\alpha/2}) = 2 \cdot \alpha/2$.

The last property is the two-sided version, and the difference between it and the one-sided definition is drawn in Figure 6.4. For a one-sided statement ('Z exceeds the cutoff') the cutoff is q_α; for a two-sided statement ('Z is beyond the cutoff in either direction') the cutoff is $q_{\alpha/2}$ because the probability α has to be split between two tails.

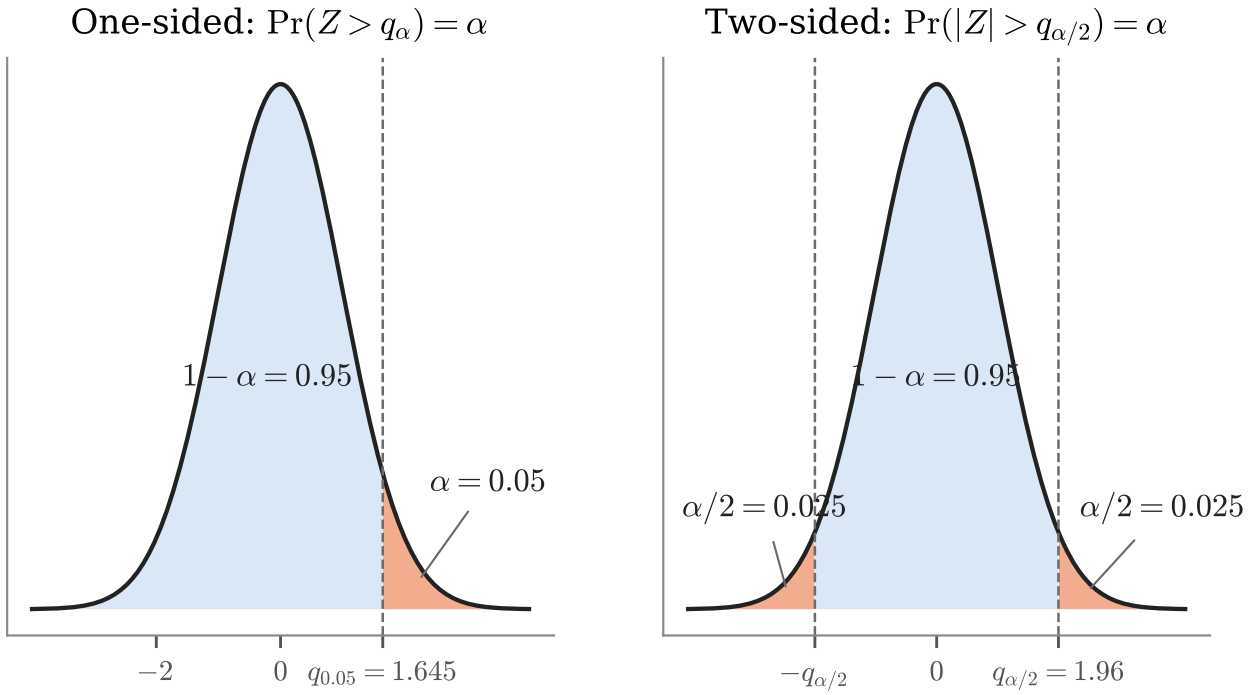


Figure 6.4: Quantiles of the standard Gaussian in this book's convention, for $\alpha = 0.05$. Left: $q_{0.05} = 1.645$ leaves 5% of the mass in the upper tail. Right: $q_{0.025} = 1.96$ leaves 2.5% in each tail, so $|Z| > 1.96$ has probability 5%. Same α, different cutoff; which one you need depends on whether the question is one-sided or two-sided.

Some important quantiles of the standard Gaussian, which recur so often that it is worth knowing them by heart:

α	10%	5%	2.5%	0.5%
q_α	1.28	1.645	1.96	2.576

Read across: $\mathbb{P}(Z > 1.28) = 10\%$, $\mathbb{P}(Z > 1.645) = 5\%$, $\mathbb{P}(Z > 1.96) = 2.5\%$, $\mathbb{P}(Z > 2.576) = 0.5\%$. And read two-sided: $\mathbb{P}(|Z| > 1.645) = 10\%$, $\mathbb{P}(|Z| > 1.96) = 5\%$, $\mathbb{P}(|Z| > 2.576) = 1\%$. The 1.96 of every 95% confidence interval is $q_{0.025}$, not $q_{0.05}$.

Pitfall: which α is that?

Two conventions collide here and both are in common use. This book's q_α has mass α *above* it: $q_{0.05} = 1.645$. Software and many textbooks index a quantile by the mass *below* it: on that convention the 0.95-quantile is 1.645 and the 0.05-quantile is -1.645. So this book's q_α is the other convention's $(1-\alpha)$-quantile, and its $q_{\alpha/2}$ for a 95% interval is the other convention's 0.975-quantile. When a formula in this book says $q_{\alpha/2}$ and you reach for a table or software that counts from below, look up $1 - \alpha/2$. Figure 6.3 is the picture to keep in mind: the subscript is the area in the upper tail.

Example 6.3 (Quantiles of the Exponential)**.** Let $X \sim \text{Expo}(\lambda)$, with CDF $F_X(x) = 1 - e^{-\lambda x}$ for $x \geq 0$. To find the number x below which the mass is p, solve $F_X(x) = p$:

$$\begin{aligned} 1 - e^{-\lambda x} &= p \\ e^{-\lambda x} &= 1 - p \\ -\lambda x &= \ln(1-p) \\ x &= -\frac{\ln(1-p)}{\lambda}. \end{aligned}$$

So, for example, the 50th percentile, or median, of an Exponential r.v. with parameter λ is $x = -\ln(1/2)/\lambda = \ln 2/\lambda \approx 0.69/\lambda$, which is less than the mean $1/\lambda$ because the distribution is skewed to the right. In this book's notation, $q_\alpha = -\ln(\alpha)/\lambda$: the subway inter-arrival time of Chapter 5, with $\hat{\lambda} = 0.16$ per minute, exceeds $q_{0.05} = -\ln(0.05)/0.16 \approx 19$ minutes about one time in twenty. Figure 6.5 shows the three numbers on the density.

6.6 What comes next

This closes the probability part of the book. Part III turns the three questions of Chapter 1 into definitions with guarantees, stated for the general model of Chapter 2: Chapter 7 says what an estimator is and how to judge one; Chapter 8 and Chapter 9 attach the two kinds of guarantee, an interval and a decision. The Gaussian quantiles of this chapter appear in every formula of the last two.

6.7 Summary

- $X \sim \text{Gauss}(\mu, \sigma^2)$ has density $\frac{1}{\sigma\sqrt{2\pi}} \exp(-(x-\mu)^2/(2\sigma^2))$. The second parameter is the variance. The CDF has no closed form; the tails are thin enough that 3σ contains 99.7% of the mass.
- Affine transformations, in particular standardization $Z = (X - \mu)/\sigma \sim \text{Gauss}(0, 1)$, preserve Gaussianity; so do sums of independent Gaussians. The sample mean of Gaussian data is exactly $\text{Gauss}(\mu, \sigma^2/n)$.

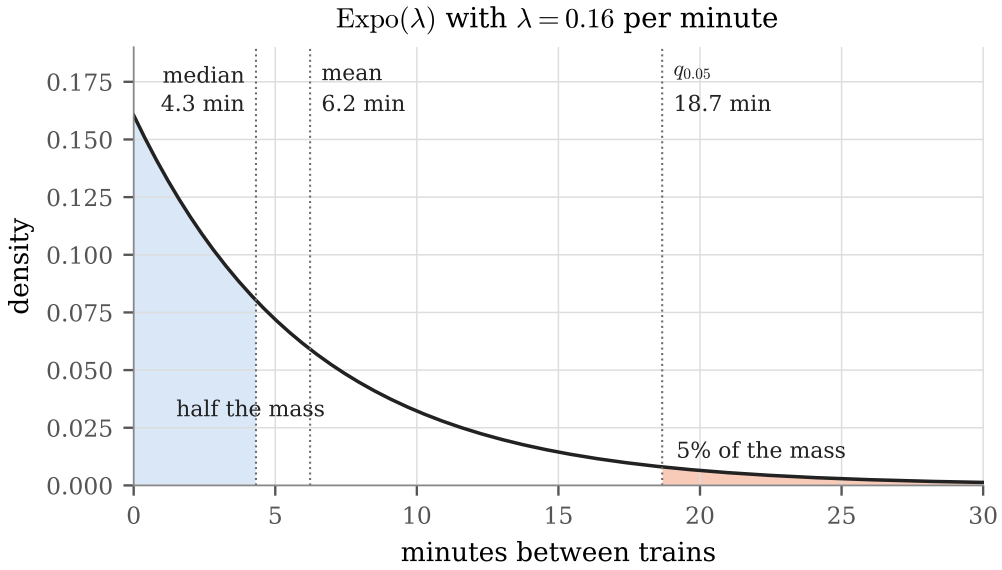


Figure 6.5: Quantiles of a skewed distribution: the Expo(λ) density of the subway inter-arrival times with $\lambda = 0.16$ per minute. The median, $\ln 2/\lambda = 4.3$ minutes, splits the mass in half and sits left of the mean, $1/\lambda = 6.2$ minutes because the long right tail pulls the mean up. The quantile $q_{0.05} = 18.7$ minutes leaves 5% of the mass above it. For the Gaussian the median and the mean coincide; for a skewed distribution they don't, and quantiles are the honest summary.

- Every Gaussian probability is a table lookup after standardization: forwards for probabilities, backwards for quantiles. $\Phi(-z) = 1 - \Phi(z)$.
- The quantile q_α of order $1 - \alpha$ satisfies $\mathbb{P}(X > q_\alpha) = \alpha$. For the standard Gaussian, $\mathbb{P}(|Z| > q_{\alpha/2}) = \alpha$. Memorize $q_{0.05} = 1.645$ and $q_{0.025} = 1.96$, and keep straight which one a question needs.

6.8 Exercises

Exercise 6.1. Let $Z \sim \text{Gauss}(0, 1)$. Compute $\mathbb{P}(Z \leq 1)$, $\mathbb{P}(Z \geq -1)$ and $\mathbb{P}(|Z| > 1)$. *(Solution on page 223.)*

Exercise 6.2. Let $X \sim \text{Gauss}(\mu, \sigma^2)$ and $Y = 2X$. Write down the PDF of Y, first from the affine property and then by the change-of-variables formula $f_Y(y) = f_X(g^{-1}(y))/|g'(g^{-1}(y))|$ for $Y = g(X)$. *(Solution on page 223.)*

Exercise 6.3. Let $X \sim \text{Gauss}(2, 2)$ (mean 2, variance 2) and let $x > 0$. Write $\mathbb{P}(X \geq -x)$ as $\Phi(g(x))$ for a function g that is positive for $x > 0$. *(Solution on page 223.)*

Exercise 6.4. Verify the three numbers of the empirical rule from a table, and find $\mathbb{P}(|X - \mu| > 2\sigma)$ for any Gaussian X. Why doesn't the answer depend on μ or σ? *(Solution on page 224.)*

Exercise 6.5. For $\alpha = 0.10$, what is the cutoff c such that $\mathbb{P}(Z > c) = 0.10$, and what is the cutoff c' such that $\mathbb{P}(|Z| > c') = 0.10$? Which of the two belongs in a 90% confidence interval of the form $\bar{X}_n \pm c\,\sigma/\sqrt{n}$? *(Solution on page 224.)*

Part III

Foundations of Inference

The three questions of Chapter 1 become definitions with guarantees, stated for the general model of Chapter 2 and computed for the kiss experiment. An estimator (Chapter 7) is a guess inside the model, judged by bias, variance and risk; a confidence interval (Chapter 8) and a test (Chapter 9) are the two ways of attaching a guarantee to the guess.

Chapter 7

Estimators: Bias, Variance and Risk

In this chapter

- What a statistic is, what an estimator is, and the sense in which a statistic throws information away (and when, by the sufficiency principle, that's fine).
- The two large-sample properties an estimator should have: consistency (it gets to the truth) and asymptotic normality (it gets there at rate $1/\sqrt{n}$ with a Gaussian error), and why the asymptotic variance is not the limit of the variance.
- The two finite-sample measures: bias and variance, with Jensen's inequality as the reason transformed estimators pick up bias.
- Quadratic risk = variance + bias^2, and the trade-off it makes visible: a biased estimator can beat an unbiased one, and no estimator wins for every value of the parameter.

The running example is again the kiss experiment, $R_1, \dots, R_n \overset{iid}{\sim} \text{Bern}(p)$ with $n = 124$, now with several competing estimators of p. Chapter 2 fixed the object inside which inference happens, and Part II supplied the three limit facts about averages; the first task inside the model is to guess θ from the data, and this chapter says what a guess is and how to judge one.

7.1 Statistics and estimators

Definition 7.1 (Statistic, estimator)**.** A **statistic** is any measurable function of the sample: $T = T(X_1, \dots, X_n)$. For example, $\bar{X}_n$, $\max_i X_i$, $X_1 + \log(1 + |X_n|)$, the sample variance, and so on. As a statistic is a function of your data, it's an r.v.

An **estimator** of θ is a statistic $\hat{\theta}_n = \hat{\theta}_n(X_1, \dots, X_n)$ whose expression doesn't depend on θ.

> **In words.** A statistic is anything you can compute from the data. An estimator is a statistic you intend as a guess for θ, and its formula may use the data and known constants only, never θ.
>
> **For the kiss data.** $\bar{R}_n$ is a statistic and an estimator of p. $\bar{R}_n - p$ is a random variable but not a statistic, since you cannot compute it. $\max_i R_i$ is a statistic (it is 1 unless every couple turned left) and a useless estimator of p.

The word 'measurable' is a technicality from probability theory (toolkit A.5); the working rule of thumb is that if you can compute it exactly once you're given the data, it's measurable. Things

that are only implicitly defined, such as a sup or inf over an infinite family (toolkit A.2), can fail the condition, but every statistic this book writes down is a formula in the observations and passes it.

The clause 'doesn't depend on θ' is the whole content of the second definition. $\bar{R}_n$ is an estimator of p in the kiss model; $\bar{R}_n - p$ is neither a statistic nor an estimator because you can't compute it from the data alone: its formula involves the thing you're trying to estimate. (It is a perfectly good random variable, and it has a name, the **estimation error**, which is the quantity the rest of this chapter measures.) The constant 0.5 is an estimator of p (a bad one, as Section 7.8 will show), and so is X_1: the definition says what an estimator is, not what a good one is. The rest of the chapter is about telling them apart.

Recall from Chapter 1 the distinction between an estimator and an estimate: the estimator $\hat{\theta}_n$ is a random variable, with a distribution, an expectation and a variance; the *estimate* $\hat{\theta}_n(x_1, \ldots, x_n)$ is the number it produces on your one sample. Every property defined in this chapter is a property of the random variable.

7.2 A statistic reduces the data

Let $\mathbf{X} = (X_1, \ldots, X_n)$ be the vector of observations and $\mathbf{x} = (x_1, \ldots, x_n)$ a realized sample. Any statistic $T(\mathbf{X})$ defines a form of data reduction. Anyone who uses only the observed value of the statistic, $T(\mathbf{x})$, rather than the entire observed sample $\mathbf{x}$, will treat as equal two samples, $\mathbf{x}$ and $\mathbf{y}$, that satisfy $T(\mathbf{x}) = T(\mathbf{y})$ even though the actual sample values may differ.

Data reduction in terms of a particular statistic can be thought of as a partition of the sample space E^n. For illustration, suppose a sample of size n is drawn and $\bar{x} = 8$. There are many points in E^n which have the same mean of 8, and you can think of them as belonging to the set $\{(x_1, \ldots, x_n) : \bar{x} = 8\}$. In fact, all points on the hyperplane

$$x_1 + \cdots + x_n = 8n$$

yield the mean $\bar{x} = 8$, and thus the hyperplane is the set. But there are of course many values that $\bar{X}_n$ can take, and thus there are many such sets. In this sense, the sample mean, or any statistic, partitions the sample space into a collection of sets (equivalence classes). Figure 7.1 draws this for $n = 2$, where the sample space is the plane and the hyperplanes are lines, next to the quite different partition induced by the maximum.

Let $\mathcal{T} = \{t = T(\mathbf{x}) \text{ for some } \mathbf{x} \in E^n\}$ be the image of E^n under T; that is, $\mathcal{T}$ is the set of all possible values of the statistic. Then T partitions the sample space into the sets $A_t = \{\mathbf{x} : T(\mathbf{x}) = t\}$ for $t \in \mathcal{T}$. Samples $\mathbf{x}_1$ and $\mathbf{x}_2$ belong to the same set if and only if $T(\mathbf{x}_1) = T(\mathbf{x}_2)$. The statistic summarizes the data in that, rather than reporting the entire sample $\mathbf{x}$, it reports only that $T(\mathbf{x}) = t$ or, equivalently, $\mathbf{x} \in A_t$.

7.2.1 The sufficiency principle

A **sufficient statistic** for a parameter θ is a statistic that, in a certain sense, captures all the information about θ contained in the sample. Any additional information in the sample, besides the value of the sufficient statistic, does not contain any more information about θ. This leads to the data reduction technique called the sufficiency principle.

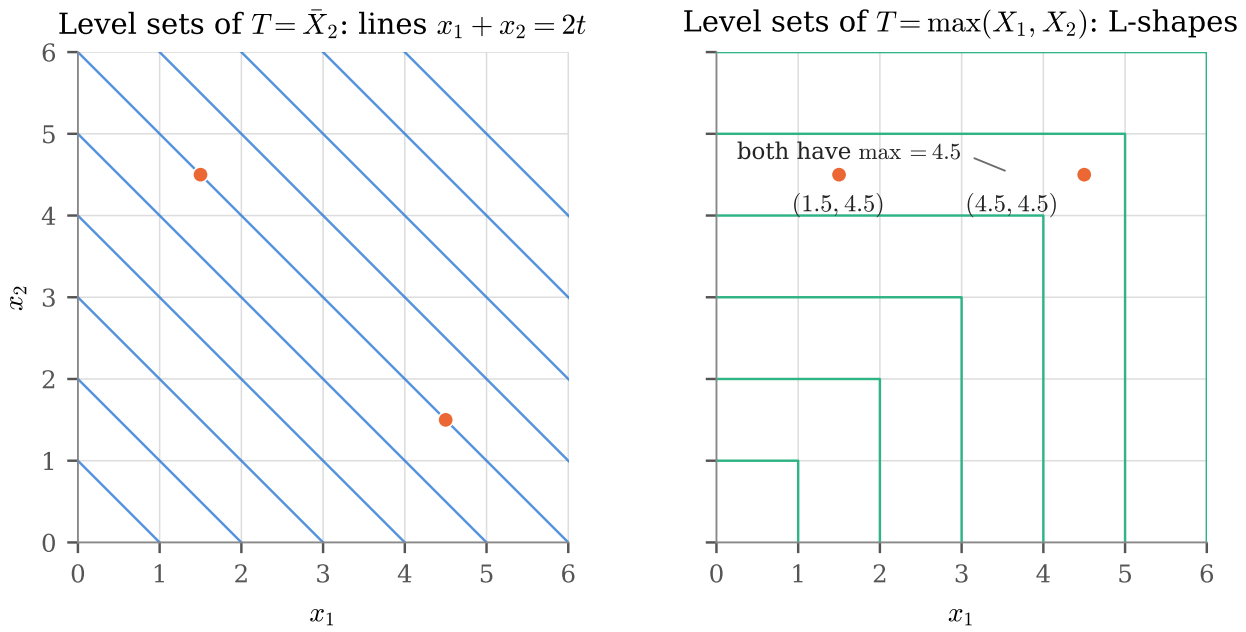


Figure 7.1: A statistic partitions the sample space. With $n = 2$ observations the sample is a point in the plane. Left: the level sets of the sample mean are the lines $x_1 + x_2 = 2t$; the samples $(1.5, 4.5)$ and $(4.5, 1.5)$ lie on the same line and are the same sample to anyone who only sees $\bar{x} = 3$. Right: the level sets of the maximum are L-shapes; $(1.5, 4.5)$ and $(4.5, 4.5)$ are the same sample to anyone who only sees max $= 4.5$. Each statistic keeps one number and discards the rest; which rest is discarded depends on the statistic.

Sufficiency principle: If $T(\mathbf{X})$ is a sufficient statistic for θ, then any inference about θ should depend on the sample $\mathbf{X}$ only through the value $T(\mathbf{X})$. That is, if $\mathbf{x}$ and $\mathbf{y}$ are two sample points such that $T(\mathbf{x}) = T(\mathbf{y})$, then the inference about θ should be the same whether $\mathbf{X} = \mathbf{x}$ or $\mathbf{X} = \mathbf{y}$ is observed.

Definition 7.2 (Sufficient statistic)**.** A statistic $T(\mathbf{X})$ is a **sufficient statistic for** θ if the conditional distribution of the sample $\mathbf{X}$ given the value of $T(\mathbf{X})$ does not depend on θ.

> **In words.** Once you know the value of T, the rest of the sample looks the same whatever θ is, so the rest carries no further information about θ.
>
> **For the kiss data.** Given that 80 of 124 couples turned right, every ordering of the 80 ones and 44 zeros is equally likely, whatever p is; the count is sufficient and the order tells you nothing more about p. The example below does the computation.

Example 7.3. In the kiss model, the number of couples turning right $S = \sum_i R_i$ is sufficient for p. Given $S = s$, every arrangement of s ones and $n - s$ zeros is equally likely:

$$\mathbb{P}_p(\mathbf{X} = \mathbf{x} \mid S = s) = \frac{p^s(1-p)^{n-s}}{\binom{n}{s} p^s(1-p)^{n-s}} = \frac{1}{\binom{n}{s}},$$

which doesn't involve p. Once you know 80 of 124 couples turned right, the order in which they did so tells you nothing further about p. This is why every method in this book for this model is a function of $\bar{R}_n$ alone: by the sufficiency principle, nothing else in the sample tells you more about p,

if the model is right. (The order of the R_i is exactly what you would look at to check whether it is, as with the wedding reception of exercise 4(b) in Chapter 1, whose couples would show up as a run of 1s.)

Sufficiency has a large theory of its own, which this book does not take up; it appears once more, in Chapter 22, where the exponential families make the sufficient statistic visible. It is here because it says precisely what the data reduction in Figure 7.1 is allowed to throw away: everything except T, whenever the conditional distribution of the sample given T is free of θ.

7.3 Consistency

The best estimator would be the one that equals the parameter of interest. Given the probabilistic nature of statistics, ascertaining rules to derive such an estimator is hopeless. But there are desirable properties you want estimators to have.

For example, as you collect more and more data you should get a better and better idea of which distribution generated the data. In the limit, as your sample size increases to infinity you should know exactly which distribution it is. Specifically, you want an estimator that converges to the true parameter θ as your sample size goes to infinity. An estimator satisfying this property is called *consistent*, and it comes in two forms, depending on whether the convergence is in probability or almost surely.

Definition 7.4 (Consistency)**.** An estimator $\hat{\theta}_n$ of θ is **weakly consistent** if

$$\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta \quad \text{(with respect to } \mathbb{P}_\theta\text{)},$$

that is, $\mathbb{P}_\theta\left(|\hat{\theta}_n - \theta| \geq \epsilon\right) \to 0$ for every $\epsilon > 0$, and **strongly consistent** if $\hat{\theta}_n \xrightarrow[n\to\infty]{\text{a.s.}} \theta$ (almost sure convergence is defined in Appendix B; the unqualified word 'consistent' is used for either, in this book and in common usage). Here $\mathbb{P}_\theta$ is the distribution that generated the data, and the statement is required to hold for every $\theta \in \Theta$.

> **In words.** With enough data the estimator lands as close to the truth as you like, with probability as high as you like, and this must be so whichever value the truth has.
>
> **For the kiss data.** $\bar{R}_n \xrightarrow[n\to\infty]{\mathbb{P}} p$ is the Law of Large Numbers, so $\bar{R}_n$ is consistent. The constant estimator 0.5 is not: it never moves, so it only gets close to p if p happens to be 0.5, and consistency has to hold for every p.

The Law of Large Numbers is the consistency of $\bar{X}_n$ for μ; its weak form (Theorem 3.5) gives weak consistency, its strong form (Theorem B.3 in Appendix B) strong consistency.

Recall. The Law of Large Numbers says that an average of independent draws settles on the number it is an average of: for any width you name, the probability that $\bar{X}_n$ sits outside a band of that width around μ goes to zero. It says nothing about how fast, and nothing about the shape of what is left; those are the Central Limit Theorem's business.

The continuous mapping theorem (Chapter 5) then makes any continuous function of a consistent estimator consistent for the corresponding function of the parameter.

Recall. The continuous mapping theorem says that applying a continuous function to a converging sequence gives a converging sequence, with the function applied to the limit: if $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$ and g is continuous at θ, then $g(\hat{\theta}_n) \xrightarrow[n\to\infty]{\mathbb{P}} g(\theta)$. It is the rule that lets you take square roots and reciprocals of estimators without starting over.

Consistency only takes a stand on where the distribution of $\hat{\theta}_n$ *concentrates*, not on the whole distribution. This is evident from its definition; consistency is defined in terms of convergence in probability or almost surely, and both are statements about the limit being a point. But consistency usually isn't enough to do inference because you want to know not just where the estimator concentrates but also the precision with which it concentrates, and how fast.

Example 7.5 (How fast is consistent?)**.** Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$. What is the smallest c such that $n^c(\bar{X}_n - p)$ does *not* converge to 0 in probability? By the CLT, $\sqrt{n}(\bar{X}_n - p)/\sigma \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$ with $\sigma = \sqrt{p(1-p)}$, so for $c < 1/2$,

$$n^c(\bar{X}_n - p) = \frac{\sigma}{n^{1/2-c}} \cdot \frac{\sqrt{n}}{\sigma}(\bar{X}_n - p) \xrightarrow[n\to\infty]{\text{(d)}} 0 \cdot \text{Gauss}(0,1) = 0$$

by Slutsky, which is convergence in probability to 0 since the limit is a constant. At $c = 1/2$ the limit is $\text{Gauss}(0, \sigma^2)$, not 0. So $c = 1/2$: the estimation error $\bar{X}_n - p$ is of exact order $1/\sqrt{n}$, and such an estimator is called $\sqrt{n}$*-consistent.* That is more than consistency says, and it's the second property.

7.4 Asymptotic normality

This is where the property of asymptotic normality enters, which unlike consistency is *defined in terms of convergence in distribution.* Asymptotic normality is a property of an estimator that, if present, gives it CLT-like behavior: a Gaussian limiting distribution. Unlike the CLT, it's not a theorem; it's a property that a given estimator may or may not have, and specifically a property of estimators, hence concerned with inference. For reasons addressed below, it involves some scaling in order to circumvent a problem with the variance of the limiting distribution.

Definition 7.6 (Asymptotic normality)**.** An estimator $\hat{\theta}_n$ of θ is **asymptotically normal** if

$$\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, \sigma^2).$$

The variance of this limiting distribution, σ^2, is called the **asymptotic variance** of $\hat{\theta}_n$.

> **In words.** Magnify the estimation error by $\sqrt{n}$ and it settles into a bell curve centered at 0. Unmagnified, the error is approximately $\text{Gauss}(0, \sigma^2/n)$: shrinking like $1/\sqrt{n}$, with a known shape.
>
> **For the kiss data.** $\sqrt{n}(\bar{R}_n - p) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, p(1-p))$ by the CLT, so $\bar{R}_n$ is asymptotically normal with asymptotic variance $p(1-p) = 0.229$ at $p = 0.645$, and its error at $n = 124$ is approximately $\text{Gauss}(0, 0.043^2)$.

NB$_1$: You might imagine the 'natural' form of the asymptotic variance would simply be the limit of $\mathbb{V}\text{ar}(\hat{\theta}_n)$ as n goes to infinity. **This is not what it is.** Why not? Think about the estimator $\bar{X}_n$. Its

variance is $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \sigma^2/n$, and hence, as $n \to \infty$, it converges to 0. It'd be silly to let 'asymptotic variance' refer to something that's always 0. Hence, the asymptotic variance σ^2 is *not* the limit of the variance of $\hat{\theta}_n$. Figure 7.2 shows the two quantities side by side for the kiss estimator.

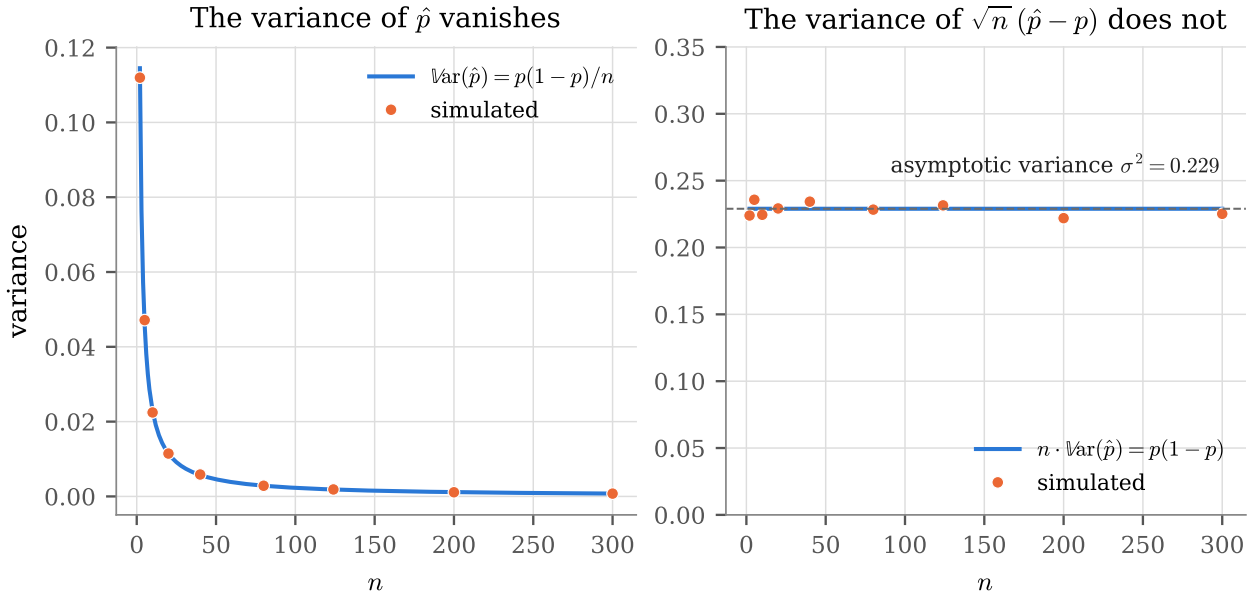


Figure 7.2: For $\hat{p} = \bar{R}_n$ with $p = 0.645$. Left: $\mathbb{V}\mathrm{ar}(\hat{p}) = p(1-p)/n$ goes to zero; its limit carries no information about precision. Right: $n \cdot \mathbb{V}\mathrm{ar}(\hat{p}) = p(1-p) = 0.229$ is constant, and it is the asymptotic variance, the variance of the limiting Gaussian of $\sqrt{n}(\hat{p}-p)$. Dots are variances estimated from 4,000 simulated samples at each n.

NB$_2$: This brings us to the rescaling in the definition of asymptotic normality. The rescaling by $\sqrt{n}$ is crucial because it holds the *spread* of the error fixed as the sample size increases, which is what allows a non-degenerate Gaussian limit to exist at all. To understand why, again think about $\bar{X}_n$: without the rescaling, $\bar{X}_n - \mu$ converges to the degenerate distribution at 0, which is consistency restated and says nothing about precision. With it, $\sqrt{n}(\bar{X}_n - \mu)$ has variance σ^2 for every n and a Gaussian limit. For the sample mean the asymptotic variance therefore coincides with $n\,\mathbb{V}\mathrm{ar}(\bar{X}_n)$; in general it is defined as the variance of the limiting Gaussian, which is the only definition that always makes sense, since convergence in distribution does not by itself say anything about variances. This is the two-scales picture of Section 3.2, and Example 7.5 says why $\sqrt{n}$ and not some other power.

NB$_3$: Note also that you subtract the unknown parameter θ from $\hat{\theta}_n$ in the definition. This focuses attention on the deviation of the estimator from its true value, $\hat{\theta}_n - \theta$, which is the estimation error, the quantity you actually care about.

Intuition: what asymptotic normality buys

For large n the error of an asymptotically normal estimator is approximately Gauss$(0, \sigma^2/n)$. That one sentence gives you a standard error, $\sigma/\sqrt{n}$; an error bar, $\hat{\theta}_n \pm 1.96\sigma/\sqrt{n}$ with probability about 95%; and a test of any hypothesis about θ. Consistency says the error vanishes; asymptotic normality says how big it is at the n you have and what shape it has. Chapter 8 is entirely about using the sentence. When the asymptotic variance depends on θ (as $p(1-p)$ does), you'll replace it by a consistent estimate and lean on Slutsky, as in Chapter 5.

7.5 Bias

Consistency and asymptotic normality are properties of the limit. The next two are properties of $\hat{\theta}_n$ at the sample size you actually have.

Definition 7.7 (Bias)**.** The **bias** of an estimator $\hat{\theta}_n$ of θ is

$$\text{Bias}(\hat{\theta}_n) = \mathop{\mathbb{E}}_{\theta}[\hat{\theta}_n] - \theta.$$

If $\text{Bias}(\hat{\theta}_n) = 0$ for every $\theta \in \Theta$, the estimator is called **unbiased**.

> **In words.** Average the estimator over all the samples you could have drawn; the bias is how far that average sits from the truth.
>
> **For the kiss data.** $\mathbb{E}[\bar{R}_n] = p$, so $\bar{R}_n$ is unbiased; the constant 0.5 has bias $0.5 - p$, which is -0.145 if $p = 0.645$.

An unbiased estimator is right on average: if you could repeat the experiment many times, the estimates would be scattered around θ with no systematic lean. It says nothing about how widely they're scattered.

Worked example: four estimators of p: bias

Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$ and consider four candidate estimators of p, the first three of which look respectable and the fourth of which does not.

1. $\hat{p} = \bar{X}_n$: $\mathbb{E}[\bar{X}_n] = p$, so the bias is 0.
2. $\hat{p} = X_1$: $\mathbb{E}[X_1] = p$, bias 0. Unbiased, and it ignores $n - 1$ observations.
3. $\hat{p} = (X_1 + X_2)/2$: $\mathbb{E} = p$, bias 0.
4. $\hat{p} = \sqrt{\mathbb{1}(X_1 = 1, X_2 = 1)}$: the indicator is 1 with probability p^2, and the square root of 0 or 1 is itself, so $\mathbb{E}[\hat{p}] = p^2$ and the bias is $p^2 - p = -p(1-p) < 0$. This estimator is systematically too low.

The fourth one is a trap for anyone who moves a square root through an expectation: $\sqrt{\mathbb{E}[\cdot]}$ is not $\mathbb{E}[\sqrt{\cdot}]$.

Worked example: a biased sample mean

Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([a, a+1])$ with a unknown. Then $\mathbb{E}[\bar{X}_n] = \mathbb{E}[X_1] = \int_a^{a+1} x \, \mathrm{d}x = a + \frac{1}{2}$, so $\bar{X}_n$ has bias $\frac{1}{2}$ as an estimator of a, for every n: more data doesn't help because the lean is built in. The fix is obvious once the bias is known: $\bar{X}_n - \frac{1}{2}$ is unbiased for a. A known bias that doesn't depend on the unknown parameter is as good as no bias at all; the dangerous biases are the ones you don't know, and the awkward ones are those that depend on the parameter, like the subway bias below.

7.5.1 Jensen's inequality

The fourth estimator above is a case of a general fact: applying a convex (or concave) function to an unbiased estimator produces a biased one, and the direction of the bias is predictable.

A function $g : \mathbb{R} \to \mathbb{R}$ is **convex** if for all $x_1 < x_2$ and all $0 \le t \le 1$,

$$g(tx_1 + (1-t)x_2) \le tg(x_1) + (1-t)g(x_2) :$$

the graph of g lies below every chord. Now read the right-hand side probabilistically. Let X take the value x_1 with probability t and x_2 with probability $1-t$. Then $tx_1 + (1-t)x_2 = \mathbb{E}[X]$ and $tg(x_1) + (1-t)g(x_2) = \mathbb{E}[g(X)]$, so convexity says $g(\mathbb{E}[X]) \le \mathbb{E}[g(X)]$ for every two-point X. It holds in general.

Theorem 7.8 (Jensen's inequality)**.** *For any r.v. X with finite mean and any convex function g,*

$$g(\mathbb{E}[X]) \le \mathbb{E}[g(X)].$$

For a concave g the inequality reverses.

The two-point argument above is the whole idea; the general case replaces the chord by the tangent line at $\mathbb{E}[X]$, which for a convex g lies below the curve everywhere (toolkit A.8); evaluating that line at X and taking expectations kills its linear term and leaves exactly $g(\mathbb{E}[X]) \le \mathbb{E}[g(X)]$.

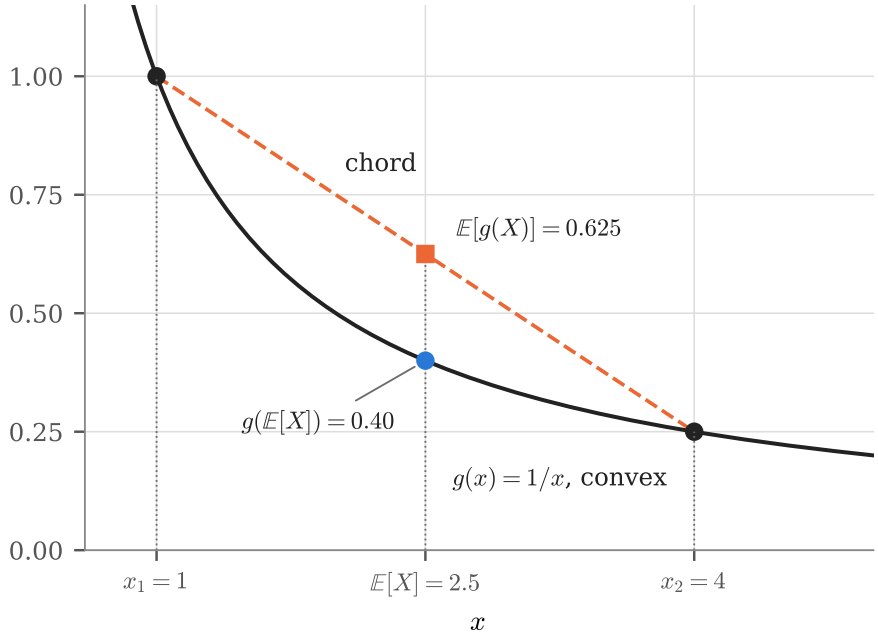


Figure 7.3: Jensen's inequality for $g(x) = 1/x$ and an X equal to 1 or 4 with probability $\frac{1}{2}$ each. The point $(\mathbb{E}[X], \mathbb{E}[g(X)])$ is the midpoint of the chord, and the chord lies above the curve, so $\mathbb{E}[g(X)] = 0.625$ exceeds $g(\mathbb{E}[X]) = 0.4$. For any convex g, averaging first and transforming second gives less than transforming first and averaging second.

Example 7.9 (The subway rate again)**.** In Chapter 5 the rate of the T was estimated by $\hat{\lambda} = 1/\bar{T}_n$, with $\mathbb{E}[\bar{T}_n] = 1/\lambda$. Since $g(x) = 1/x$ is convex on $(0, \infty)$, Jensen gives $\mathbb{E}[1/\bar{T}_n] \ge 1/\,\mathbb{E}[\bar{T}_n] = \lambda$: the estimator is biased upward, for every n. The bias does vanish as n grows, but not because the estimator is consistent and asymptotically normal (neither property says anything about the bias); it vanishes because an exact computation gives $\mathbb{E}[1/\bar{T}_n] = \frac{n}{n-1}\lambda$, a bias of $\lambda/(n-1)$, which is of order $1/n$ and so is swamped by the $1/\sqrt{n}$ standard error. The exercises take this up. Biased and consistent is a common combination, and nothing is wrong with it.

From undergraduate statistics: why the sample variance divides by $n-1$

Every introductory course divides by $n-1$ and most students never learn why. It's a bias calculation. For $X_1, \dots, X_n$ i.i.d. with mean μ and variance σ^2, the natural estimator $\frac{1}{n}\sum(X_i - \bar{X}_n)^2$ has expectation

$$\mathbb{E}\left[\frac{1}{n}\sum_{i=1}^{n}(X_i - \bar{X}_n)^2\right] = \mathbb{E}\left[\frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2\right] - \mathbb{E}\left[(\bar{X}_n - \mu)^2\right] = \sigma^2 - \frac{\sigma^2}{n} = \frac{n-1}{n}\sigma^2,$$

using $\sum(X_i - \bar{X}_n)^2 = \sum(X_i - \mu)^2 - n(\bar{X}_n - \mu)^2$. Its bias is $-\sigma^2/n$: the deviations from $\bar{X}_n$ are systematically smaller than the deviations from μ because $\bar{X}_n$ was chosen to be in the middle of this sample. Multiplying by $n/(n-1)$ removes the bias exactly, which is the $S^2 = \frac{1}{n-1}\sum(X_i - \bar{X}_n)^2$ you were taught. Both versions are consistent; the correction matters at $n = 5$ and not at $n = 500$.

7.6 Variance

An estimator is a random variable, so it has a variance, $\mathbb{V}\mathrm{ar}(\hat{\theta}_n) = \mathbb{E}[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2]$, which measures how widely it scatters around its own mean (not around θ). Small variance means the estimate you got is close to the estimate you would have got from a different sample; it says nothing about whether either is close to θ.

Worked example: four estimators of p: variance

For the same four estimators, with $\mathbb{V}\mathrm{ar}(X_i) = p(1-p)$:

1. $\bar{X}_n$: $\mathbb{V}\mathrm{ar} = p(1-p)/n$, since the variance of a sum of independent terms is the sum of the variances.
2. X_1: $\mathbb{V}\mathrm{ar} = p(1-p)$. Unbiased, but as variable as a single coin flip, however large n is.
3. $(X_1 + X_2)/2$: $\mathbb{V}\mathrm{ar} = p(1-p)/2$.
4. $\sqrt{\mathbb{1}(X_1 = 1, X_2 = 1)}$: a Bernoulli with parameter p^2, so $\mathbb{V}\mathrm{ar} = p^2(1-p^2)$.

The first three are unbiased, so among them the comparison is purely on variance, and $\bar{X}_n$ wins by a factor of n over X_1. This is the sense in which using all the data matters: an unbiased estimator that ignores observations is paying in variance for no return in bias. Figure 7.4 draws all four sampling distributions on one axis.

For the uniform example, $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \mathbb{V}\mathrm{ar}(X_1)/n = \frac{1}{12n}$, and the bias-corrected $\bar{X}_n - \frac{1}{2}$ has the same variance, since subtracting a constant changes nothing.

7.7 Quadratic risk

It may not be possible to find an estimator that is both unbiased and has the smallest variance. Is there a compromise that can be made between a biased estimator and a high-variance one? That is, can you find an estimator that simply has a small bias and a small variance? That's where the quadratic risk comes in: a single number that charges for both.

Definition 7.10 (Quadratic risk)**.** The **quadratic risk** (or **risk**) of an estimator $\hat{\theta}_n \in \mathbb{R}$ is

$$R(\hat{\theta}_n) = \mathop{\mathbb{E}}_{\theta}\left[|\hat{\theta}_n - \theta|^2\right],$$

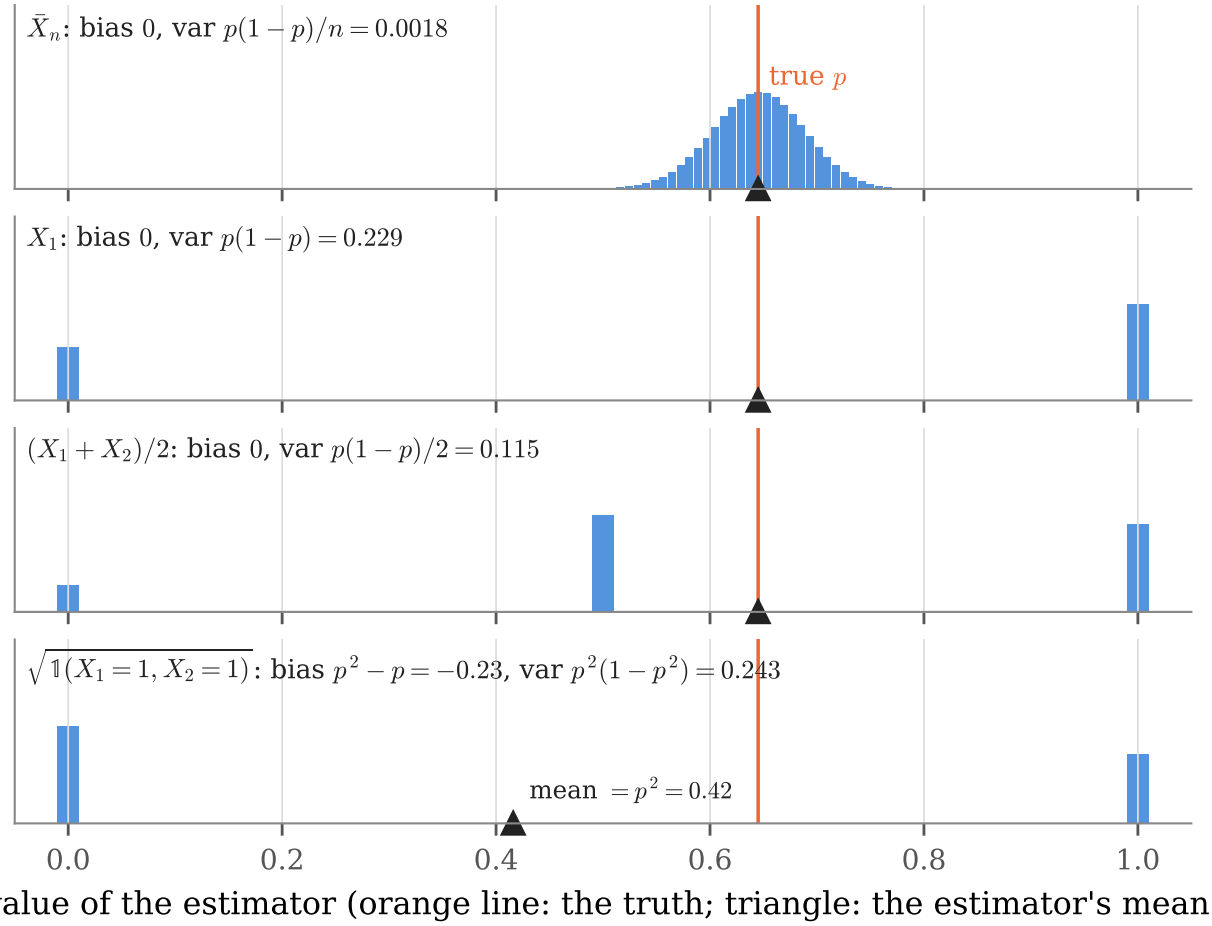


Figure 7.4: The four estimators of p at $n = 124$, $p = 0.645$: the exact sampling distribution of each. $\bar{X}_n$ is a narrow spike at p. X_1 is two bars, at 0 and 1: unbiased, since its mean is p, and never anywhere near p. $(X_1 + X_2)/2$ is three bars. The square-root indicator is two bars with its mean at $p^2 = 0.42$, off the truth. Bias is where the triangle sits relative to the orange line; variance is how spread out the bars are; only the first estimator does well on both.

the expected squared distance from the estimator to the truth. It is more commonly called the **mean squared error**, and the two names mean the same thing.

> **In words.** Square the miss, average over all possible samples. One number that punishes both a systematic lean and a wide scatter.
>
> **For the kiss data.** $R(\bar{R}_n) = p(1-p)/n = 0.00185$ at $p = 0.645$, $n = 124$; a typical miss is its square root, 0.043.

Theorem 7.11 (Bias-variance decomposition)**.**

$$R(\hat{\theta}_n) = \mathbb{V}ar(\hat{\theta}_n) + Bias(\hat{\theta}_n)^2.$$

Proof strategy (core). Add and subtract $\mathbb{E}[\hat{\theta}_n]$ inside the square, expand, and watch the cross term vanish because $\mathbb{E}[\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]] = 0$.

Proof. First note that $|\hat{\theta}_n - \theta|^2 = (\hat{\theta}_n - \theta)^2$. Add and subtract $\mathbb{E}[\hat{\theta}_n]$ inside the square and expand:

$$\begin{aligned} R(\hat{\theta}_n) &= \mathbb{E}\left[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n] + \mathbb{E}[\hat{\theta}_n] - \theta)^2\right] \\ &= \mathbb{E}\left[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])^2\right] + \mathbb{E}\left[(\mathbb{E}[\hat{\theta}_n] - \theta)^2\right] + 2\,\mathbb{E}\left[(\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n])(\mathbb{E}[\hat{\theta}_n] - \theta)\right] \\ &= \mathbb{V}\mathrm{ar}(\hat{\theta}_n) + (\mathbb{E}[\hat{\theta}_n] - \theta)^2 + 2(\mathbb{E}[\hat{\theta}_n] - \theta)\ \mathbb{E}\left[\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]\right] \\ &= \mathbb{V}\mathrm{ar}(\hat{\theta}_n) + \mathrm{Bias}(\hat{\theta}_n)^2 + 2(\mathbb{E}[\hat{\theta}_n] - \theta) \cdot 0. \end{aligned}$$

In the second term, $\mathbb{E}[\hat{\theta}_n] - \theta$ is a number, so its expectation is itself; in the third, the same number comes out of the expectation, and what remains, $\mathbb{E}[\hat{\theta}_n - \mathbb{E}[\hat{\theta}_n]]$, is zero. □

The decomposition is the reason a biased estimator can be the better one: if accepting a little bias buys a large reduction in variance, the risk goes down. It also makes a point that is easy to miss when you only ever compute bias and variance separately: the risk is a function of θ, through both terms, and an estimator that is best for one value of θ need not be best for another. Figure 7.5 shows both facts in the kiss model.

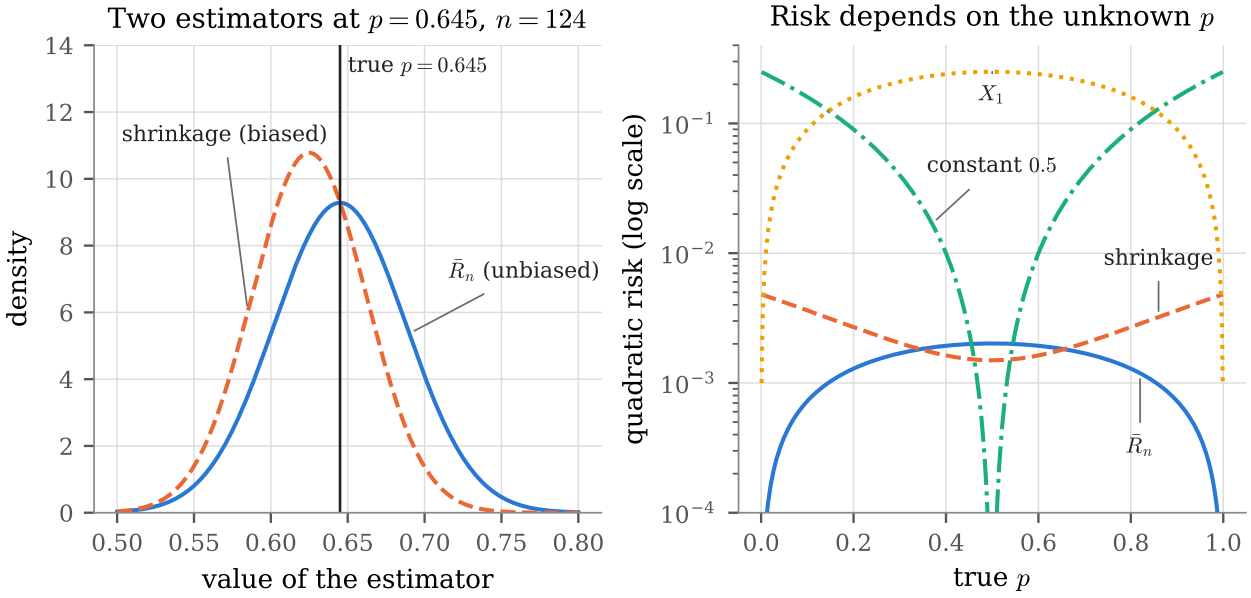


Figure 7.5: The trade-off in the kiss model, $n = 124$. Left: at $p = 0.645$, the sampling distributions (drawn as their Gaussian approximations) of the unbiased $\bar{R}_n$ (standard deviation 0.043) and of the shrinkage estimator $(n\bar{R}_n + 10)/(n + 20)$, which pulls the estimate toward $\frac{1}{2}$ as if 20 extra coin flips had come up half heads. The shrinkage estimator is biased by -0.020 but has standard deviation 0.037, and its risk, 0.00177, is smaller than $\bar{R}_n$'s 0.00185. Right: risk against the true p for four estimators, on a log scale. Shrinkage beats $\bar{R}_n$ for p between 0.34 and 0.66 and loses outside; the constant 0.5 is unbeatable at $p = 0.5$ and terrible elsewhere; X_1 is unbiased and hopeless everywhere. No curve lies below all the others, so 'best estimator' is not a well-posed question without saying for which p.

Intuition: bias and variance as a target

Think of the truth θ as the bull's-eye and each estimate as a shot. Variance is how tight the group is; bias is how far the center of the group is from the bull's-eye. The risk is the average squared distance of a shot from the bull's-eye, and the decomposition says it's the squared distance of the group's center plus the group's spread: a tight group off to one side can be closer on average than a loose group centered on the target. The shrinkage estimator in Figure 7.5 is the tight group off to one side. What the target picture doesn't show is that moving the bull's-eye (changing p) changes which group is closer, and that's the right-hand panel. Figure 7.6 draws the four combinations.

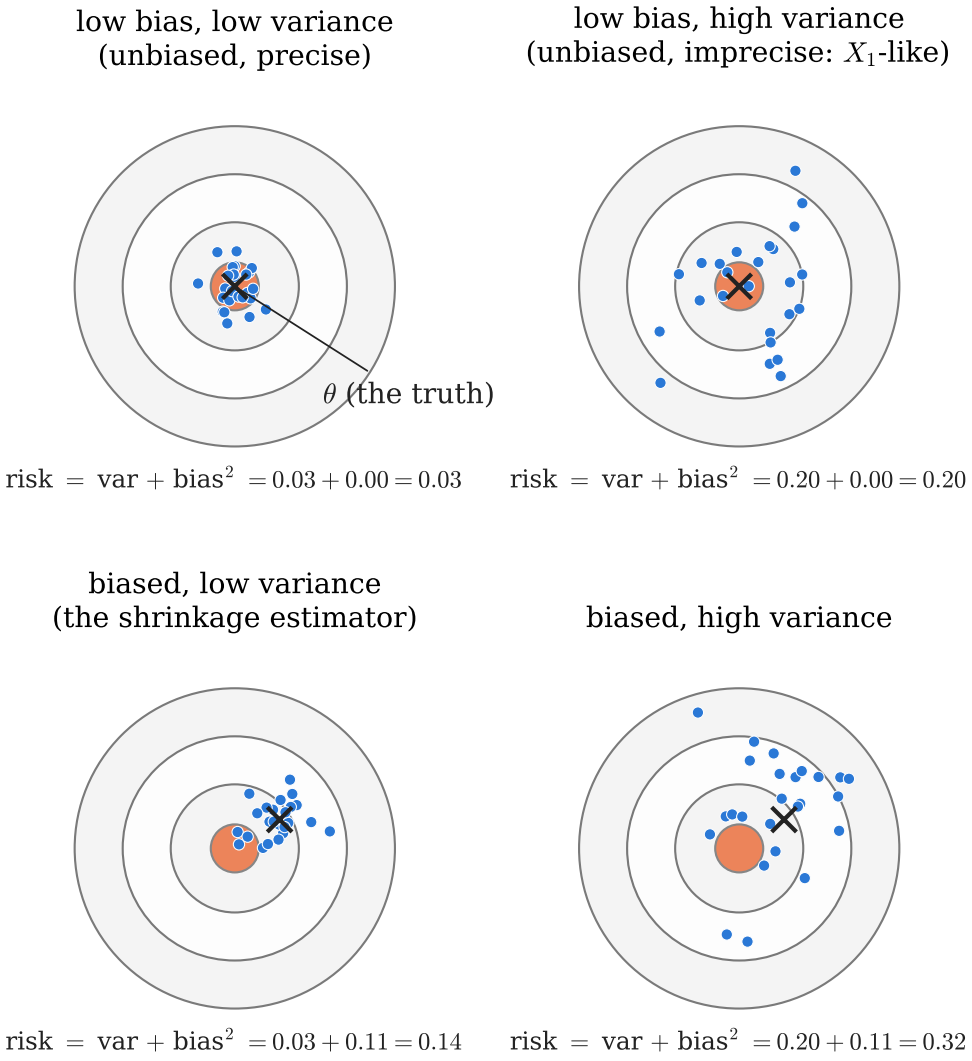


Figure 7.6: Bias and variance as a target. The bull's-eye is the truth θ; each dot is the estimate from one sample; the cross is the center of the group, $\mathbb{E}[\hat{\theta}_n]$. Bias is the distance from the cross to the bull's-eye, variance is the spread of the dots around the cross, and the risk is the average squared distance from a dot to the bull's-eye, which the decomposition splits into the two. Lower left is the shrinkage estimator of Figure 7.5: a tight group slightly off center, closer on average than the loose group centered on the target at upper right.

From undergraduate statistics: the $n-1$ story, finished

The unbiased S^2 with $n-1$ is not the minimum-risk estimator of σ^2. For Gaussian data, dividing by $n+1$ instead gives a biased estimator with strictly smaller quadratic risk than both the n and $n-1$ versions. Nobody uses it because unbiasedness is a convention that makes results comparable across studies and because the difference is small, but it shows that 'unbiased' and 'best' are different properties, and that an unbiasedness calculation settles only the first.

7.8 Strengths and weaknesses of estimators

Everything in this chapter can be applied at once to three estimators of p for $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$, n large: $\bar{X}_n$, the constant 0.5, and X_1. Four criteria: bias, consistency, whether the estimator is efficiently computable (in this section, whether there is an explicit formula for it; the maximum likelihood estimators of Chapter 23 have no closed form and are found by an iterative algorithm instead), and whether the quadratic risk tends to 0 as $n \to \infty$.

Estimator	Bias	Consistent?	Explicit formula?	Risk $\to 0$?
$\bar{X}_n$	0	yes (LLN)	yes	yes: $p(1-p)/n$
0.5	$0.5 - p$, nonzero unless $p = \frac{1}{2}$	no, unless $p = \frac{1}{2}$	yes	no: $(0.5-p)^2$ for every n
X_1	0	no: it's 0 or 1, never near p	yes	no: $p(1-p)$ for every n

The constant 0.5 uses none of the data, so nothing can be learned from it; its only virtue is that it's exactly right when $p = \frac{1}{2}$. X_1 is unbiased, which sounds good, and uses one observation out of n, which is the actual problem: its risk doesn't shrink, so its accuracy is limited no matter how many samples are collected. $\bar{X}_n$ has every property on the list. Unbiasedness is desirable and may be the first thing to check, but the comparison with X_1 shows why it isn't enough: there are estimators that are unbiased and have a lot of variability, and risk is the number that catches them.

7.9 What comes next

Asymptotic normality hands you a sentence, 'the error of $\hat{\theta}_n$ is approximately Gauss$(0, \sigma^2/n)$', and the next two chapters are the two ways of using it. Chapter 8 turns it into a range of plausible values of θ with a guaranteed capture rate; Chapter 9 turns it into a yes/no decision about θ with a guaranteed error rate. Both need the quantiles of Chapter 6 and the plug-in step of Chapter 5, since σ^2 usually depends on θ.

7.10 Summary

- A statistic is a function of the data; an estimator of θ is a statistic whose formula doesn't involve θ. Both are random variables. A statistic partitions the sample space and keeps one number per cell; a sufficient statistic keeps everything the sample says about θ.

- Consistent: $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$ under $\mathbb{P}_\theta$, for every θ. The sample mean is $\sqrt{n}$-consistent: its error is of exact order $1/\sqrt{n}$ (Example 7.5). Asymptotically normal: $\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, \sigma^2)$; σ^2 is the asymptotic variance, the variance of the limiting Gaussian, which for the sample mean equals $n\,\mathbb{V}\text{ar}(\bar{X}_n)$ and is never the limit of $\mathbb{V}\text{ar}(\hat{\theta}_n)$, which for any estimator whose variance shrinks is 0 and carries no information.
- Bias: $\mathbb{E}[\hat{\theta}_n] - \theta$. Unbiased means right on average. Convex or concave functions of unbiased estimators are biased, in the direction Jensen's inequality says: $g(\mathbb{E}[X]) \le \mathbb{E}[g(X)]$ for convex g.
- Variance: spread of $\hat{\theta}_n$ around its own mean. Using more of the data lowers it.
- Quadratic risk: $\mathbb{E}[(\hat{\theta}_n - \theta)^2] = \mathbb{V}\text{ar} + \text{Bias}^2$. It depends on θ; a biased estimator can have lower risk than an unbiased one, and no estimator has the lowest risk for every θ.
- Of $\bar{X}_n$, the constant $\frac{1}{2}$ and X_1, only $\bar{X}_n$ is unbiased, consistent and has risk tending to 0. Unbiasedness alone, as X_1 shows, is not enough.

7.11 Exercises

Exercise 7.1. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([a, a+1])$. Find $\mathbb{E}[\bar{X}_n]$, the bias of $\bar{X}_n$ for a, $\mathbb{V}\text{ar}(\bar{X}_n)$, and the quadratic risk of $\bar{X}_n - \frac{1}{2}$. *(Solution on page 224.)*

Exercise 7.2. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$ and $\hat{p} = \sqrt{\mathbb{1}(X_1 = 1, X_2 = 1)}$. Find its bias, variance and quadratic risk, and say whether it is consistent. *(Solution on page 224.)*

Exercise 7.3. In the kiss model with $n = 124$, compute the bias, variance and risk of the shrinkage estimator $\tilde{p}_n = (n\bar{R}_n + 10)/(n + 20)$ at $p = 0.645$ and at $p = 0.9$, and compare each with $\bar{R}_n$. *(Solution on page 224.)*

Exercise 7.4. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Expo}(\lambda)$ and $\hat{\lambda} = 1/\bar{X}_n$. Use Jensen's inequality to determine the sign of the bias, and use the fact (which you may take as given) that $\bar{X}_n$ has the Gamma distribution with $\mathbb{E}[1/\bar{X}_n] = \frac{n}{n-1}\lambda$ for $n \ge 2$ to find the bias exactly. Is there an unbiased estimator of λ of the form $c/\bar{X}_n$? *(Solution on page 224.)*

Exercise 7.5. Show that an estimator whose bias and variance both tend to 0 is consistent. Is the converse true? *(Solution on page 224.)*

Chapter 8

Confidence Intervals

In this chapter

- What a confidence interval is: a *random* interval, computable from the data alone, that traps the fixed parameter with a guaranteed probability. The two words that carry the definition are 'random' and 'computable'.
- The recipe that produces every interval in this book: find a pivot (a quantity whose distribution you know), trap it between two quantiles, solve for the parameter. An exact pivot gives an exact interval; the CLT gives an asymptotic one.
- Three ways around a pivot whose scale involves the unknown parameter: conservative, solve, plug-in. One picture explains all three and shows how they nest.
- What drives the width (the level, through a quantile; the sample size, through $1/\sqrt{n}$), and what the 95% does and does not mean.

The running examples are the kiss experiment ($n = 124$, $\bar{R}_n = 0.645$), for which Chapter 1 promised the interval $[0.56, 0.73]$, and the subway inter-arrival times of Chapter 5 ($n = 64$, $\bar{T}_n = 6.23$).

8.1 What a confidence interval is

Chapter 7 produced a number, $\hat{p} = 0.645$, and a sampling distribution around it. A confidence interval converts the two into a statement with a guarantee: a range of values of p, and a probability that the range is right.

Definition 8.1 (Confidence interval)**.** Let $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ be a statistical model based on observations $X_1, \dots, X_n$, with $\Theta \subseteq \mathbb{R}$, and let $\alpha \in (0, 1)$.

A **confidence interval of level** $1-\alpha$ for θ is a random interval $\mathcal{I} = [L(X_1, \dots, X_n),\, U(X_1, \dots, X_n)]$ whose endpoints are statistics (their formulas don't involve θ), such that

$$\mathbb{P}_\theta\left(\mathcal{I} \ni \theta\right) \geq 1 - \alpha \quad \text{for all } \theta \in \Theta.$$

A confidence interval of **asymptotic level** $1 - \alpha$ is a random interval $\mathcal{I} = \mathcal{I}_n$ of the same kind such that

$$\lim_{n\to\infty} \mathbb{P}_\theta\left(\mathcal{I}_n \ni \theta\right) \geq 1 - \alpha \quad \text{for all } \theta \in \Theta.$$

In words. A rule that turns a sample into an interval, computable from the sample alone, such that the interval it produces contains the true parameter at least $100(1-\alpha)\%$ of the time, whatever the true parameter is. 'Asymptotic' means the guarantee holds in the limit of large samples.

For the kiss data. The rule '$\bar{R}_n \pm 1.96\sqrt{\bar{R}_n(1-\bar{R}_n)/n}$' is a confidence interval of asymptotic level 95%; on this sample it produced $[0.56, 0.73]$. The rule '$\bar{R}_n \pm 1.96\sqrt{p(1-p)/n}$' is not one, since p is in its formula.

The notation $\mathcal{I} \ni \theta$ says '$\mathcal{I}$ contains θ'. It means the same as $\theta \in \mathcal{I}$, and this book writes it backwards on purpose: the random object, the thing the probability is about, is the interval on the left, not the parameter on the right. Both appearances of θ in $\mathbb{P}_\theta(\mathcal{I} \ni \theta)$ are the same fixed number doing two jobs: as the subscript it fixes which distribution the sample is drawn from, and inside the parentheses it is the target that the interval either contains or misses. Nothing in the expression is random except $\mathcal{I}$.

Three things in the definition deserve a second look.

1. *The interval is random; the parameter is not.* L and U are functions of the sample, so they're random variables, and $\mathcal{I}$ moves from sample to sample. θ is a fixed number that doesn't move. The probability is the probability that the moving interval lands on the fixed point.
2. *The endpoints must be computable.* An interval whose formula contains θ is not a confidence interval, however nice its coverage probability: you can't compute it. This clause is what the three 'solutions' of Section 8.4 are solutions to.
3. *The guarantee holds for every θ.* You don't know θ, so a guarantee that holds for some values of it and not others is no guarantee. This clause is what makes one of the intervals in Exercise 8.2 fail.

The inequality $\geq 1-\alpha$ has a consequence worth stating once: a confidence interval of level 95% is also a confidence interval of level 90%, of level 50%, of any level below 95%. It is not, in general, one of level 98%. 'Level' is a floor on the coverage, not its exact value.

The word 'level' is used for tests too (Chapter 9), where it denotes α, the probability of rejecting a true hypothesis; for an interval the level is $1-\alpha$, the probability of a correct capture. Same α, complementary conventions, and the context always says which is meant.

Intuition: a procedure, not a number

The confidence interval is the *rule* that turns a sample into an interval, $(x_1, \ldots, x_n) \mapsto [L, U]$, and the level is a property of the rule: the fraction of samples on which the rule's output contains θ. Once you apply the rule to your one sample you get two numbers, $[0.56, 0.73]$. Those numbers are not random and the parameter is not random, so the pair has no probability attached: $[0.56, 0.73]$ either contains p or it doesn't, and you will never know which. What you know is that the rule that produced it is right 95% of the time. Figure 8.1 is this sentence as a picture.

8.2 The recipe, on an exact case

Every interval in this book is built the same way, and the cleanest place to see the construction is the one case where nothing is asymptotic.

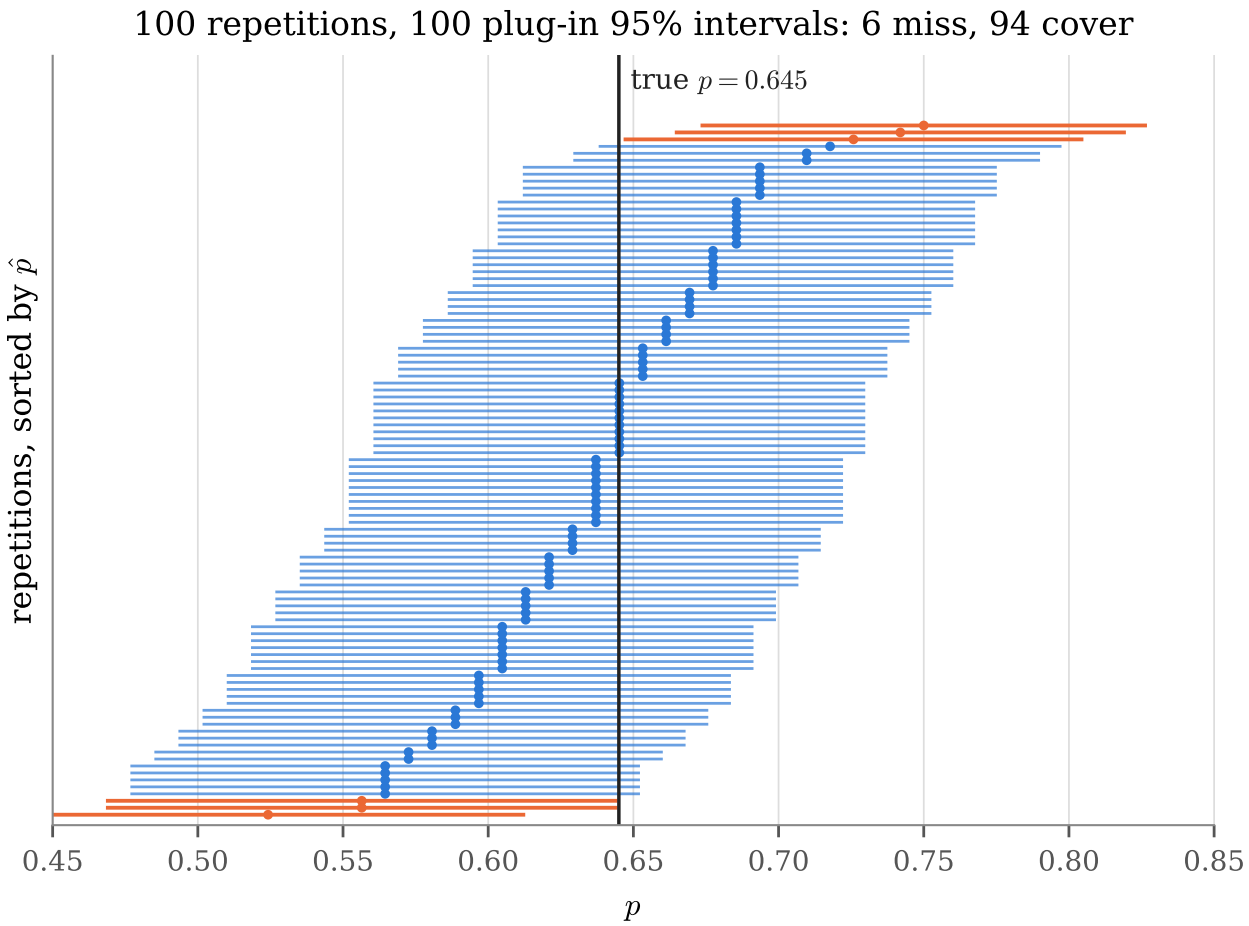


Figure 8.1: The frequentist meaning of '95%'. The kiss experiment was simulated 100 times with $n = 124$ and $p = 0.645$, and the plug-in interval of Section 8.4 was computed each time. The parameter is the fixed vertical line; the intervals move, and each is centered at its own $\hat{p}$ (dot). Sorted by $\hat{p}$, the six that miss are the samples whose proportion came out unusually low or high. The level 95% is the long-run hit rate of the rule; it is not a statement about any one of the hundred intervals, each of which either contains 0.645 or doesn't.

Worked example: Gaussian data, known variance

Let $X_1, \dots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$ with σ^2 known and μ unknown. By Chapter 6, $\bar{X}_n \sim \text{Gauss}(\mu, \sigma^2/n)$ exactly, for every n, so

$$Z = \sqrt{n}\,\frac{\bar{X}_n - \mu}{\sigma} \sim \text{Gauss}(0,1) \quad \text{for every } n \text{ and every } \mu.$$

A quantity like Z, a function of the data and the parameter whose distribution doesn't depend on the parameter, is called a **pivot**. The recipe has three steps.

1. *Trap the pivot.* With $q_{\alpha/2}$ the quantile of Section 6.5, $\mathbb{P}(-q_{\alpha/2} \le Z \le q_{\alpha/2}) = 1 - \alpha$.
2. *Rewrite the event.* $-q_{\alpha/2} \le \sqrt{n}(\bar{X}_n - \mu)/\sigma \le q_{\alpha/2}$ is the same event as $|\bar{X}_n - \mu| \le q_{\alpha/2}\,\sigma/\sqrt{n}$.

3. *Solve for the parameter.* That event is the same as

$$\mu \in \left[\bar{X}_n - q_{\alpha/2}\frac{\sigma}{\sqrt{n}},\ \bar{X}_n + q_{\alpha/2}\frac{\sigma}{\sqrt{n}}\right] =: \mathcal{I}.$$

The event in step 3 is the event in step 1, written differently, so $\mathbb{P}_\mu(\mathcal{I} \ni \mu) = 1-\alpha$ for every μ and every n. The endpoints involve $\bar{X}_n$, σ (known), n and a quantile, none of which is μ, so $\mathcal{I}$ is a confidence interval of level $1-\alpha$, with no limit anywhere. For $\alpha = 0.05$ it is $\bar{X}_n \pm 1.96\,\sigma/\sqrt{n}$.

From undergraduate statistics: estimate ± margin of error

The interval you learned is 'point estimate ± (critical value) × (standard error)'. The recipe says where each piece comes from: the point estimate is the center because the pivot is centered at the parameter; the critical value is the quantile that traps the pivot; the standard error is the pivot's scale, $\sigma/\sqrt{n}$. When σ is unknown, an introductory course replaces it by s and switches to a t critical value: that is a different pivot, $\sqrt{n}(\bar{X}_n - \mu)/s$, with an exact Student distribution for Gaussian data, and Chapter 14 derives it. For large n and any parent distribution, the replacement is justified by Slutsky instead (Section 5.2), which is the route this chapter takes.

One-sided intervals. Nothing forces step 1 to trap the pivot symmetrically. $\mathbb{P}(Z \le q_\alpha) = 1-\alpha$ as well, and solving $\sqrt{n}(\bar{X}_n - \mu)/\sigma \le q_\alpha$ for μ gives

$$\mu \in \left[\bar{X}_n - q_\alpha\frac{\sigma}{\sqrt{n}},\ \infty\right),$$

a **lower confidence bound**: 'μ is at least this much', with level $1-\alpha$. Trapping $Z \ge -q_\alpha$ gives the upper bound $(-\infty, \bar{X}_n + q_\alpha\sigma/\sqrt{n}]$. Note the quantile: q_α (1.645 at 95%), not $q_{\alpha/2}$ (1.96) because all of the 5% of misses is spent on one side. One-sided intervals are the confidence-interval face of one-sided tests, and Exercise 8.6 makes the correspondence exact.

8.3 The kiss experiment: an interval that isn't one

Now $R_1, \dots, R_n \overset{iid}{\sim} \text{Bern}(p)$, $\hat{p} = \bar{R}_n$, and there is no exact pivot in sight: the distribution of $\bar{R}_n$ is Binomial, scaled, and it depends on p in a way that no simple transformation removes. What there is instead is the CLT (Theorem 4.8).

Recall. The Central Limit Theorem says that the error of an average, magnified by $\sqrt{n}$ and divided by the standard deviation of one observation, has a distribution that settles down to Gauss(0, 1), whatever the observations were drawn from. You do not need to know that distribution to use the theorem; you need only its mean and variance.

$$\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{p(1-p)}} \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1).$$

The left side is an *asymptotic* pivot: its distribution depends on p for every finite n, but its limit doesn't. Convergence in distribution means convergence of the CDF at every continuity point of the limit (Section 4.1), and the Gaussian CDF is continuous everywhere, so

$$\lim_{n\to\infty} \mathbb{P}_p\left(-q_{\alpha/2} \le \sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{p(1-p)}} \le q_{\alpha/2}\right) = \mathbb{P}(-q_{\alpha/2} \le Z \le q_{\alpha/2}) = 1-\alpha.$$

Steps 2 and 3 of the recipe then give

$$\lim_{n\to\infty} \mathbb{P}_p\left(\left[\bar{R}_n - q_{\alpha/2}\sqrt{\frac{p(1-p)}{n}},\ \bar{R}_n + q_{\alpha/2}\sqrt{\frac{p(1-p)}{n}}\right] \ni p\right) = 1-\alpha. \tag{8.1}$$

This looks like the Gaussian case with σ replaced by $\sqrt{p(1-p)}$, and it has the right asymptotic coverage, *but it is not a confidence interval*: the endpoints contain p. You can't compute $\bar{R}_n \pm 1.96\sqrt{p(1-p)/124}$ without knowing the thing you're estimating. The pivot's scale depends on the parameter, and the recipe's step 3 is stuck.

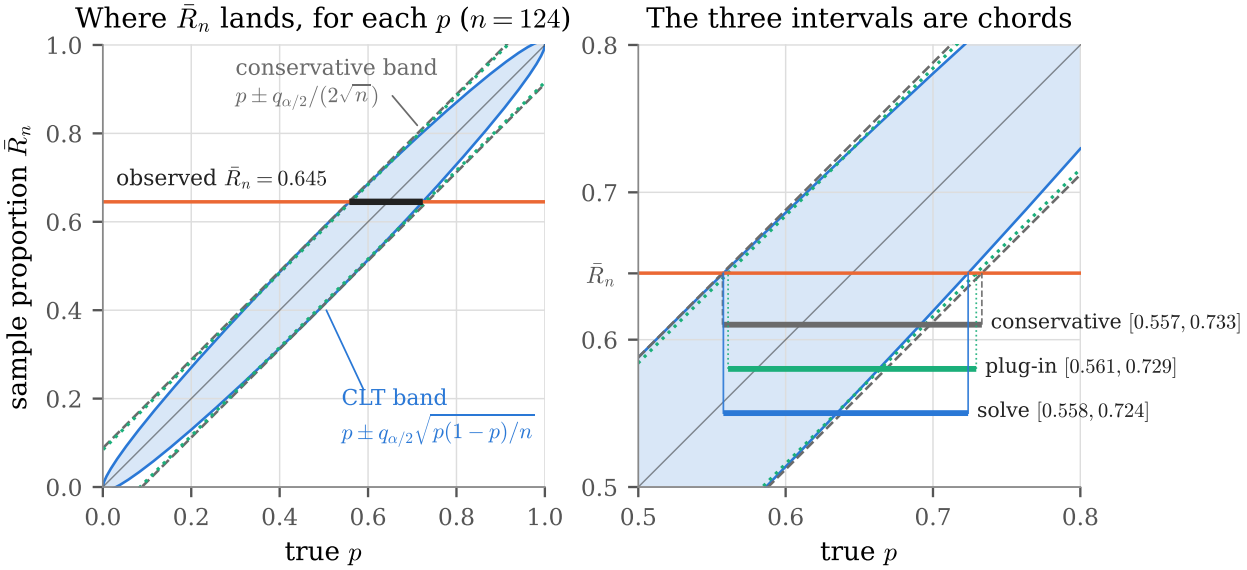


Figure 8.2: Where the three intervals come from. Horizontal axis: the true p. Vertical axis: the sample proportion. For each p, Equation (8.1) says $\bar{R}_n$ lands in the band $p \pm q_{\alpha/2}\sqrt{p(1-p)/n}$ about 95% of the time; the band is a lens, widest at $p = \frac{1}{2}$ and pinched at the ends. We observed $\bar{R}_n = 0.645$: the horizontal line. The set of p whose band contains the observed value is the chord the line cuts out of the lens, and that chord is the solve interval. The conservative interval replaces the lens by the widest band it fits in (gray dashes), and gets a longer chord. The plug-in interval replaces the lens by a band of constant width $q_{\alpha/2}\sqrt{\hat{p}(1-\hat{p})/n}$ (dotted), the width the lens has at $p = \hat{p}$, and gets a chord of nearly the same length, centered at $\hat{p}$. Right: the three chords, to scale, with $n = 124$.

8.4 Three ways to get unstuck

Figure 8.2 is the picture to hold on to. For each candidate value of p on the horizontal axis, the CLT band is the set of sample proportions that value would plausibly produce. Reading the picture downwards, from a p to its band, is probability. Reading it sideways, from the observed $\bar{R}_n$ to the set of p whose bands reach it, is statistics, and that set is the confidence interval. The three solutions differ in which band they read.

Solution 1: conservative

Whatever p is, $p(1-p) \le \frac{1}{4}$, with equality at $p = \frac{1}{2}$. So the event in Equation (8.1) is contained in a larger event with a computable boundary:

$$|\bar{R}_n - p| \le q_{\alpha/2}\sqrt{\frac{p(1-p)}{n}} \quad \Longrightarrow \quad |\bar{R}_n - p| \le \frac{q_{\alpha/2}}{2\sqrt{n}}.$$

Define

$$\mathcal{I}_{\text{conserv}} = \left[\bar{R}_n - \frac{q_{\alpha/2}}{2\sqrt{n}},\ \bar{R}_n + \frac{q_{\alpha/2}}{2\sqrt{n}}\right].$$

Its coverage is at least that of the CLT event, so $\lim_{n\to\infty} \mathbb{P}_p(\mathcal{I}_{\text{conserv}} \ni p) \ge 1-\alpha$ for every p: an asymptotic confidence interval, with the inequality of the definition used in earnest. The coverage is exactly $1-\alpha$ in the limit only at $p = \frac{1}{2}$; elsewhere the interval is wider than it needs to be, which is what 'conservative' means. In the figure it reads the widest band instead of the lens.

With $n = 124$: $q_{0.025}/(2\sqrt{124}) = 1.96/22.27 = 0.088$, so $\mathcal{I}_{\text{conserv}} = [0.557, 0.733]$.

Solution 2: solve

Keep the lens and do step 3 properly. The event $|\bar{R}_n - p| \le q_{\alpha/2}\sqrt{p(1-p)/n}$ is, after squaring,

$$(\bar{R}_n - p)^2 \le \frac{q_{\alpha/2}^2}{n}\,p(1-p), \qquad \text{that is,} \qquad \left(1 + \frac{q_{\alpha/2}^2}{n}\right)p^2 - \left(2\bar{R}_n + \frac{q_{\alpha/2}^2}{n}\right)p + \bar{R}_n^2 \le 0.$$

The left side is a convex quadratic in p, so the inequality holds exactly between its two roots $p_1 < p_2$, and

$$\mathcal{I}_{\text{solve}} = [p_1, p_2]$$

is the *same event* as the one in Equation (8.1), now with computable endpoints: $\lim_{n\to\infty} \mathbb{P}_p(\mathcal{I}_{\text{solve}} \ni p) = 1-\alpha$, with equality. The quadratic formula gives the roots in closed form,

$$p_{1,2} = \frac{1}{1 + q_{\alpha/2}^2/n}\left(\bar{R}_n + \frac{q_{\alpha/2}^2}{2n} \mp q_{\alpha/2}\sqrt{\frac{\bar{R}_n(1-\bar{R}_n)}{n} + \frac{q_{\alpha/2}^2}{4n^2}}\right), \tag{8.2}$$

an interval known as the Wilson score interval. The word 'score' is explained in Section 9.8: this interval is what you get by inverting a standard test of $H_0 : p = p_0$, the score test, which measures distance from p_0 in units of the standard error computed at p_0. Note that it is *not* centered at $\bar{R}_n$: its center is pulled toward $\frac{1}{2}$, by an amount of order $1/n$ because the lens is wider there.

With $n = 124$, $\bar{R}_n = 80/124$ and $q_{0.025} = 1.96$ the quadratic is $1.0310\,p^2 - 1.3213\,p + 0.4162 \le 0$, with roots $p_1 = 0.558$ and $p_2 = 0.724$: $\mathcal{I}_{\text{solve}} = [0.558, 0.724]$.

Pitfall: rounding the quadratic before solving it

Solve for p_1 and p_2 from the coefficients as computed, not from rounded versions of them. A quadratic's roots are sensitive to its coefficients, and this quadratic is a bad case. Its constant term is 0.4162; carrying that to two decimals as 0.41 moves the roots three points each way, to $[0.53, 0.75]$, and rounding it up to 0.42 instead gives $[0.59, 0.69]$, wrong in the other direction. Two decimals are not enough here either way, and the error is large enough to change the

conclusion: the $[0.53, 0.75]$ that the truncated coefficients produce is wider than the conservative interval $[0.557, 0.733]$, while the true solve interval $[0.558, 0.724]$ is the narrowest of the three. The solve interval can never be the widest, since Proposition 8.2 shows it always sits inside the conservative interval, so an interval that comes out wider than the conservative one is an arithmetic error and nothing else.

Solution 3: plug-in

Replace the unknown scale $\sqrt{p(1-p)}$ by its estimate $\sqrt{\hat{p}(1-\hat{p})}$, and let Slutsky's theorem pay for it.

Recall. Slutsky's theorem is the rule that lets you substitute an estimate for an unknown constant inside a limit: if one quantity settles into a distribution and a second converges in probability to a constant, their product and their sum behave as if the second simply *were* that constant. It is what makes every 'estimate $\pm$ standard error' formula legitimate.

This is the worked example of Section 5.2: since $\hat{p} \xrightarrow[n\to\infty]{\mathbb{P}} p$ (LLN) and $x \mapsto \sqrt{p(1-p)/(x(1-x))}$ is continuous at p (CMT),

$$\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{\hat{p}(1-\hat{p})}} = \underbrace{\sqrt{n}\,\frac{\bar{R}_n - p}{\sqrt{p(1-p)}}}_{\xrightarrow[n\to\infty]{(\mathrm{d})}\,\mathrm{Gauss}(0,1)} \cdot \underbrace{\sqrt{\frac{p(1-p)}{\hat{p}(1-\hat{p})}}}_{\xrightarrow[n\to\infty]{\mathbb{P}}\,1} \xrightarrow[n\to\infty]{(\mathrm{d})} \mathrm{Gauss}(0,1).$$

The left side is a new asymptotic pivot whose scale no longer involves p, so the recipe runs to the end:

$$\mathcal{I}_{\text{plug-in}} = \left[\bar{R}_n - q_{\alpha/2}\sqrt{\frac{\hat{p}(1-\hat{p})}{n}},\ \bar{R}_n + q_{\alpha/2}\sqrt{\frac{\hat{p}(1-\hat{p})}{n}}\right], \qquad \lim_{n\to\infty} \mathbb{P}_p(\mathcal{I}_{\text{plug-in}} \ni p) = 1-\alpha.$$

In the figure it reads a band of constant width, the lens's width at $\hat{p}$. The price is a second asymptotic approximation on top of the CLT: Slutsky's theorem is a statement about $n \to \infty$, and at finite n the estimated scale is itself noisy.

With $n = 124$: $1.96\sqrt{0.645 \times 0.355/124} = 0.084$, so $\mathcal{I}_{\text{plug-in}} = [0.561, 0.729]$.

The three side by side

Interval	95%, $n = 124$	Reads	Asymptotic coverage
conservative	$[0.557, 0.733]$	the widest band	$\geq 1-\alpha$; $= 1-\alpha$ only at $p = \frac{1}{2}$
solve	$[0.558, 0.724]$	the lens itself	$= 1-\alpha$; one approximation (CLT)
plug-in	$[0.561, 0.729]$	a band of half-width $q_{\alpha/2}\sqrt{\hat{p}(1-\hat{p})/n}$	$= 1-\alpha$; two approximations (CLT, Slutsky)

The conservative and plug-in intervals round to the $[0.56, 0.73]$ that Chapter 1 promised, the solve interval to $[0.56, 0.72]$, and Chapter 1's drug example is the same computation: 78 of 100 cured gives $\hat{p} = 0.78$ and $\mathcal{I}_{\text{plug-in}} = 0.78 \pm 1.96\sqrt{0.78 \times 0.22/100} = [0.699, 0.861]$, the $[0.70, 0.86]$ quoted there. There are more intervals still; the exact interval built from the Binomial distribution of $n\bar{R}_n$ rather than from the CLT (the Clopper–Pearson interval mentioned in Section 4.4) gives $[0.554, 0.729]$ for the kiss data, and Newcombe (1998) compares seven of them on coverage and width.[1] That they all agree to within a point is the CLT doing its job at $n = 124$.

Proposition 8.2 (How the three nest)**.** *For every sample and every n: $\mathcal{I}_{solve} \subseteq \mathcal{I}_{conserv}$ and $\mathcal{I}_{plug\text{-}in} \subseteq \mathcal{I}_{conserv}$, and both $\mathcal{I}_{solve}$ and $\mathcal{I}_{plug\text{-}in}$ contain $\bar{R}_n$.*

Proof strategy (core). Everything follows from $p(1-p) \le \frac{1}{4}$: a p that passes the solve test passes the conservative one, and the plug-in half-width is at most the conservative half-width.

Proof. If $p \in \mathcal{I}_{\text{solve}}$ then $|\bar{R}_n - p| \le q_{\alpha/2}\sqrt{p(1-p)/n} \le q_{\alpha/2}/(2\sqrt{n})$, so $p \in \mathcal{I}_{\text{conserv}}$. The plug-in half-width is $q_{\alpha/2}\sqrt{\hat{p}(1-\hat{p})/n} \le q_{\alpha/2}/(2\sqrt{n})$, the conservative one. Finally $p = \bar{R}_n$ satisfies the solve inequality trivially ($0 \le$ something non-negative), and the plug-in interval is centered there. □

In the picture: the lens lies inside the widest band, so any chord of the lens lies inside the corresponding chord of the band. Solve and plug-in need not nest with each other; here the solve interval is shifted toward $\frac{1}{2}$ and slightly shorter.

Intuition: which one to use

Use the solve interval for a proportion, the plug-in interval for everything else. The defense: the solve interval has the exact asymptotic level with one approximation, has a closed form (Equation (8.2)) and, as Figure 8.3 shows, holds up at finite n far better than the plug-in interval near $p = 0$ or 1, where the latter breaks; the only reason it isn't universal is that 'solve for the parameter' is a quadratic here and may be intractable elsewhere. The plug-in interval is the one every formula sheet gives, is what the rest of this book uses, and is fine whenever $n\hat{p}(1-\hat{p})$ is comfortably large. The conservative interval is for two jobs: a guaranteed level, and planning sample sizes before any data exist (Section 8.7), since it is the only one of the three whose width you know in advance.

8.5 What 'asymptotic level' buys at a finite n

Asymptotic level $1-\alpha$ is a promise about $n \to \infty$. For a proportion you can check what it delivers at the n you have because the distribution of $n\bar{R}_n$ is Binomial and the coverage probability is a finite sum:

$$\mathbb{P}_p(\mathcal{I} \ni p) = \sum_{k=0}^{n} \binom{n}{k} p^k (1-p)^{n-k} \, \mathbb{1}\big(\mathcal{I}(k) \ni p\big),$$

where $\mathcal{I}(k)$ is the interval the rule produces when k of the n couples turn right. Figure 8.3 plots this against p for the three intervals, at $n = 124$ and at $n = 20$.

[1] R. Newcombe, 'Two-sided confidence intervals for the single proportion: comparison of seven methods', *Statistics in Medicine* 17 (1998).

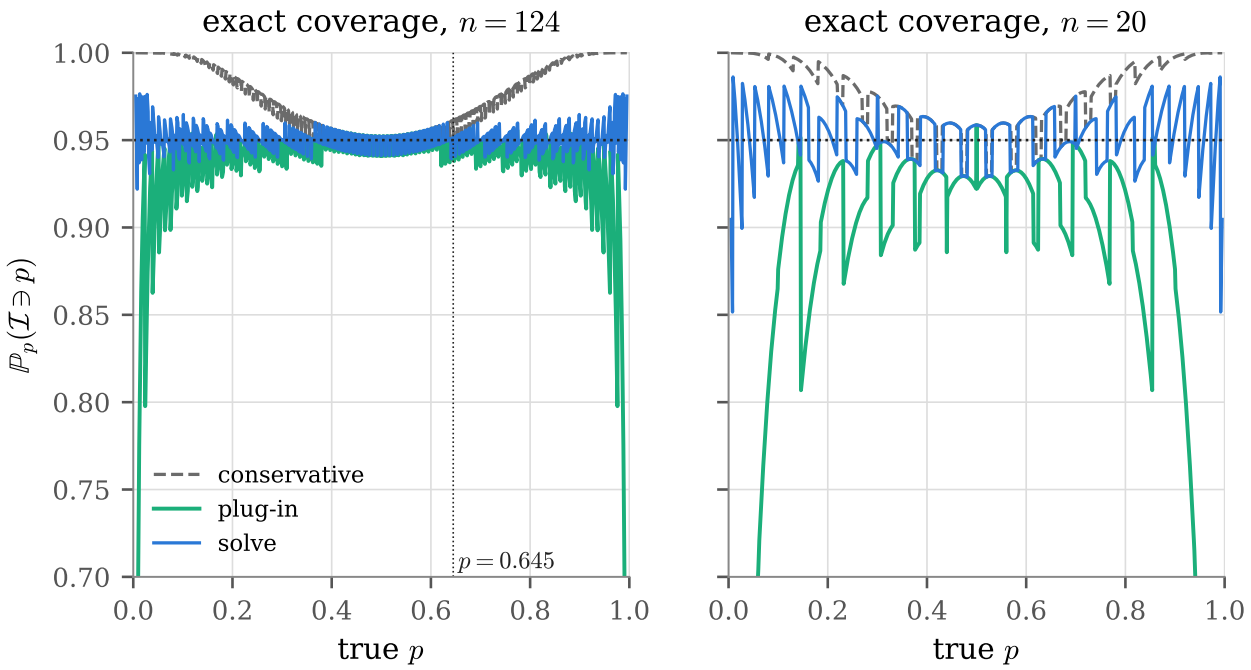


Figure 8.3: Exact coverage probability $\mathbb{P}_p(\mathcal{I} \ni p)$ of the three 95% intervals, as a function of the true p, from the Binomial distribution of $n\bar{R}_n$. Left, $n = 124$; right, $n = 20$. The conservative interval over-covers away from $p = \frac{1}{2}$ (its coverage is 0.99995 at $p = 0.05$). The plug-in interval under-covers, mildly at $n = 124$ (0.939 at $p = 0.645$) and badly at $n = 20$ near the edges. The solve interval stays near 0.95 at $n = 124$ (0.952 at $p = 0.645$) and above 0.92 at $n = 20$ for p between 0.01 and 0.99, where the plug-in interval falls as low as 0.33 (below the bottom of the plot). The saw teeth are real, not noise: $\bar{R}_n$ takes only $n + 1$ values, so there are only $n + 1$ possible intervals, and the coverage jumps whenever p crosses an endpoint of one of them.

Three lessons from the figure. The asymptotic promise is kept at $n = 124$ to within a point or two for p in the middle of $(0, 1)$, which is where the kiss data put it. The plug-in interval is the one that fails first as n shrinks or p approaches the edges because its width is $\sqrt{\hat{p}(1-\hat{p})/n}$ and $\hat{p}$ near 0 or 1 makes it think it is very precise (at $\hat{p} = 0$ its width is zero). And the conservative interval is not free: its extra width buys coverage you didn't ask for, and that over-coverage is wasted precision.

This is the chapter's instance of the book's stance on asymptotics (the intuition box that opens Chapter 3, and Section 4.4): the theorem licenses the Gaussian replacement and is silent about the error at your n. For a proportion the error can be computed, as here. For most models it can't, and the plug-in interval is used on trust and on the experience that n in the dozens with a parameter away from the boundary is usually enough.

8.6 The subway: when the scale depends on the parameter in a new way

The three solutions were introduced for a proportion, where the scale $\sqrt{p(1-p)}$ is bounded. The inter-arrival-time example of Section 5.4 shows what changes when it isn't. $T_1, \ldots, T_n \overset{iid}{\sim} \text{Expo}(\lambda)$, $\hat{\lambda} = 1/\bar{T}_n$, and the delta method gave

Recall. The delta method says what happens to a transformed estimator: if $\hat{\theta}_n$ is approximately Gaussian around θ with standard deviation $\sigma/\sqrt{n}$, then $g(\hat{\theta}_n)$ is approximately Gaussian around $g(\theta)$ with standard deviation $|g'(\theta)|\sigma/\sqrt{n}$. The slope of g at θ stretches or compresses the error.

$$\sqrt{n}\,(\hat{\lambda}-\lambda) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,\lambda^2), \qquad \text{so} \qquad \lim_{n\to\infty} \mathbb{P}_\lambda\left(|\hat{\lambda}-\lambda| \le q_{\alpha/2}\frac{\lambda}{\sqrt{n}}\right) = 1-\alpha.$$

Again the half-width involves the parameter, now as a factor λ rather than $\sqrt{p(1-p)}$, and Figure 8.4 shows the consequence: the band is a wedge, not a lens.

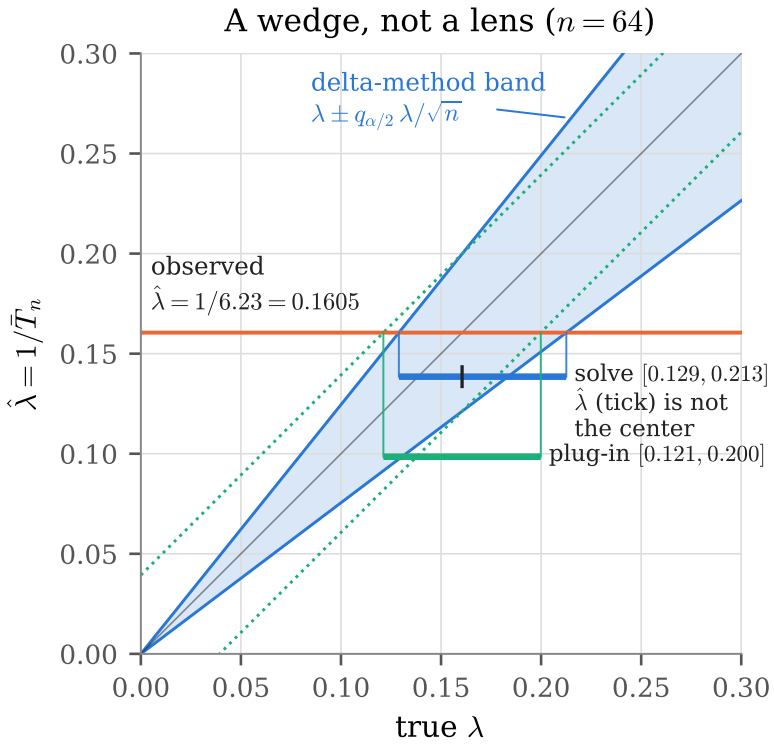


Figure 8.4: The inversion picture for the subway rate. For each true λ, the delta method says $\hat{\lambda}$ lands in $\lambda \pm q_{\alpha/2}\lambda/\sqrt{n}$ about 95% of the time: a wedge that keeps opening as λ grows, with no widest point. The observed $\hat{\lambda} = 1/6.23 = 0.1605$ is the horizontal line, and the solve interval is the chord it cuts from the wedge, which is not centered at $\hat{\lambda}$ because the wedge is wider to the right. The plug-in interval uses a band of constant width $q_{\alpha/2}\hat{\lambda}/\sqrt{n}$ (dotted) and is centered. There is no conservative interval: no band of constant width contains the wedge.

1. *Conservative:* there is nothing to bound. λ ranges over $(0,\infty)$, so $\sup_\lambda q_{\alpha/2}\lambda/\sqrt{n} = \infty$ and the 'conservative interval' is $(0,\infty)$, which has level 100% and says nothing.
2. *Solve:* the inequality is linear in λ this time. $|\hat{\lambda}-\lambda| \le q_{\alpha/2}\lambda/\sqrt{n}$ is the same as $\lambda(1-q_{\alpha/2}/\sqrt{n}) \le \hat{\lambda} \le \lambda(1+q_{\alpha/2}/\sqrt{n})$, and if $n > q_{\alpha/2}^2$ (so that $1-q_{\alpha/2}/\sqrt{n} > 0$) this solves to

$$\mathcal{I}_{\text{solve}} = \left[\frac{\hat{\lambda}}{1+q_{\alpha/2}/\sqrt{n}},\ \frac{\hat{\lambda}}{1-q_{\alpha/2}/\sqrt{n}}\right].$$

3. *Plug-in:* replace λ in the half-width by $\hat{\lambda}$ (Slutsky, since $\hat{\lambda} \xrightarrow[n\to\infty]{\mathbb{P}} \lambda$):

$$\mathcal{I}_{\text{plug-in}} = \left[\hat{\lambda}\left(1 - \frac{q_{\alpha/2}}{\sqrt{n}}\right),\ \hat{\lambda}\left(1 + \frac{q_{\alpha/2}}{\sqrt{n}}\right)\right].$$

With $n = 64$ and $\bar{T}_n = 6.23$ minutes: $\hat{\lambda} = 0.1605$ per minute and $q_{0.025}/\sqrt{64} = 0.245$, so

$$\mathcal{I}_{\text{solve}} = [0.129, 0.213] \approx [0.13, 0.21], \qquad \mathcal{I}_{\text{plug-in}} = [0.121, 0.200] \approx [0.12, 0.20].$$

The two differ more than the kiss intervals did because $n = 64$ is smaller and $q/\sqrt{n} = 0.245$ is not small: $1/(1 - 0.245) = 1.32$ against $1 + 0.245 = 1.245$. The solve interval reaches further to the right, where the wedge is wider: large values of λ produce more variable $\hat{\lambda}$, so more of them are compatible with the observed 0.16.

Remark. There is a second way to see the solve interval that explains its asymmetry. Build the plug-in interval for the *mean* $1/\lambda$ first: the CLT gives $\sqrt{n}(\bar{T}_n - 1/\lambda) \xrightarrow[n\to\infty]{(\text{d})} \text{Gauss}(0, 1/\lambda^2)$, and plugging in gives $\bar{T}_n(1 \pm q_{\alpha/2}/\sqrt{n}) = [4.70, 7.76]$ minutes. Now map it through $x \mapsto 1/x$, which is decreasing: $\lambda \in [1/7.76, 1/4.70] = [0.129, 0.213]$, which is $\mathcal{I}_{\text{solve}}$ exactly. A confidence interval for θ becomes a confidence interval for $g(\theta)$, with the same level, by applying any continuous and strictly monotone g to its endpoints, in the right order ($[g(L), g(U)]$ for increasing g, $[g(U), g(L)]$ for decreasing g) because $\theta \in [L, U]$ and $g(\theta) \in g([L, U])$ are the same event. The delta method is needed when g is not monotone (the odds of Chapter 5 are, $p(1-p)$ is not), or when you want the interval for $g(\theta)$ to be symmetric about $g(\hat{\theta}_n)$, and symmetry is not a virtue in itself.

8.7 What makes an interval wide

The half-width of the conservative and plug-in intervals is exactly

$$\text{half-width} = q_{\alpha/2} \times \frac{\widehat{\sigma}}{\sqrt{n}},$$

a quantile times an estimated standard error, and the solve intervals have the same size up to terms of smaller order. The two factors are the two things you control.

The level. Raising the confidence level widens the interval, through the quantile, and the cost accelerates (Figure 8.5, left). Going from 90% to 95% costs 19% more width; from 95% to 99%, 31% more; and 100% confidence costs the whole parameter space, since $q_{\alpha/2} \to \infty$ as $\alpha \to 0$. This is why 95% is a convention and not a law: it is a point on a curve where the interval is still informative and the miss rate is still small. A 98% interval of the same type always contains the 95% one because $q_{0.01} > q_{0.025}$; but a 98% interval need not contain a 95% interval built differently. $[0,\ \bar{R}_n + q_{0.02}\,\widehat{\sigma}/\sqrt{n}]$ has asymptotic level 98% and does not contain the 95% lower-bound interval $[\bar{R}_n - q_{0.05}\,\widehat{\sigma}/\sqrt{n},\ 1]$, nor the other way around: the 98% interval reaches down to 0 but stops short of 1, the 95% one reaches up to 1 but stops short of 0, so each contains points the other does not. Nesting holds within one family of intervals indexed by α, not between two intervals of different shapes whose levels happen to differ.

The sample size. The width shrinks like $1/\sqrt{n}$ (Figure 8.5, right), which is the rate of the CLT and the rate of every interval in this book built from a sample mean. The practical form of $1/\sqrt{n}$: to halve the width you need four times the data; to get one more decimal place you need a hundred

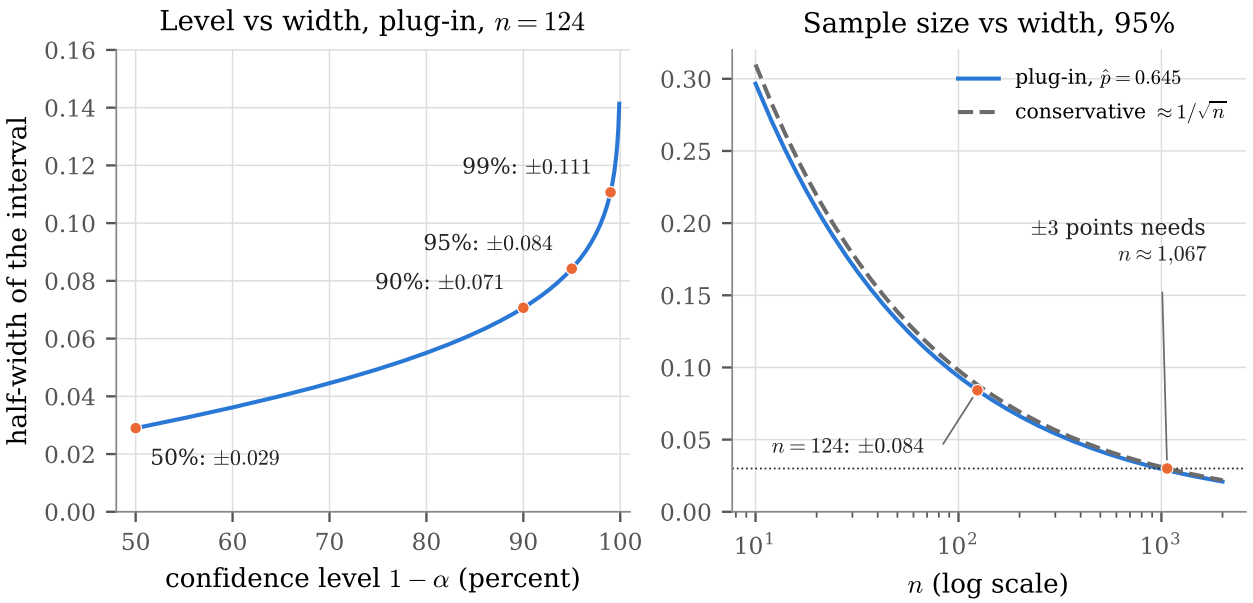


Figure 8.5: Left: the half-width of the plug-in interval for the kiss data as a function of the confidence level. The quantile $q_{\alpha/2}$ does the work: 0.67 at 50%, 1.645 at 90%, 1.96 at 95%, 2.58 at 99%, 3.29 at 99.9%, and infinite at 100%. Right: the half-width at 95% as a function of n, for the plug-in interval at $\hat{p} = 0.645$ and the conservative interval, whose half-width $q_{0.025}/(2\sqrt{n}) = 0.98/\sqrt{n}$ is the pollster's 'about $1/\sqrt{n}$'. The curve is $1/\sqrt{n}$: four times the data halves the width, and a margin of ±3 points needs $n \approx 1{,}067$.

times the data. Intervals that shrink faster exist when the model has special structure (Exercise 8.5 has one that shrinks like $1/n$), but they are built on the model's exact distribution rather than on the CLT, and they are fragile when the model is wrong.

From undergraduate statistics: the pollster's ±3 points

National polls report a 'margin of error of ±3 percentage points' with about 1,000 respondents, whatever the question. That is the conservative interval: its half-width $q_{0.025}/(2\sqrt{n}) = 0.98/\sqrt{n} \approx 1/\sqrt{n}$ doesn't depend on p, so a pollster can state it before seeing a single answer. Setting $0.98/\sqrt{n} = 0.03$ gives $n = (1.96/0.06)^2 = 1{,}067$. The same arithmetic plans any proportion study: the sample size for a margin m at 95% is $n = (0.98/m)^2$, and the conservative bound is the right one to use because at the planning stage there is no $\hat{p}$ to plug in.

8.8 What the 95% means, and three things it doesn't

Take $\mathcal{I}_{\text{plug-in}} = [0.12, 0.20]$ for the subway rate. The sentence '$[0.12, 0.20]$ is a confidence interval of asymptotic level 95% for λ' means: *if you were to repeat the experiment*, collecting 64 new inter-arrival times and computing the interval each time, the parameter λ would be inside the resulting interval about 95% of the time, with the approximation improving as n grows. This is the frequentist interpretation, and Figure 8.1 is a hundred such repetitions. Three readings that sound the same and are not:

Pitfall: three wrong readings

1. *'The probability that* $\lambda \in [0.12, 0.20]$ *is* 0.95*.'* No. λ is a fixed number and $[0.12, 0.20]$ is a fixed interval; $\mathbb{P}(\lambda \in [0.12, 0.20])$ is 1 or 0, and the theory doesn't tell you which. The 0.95 belongs to the procedure. (The sentence 'you are 95% confident that λ is in $[0.12, 0.20]$', as in the drug example of Chapter 1, is the accepted shorthand for the procedure statement, and is fine as long as it's read that way.)
2. *'*95% *of the inter-arrival times fall in* $[0.12, 0.20]$*.'* No. The interval is for the parameter λ, not for the data; it shrinks to a point as $n \to \infty$ while the data's spread stays what it is.
3. *'*95% *of the future estimates* $\hat{\lambda}$ *will fall in* $[0.12, 0.20]$*.'* No. A future $\hat{\lambda}$ is centered at λ, not at this sample's 0.16, and lands in this particular interval about 95% of the time only if this sample's 0.16 happens to be close to λ, and less the further it is off.

The first reading is the one everyone wants, a probability statement about the parameter given the data. It is available, but at a price: you must treat λ as random, with a distribution before the data (a prior), and compute its distribution after the data (the posterior). That is the Bayesian approach of Chapter 18, and the interval it produces, which Chapter 19 calls a Bayesian confidence region, does carry the probability 0.95 in the sense of reading 1. For large n the two kinds of interval nearly coincide, which is the Bernstein–von Mises result mentioned in Section 5.5; what differs is what the number 0.95 is a probability *of*.

8.9 What comes next

A confidence interval answers 'which values of θ are compatible with the data?'. A test answers a narrower question, 'is this particular value compatible?', and Exercise 8.6 already showed that the two are the same construction: the value $\frac{1}{2}$ is rejected exactly when it falls outside a one-sided interval. Chapter 9 starts from the decision rather than the interval, which brings in two new things the interval didn't need: two kinds of error, with a level for one and a power for the other, and the p-value, which reports a test without fixing α in advance. The thresholds 72 and 74 that Chapter 1 quoted are derived there.

8.10 Summary

- A confidence interval of level $1 - \alpha$ is a random interval with computable endpoints such that $\mathbb{P}_\theta(\mathcal{I} \ni \theta) \geq 1 - \alpha$ for every θ; asymptotic level replaces the probability by its limit. Level is a floor: a 95% interval is also a 90% one.
- The recipe: find a pivot, trap it between quantiles ($q_{\alpha/2}$ two-sided, q_α one-sided), solve for the parameter. Gaussian data with known variance give an exact pivot and an exact interval; the CLT gives an asymptotic pivot and an asymptotic interval.
- When the pivot's scale involves the parameter, the interval it suggests isn't computable. Three fixes: bound the scale (conservative, level $\geq 1 - \alpha$, widest), solve the inequality for the parameter (solve, exact asymptotic level, one approximation), or estimate the scale and invoke Slutsky (plug-in, exact asymptotic level, two approximations). In the inversion picture they read the widest band, the lens, and a band of constant width.
- Kiss experiment, 95%: conservative $[0.557, 0.733]$, solve $[0.558, 0.724]$, plug-in $[0.561, 0.729]$; they agree to within a point, and here the solve interval is the narrowest, not the widest.

Subway: solve $[0.13, 0.21]$, plug-in $[0.12, 0.20]$, no conservative interval.

- At finite n the three behave differently: for a proportion the exact coverage can be computed from the Binomial, and at $n = 124$ the plug-in interval covers 0.939 and the solve interval 0.952 at $p = 0.645$, while at $n = 20$ the plug-in interval fails near the edges and the solve interval holds. Use solve for a proportion, plug-in otherwise.
- A continuous, strictly monotone map carries a confidence interval for θ to one for $g(\theta)$ with the same level; the subway solve interval is the plug-in interval for $1/\lambda$ mapped through $x \mapsto 1/x$.
- Half-width $= q_{\alpha/2}\,\widehat{\sigma}/\sqrt{n}$: the level acts through the quantile (and 100% costs everything), the sample size through $1/\sqrt{n}$ (four times the data, half the width; ± 3 points needs about 1,000).
- The level is the long-run hit rate of the procedure. A realized interval has no probability attached; the statement '$\mathbb{P}(p \in [0.56, 0.73]) = 0.95$' is Bayesian, and Chapter 18 says what it costs.

8.11 Exercises

Exercise 8.1. For the kiss data, find a 98% asymptotic confidence interval for p of the conservative and of the plug-in type. If $[0.34, 0.57]$ is a 95% confidence interval for some proportion, is it also a 98% one? A 90% one? *(Solution on page 224.)*

Exercise 8.2. A second study observes 150 couples, 75 turning each way, so $\bar{R}_n = 0.5$. Which of the following are realizations of 50% asymptotic confidence intervals for p: (1) $[0, 0.5]$; (2) $[0.5, 1]$; (3) $[0.466, 0.533]$; (4) $[0.48, 0.52]$? *(Solution on page 225.)*

Exercise 8.3. Prove that $\mathcal{I}_{\text{plug-in}} \subseteq \mathcal{I}_{\text{conserv}}$ always, and find the sample for which the two coincide. Then show that $\mathcal{I}_{\text{solve}} = \mathcal{I}_{\text{plug-in}}$ never happens unless $\bar{R}_n = \frac{1}{2}$, by comparing centers. *(Solution on page 225.)*

Exercise 8.4. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(\theta, \theta)$ for an unknown $\theta > 0$: the mean and the variance are the same unknown number. (a) Find an exact pivot based on $\bar{X}_n$. (b) Write the interval $\mathcal{I}_\theta$, depending on θ, with exact coverage 90%. (c) Find the plug-in interval of asymptotic level 90%. (d) Find the solve interval, and say whether its level is exact or asymptotic. *(Solution on page 225.)*

Exercise 8.5. Let $X_1, \ldots, X_n$ be i.i.d. with density $f(x) = e^{-(x-a)}\mathbb{1}(x \geq a)$, a unit Exponential shifted by an unknown $a \in \mathbb{R}$. (a) Show that $\hat{a}_1 = \bar{X}_n - 1$ is consistent and find its asymptotic distribution and the two-sided plug-in interval of asymptotic level $1 - \alpha$. (b) Let $\hat{a}_2 = \min_i X_i$. Show that $n(\hat{a}_2 - a) \sim \text{Expo}(1)$ exactly, and build a one-sided interval of exact level $1 - \alpha$ from it. (c) Compare the two intervals' widths and their robustness. *(Solution on page 225.)*

Exercise 8.6. For the kiss data, build a 95% lower confidence bound for p of the plug-in type and of the conservative type. Then show that '$\frac{1}{2}$ is below the conservative lower bound' is the same event as '$n\bar{R}_n \geq 72$', the one-sided threshold that Chapter 1 quoted and Chapter 9 derives. *(Solution on page 225.)*

Exercise 8.7. The subway data have an exact interval that uses no limit theorem. If $T_1, \ldots, T_n \overset{iid}{\sim} \text{Expo}(\lambda)$ then $\lambda \sum_i T_i = \lambda n \bar{T}_n$ has the Gamma distribution with shape n and rate 1, whatever λ is (you may take this as given). Use it as a pivot to build a confidence interval for λ of exact level 95%, and compute it for $n = 64$, $\bar{T}_n = 6.23$, given that the 2.5% and 97.5% points of that Gamma distribution are 49.29 and 80.60. *(Solution on page 226.)*

Chapter 9

Hypothesis Tests: Errors, Level, Power and p-Values

In this chapter

- What a test is: a yes/no rule applied to the data, written as an indicator of a rejection region. The null hypothesis is the claim the data are allowed to overturn; the alternative is what you conclude when they do.
- The two ways a test can be wrong, why they are not treated alike, and the two numbers that describe a test: its level, a cap on one error, and its power, one minus the other.
- The recipe behind every test in this book: a test statistic whose distribution under the null hypothesis you know, and a quantile of that distribution as the threshold. One-sided and two-sided versions, and why the null hypothesis makes the unknown parameter known.
- The test for the kiss data, done exactly and done by the CLT: thresholds 72 and 74 one way, 72 and 73 the other, and why they differ.
- The p-value: the smallest level at which the data reject, and the one number a test reports.
- Tests and confidence intervals are the same construction seen from two sides.

The running example is the kiss experiment: $n = 124$ couples, 80 turning right, and the question Chapter 1 opened with, whether 80 is more than chance would produce if couples had no preference for either side.

9.1 A yes/no question about a parameter

A confidence interval reports a range. Often what is wanted is a decision. The examples below are the kind a manager or a scientist actually faces, and each is a yes/no question about a parameter of a model.

- **Boarding a plane.** Two airlines board differently: one from the rear of the plane to the front ('R2F'), the other by window seats first, then middle, then aisle ('WilMA'). Which is faster? The data are $n = 72$ R2F flights averaging 24.2 minutes with a standard deviation of 2.1, and $m = 56$ WilMA flights averaging 15.9 minutes with a standard deviation of 1.3. With boarding times modeled as $X \sim \text{Gauss}(\mu_1, \sigma_1^2)$ and $Y \sim \text{Gauss}(\mu_2, \sigma_2^2)$ and the two

samples independent, the question is: is $\mu_1 = \mu_2$, or is $\mu_1 > \mu_2$? Two samples, one question: a **two-sample test**.

- **Waiting in the emergency room.** The national average wait is 30 minutes. Patients at one hospital say its wait is longer. With $X_1, \dots, X_n$ the waits of n random patients there and $\mu = \mathbb{E}[X_1]$, the question is: is $\mu > 30$? One sample: a **one-sample test**.
- **The kiss experiment.** If couples have no preference for either side, $p = \frac{1}{2}$. Of 124 couples, 80 turned right. Is $p = \frac{1}{2}$, or not?

The obvious first answer to the boarding question is '$\bar{X}_n = 24.2 > 15.9 = \bar{Y}_m$, so $\mu_1 > \mu_2$'. What is wrong with it is this: by bad luck the 72 R2F flights might all have been full and the 56 WilMA flights half empty, and sample means move from sample to sample. A better rule puts a buffer around each mean and asks whether the buffers still separate:

$$\text{conclude } \mu_1 > \mu_2 \text{ if } \bar{X}_n - \text{buffer} > \bar{Y}_m + \text{buffer}.$$

For the one-sample version, conclude $\mu > 30$ if $\bar{X}_n > 30 + \text{buffer}$. What should the buffer be? Something that measures the size of the random fluctuations of $\bar{X}_n$, and Chapter 4 says what that is: a few multiples of $\sigma/\sqrt{n}$. Figure 9.1 draws the heuristic for the boarding data and for two one-sample examples. With $n = 72$ and $m = 56$ the buffers around 24.2 and 15.9 are a fraction of a minute wide, and the conclusion is not in doubt. For a coin tossed 30 times with 13 heads, the buffer around $\hat{p} = 0.43$ reaches well past $\frac{1}{2}$, and the heuristic says nothing.

The rest of the chapter makes the heuristic precise: what counts as a buffer, what guarantee it comes with, and what to report when the data land near the edge.

From undergraduate statistics: statistically significant

Every one of these questions was answered in an introductory course with a z-test or a t-test and the phrase 'the difference is statistically significant at the 5% level'. This chapter is that procedure, written so that each piece has a name and a guarantee: the hypotheses, the rejection region, the two errors, the level, the power and the p-value. 'Significant at 5%' will turn out to mean 'the test of level 5% rejected', and '$p < 0.05$' will be a theorem rather than a rule of thumb.

9.2 Null and alternative hypotheses

Let $X_1, \dots, X_n$ be a sample of i.i.d. r.v.s and $(E, (\mathbb{P}_\theta)_{\theta \in \Theta})$ a statistical model (Chapter 2). Let Θ_0 and Θ_1 be disjoint subsets of Θ. The two hypotheses are

$$H_0 : \theta \in \Theta_0 \qquad \text{and} \qquad H_1 : \theta \in \Theta_1.$$

H_0 is the **null hypothesis**, H_1 the **alternative hypothesis**. If you believe the true θ is in one of the two sets, you test H_0 against H_1: you look in the data for evidence against H_0, and you decide whether to **reject** H_0.[1]

[1] Some treatments require Θ_0 and Θ_1 to partition Θ, and others, this book among them, ask only that the two sets be disjoint. The difference matters when values of θ in neither set are possible, in which case the test is silent about them. In every example in this book the two sets partition Θ.

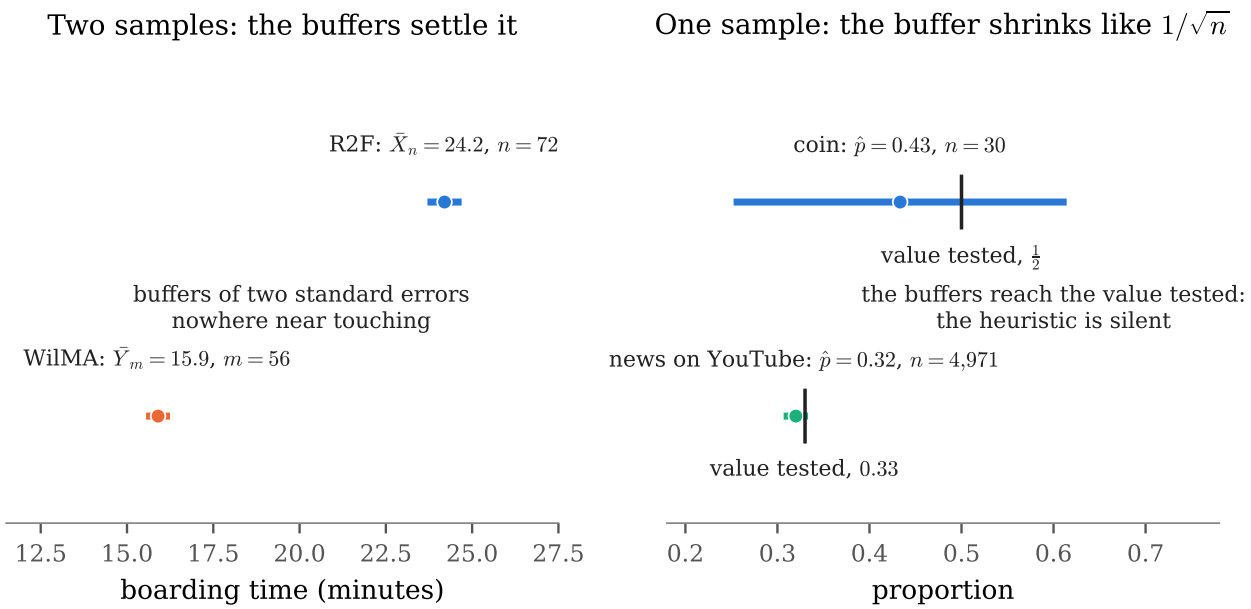


Figure 9.1: A first heuristic: put a buffer of two standard errors around each estimate and see whether it reaches the value being tested. Left: the boarding data; the buffers around 24.2 and 15.9 minutes reach about half a minute (R2F) and a third of a minute (WilMA) on each side and are nowhere near each other, so the two methods differ. Right: two one-sample examples, each with the value being tested drawn as a tick. A coin with 13 heads in 30 tosses has a buffer 0.43 ± 0.18 that reaches $\frac{1}{2}$ easily; a survey of 4,971 Americans in which 32% get news on YouTube has a buffer 0.32 ± 0.013 that just reaches the 33% being tested. In both the heuristic is silent. The buffer is $2\sqrt{\hat{p}(1-\hat{p})/n}$, and shrinks like $1/\sqrt{n}$.

The two hypotheses do not play symmetric roles, and this asymmetry is the first thing to internalize. The data are used only to try to disprove H_0. If they fail to, that is not evidence that H_0 is true; it is a lack of evidence against it. The standard phrase is 'innocent until proven guilty': the null hypothesis is the defendant, the data are the evidence, and a verdict of not guilty is not a finding of innocence. So the conclusions of a test are 'reject H_0' and 'fail to reject H_0', never 'accept H_0'.

Which claim goes in the null? The one that gets the benefit of the doubt: the status quo, the claim of no effect, the value you would act on if the data said nothing. For the kiss experiment the null is $p = \frac{1}{2}$, no preference; for the emergency room it is 'the wait is not longer than the national average'; for the boarding methods it is 'no difference'. The alternative is the claim that needs evidence to be adopted, and it is the one the study was run to look for.

Simple and composite. A hypothesis is **simple** if its set is a single point, $\Theta_k = \{\theta_k\}$, and **composite** otherwise, typically a half-line $\{\theta > \theta_k\}$ or $\{\theta < \theta_k\}$ or everything but a point, $\{\theta \neq \theta_k\}$. When $\Theta \subseteq \mathbb{R}$ and H_0 is simple, $H_0 : \theta = \theta_0$, the terminology is:

- **two-sided test**: $H_1 : \theta \neq \theta_0$, departures in either direction count as evidence;
- **one-sided test**: $H_1 : \theta > \theta_0$ or $H_1 : \theta < \theta_0$, only one direction counts.

One-sided tests are also written with a composite null, $H_0 : \theta \leq \theta_0$ against $H_1 : \theta > \theta_0$; Section 9.4 explains why the two forms of the null lead to the same test. For the examples above: the boarding question is one-sided ($\mu_1 > \mu_2$); the emergency room is one-sided ($\mu > 30$); the fair coin is two-sided ($p \neq \frac{1}{2}$); and the test for the kiss data can be put either way, two-sided as $H_0 : p = \frac{1}{2}$ against

$H_1 : p \neq \frac{1}{2}$, or one-sided as Chapter 1 asked it, $H_1 : p > \frac{1}{2}$. Both are done below.

Pitfall: one-sided is a modeling decision, not a reading of the data

Writing $H_1 : \mu_1 > \mu_2$ for the boarding methods says that $\mu_1 < \mu_2$ is not a possibility you entertain. The modeling assumption reduces the number of ways the hypothesis $\mu_1 = \mu_2$ can be rejected: $\mu_1 < \mu_2$ is not allowed, so a sample pointing that way counts as no evidence against the null at all. That decision has to be made before looking at the data. Choosing the side after seeing which way $\bar{X}_n - \bar{Y}_m$ came out is a way of rejecting more often than the level promises: a nominal 5% becomes an actual 10%, since a two-sided question is then being answered with a one-sided threshold. Section 9.6 shows the thresholds differ (72 against 74 right turns for the kiss data), which is exactly the room this mistake exploits.

9.3 What a test is

Definition 9.1 (Test, rejection region)**.** A **test** is a statistic $\psi : E^n \to \{0, 1\}$: a function of the sample $X_1, \ldots, X_n$ taking the values 0 and 1, such that

- if $\psi(X_1, \ldots, X_n) = 0$, H_0 is not rejected;
- if $\psi(X_1, \ldots, X_n) = 1$, H_0 is rejected.

The **rejection region** of ψ is the set of samples on which it rejects,

$$R_\psi = \{(x_1, \ldots, x_n) \in E^n : \psi(x_1, \ldots, x_n) = 1\},$$

so that $\psi = \mathbb{1}\{R_\psi\}$ (toolkit A.4).

> **In words.** A test is a rule, fixed before the data arrive, that sorts every possible sample into two bins: reject or don't. The rejection region is the 'reject' bin, and the test is its indicator. Because the rule is applied to a random sample, ψ is a random variable, a Bernoulli one, and asking how often it equals 1 is a probability question, like asking how often an interval covers.
>
> **For the kiss data.** 'Reject $p = \frac{1}{2}$ if 72 or more of the 124 couples turn right' is the test $\psi = \mathbb{1}\{n\bar{R}_n \geq 72\}$. Its rejection region is every sequence of 124 zeros and ones with at least 72 ones. On the sample we have, $n\bar{R}_n = 80$ and $\psi = 1$.

The definition contains no probability: a test is a partition of the sample space E^n of Section 2.3 into two pieces. Probability enters when you ask how the test behaves under each candidate $\mathbb{P}_\theta$, which is the next section.

In practice every test in this book has the form

$$\psi = \mathbb{1}\{T_n > c\}$$

for a statistic $T_n = T_n(X_1, \ldots, X_n)$, called the **test statistic**, and a number c, the **threshold**; the rejection region is then $\{T_n > c\}$ (for a count, '≥ 72' is '> 71', so the strict inequality is only a matter of where c sits). The test statistic compresses the sample into one number that measures how far the data are from what H_0 predicts, and the threshold says how far is too far. Everything a test does is in those two choices, and the sections that follow are about making them.

From undergraduate statistics: the test statistic you already know

In an introductory course the test statistic was $z = (\bar{x} - \mu_0)/(\sigma/\sqrt{n})$, or t with s in place of σ, and the threshold was a number from a table, 1.96 or 1.645 or a t critical value. That is $\psi = \mathbb{1}\{|T_n| > c\}$ with T_n the standardized mean. What this chapter adds is the function ψ itself as an object, so that 'the test rejects with probability α when H_0 is true' is a statement about a random variable with a definition, and so that tests with other statistics (likelihood ratios in Chapter 15, chi-squared distances in Chapter 16) are the same object with a different T_n.

9.4 Two kinds of error, level and power

A test can be wrong in two ways, and the standard table lays them out with the truth in the rows and the decision in the columns:

	not rejected ($\psi = 0$)	rejected ($\psi = 1$)
H_0 true ($\theta \in \Theta_0$)	correct	**type 1 error**
H_1 true ($\theta \in \Theta_1$)	**type 2 error**	correct

A **type 1 error** is rejecting H_0 when it is true: convicting the innocent. A **type 2 error** is failing to reject H_0 when H_1 is true: acquitting the guilty. Since the truth is one fixed θ and the decision is random, the natural measure of each error is a probability, and since the probability depends on which θ is true, each error is a *function* of θ on its own hypothesis.

Definition 9.2 (Type 1 error, type 2 error, power)**.** For a test ψ of $H_0 : \theta \in \Theta_0$ against $H_1 : \theta \in \Theta_1$:

- the **type 1 error** of ψ is the function

$$\alpha_\psi : \Theta_0 \to [0,1], \qquad \theta \mapsto \mathbb{P}_\theta(\psi = 1);$$

- the **type 2 error** of ψ is the function

$$\beta_\psi : \Theta_1 \to [0,1], \qquad \theta \mapsto \mathbb{P}_\theta(\psi = 0);$$

- the **power** of ψ is the number (toolkit A.2)

$$\pi_\psi = \inf_{\theta \in \Theta_1} \left(1 - \beta_\psi(\theta)\right).$$

In words. Fix a truth θ and imagine rerunning the study many times. If θ is in the null set, the fraction of reruns on which the test rejects is the type 1 error at that θ; if θ is in the alternative, the fraction on which it fails to reject is the type 2 error at that θ, and the fraction on which it rejects is the power at that θ. The power of the test is the worst case of that last fraction over the whole alternative: the least the test can be relied on to detect a departure from H_0.

For the kiss data. For $\psi = \mathbb{1}\{n\bar{R}_n \geq 72\}$ and the one-sided hypotheses $H_0 : p = \frac{1}{2}$, $H_1 : p > \frac{1}{2}$: the type 1 error is the single number $\mathbb{P}_{1/2}(n\bar{R}_n \geq 72) = 0.044$. The type 2 error at $p = 0.645$ is $\mathbb{P}_{0.645}(n\bar{R}_n \leq 71) = 0.057$, so the test detects a true p of 0.645 with probability 0.943. The power in the definition's sense is $\inf_{p>1/2}$ of that detection probability, which is 0.044, approached as $p \downarrow \frac{1}{2}$: no test can reliably detect a p just above $\frac{1}{2}$ with 124 couples, and the infimum records that.

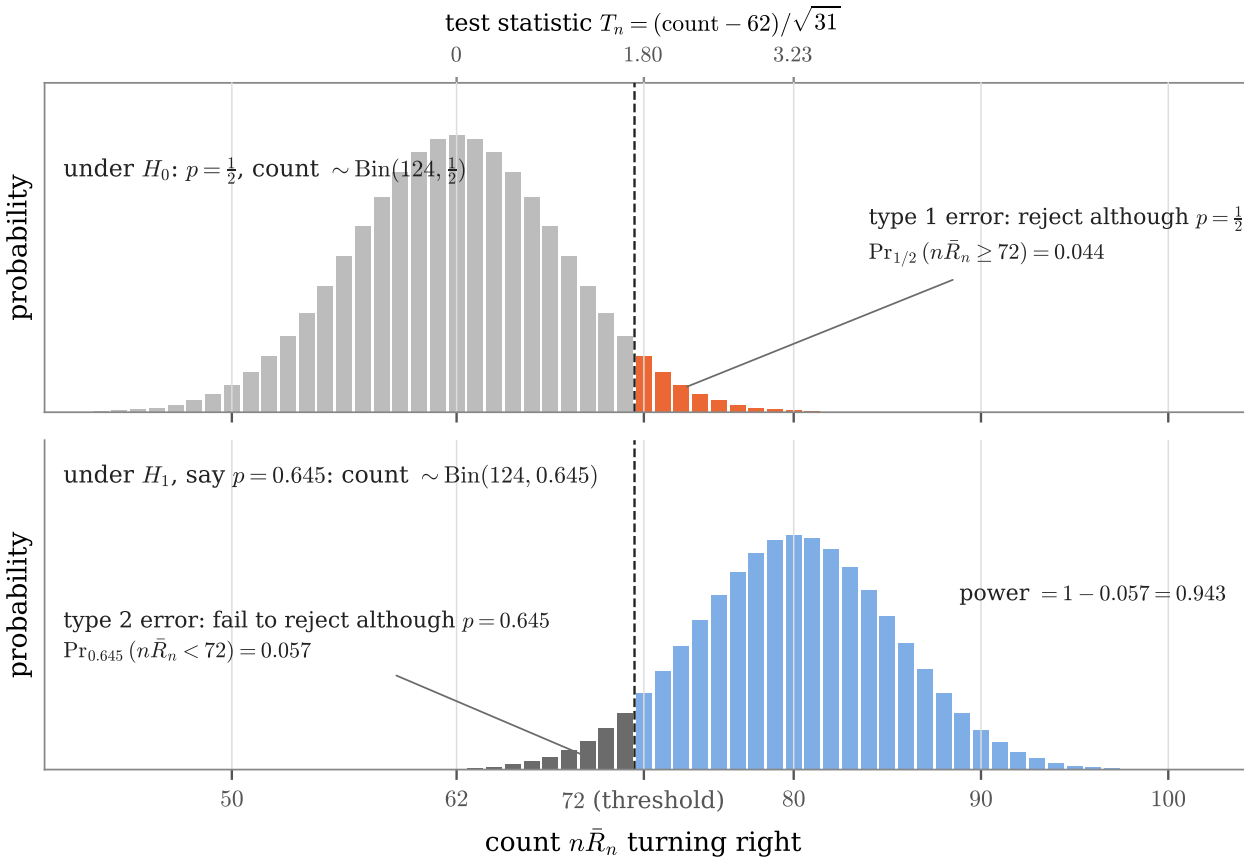


Figure 9.2: One threshold, two distributions, two errors, for the one-sided test on the kiss data 'reject $p = \frac{1}{2}$ if the count is 72 or more'. Top: if $p = \frac{1}{2}$ the count is $\text{Bin}(124, \frac{1}{2})$, centered at 62, and the orange mass at 72 and above is the type 1 error, 0.044. Bottom: if instead $p = 0.645$ the count is $\text{Bin}(124, 0.645)$, centered at 80; the gray mass below 72 is the type 2 error at that p, 0.057, and the blue mass is the power there, 0.943. Moving the threshold right shrinks the orange and grows the gray. The scale along the top is the test statistic $T_n = (\text{count} - 62)/\sqrt{31}$ of Section 9.5: the threshold 72 is $T_n = 1.80$, and the observed 80 is 3.23.

Figure 9.2 is the picture to hold on to. The same threshold cuts two different distributions of the count: the one H_0 predicts and one of the many H_1 allows. What falls on the wrong side of the cut under each is the corresponding error. And the picture shows the trade: sliding the threshold to the right makes the type 1 error smaller and the type 2 error larger, sliding it left does the reverse, and at a fixed n you cannot shrink both. The two distributions pull apart as n grows (their standard deviations are $\sqrt{p(1-p)/n}$ on the proportion scale), and that is the only lever that lowers both errors at once.

Rather than two functions on two sets, it is often clearer to look at one function on all of Θ, the probability of rejecting,

$$\theta \mapsto \mathbb{P}_\theta(\psi = 1),$$

which is α_ψ on Θ_0 and $1 - \beta_\psi$ on Θ_1.[2] Figure 9.3 draws it for the kiss data. A good test has a

[2]Two conventions for the letter β are in common use. This book writes β_ψ for the type 2 error, as in Definition 9.2. Elsewhere $\beta(\theta)$ names this rejection probability itself, the power function, which is $1 - \beta_\psi(\theta)$ on Θ_1 and $\alpha_\psi(\theta)$ on Θ_0. Check which one a source means before carrying a formula across.

rejection curve that is low on Θ_0 and climbs quickly to 1 on Θ_1; the level, defined next, caps the low part, and the power is the lowest point of the high part.

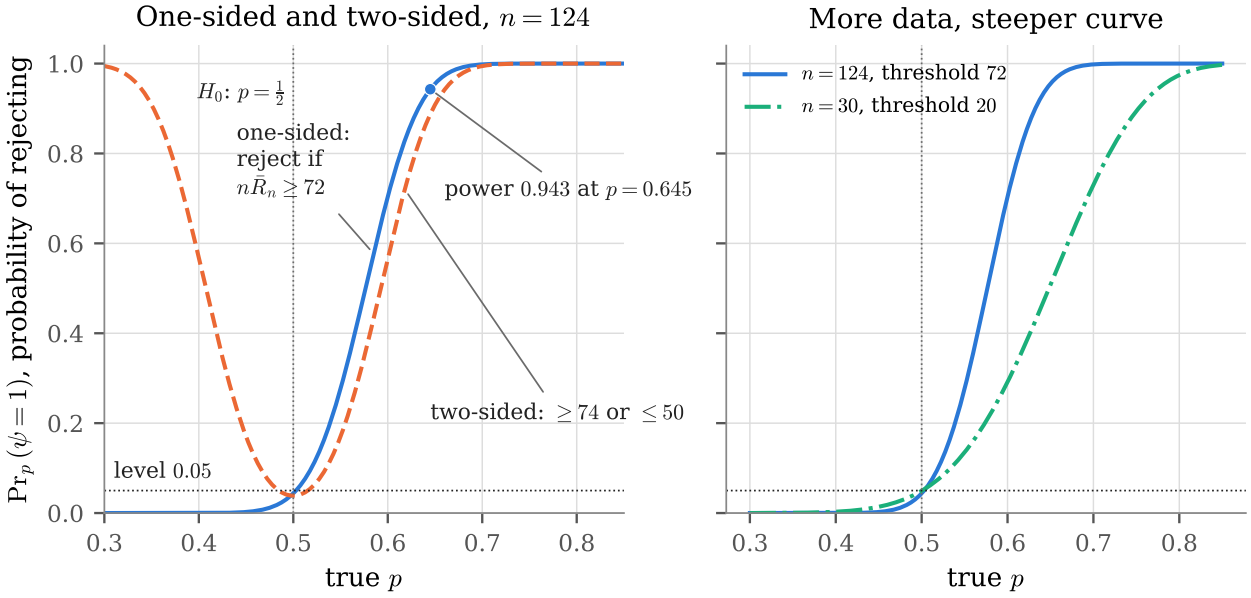


Figure 9.3: The probability of rejecting as a function of the true p, computed exactly from the Binomial distribution of the count. Left: the one-sided test (reject if the count is ≥ 72) and the two-sided test (reject if ≥ 74 or ≤ 50) at $n = 124$. At $p = \frac{1}{2}$ both curves are at or below 0.05 (they are 0.044 and 0.038); to the right both climb to 1, the one-sided test faster since it spends its whole 5% on that side; to the left the two-sided test climbs and the one-sided test, which does not watch that side, stays at 0. Right: the one-sided test at $n = 30$ (threshold 20) and $n = 124$. More data makes the curve steeper: both errors fall at once.

From undergraduate statistics: α, β and the power calculation

The introductory course called the type 1 error probability α and the type 2 error probability β, said 'power $= 1 - \beta$', and may have made you compute the sample size needed for 80% power against a specific alternative. All of that is here. The only addition is that α and β are functions of θ, not single numbers: a type 2 error of 0.057 is a type 2 error *at* $p = 0.645$, and at $p = 0.55$ the same test has a type 2 error of 0.72, which is why failing to reject is not accepting: this test misses a true p of 0.55 nearly three times in four. The power calculation of an introductory course picks one alternative and reads one point off Figure 9.3's curve; the power π_ψ of Definition 9.2 takes the infimum over all of them, which for a composite alternative that touches Θ_0 is the rejection probability at the edge of Θ_0, 0.044 here: at most the level, and useless as a summary of what the test can detect. The one-point version is the useful one in planning a study, and the infimum is the one that describes the test over the whole alternative.

Definition 9.3 (Level, asymptotic level)**.** Let $\alpha \in (0, 1)$. A test ψ has **level** α if

$$\alpha_\psi(\theta) \leq \alpha \quad \text{for all } \theta \in \Theta_0,$$

and **asymptotic level** α if

$$\lim_{n\to\infty} \alpha_{\psi_n}(\theta) \leq \alpha \quad \text{for all } \theta \in \Theta_0,$$

where ψ_n is the test built on n observations (toolkit A.1).

In words. A test of level 5% is one that convicts the innocent at most 5% of the time, whichever innocent θ is the truth. The level is a cap on the worst case of the type 1 error over the whole null set, not the type 1 error at one point. 'Asymptotic' means the cap is guaranteed only in the limit of large samples, which is what a CLT-based threshold delivers.

For the kiss data. $\psi = \mathbb{1}\{n\bar{R}_n \geq 72\}$ has level 5% for $H_0 : p = \frac{1}{2}$ exactly, since $0.044 \leq 0.05$; it also has level 5% for the composite null $H_0 : p \leq \frac{1}{2}$, since $\mathbb{P}_p(n\bar{R}_n \geq 72)$ only decreases as p falls below $\frac{1}{2}$. The threshold 71 would not do: $\mathbb{P}_{1/2}(n\bar{R}_n \geq 71) = 0.063$.

Two remarks on the definition, both consequences of the asymmetry of Section 9.2. First, the level is a guarantee about the type 1 error only. A test of level 5% promises nothing about its type 2 error; $\psi \equiv 0$, the test that never rejects, has level 5% (its type 1 error is 0) and is useless. The standard way of choosing a test is to fix the level first and then, among tests of that level, prefer the one with the most power. Second, the level is a cap, so a test of level 5% is also a test of level 10%: 'level 5%' is a one-directional bound, as 'level 95%' was for an interval in Chapter 8, the worst the procedure can do rather than what it does.

Pitfall: level α for a test, level $1 - \alpha$ for an interval

The same word, the same α, complementary conventions. For a confidence interval the level is $1 - \alpha$, the probability of a correct capture. For a test the level is α, the probability of an incorrect rejection. Both conventions describe the procedure by the guarantee it gives, and the guarantee for a test is about the error it controls. Section 9.8 shows the two procedures are one object, which is why the same α appears in both.

Intuition: why the worst case over Θ_0

For a one-sided question such as 'can we conclude that at most a third of Americans get news on YouTube?' you would write $H_0 : p \geq 0.33$ against $H_1 : p < 0.33$, a composite null. Which p in Θ_0 should the threshold be computed from? The definition answers: the one that makes rejection most likely, since the level must hold for every p in Θ_0. The test rejects when $\hat{p}$ is small, and a small $\hat{p}$ is likeliest when p is as small as Θ_0 allows, that is, at the boundary $p = 0.33$. So the composite null $p \geq 0.33$ is tested with the threshold of the simple null $p = 0.33$, and no separate computation is needed. The same reasoning gives the threshold 72 of the kiss data for $H_0 : p \leq \frac{1}{2}$. This is why one-sided tests are freely written either way.

9.5 Building a test from the Central Limit Theorem

The general-purpose test is the standardized mean, and, as in the plug-in interval of Chapter 8, all three limit tools of Part II work together in its one argument.

Let $X_1, \ldots, X_n$ be i.i.d. with mean μ and variance σ^2, both finite, and let μ_0 be a fixed number. To test $H_0 : \mu = \mu_0$ against $H_1 : \mu \neq \mu_0$, take

$$T_n = \sqrt{n}\,\frac{\bar{X}_n - \mu_0}{\sigma}.$$

If H_0 is true then $\mu = \mu_0$ and the CLT (Theorem 4.8) says $T_n \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$. So under H_0 the probability that $|T_n|$ exceeds $q_{\alpha/2}$ tends to $\mathbb{P}(|Z| > q_{\alpha/2}) = \alpha$, and the test

$$\psi_\alpha = \mathbb{1}\{|T_n| > q_{\alpha/2}\}$$

has asymptotic level α. The one-sided versions use one tail: for $H_1 : \mu > \mu_0$, $\psi_\alpha = \mathbb{1}\{T_n > q_\alpha\}$; for $H_1 : \mu < \mu_0$, $\psi_\alpha = \mathbb{1}\{T_n < -q_\alpha\}$. Figure 6.4 in Chapter 6 is the picture: $q_{\alpha/2}$ leaves $\alpha/2$ in each tail, q_α leaves α in one.

Proposition 9.4 (The standardized-mean test)**.** *Under the assumptions above, with σ known or replaced by any consistent estimator $\hat{\sigma}_n$:*

1. *$\psi_\alpha = \mathbb{1}\{|T_n| > q_{\alpha/2}\}$ has asymptotic level α for $H_0 : \mu = \mu_0$, and $\mathbb{1}\{T_n > q_\alpha\}$ has asymptotic level α for $H_0 : \mu \le \mu_0$.*
2. *Under any $\mu \neq \mu_0$ in the alternative, $\mathbb{P}_\mu(\psi_\alpha = 1) \to 1$: the type 2 error at every fixed alternative tends to 0.*

Proof strategy (core). Under H_0 the CLT gives the limit of the rejection probability, and it is a limit of numbers, computed from Φ. With $\hat{\sigma}_n$ in place of σ, Slutsky supplies the same limit. Under H_1 the LLN sends $\bar{X}_n - \mu_0$ to a nonzero constant, so T_n runs off to $\pm\infty$.

Proof. (1) Under $\mathbb{P}_{\mu_0}$, $T_n \xrightarrow[n\to\infty]{(\mathrm{d})} Z \sim \mathrm{Gauss}(0,1)$ by the CLT, and the set $\{|z| > q_{\alpha/2}\}$ has boundary points $\pm q_{\alpha/2}$ at which Z has no mass, so convergence in distribution gives convergence of the probabilities (Proposition 4.3):

$$\alpha_{\psi_n}(\mu_0) = \mathbb{P}_{\mu_0}(|T_n| > q_{\alpha/2}) \to \mathbb{P}(|Z| > q_{\alpha/2}) = \alpha.$$

This is a limit of a sequence of numbers, one per n, and not a mode of convergence of random variables. The random variable T_n is the thing that converges in distribution; the probability it assigns to the fixed set $\{|z| > q_{\alpha/2}\}$ is a number for each n, and those numbers converge in the ordinary calculus sense (toolkit A.1). If σ is replaced by $\hat{\sigma}_n \xrightarrow[n\to\infty]{\mathbb{P}} \sigma$, then $\sqrt{n}(\bar{X}_n - \mu_0)/\hat{\sigma}_n = \big(\sqrt{n}(\bar{X}_n - \mu_0)/\sigma\big) \cdot (\sigma/\hat{\sigma}_n)$ is a product of something converging in distribution and something converging in probability to the constant 1, and Slutsky's theorem (Theorem 5.1) gives the same limit. For the one-sided test with composite null $\mu \le \mu_0$: at $\mu = \mu_0$ the same argument gives $\mathbb{P}_{\mu_0}(T_n > q_\alpha) \to \mathbb{P}(Z > q_\alpha) = \alpha$; at $\mu < \mu_0$, write the event $\{T_n > q_\alpha\}$ as $\{(\bar{X}_n - \mu_0)/\sigma > q_\alpha/\sqrt{n}\}$: the left side converges in probability to the negative number $(\mu - \mu_0)/\sigma$ by the LLN, while the right side tends to 0, so the event's probability tends to $0 \le \alpha$ (once $q_\alpha/\sqrt{n}$ is within ϵ of 0, the event is contained in $\{|(\bar{X}_n - \mu_0)/\sigma - (\mu - \mu_0)/\sigma| > |\mu - \mu_0|/\sigma - \epsilon\}$, whose probability tends to 0).

(2) Under $\mu \neq \mu_0$, write $\{|T_n| > q_{\alpha/2}\}$ as $\{|(\bar{X}_n - \mu_0)/\sigma| > q_{\alpha/2}/\sqrt{n}\}$: the left side converges in probability to the positive number $|\mu - \mu_0|/\sigma$ while the right side tends to 0, so the probability tends to 1 by the same containment argument. For the one-sided test, $\mu > \mu_0$ sends $(\bar{X}_n - \mu_0)/\sigma$ to a positive number and $\mathbb{P}_\mu(T_n > q_\alpha) \to 1$ the same way. □

Part (2) is the reason for liking this test: against any fixed departure from H_0, enough data will detect it. A test with this property is called **consistent**, by analogy with a consistent estimator. Part (1) is the promise Section 5.5 made: the level of a test built on a standardized average rests on the CLT, and on Slutsky as soon as the scale is estimated.

Intuition: under the null the parameter becomes known

The reason tests are easier than intervals is that under the null hypothesis the value of the unknown parameter becomes known, so there is nothing to plug in. A confidence interval for p had to deal with the unknown p in the standard error $\sqrt{p(1-p)/n}$, which is what the

three solutions of Section 8.4 were for. A test of $H_0 : p = \frac{1}{2}$ computes everything *as if* $p = \frac{1}{2}$ because the level is a statement about what happens when H_0 is true. So the standard error is $\sqrt{\frac{1}{2} \cdot \frac{1}{2}/n} = 1/(2\sqrt{n})$, a number, and the test statistic is

$$T_n = \sqrt{n}\,\frac{\bar{R}_n - \frac{1}{2}}{\sqrt{\frac{1}{2} \cdot \frac{1}{2}}} = 2\sqrt{n}\left(\bar{R}_n - \tfrac{1}{2}\right),$$

with nothing unknown in it. The alternative only enters through the power, which is computed separately and does not affect the threshold.

Worked example: three tests of a proportion

Each is the two-sided or one-sided test above with σ the null standard deviation $\sqrt{p_0(1-p_0)}$.

- **The kiss experiment**, $H_0 : p = \frac{1}{2}$ against $H_1 : p \neq \frac{1}{2}$. $T_n = \sqrt{124}\,(0.645 - 0.5)/0.5 = 3.23$, and $3.23 > 1.96 = q_{0.025}$: reject at asymptotic level 5%. The one-sided version compares with $q_{0.05} = 1.645$ and rejects too.
- **The fair coin**, 13 heads in 30 tosses, $H_0 : p = \frac{1}{2}$ against $H_1 : p \neq \frac{1}{2}$. $T_n = \sqrt{30}\,(13/30 - 0.5)/0.5 = -0.73$, and $0.73 < 1.96$: do not reject. A value of -0.73 is a perfectly plausible draw from a standard Gaussian: about 47% of draws fall further from zero than that.
- **News on YouTube**, $\hat{p} = 0.32$ of $n = 4{,}971$, $H_0 : p \geq 0.33$ against $H_1 : p < 0.33$. A one-sided test with the threshold computed at the boundary $p = 0.33$: $T_n = \sqrt{4971}\,(0.32 - 0.33)/\sqrt{0.33 \times 0.67} = -1.50$, and $-1.50 > -1.645 = -q_{0.05}$: do not reject at level 5%. The data lean the alternative's way but not far enough, which is the same thing Figure 9.1 showed with the buffer just reaching 0.33.

Note that the three conclusions came from the same kind of comparison, T_n against a Gaussian quantile; what differed was the quantile, $q_{0.025}$ for a two-sided question and $q_{0.05}$ for a one-sided one, and, for the one-sided test, that only one sign counts. It is easier to reject one-sided than two-sided because a one-sided test puts its whole α in the tail it watches.

Example 9.5 (A Poisson rate)**.** Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$, so $\mathbb{E}[X_i] = \mathbb{V}\text{ar}(X_i) = \lambda$, and let $\hat{\lambda} = \bar{X}_n$. The CLT gives $\sqrt{n}(\hat{\lambda} - \lambda)/\sqrt{\lambda} \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1)$. To test $H_0 : \lambda = 2$ against $H_1 : \lambda \neq 2$ at asymptotic level 5%, use $T_n = \sqrt{n}(\hat{\lambda} - 2)/\sqrt{2}$, with the null value 2 in the variance as well as the mean, and reject when $|T_n| > 1.96$. To test $H_0 : \lambda \geq 2$ against $H_1 : \lambda < 2$, reject when $T_n < -1.645$; at any $\lambda > 2$ the LLN sends T_n to $+\infty$ and the rejection probability to 0, so the level is set at the boundary $\lambda = 2$ as in the intuition box above. A fourth pair of hypotheses, $H_0 : |\lambda - 2| \leq 1$ against $H_1 : |\lambda - 2| > 1$, has a null set that is an interval, $[1, 3]$; the worst case over it is at the two endpoints, 1 and 3, the test has to be sized at both, and Exercise 9.5 works it out.

9.6 The kiss experiment, exactly and by the CLT

Chapter 1 quoted two thresholds, 72 for the one-sided question and 74 for the two-sided one, and said both come from the exact Binomial distribution of the count. The CLT test of the previous

section gives 72 and 73. Here are both computations and the reason for the gap.

Exactly. Under $H_0 : p = \frac{1}{2}$ the count $n\bar{R}_n$ is exactly Bin$(124, \frac{1}{2})$, with no approximation. A one-sided test of level 5% rejects for counts $\geq k$ with k the smallest integer such that $\mathbb{P}_{1/2}(n\bar{R}_n \geq k) \leq 0.05$. The tail probabilities are 0.063 at $k = 71$, 0.044 at $k = 72$, so $k = 72$. A two-sided test of level 5% rejects for counts $\geq k$ or $\leq 124 - k$, and by symmetry needs $\mathbb{P}_{1/2}(n\bar{R}_n \geq k) \leq 0.025$: the tail is 0.029 at $k = 73$ and 0.019 at $k = 74$, so $k = 74$, with the mirror threshold 50. These are the **exact Binomial tests** that Section 4.4 promised: the finite-sample alternatives to the CLT that exist whenever the parent distribution is known. Their type 1 errors are 0.044 and 0.038, not 0.05: with a discrete statistic the level can only be met from below, since the tail probability jumps at each integer and 0.05 is not one of the values it takes.

By the CLT. The one-sided test rejects when $2\sqrt{n}(\bar{R}_n - \frac{1}{2}) > 1.645$, that is, when $n\bar{R}_n > 62 + 1.645 \times \sqrt{31} = 71.2$, that is, for counts ≥ 72. The two-sided test rejects when $|n\bar{R}_n - 62| > 1.96 \times \sqrt{31} = 10.9$, that is, for counts ≥ 73 or ≤ 51. So the one-sided thresholds agree and the two-sided ones differ by one count.

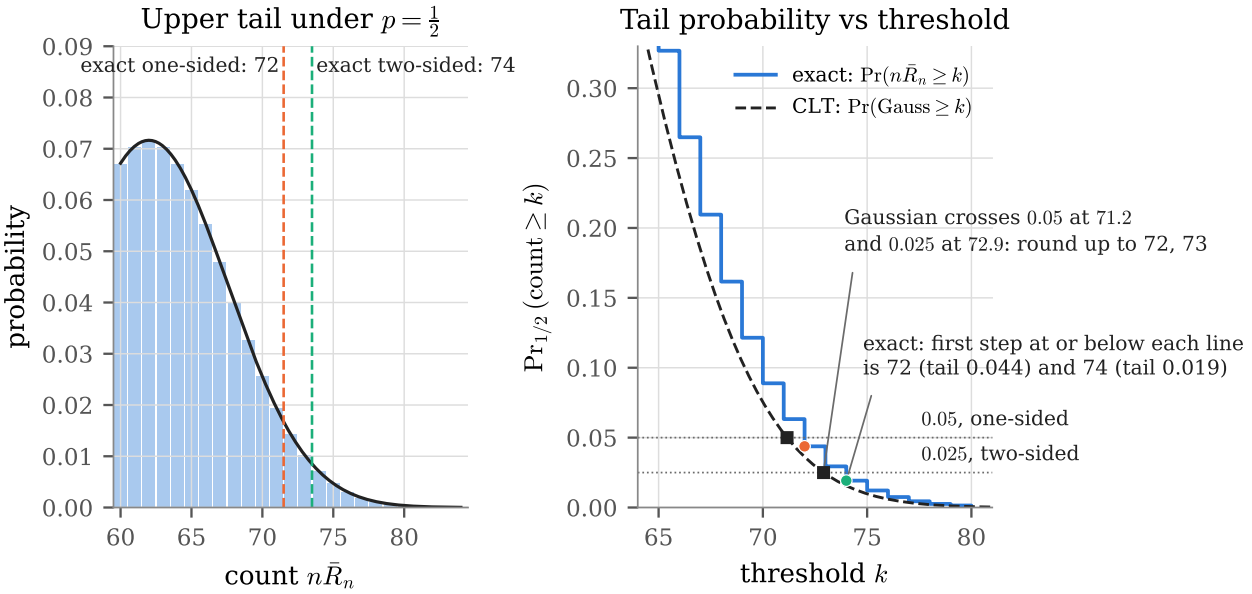


Figure 9.4: Why the exact and CLT thresholds can differ. Left: the upper tail of Bin$(124, \frac{1}{2})$ (bars) and the Gauss$(62, 31)$ density the CLT replaces it with (curve), with the two exact thresholds marked. Right: the tail probability $\mathbb{P}_{1/2}(\text{count} \geq k)$ as a function of the threshold k, a staircase for the exact distribution and a smooth curve for the Gaussian. The exact threshold is the first step at or below the level line (72 for 0.05, 74 for 0.025); the CLT threshold is where the curve crosses it (71.2 and 72.9, rounded up to 72 and 73). At $k = 73$ the exact tail is 0.029 and the Gaussian tail 0.024: one is above the 0.025 line and the other below, and that is the whole disagreement.

Why they differ. Figure 9.4 shows the two tail probabilities side by side. The Gaussian curve is a good approximation of the staircase, off by about 0.005 at $k = 73$, but 0.005 is enough to put the exact tail (0.029) above the line 0.025 and the Gaussian tail (0.024) below it. The gap is mostly discreteness, not the CLT's inaccuracy: the Gaussian approximates the probability of the count being ≥ 73 better if the threshold is placed at 72.5 (the **continuity correction** of an introductory course), and with it the two-sided threshold becomes $62 + 0.5 + 10.9 = 73.4$, rounding up to 74, in agreement with the exact test. This book does not use the correction. For a proportion, whenever

the Gaussian approximation lands close enough to the line that the correction would change the decision, the exact Binomial test is available and settles the question outright. An asymptotic level 5% is not an exact level 5%, and at a threshold where the exact tail is near the line the two can round in different directions. For 80 right turns, far from either threshold, the decision is the same every way: reject.

From undergraduate statistics: the binomial test and the z-test for a proportion

An introductory course presents these as two tests: the 'exact binomial test' when n is small, the 'z-test for a proportion' when np_0 and $n(1-p_0)$ are both at least 10. They test the same hypothesis with the same statistic, the count, and differ only in whether the count's distribution under H_0 is used as it is or replaced by its Gaussian limit. The rule of thumb is a statement about when Figure 9.4's staircase and curve are close enough that the thresholds agree.

9.7 p-values

A test as defined reports one bit, reject or not, at one level fixed in advance. That loses information: a T_n of 1.97 and a T_n of 3.23 both reject at 5%, and the second is much stronger evidence. It also forces the reader to accept the writer's α. The p-value fixes both.

Definition 9.6 (p-value)**.** Let $(\psi_\alpha)_{\alpha\in(0,1)}$ be a family of tests of H_0 indexed by their level, nested so that rejecting at level α implies rejecting at every level above α. The **(asymptotic) p-value** of the family on a given sample is the smallest (asymptotic) level α at which ψ_α rejects H_0 on that sample.[3] It is random: it depends on the sample.

> **In words.** Run the test at every level at once. At large α it rejects easily, at small α it is hard to satisfy; the p-value is the level where it switches. Equivalently, for the standardized-mean tests of this chapter, it is the probability under H_0 of a test statistic at least as extreme as the one observed: the tail area beyond the observed value.
>
> **For the kiss data.** $T_n = 3.23$. Two-sided, the p-value is $\mathbb{P}(|Z| > 3.23) = 2(1-\Phi(3.23)) = 0.0012$; one-sided, $\mathbb{P}(Z > 3.23) = 0.0006$. The exact Binomial tails give 0.0016 and 0.0008. Under $p = \frac{1}{2}$, 80 or more right turns out of 124 is a one-in-a-thousand event.

The two readings in Figure 9.5 are the same number because of how the thresholds are built. For the two-sided test, ψ_α rejects when $|T_n| > q_{\alpha/2}$, and $q_{\alpha/2}$ decreases as α grows, so the levels at which the test rejects are exactly those with $q_{\alpha/2} < |t|$ for the observed t, that is, those with $\alpha/2 > \mathbb{P}(Z > |t|)$, that is, $\alpha > \mathbb{P}(|Z| > |t|)$. The smallest such α is $\mathbb{P}(|Z| > |t|)$, the tail area. For a one-sided test with $H_1 : \mu > \mu_0$ the same argument gives $\mathbb{P}(Z > t)$. Since $\mathbb{P}(|Z| > |t|) = 2\mathbb{P}(Z > |t|)$, a two-sided p-value is twice the one-sided one when the data lean the alternative's way; when they lean the other way the one-sided p-value is above $\frac{1}{2}$ and the one-sided test cannot reject. That is the one-sided test's advantage of Section 9.5 seen from the other end: it spends its whole α on one side and is blind to the other.

The golden rule. The consequence of the definition is the rule worth memorizing:

$$p\text{-value} \le \alpha \quad \Longleftrightarrow \quad H_0 \text{ is rejected by } \psi_\alpha \text{ at (asymptotic) level } \alpha.$$

[3] Strictly an infimum: at the boundary level the threshold equals the observed statistic and the strict inequality in ψ_α fails. The boundary case has probability zero, and the golden rule's $\le$ settles it by convention.

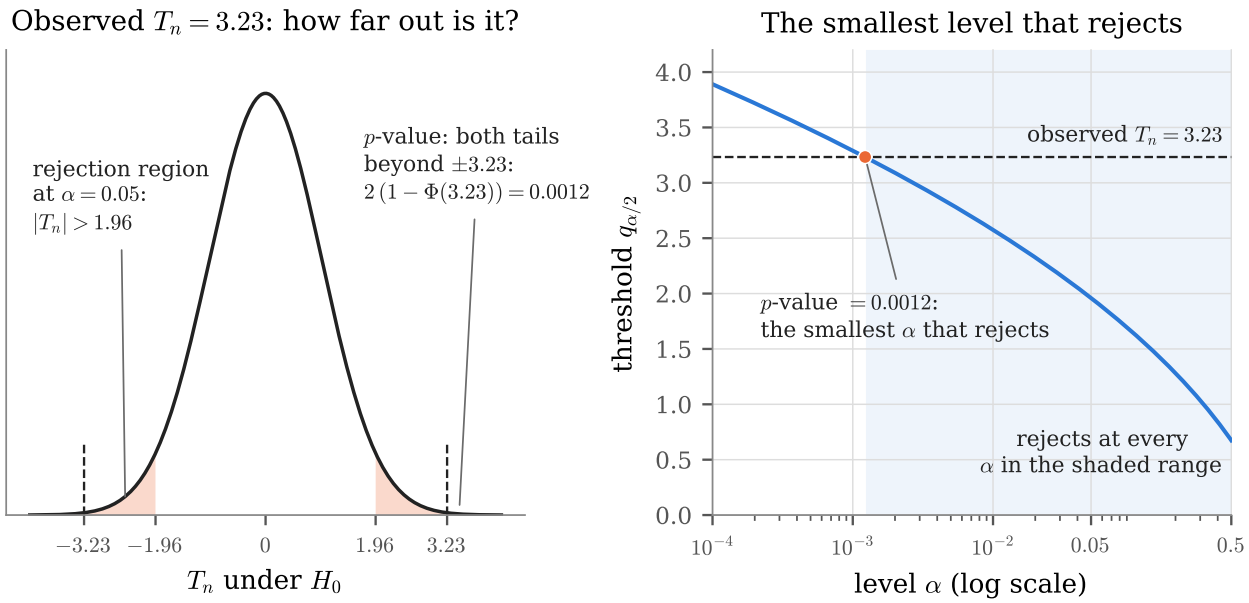


Figure 9.5: The p-value two ways, for $T_n = 3.23$. Left: the density of T_n under H_0 (standard Gaussian by the CLT), the 5% rejection region $|T_n| > 1.96$ shaded lightly, and the observed value marked. The p-value is the area beyond ± 3.23, too small to see at this scale, the probability under H_0 of something at least this extreme: 0.0012. Right: the threshold $q_{\alpha/2}$ as a function of the level, with the observed value as a horizontal line. The test rejects at every level whose threshold lies below the line; the smallest such level, where the curve crosses, is the p-value. Left and right say the same thing: 3.23 is far out.

So one number replaces the whole family of tests. A reader who wants level 5% compares the p-value with 0.05; one who wants 1% compares with 0.01; and the smaller the p-value, the more levels reject, the more confidently H_0 is rejected.

From undergraduate statistics: $p < 0.05$

'Reject if $p < 0.05$' was the intro course's rule, and the golden rule is its justification: the p-value is the smallest level that rejects, so a p-value below 0.05 means the level-5% test rejects, no more and no less. Two habits from that course transfer intact. The p-value is computed from the null distribution of the statistic, which is why it needed μ_0 (or p_0) and not any guess at the alternative. And a one-sided p-value is half the two-sided one for a symmetric statistic when the data lean the alternative's way, so the choice of sides made in Section 9.2 before seeing the data is the choice that fixes the p-value.

Pitfall: what a p-value is not

A p-value of 0.0012 is the probability, computed as if H_0 were true, of data at least as extreme as these. It is not the probability that H_0 is true; in this part of the book H_0 is true or false and has no probability (Chapter 18 is where hypotheses get probabilities, and the number it produces is a different one). It is not the probability that the result 'was due to chance'. And it is not a measure of the size of the effect: with n large enough, a true p of 0.51 produces a tiny p-value (at $n = 100{,}000$, $T_n = 2\sqrt{n} \times 0.01 = 6.3$), since by Proposition 9.4 the test detects any fixed departure eventually. The p-value answers one question, how surprising the data would be

under H_0, and the confidence interval of Chapter 8 answers the other, how big the departure is. Report both.

Intuition: the p-value is a statistic, and under H_0 it is uniform

The p-value is a function of the sample, so it is a random variable, and it has a distribution under each θ. Under a simple H_0, for a test statistic with a continuous distribution, that distribution is uniform on $[0, 1]$: a p-value below 0.05 happens 5% of the time, below 0.01 1% of the time, which is just the level restated (Exercise 9.6 proves it). Under H_1 the distribution piles up near 0, and the more so the larger n. So a small p-value is not proof; it is an event that is rare under H_0 and common under H_1.

9.8 Tests and confidence intervals are one construction

Exercise 8.6 in Chapter 8 found that '$\frac{1}{2}$ is below the conservative lower confidence bound' is the same event as '$n\bar{R}_n \geq 72$', the one-sided test's rejection. That was not a coincidence of the kiss data.

Proposition 9.7 (Duality of tests and confidence intervals)**.** *Let $\Theta \subseteq \mathbb{R}$.*

1. *Suppose that for every $\theta_0 \in \Theta$ there is a test $\psi(\theta_0)$ of $H_0 : \theta = \theta_0$ with level α. Then the set of values not rejected,*
$$\mathcal{I} = \{\theta_0 \in \Theta : \psi(\theta_0) = 0\},$$
is a confidence set of level $1 - \alpha$ for θ.
2. *Conversely, if $\mathcal{I}$ is a confidence interval of level $1 - \alpha$ for θ, then for every θ_0 the test $\psi(\theta_0) = \mathbb{1}\{\theta_0 \notin \mathcal{I}\}$ has level α for $H_0 : \theta = \theta_0$.*

The same holds with 'asymptotic' in front of both levels.

Proof strategy (core). Both directions are one line: the event 'the true θ is not in $\mathcal{I}$' is the event 'the test of the true θ rejects', whose probability is the type 1 error.

Proof. (1) For any true θ, $\mathbb{P}_\theta(\mathcal{I} \not\ni \theta) = \mathbb{P}_\theta(\psi(\theta) = 1) = \alpha_{\psi(\theta)}(\theta) \leq \alpha$, since $\psi(\theta)$ has level α for the null $\theta = \theta$, which is true. So $\mathbb{P}_\theta(\mathcal{I} \ni \theta) \geq 1 - \alpha$ for every θ: the coverage condition of Definition 8.1, applied to a set rather than an interval. (2) Under $H_0 : \theta = \theta_0$ the truth is θ_0, and $\mathbb{P}_{\theta_0}(\psi(\theta_0) = 1) = \mathbb{P}_{\theta_0}(\mathcal{I} \not\ni \theta_0) \leq \alpha$ by the coverage guarantee. The asymptotic version replaces each probability by its limit. □

Figure 9.6 draws part (1) for the kiss data, running the test for every p_0 at once. The test being inverted is
$$\psi(p_0) = \mathbb{1}\left\{|\bar{R}_n - p_0| > q_{\alpha/2}\sqrt{p_0(1-p_0)/n}\right\},$$
which measures the distance from p_0 in units of the standard error computed *at* p_0, the value the null hypothesis asserts. A test built that way, from the standard error under the null rather than from an estimate, is called a **score test**. The candidates that survive are those with $|\bar{R}_n - p_0| \leq q_{\alpha/2}\sqrt{p_0(1-p_0)/n}$, which is the inequality that Section 8.4 solved to get the solve interval $[0.558, 0.724]$. So the solve interval is the set of null values the score test does not reject, and the score test of $p = \frac{1}{2}$ rejects exactly because $\frac{1}{2}$ is outside the interval. This is why Section 8.4

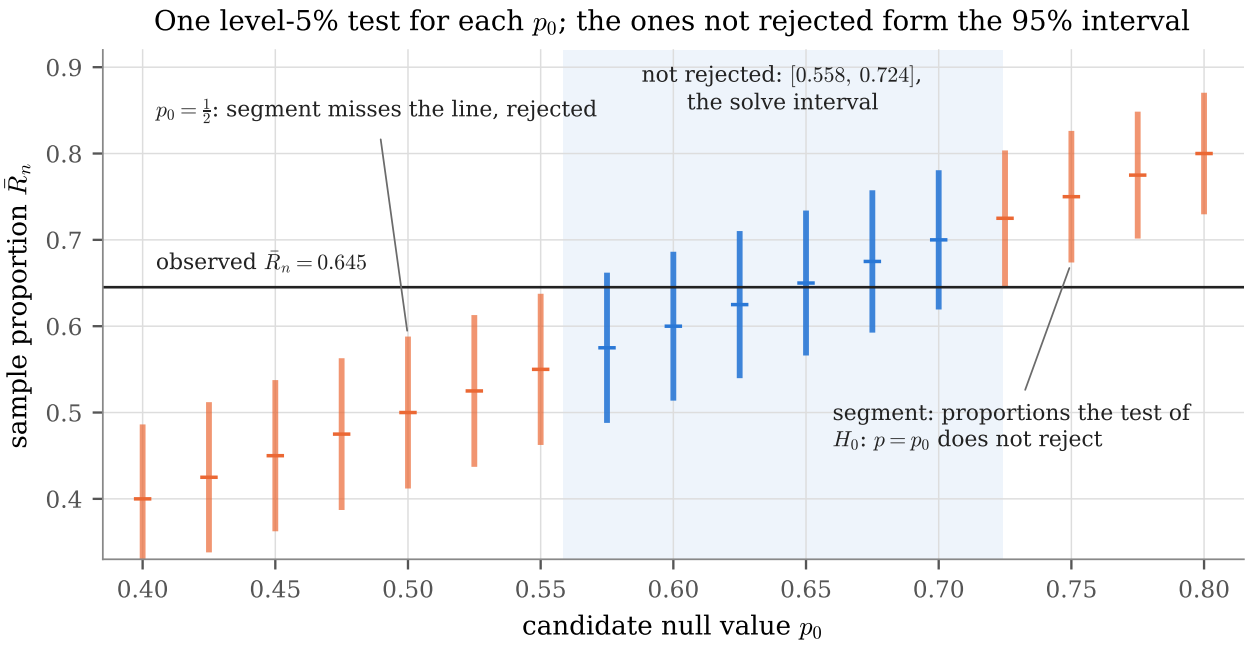


Figure 9.6: The duality for the kiss data. For each candidate p_0 on the horizontal axis, the vertical segment is the set of sample proportions the level-5% test of $H_0 : p = p_0$ does not reject, $p_0 \pm 1.96\sqrt{p_0(1-p_0)/124}$. The horizontal line is the observed $\bar{R}_n = 0.645$. A candidate whose segment crosses the line is not rejected (blue); one whose segment misses it is rejected (orange). The set of candidates not rejected (shaded) is $[0.558, 0.724]$, the solve interval of Chapter 8 exactly, and $\frac{1}{2}$ is orange.

called the result the Wilson score interval: it is the score test, inverted. Inverting the test built from the *estimated* standard error $\sqrt{\bar{R}_n(1-\bar{R}_n)/n}$ instead gives the plug-in interval $[0.561, 0.729]$. That one is **Wald's test**, which Chapter 15 treats in general. Score and Wald are the same idea evaluated at two different places, the null value and the estimate, and they are the two tests you will meet most often. And the one-sided test of Exercise 8.6 inverts to a one-sided interval, a lower confidence bound. Every interval in Chapter 8 is a test turned inside out, and every test here is an interval asked about one point.

From undergraduate statistics: if the confidence interval excludes the null value, reject

An introductory course states this as a shortcut: 'the 95% confidence interval is $[0.56, 0.73]$, which does not contain 0.5, so the difference is significant at 5%'. Proposition 9.7 is why the shortcut is legitimate, and it also says when it is exact and when it is approximate. The shortcut reproduces the test's decision exactly when the test and the interval are built from the same standard error, the solve interval for the score test and the plug-in interval for Wald's, and only approximately otherwise. Using the plug-in interval to decide a score test can disagree with it on a borderline sample, since the two standard errors differ: $\sqrt{\bar{R}_n(1-\bar{R}_n)/n} = 0.043$ against $\sqrt{\frac{1}{2}\cdot\frac{1}{2}/n} = 1/(2\sqrt{n}) = 0.045$ here. Far from the border, as with 80 picks, every version agrees.

9.9 What comes next

This chapter built one test, the standardized mean against a Gaussian quantile, and the ideas that describe any test. The remaining tests in this book are the same object with a different statistic and a different null distribution. Chapter 14 replaces the Gaussian quantile by a Student quantile when n is small and the data are Gaussian, which needs the chi-squared distribution and Cochran's theorem for the sample variance; it also finishes the boarding example as a two-sample test. Chapter 15 gives two general recipes for a test statistic in any parametric model, Wald's test (the standardized maximum likelihood estimator) and the likelihood ratio test, with chi-squared limits in place of Gaussian ones. Chapter 16 and Chapter 17 test the model itself rather than a parameter in it, with the same vocabulary: a statistic, a null distribution, a threshold, a p-value. In all of them the level is the cap on the type 1 error, the power is what the test can detect, and the p-value is what gets reported.

9.10 Summary

- Hypotheses: $H_0 : \theta \in \Theta_0$ against $H_1 : \theta \in \Theta_1$, disjoint. The data are used only to try to reject H_0; failing to reject is not accepting. Simple ($\{\theta_0\}$) or composite; one-sided or two-sided, and the side is chosen before the data are seen.
- A test is $\psi : E^n \to \{0, 1\}$, the indicator of its rejection region; in practice $\psi = \mathbb{1}\{T_n > c\}$ for a test statistic T_n and a threshold c.
- Type 1 error $\alpha_\psi(\theta) = \mathbb{P}_\theta(\psi = 1)$ on Θ_0; type 2 error $\beta_\psi(\theta) = \mathbb{P}_\theta(\psi = 0)$ on Θ_1; power $\pi_\psi = \inf_{\Theta_1}(1 - \beta_\psi)$. At fixed n the two errors trade off; only n lowers both.
- Level α: $\alpha_\psi \leq \alpha$ on all of Θ_0, a cap on the worst case; asymptotic level: the same in the limit. For a composite null the worst case is at the boundary, so $H_0 : \theta \leq \theta_0$ and $H_0 : \theta = \theta_0$ get the same test. For a test the level is α; for an interval it is $1 - \alpha$.
- The standardized-mean test: $T_n = \sqrt{n}(\bar{X}_n - \mu_0)/\sigma$, reject if $|T_n| > q_{\alpha/2}$ (two-sided) or $T_n > q_\alpha$ (one-sided). Asymptotic level α by the CLT, plus Slutsky if σ is estimated; consistent (power $\to 1$ at every fixed alternative) by the LLN. Under H_0 the parameter is known, so no plug-in is needed.
- Kiss experiment: $T_n = 3.23$, reject at 5% either way. Exact Binomial thresholds 72 (one-sided) and 74 (two-sided); CLT thresholds 72 and 73. They differ because the exact tail is a staircase and 0.025 falls between two of its steps.
- p-value: the smallest level at which the data reject; for these tests, the tail area beyond the observed statistic under H_0. Golden rule: p-value $\leq \alpha$ if and only if ψ_α rejects. Kiss experiment: 0.0012 two-sided, 0.0006 one-sided. Not the probability that H_0 is true.
- Duality: the null values a level-α test does not reject form a level-$(1 - \alpha)$ confidence set, and a level-$(1 - \alpha)$ interval rejects every null value outside it at level α. The solve interval inverts the score test, which uses the standard error at the null value; the plug-in interval inverts Wald's test, which uses the estimated standard error.

9.11 Exercises

Exercise 9.1. For 80 of 124 couples turning right, formulate the two-sided hypothesis problem, compute the p-value by the CLT and exactly, and conclude. Then do the one-sided version. *(Solution*

on page 226.)

Exercise 9.2. (From OpenIntro Statistics.) Students count the chocolate chips in 32 cookies and find a mean of 14.77 with a standard deviation of 4.37. The packaging claims at least 20 chips per cookie. One student says the claim is unreasonable; another says the difference might be chance. Formulate the hypotheses, compute a p-value and decide, saying what approximation you used. *(Solution on page 226.)*

Exercise 9.3. A machine learning procedure claims to reduce the rate of surgery on benign tumors below the state of the art, 95%, while keeping the true positive rate. On 297 benign tumors it recommended surgery for 206. Let p be the proportion of benign tumors on which it prescribes surgery. Formulate the hypotheses, compute the p-value and conclude. *(Solution on page 226.)*

Exercise 9.4. For the fair coin (13 heads in 30 tosses), compute the two-sided p-value by the CLT and exactly, and comment on the gap. *(Solution on page 226.)*

Exercise 9.5. For $X_1, \dots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$, build a test of asymptotic level 5% for $H_0 : |\lambda - 2| \leq 1$ against $H_1 : |\lambda - 2| > 1$. *(Solution on page 226.)*

Exercise 9.6. Let T be a test statistic with a continuous, strictly increasing CDF F_0 under H_0, and let the test reject for large T, so the p-value is $P = 1 - F_0(T)$. Show that under H_0, P is uniform on $[0, 1]$, and deduce that $\mathbb{P}_{H_0}(P \leq \alpha) = \alpha$ for every α: the golden rule gives exact level α. *(Solution on page 227.)*

Exercise 9.7. Show that the set of $p_0 \in (0, 1)$ not rejected by the test $\psi(p_0) = \mathbb{1}\{|\bar{R}_n - p_0| > q_{\alpha/2}\sqrt{p_0(1-p_0)/n}\}$ is exactly the solve interval of Section 8.4, and that inverting the test with the plug-in standard error gives the plug-in interval. *(Solution on page 227.)*

Part IV

Methods of Estimation

Part III judged estimators; this part produces them. The distances between distributions of Chapter 10 say what 'close to the truth' means; maximum likelihood (Chapters 11 and 12) is the main method, with its consistency and asymptotic normality through the Fisher information; the method of moments and M-estimation (Chapter 13) are the two generalizations.

Chapter 10

Distances Between Distributions: Total Variation and Kullback–Leibler

In this chapter

- Why 'the estimate is close to the truth' should be measured between distributions and not between parameters, and an example where a parameter difference of 0.01 means almost nothing in one place and a great deal in another.
- Total variation: the largest disagreement between two distributions over all events, equal to the area between their densities. What it says about tests: the best possible test between two distributions has $\alpha + \beta = 1 - \mathrm{TV}$.
- Why total variation, for all its good properties, can't be estimated by an average of the data.
- The Kullback–Leibler divergence: the average evidence one observation supplies against the wrong distribution. Its properties, its asymmetry, why it adds up over independent observations, and why near the truth it is a parabola.
- Estimating the KL divergence by the Law of Large Numbers, and the discovery that minimizing the estimate is maximizing the likelihood.
- What an estimator estimates when the model is wrong: the member of the family closest to the truth in KL. For the Gaussian family that is the member with the right mean and variance.

Part III judged estimators that were handed to you. This chapter is the first step toward producing them: it says what 'close' means, and the answer leads straight to maximum likelihood in Chapter 11.

10.1 What 'close to the truth' should mean

Go back to the setup of Chapter 2. You have a statistical model $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$, data $X_1, \dots, X_n \overset{iid}{\sim} \mathbb{P}_{\theta^*}$ from a true parameter θ^* that you don't know, and an estimator $\hat{\theta}_n$. What would it mean for $\hat{\theta}_n$ to be a good estimator?

The answer of Chapter 7 was that $\hat{\theta}_n$ should be close to θ^* as numbers: small bias, small variance, small quadratic risk $\mathbb{E}[(\hat{\theta}_n - \theta^*)^2]$. That is a fine answer when the parameter is the thing you care about, like the share of couples who turn right. But a parameter is a label for a distribution. What

you actually use a fitted model for is its distribution: the probability it assigns to an interval, a tail, a range of outcomes. So the question you'd really like answered is whether the fitted distribution $\mathbb{P}_{\hat{\theta}_n}$ is close to the true one $\mathbb{P}_{\theta^*}$, in the sense that for every event A,

$$\mathbb{P}_{\hat{\theta}_n}(A) \approx \mathbb{P}_{\theta^*}(A).$$

If that holds, every probability you compute from the fitted model is about right, and that is the whole point of fitting it.

Intuition: a parameter difference has no fixed meaning

Is a difference of 0.01 between two parameters small? It depends on where in the parameter space it happens. Two Gaussians with standard deviation 1 whose means differ by 0.01 are, for every practical purpose, the same distribution: no event has probabilities under the two that differ by more than 0.004. Two Gaussians with mean 0 whose standard deviations are 0.01 and 0.02 also have parameters 0.01 apart, and they are obviously different distributions: one is twice as wide as the other, and there is an event whose probability is 0.32 higher under one than under the other. Figure 10.1 draws both.

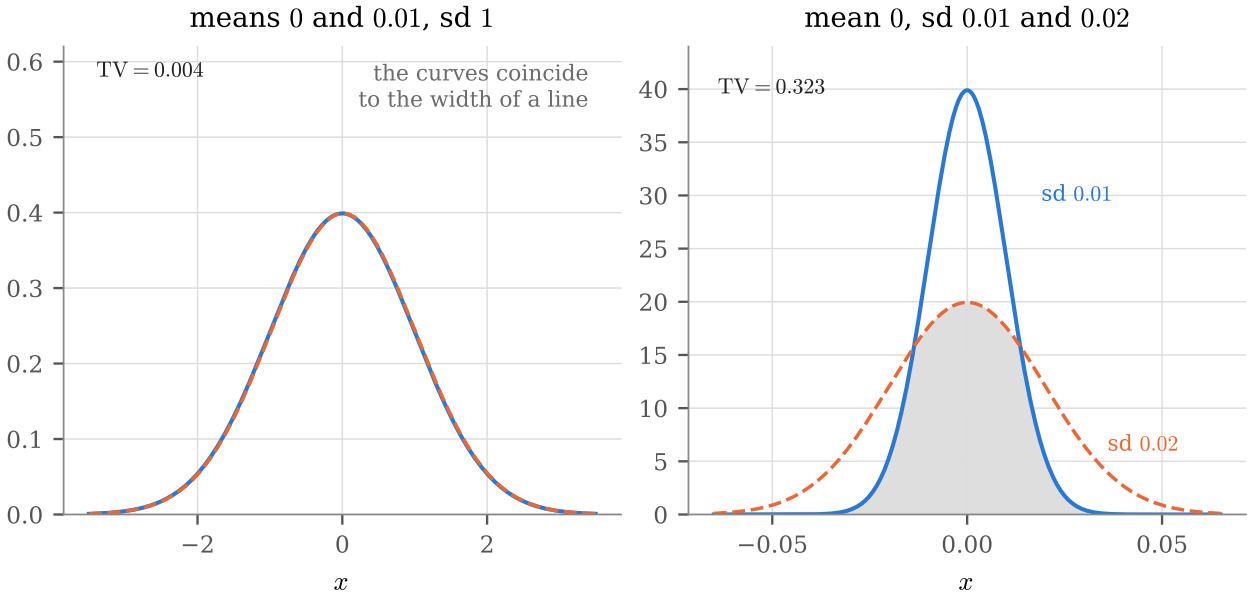


Figure 10.1: Two pairs of Gaussians whose parameters differ by 0.01. Left: means 0 and 0.01, standard deviation 1; the curves coincide to the width of a line and the total variation distance is 0.004. Right: mean 0, standard deviations 0.01 and 0.02; the curves share only the gray area, 0.677, so the total variation distance is 0.323.

There is a second reason to measure closeness between distributions, and it matters as much as the first. Section 2.4 asked whether the model is well-specified, whether the true distribution is in the family at all. If it isn't, there is no θ^*, and 'the distance from $\hat{\theta}_n$ to θ^*' has no meaning. A distance between distributions still has one: you can ask which member of the family is closest to the truth, even when none of them is the truth. Section 10.7 answers that question.

The chapter builds two such measures. The first, total variation, is the natural one, and it is a true distance. The second, the Kullback–Leibler divergence, is not a true distance, and it is the one statistics uses, for a reason that takes Section 10.4 to explain.

10.2 Total variation

The requirement '$\mathbb{P}(A) \approx \mathbb{Q}(A)$ for every event A' already contains a definition. Look at all the events, find the one on which the two distributions disagree most, and call that disagreement the distance.

Definition 10.1 (Total variation distance)**.** Let $\mathbb{P}$ and $\mathbb{Q}$ be two probability distributions on the same sample space E. The **total variation distance** between them is

$$\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = \max_{A \subset E} |\mathbb{P}(A) - \mathbb{Q}(A)|,$$

the largest difference between the probabilities they assign to one event.

> **In words.** Run through every question you could ask about one draw (is it positive? is it between 2 and 3? is it even?) and record how differently the two distributions answer. The total variation distance is the worst case. If it is 0.01, no question you could ever ask gets answers more than one percentage point apart. That is exactly the guarantee you want from a fitted model: every probability it reports is within TV of the true one.
>
> **For the kiss data.** One couple either turns right or doesn't, so there are four events: nothing, 'right', 'left' and everything. Compare Bern(0.645), the fitted distribution, with Bern($\frac{1}{2}$), 'no preference'. The empty event and the whole space get probabilities 0 and 1 from both. 'Right' gets 0.645 and 0.5; 'left' gets 0.355 and 0.5. The largest difference is 0.145, so TV = 0.145. As a natural frequency: of 1,000 couples, the fitted model predicts 645 right turns and 'no preference' predicts 500; the two descriptions disagree about 145 couples in a thousand.

For a sample space with four events the maximum is easy. For a sample space like $\mathbb{R}$ there are uncountably many events, and searching over them is not a calculation anyone wants to do. The next result turns the maximum into a sum or an integral.

Proposition 10.2 (Total variation as an area)**.** *If $\mathbb{P}$ and $\mathbb{Q}$ have PMFs p and q on a discrete E,*

$$\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = \frac{1}{2} \sum_{x \in E} |p(x) - q(x)|;$$

if they have densities p and q,

$$\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = \frac{1}{2} \int_E |p(x) - q(x)| \, \mathrm{d}x.$$

The maximum in Definition 10.1 is attained at the event $A^\star = \{x : p(x) > q(x)\}$, where $\mathbb{P}$ puts more mass than $\mathbb{Q}$.

Proof strategy (core). The event that favors $\mathbb{P}$ most collects every point where p beats q and no other point. Its advantage is the area of the region where p is above q; the region where q is above p has the same area, since both densities have total mass 1. Half the total area between the curves is therefore the area of either piece.

Proof. The discrete case; the continuous one replaces sums by integrals. Split E into $A^\star = \{p > q\}$ and its complement $\{p \le q\}$. For any event A,

$$\mathbb{P}(A) - \mathbb{Q}(A) = \sum_{x \in A} \big(p(x) - q(x)\big) \le \sum_{x \in A \cap A^\star} \big(p(x) - q(x)\big) \le \sum_{x \in A^\star} \big(p(x) - q(x)\big) = \mathbb{P}(A^\star) - \mathbb{Q}(A^\star).$$

The first inequality drops the points of A where $p - q \le 0$, which can only increase the sum; the second adds the points of $A^\star$ that A missed, where $p - q > 0$, which can only increase it again. The same argument with $\mathbb{P}$ and $\mathbb{Q}$ exchanged gives $\mathbb{Q}(A) - \mathbb{P}(A) \le \mathbb{Q}(A^{\star c}) - \mathbb{P}(A^{\star c})$. Both p and q sum to 1, so

$$\sum_{x \in A^\star} \big(p(x) - q(x)\big) + \sum_{x \notin A^\star} \big(p(x) - q(x)\big) = 1 - 1 = 0,$$

which says that $\mathbb{P}(A^\star) - \mathbb{Q}(A^\star) = \mathbb{Q}(A^{\star c}) - \mathbb{P}(A^{\star c})$: the area where p is above q equals the area where q is above p. Call this common value D. Every event has $|\mathbb{P}(A) - \mathbb{Q}(A)| \le D$, and $A^\star$ attains it, so $\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = D$. Finally $\sum_x |p(x) - q(x)|$ is the sum of the two pieces, $D + D = 2D$, and so $D = \frac{1}{2}\sum_x |p(x) - q(x)|$. □

Figure 10.2 is the proposition as a picture. Lay the two densities on top of each other. Where they overlap, the gray region, they agree. Where they don't, the area splits into two pieces of equal size, one where p is above and one where q is above, and the total variation distance is the area of either piece. Equivalently, $1 - \mathrm{TV}$ is the area of the overlap, the probability mass the two distributions share.

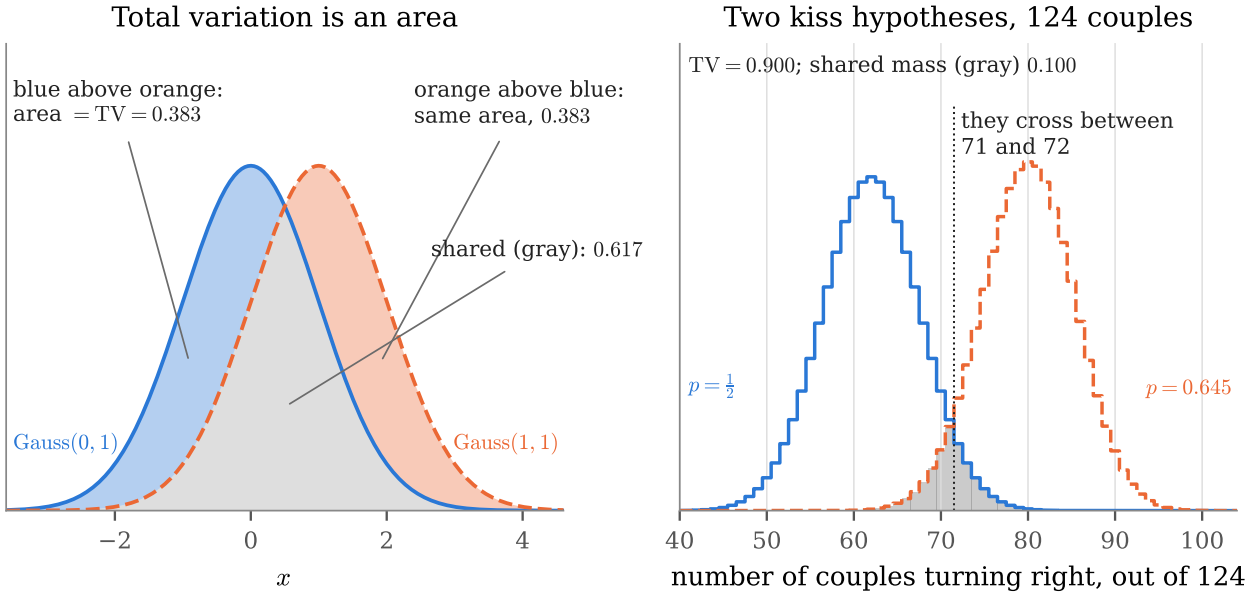


Figure 10.2: Total variation as an area. Left: Gauss$(0, 1)$ and Gauss$(1, 1)$. The blue region, where the first density is above the second, and the orange region, where the second is above the first, have the same area, 0.383, which is the total variation distance; the gray overlap has area $1 - 0.383 = 0.617$. Right: the distributions of the number of right turns among 124 couples when $p = \frac{1}{2}$ and when $p = 0.645$. They share only 0.100 of their mass (gray), so TV $= 0.900$, and the two PMFs cross between 71 and 72.

Worked example: four total variation distances

- **Two Bernoullis.** For Bern(p) and Bern(q) the PMFs differ by $|p - q|$ at 1 and by $|(1 - p) - (1 - q)| = |p - q|$ at 0, so

  $$\mathrm{TV}(\mathrm{Bern}(p), \mathrm{Bern}(q)) = \tfrac{1}{2}\big(|p - q| + |p - q|\big) = |p - q|.$$

  For the kiss data, $|0.645 - 0.5| = 0.145$, as the in-words box found by listing events. For

Bern($\frac{1}{2}$) against Bern($\frac{1}{10}$), 0.4.

- **Two Gaussians with the same variance.** The densities of Gauss$(0,1)$ and Gauss$(\mu,1)$, $\mu > 0$, cross at the midpoint $\mu/2$, and $A^\star = \{x < \mu/2\}$ is where the first is above. So $\mathrm{TV} = \mathbb{P}(Z < \mu/2) - \mathbb{P}(Z + \mu < \mu/2) = \Phi(\mu/2) - \Phi(-\mu/2) = 2\Phi(\mu/2) - 1$ for $Z \sim \mathrm{Gauss}(0,1)$. With $\mu = 1$ this is 0.383, the left of Figure 10.2; with $\mu = 0.01$ it is 0.004, the left of Figure 10.1.
- **An exponential and a uniform.** Expo(1) has density e^{-x} on $[0,\infty)$, and Unif$([0,1])$ has density 1 on $[0,1]$. The two need a common sample space, and $[0,\infty)$ contains both. On $[0,1]$, $1 \geq e^{-x}$; beyond 1, only the exponential has mass. So

$$\mathrm{TV} = \frac{1}{2}\left(\int_0^1 (1 - e^{-x})\,\mathrm{d}x + \int_1^\infty e^{-x}\,\mathrm{d}x\right) = \frac{1}{2}\left((1 - 1 + e^{-1}) + e^{-1}\right) = e^{-1} = 0.368.$$

- **Number of right turns.** Bin$(124, \frac{1}{2})$ against Bin$(124, 0.645)$: half the sum of 125 absolute differences, computed by machine, is 0.900. One couple's distributions are 0.145 apart; 124 couples' are 0.900 apart. Section 10.3 says what that growth means.

Proposition 10.3 (Total variation is a distance)**.** *For distributions $\mathbb{P}$, $\mathbb{Q}$, $\mathbb{V}$ on the same E:*

1. ***symmetry:*** $\mathrm{TV}(\mathbb{P},\mathbb{Q}) = \mathrm{TV}(\mathbb{Q},\mathbb{P})$*;*
2. ***bounds:*** $0 \leq \mathrm{TV}(\mathbb{P},\mathbb{Q}) \leq 1$*, with* $\mathrm{TV} = 1$ *exactly when there is an event that has probability* 1 *under one and* 0 *under the other;*
3. ***definiteness:*** $\mathrm{TV}(\mathbb{P},\mathbb{Q}) = 0$ *only if* $\mathbb{P} = \mathbb{Q}$*;*
4. ***triangle inequality:*** $\mathrm{TV}(\mathbb{P},\mathbb{Q}) \leq \mathrm{TV}(\mathbb{P},\mathbb{V}) + \mathrm{TV}(\mathbb{V},\mathbb{Q})$.

Proof strategy (for completeness; skip on a first reading). Each property holds for every single event, and taking the maximum over events keeps it.

Proof. (1) $|\mathbb{P}(A) - \mathbb{Q}(A)| = |\mathbb{Q}(A) - \mathbb{P}(A)|$. (2) Every $|\mathbb{P}(A) - \mathbb{Q}(A)|$ is a difference of two numbers in $[0,1]$, so it lies in $[0,1]$. It equals 1 only if one probability is 1 and the other is 0. (3) If the maximum over events is 0, every event has the same probability under both, and that is what $\mathbb{P} = \mathbb{Q}$ means. (4) For each event, $|\mathbb{P}(A) - \mathbb{Q}(A)| \leq |\mathbb{P}(A) - \mathbb{V}(A)| + |\mathbb{V}(A) - \mathbb{Q}(A)|$ by the triangle inequality for numbers (toolkit A.7), and each term on the right is at most the corresponding total variation. Take the maximum over A on the left. □

A total variation of 1 means the two distributions live on separate parts of the sample space: one draw tells you, with certainty, which one it came from. A total variation of 0 means they are the same. Everything else is in between, and Section 10.3 makes 'in between' precise.

From undergraduate statistics: effect sizes and overlap

An introductory course measured how far apart two groups are with Cohen's d, the difference in means in units of the standard deviation, and called 0.2 small, 0.5 medium and 0.8 large. For two Gaussians with a common standard deviation, d is a parameter distance, and the second item of the worked example converts it to a distance between distributions: $\mathrm{TV} = 2\Phi(d/2) - 1$. Figure 10.3 draws the three conventional sizes. A 'large' effect still leaves the two distributions sharing 69% of their mass, and a 'small' one leaves them sharing 92%. The shared area $1 - \mathrm{TV}$

is what some texts call the overlapping coefficient, and TV is the share of each group that has no counterpart in the other. With equal numbers in the two groups, the best rule for guessing which group one person came from is right with probability $(1 + \text{TV})/2$, which at $d = 0.8$ is 0.66.

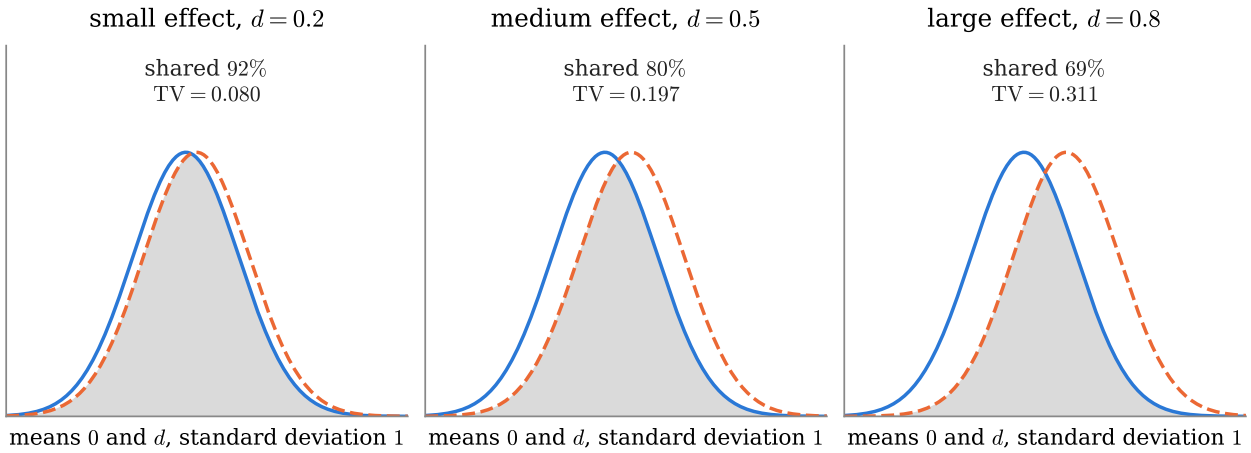


Figure 10.3: Cohen's three conventional effect sizes as total variation distances between Gauss$(0, 1)$ (blue, solid) and Gauss$(d, 1)$ (orange, dashed). The shared mass (gray) is 92%, 80% and 69%; the total variation distances are 0.080, 0.197 and 0.311.

Pitfall: total variation does not know which outcomes are near each other

Let $X \sim \text{Bern}(\frac{1}{2})$ and $a = 0.001$. The random variable $X + a$ is X shifted by a thousandth, and any reasonable notion of 'close' should call the two distributions close. Total variation calls them as far apart as two distributions can be: X lives on $\{0, 1\}$, $X + a$ on $\{0.001, 1.001\}$, the event $\{0, 1\}$ has probability 1 under the first and 0 under the second, and $\text{TV} = 1$. Total variation compares probabilities point by point; it has no idea that 0 and 0.001 are neighbors. The same thing happens with the CLT.

Recall. The Central Limit Theorem (Theorem 4.8): for iid data with mean μ and variance σ^2, $\sqrt{n}(\bar{X}_n - \mu)/\sigma \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1)$.

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(\frac{1}{2})$, $2\sqrt{n}(\bar{X}_n - \frac{1}{2}) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1)$, but $2\sqrt{n}(\bar{X}_n - \frac{1}{2})$ takes only the $n+1$ values $2\sqrt{n}(k/n - \frac{1}{2})$, a set that has probability 1 under its own distribution and probability 0 under any continuous one. So the total variation distance between the standardized mean and its Gaussian limit is 1 for every n. Convergence in distribution (Definition 4.1) is not convergence in total variation, and it can't be: a discrete distribution is never close to a continuous one in total variation, however finely its mass is spread.

10.3 What total variation says about tests

Total variation has an operational meaning that makes 'in between' concrete: it measures how well data from one distribution can be told apart from data from the other.

Suppose you have to decide, from one draw X, whether it came from $\mathbb{P}$ or from $\mathbb{Q}$. This is the test

of Chapter 9 with $H_0 : X \sim \mathbb{P}$ and $H_1 : X \sim \mathbb{Q}$, both simple. A test is an event R, the rejection region, and its two errors are $\alpha = \mathbb{P}(R)$, rejecting when $\mathbb{P}$ is true, and $\beta = \mathbb{Q}(R^c)$, not rejecting when $\mathbb{Q}$ is true. How small can the two errors be together?

Proposition 10.4 (The best test and total variation)**.** *For every test R of $\mathbb{P}$ against $\mathbb{Q}$,*

$$\alpha + \beta = \mathbb{P}(R) + \mathbb{Q}(R^c) \geq 1 - \mathrm{TV}(\mathbb{P}, \mathbb{Q}),$$

with equality for the test that rejects where $\mathbb{Q}$'s density is larger, $R = \{x : q(x) > p(x)\}$.

Proof strategy (core). Rewrite $\alpha + \beta$ as 1 minus the advantage $\mathbb{Q}(R) - \mathbb{P}(R)$ the rejection region gives the alternative. That advantage can never exceed TV, and the region where q beats p achieves it.

Proof. $\mathbb{Q}(R^c) = 1 - \mathbb{Q}(R)$, so $\alpha + \beta = 1 - \big(\mathbb{Q}(R) - \mathbb{P}(R)\big)$. By Definition 10.1, $\mathbb{Q}(R) - \mathbb{P}(R) \leq \mathrm{TV}(\mathbb{P}, \mathbb{Q})$ for every event R, which gives the inequality. By Proposition 10.2 with the roles of $\mathbb{P}$ and $\mathbb{Q}$ exchanged, the event $\{q > p\}$ attains the maximum, which gives the equality. □

Intuition: total variation is distinguishability

Read Proposition 10.4 at its two ends. If TV = 0, the two distributions are the same, and every test has $\alpha + \beta = 1$ exactly, no better than a test that ignores the data. If TV = 1, the two live on separate parts of E, and the test 'reject if X lands in $\mathbb{Q}$'s part' makes no errors at all. In between, $1 - \mathrm{TV}$ is the irreducible error: the total of the two error rates that no test, however clever, can get below. In Figure 10.2 it is the gray overlap. A draw that lands in the overlap could have come from either distribution, and that is where every test makes its mistakes.

Worked example: how distinguishable is 'no preference' from the kiss data?

For one couple, $\mathrm{TV}(\mathrm{Bern}(\frac{1}{2}), \mathrm{Bern}(0.645)) = 0.145$, so the best test based on a single couple has $\alpha + \beta \geq 0.855$: almost no better than a coin flip.

For 124 couples, the relevant distributions are those of the count of right turns, $\mathrm{Bin}(124, \frac{1}{2})$ and $\mathrm{Bin}(124, 0.645)$; the order in which the couples turned carries no information about p, so the count is all that matters. Their total variation is 0.900, so the best test has $\alpha + \beta = 0.100$. The right of Figure 10.2 shows where the best test rejects: where the 0.645 PMF is above the $\frac{1}{2}$ PMF, which is from 72 right turns up. Its two errors are

$$\alpha = \mathbb{P}_{1/2}(\text{count} \geq 72) = 0.0438, \qquad \beta = \mathbb{P}_{0.645}(\text{count} \leq 71) = 0.0567,$$

and they add up to 0.100. In natural frequencies: of 1,000 studies of 124 couples in a world with no preference, the best test cries 'preference' in about 44; of 1,000 studies in a world where $p = 0.645$, it misses the preference in about 57. No test of 124 couples does better on both counts at once.

That threshold, 72, is the one Section 9.6 found for the one-sided exact test at level 5%. Both tests reject for large counts, which is no accident; that they pick the same cutoff is a coincidence of these particular numbers: the level-5% test fixes α near 0.05 and accepts whatever β results, while the best test of Proposition 10.4 minimizes the sum and accepts whatever split results.

Does this make sense? More data, more distinguishable

Going from one couple to 124 took the total variation from 0.145 to 0.900. That is the direction it has to go: more data can only make two hypotheses easier to tell apart, since you could always ignore the extra data. But it doesn't go up by a factor of 124, and it couldn't: total variation is capped at 1. Total variation doesn't add over independent observations, and it has no simple formula for n observations in terms of one. The divergence of Section 10.5 does, which is one of several reasons it is the one that statistics uses.

10.4 Why total variation can't be estimated by an average

Here is the plan the chapter started with. Total variation measures how close $\mathbb{P}_\theta$ is to the truth $\mathbb{P}_{\theta^*}$. So estimate the function

$$\theta \mapsto \mathrm{TV}(\mathbb{P}_\theta, \mathbb{P}_{\theta^*})$$

from the data, and take as $\hat{\theta}_n$ the θ that makes the estimate smallest. If the estimate is good, its minimizer should be near the true minimizer, which is θ^* (the only θ at distance 0, if the model is identifiable).

The obstacle is the word 'estimate'. The one general-purpose estimation tool in this book is the Law of Large Numbers.

Recall. The Law of Large Numbers (Theorem 3.5): if $X_1, \ldots, X_n \overset{iid}{\sim} \mathbb{P}_{\theta^*}$ and h is a function with $\mathbb{V}\mathrm{ar}_{\theta^*}(h(X)) < \infty$, then $\frac{1}{n}\sum_{i=1}^n h(X_i) \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}_{\theta^*}[h(X)]$. An expectation under the truth is estimated by an average over the data.

So the question is whether $\mathrm{TV}(\mathbb{P}_\theta, \mathbb{P}_{\theta^*})$ is an expectation under $\mathbb{P}_{\theta^*}$ of something you can compute. It can be written as one. Factor $p_{\theta^*}(x)$ out of the integrand of Proposition 10.2:

$$\mathrm{TV}(\mathbb{P}_\theta, \mathbb{P}_{\theta^*}) = \frac{1}{2}\int \left|\frac{p_\theta(x)}{p_{\theta^*}(x)} - 1\right| p_{\theta^*}(x)\,\mathrm{d}x = \mathop{\mathbb{E}}_{\theta^*}\left[\frac{1}{2}\left|\frac{p_\theta(X)}{p_{\theta^*}(X)} - 1\right|\right].$$

But the function inside the expectation contains p_{θ^*}, the density you don't know, and it sits inside the absolute value tangled up with p_θ. To average it over the data you would need to evaluate it at each X_i, and to evaluate it you would need θ^*, so no average of the data estimates it. For a sample space with a few outcomes there is a way around this: plug in the observed frequencies. For the kiss data, $\mathrm{TV}(\mathrm{Bern}(q), \mathrm{Bern}(p^*)) = |q - p^*|$ is estimated by $|q - \hat{p}|$, whose minimizer is $\hat{p}$. For continuous data the plug-in fails outright: the observed frequencies form a discrete distribution, and by the pitfall of Section 10.2 a discrete distribution sits at total variation 1 from every continuous candidate.

What would rescue the plan is a measure of distance in which the unknown p_{θ^*} separates from p_θ, so that the part involving θ^* is a constant that doesn't depend on θ. A constant can be left unestimated: it doesn't move the minimizer. The logarithm turns a ratio into a difference, $\log(p_{\theta^*}/p_\theta) = \log p_{\theta^*} - \log p_\theta$, and that is how the next divergence is built.

10.5 The Kullback–Leibler divergence

Here is a second way to compare two distributions, built from a question about evidence. Suppose the truth is $\mathbb{P}$, with PMF or density p, and you are weighing it against an alternative $\mathbb{Q}$, with q. One

observation x arrives. How much does it favor $\mathbb{P}$ over $\mathbb{Q}$? A natural measure is the ratio $p(x)/q(x)$, how many times more likely x is under $\mathbb{P}$ than under $\mathbb{Q}$, and its logarithm $\log\big(p(x)/q(x)\big)$ turns the product of the ratios for independent observations into a sum. This log-ratio is positive when x favors $\mathbb{P}$, negative when it favors $\mathbb{Q}$ and 0 when it favors neither. A single observation can point the wrong way. On average, over observations drawn from the truth $\mathbb{P}$, it should point the right way. That average is the divergence.

Definition 10.5 (Kullback–Leibler divergence)**.** Let $\mathbb{P}$ and $\mathbb{Q}$ be distributions on E with PMFs or densities p and q. The **Kullback–Leibler (KL) divergence** of $\mathbb{Q}$ from $\mathbb{P}$ is

$$\mathrm{KL}(\mathbb{P},\mathbb{Q}) = \mathop{\mathbb{E}}_{\mathbb{P}}\left[\log\frac{p(X)}{q(X)}\right] = \begin{cases} \displaystyle\sum_{x\in E} p(x)\log\frac{p(x)}{q(x)} & \text{(discrete)},\\[2ex] \displaystyle\int_E p(x)\log\frac{p(x)}{q(x)}\,\mathrm{d}x & \text{(continuous)}, \end{cases}$$

where the sum or integral runs over the points with $p(x) > 0$, and $\mathrm{KL}(\mathbb{P},\mathbb{Q}) = +\infty$ if $q(x) = 0$ at some point where $p(x) > 0$. It is also written $\mathrm{KL}(\mathbb{P}\,\|\,\mathbb{Q})$, and called the **relative entropy**.

In words. Draw X from the truth $\mathbb{P}$ and compute how much more likely X is under $\mathbb{P}$ than under the alternative $\mathbb{Q}$, on the log scale. Average that over many draws. The result is the average evidence one observation supplies for the truth against $\mathbb{Q}$. The divergence is large when $\mathbb{Q}$ makes the observations that actually happen look improbable, and it is infinite when $\mathbb{Q}$ calls something impossible that $\mathbb{P}$ produces: one such observation refutes $\mathbb{Q}$ outright. The logarithm is natural (base e), and the unit is the *nat*; with base-2 logarithms the unit is the *bit*, and 1 nat is 1.44 bits.

For the kiss data. Take the truth to be the fitted $\mathrm{Bern}(\hat{p})$, $\hat{p} = 80/124 = 0.6452$, and the alternative $\mathrm{Bern}(\frac{1}{2})$. A couple turning right favors $\hat{p}$ by $\log(0.6452/0.5) = 0.2549$; a couple turning left favors $\frac{1}{2}$ by $\log(0.5/0.3548) = 0.3430$. Right turns happen 64.5% of the time, so on average

$$\mathrm{KL}(\mathrm{Bern}(\hat{p}),\mathrm{Bern}(\tfrac{1}{2})) = 0.6452\times 0.2549 - 0.3548\times 0.3430 = 0.0428 \text{ nats}.$$

Each couple, on average, supplies 0.043 nats of evidence that $p = 0.645$ rather than $\frac{1}{2}$. Left turns are individually stronger evidence the other way; right turns are more frequent, and frequency wins.

Two conventions in the definition need a word. Points where $p(x) = 0$ never occur when $X \sim \mathbb{P}$, so they contribute nothing to an expectation under $\mathbb{P}$; this is also the convention $0\log(0/q) = 0$, which agrees with the limit $t\log t \to 0$ as $t \to 0^+$. Points where $q(x) = 0$ but $p(x) > 0$ are the opposite: they occur, and when they do, the log-ratio is $+\infty$. An expectation that puts positive probability on $+\infty$ is $+\infty$.

Worked example: KL divergences you can compute by hand

- **Two Bernoullis.** With two outcomes the sum has two terms:

$$\mathrm{KL}(\mathrm{Bern}(p),\mathrm{Bern}(q)) = p\log\frac{p}{q} + (1-p)\log\frac{1-p}{1-q}.$$

  Figure 10.4 plots it with $p = 0.645$ fixed. Two values to keep: $\mathrm{KL}(\mathrm{Bern}(\frac{1}{2}),\mathrm{Bern}(\frac{1}{10})) = 0.511$ and $\mathrm{KL}(\mathrm{Bern}(\frac{1}{10}),\mathrm{Bern}(\frac{1}{2})) = 0.368$.

- **Two Gaussians with the same variance.** For $\mathbb{P} = \mathrm{Gauss}(a,\sigma^2)$ and $\mathbb{Q} = \mathrm{Gauss}(b,\sigma^2)$

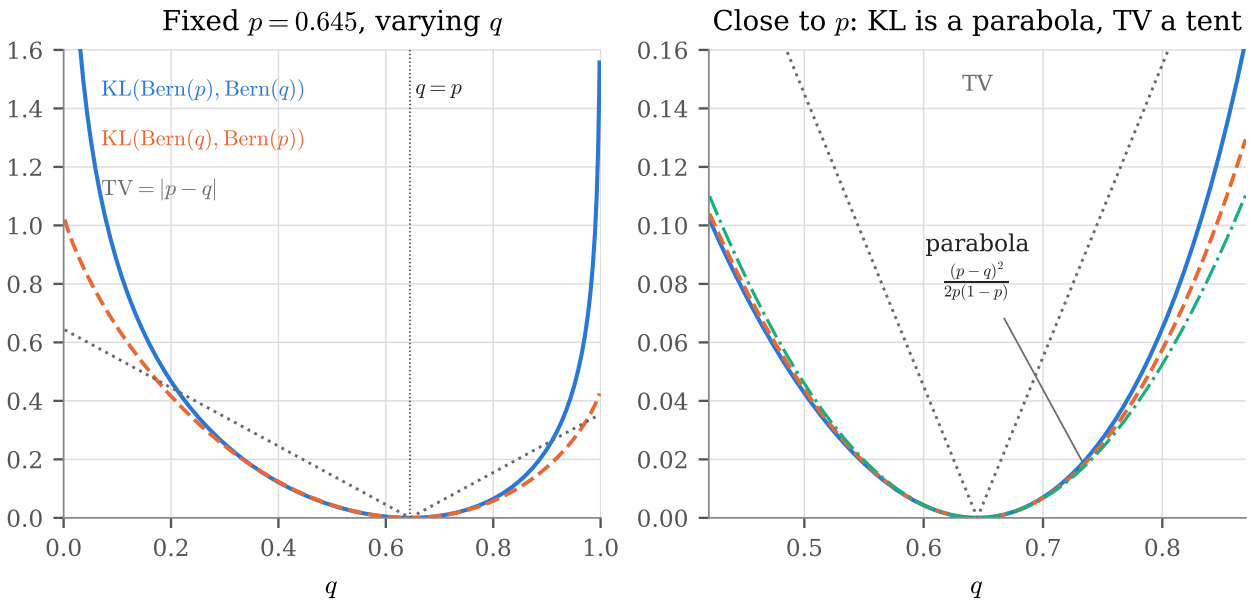


Figure 10.4: KL divergence and total variation between Bern(p) and Bern(q), with $p = 0.645$ fixed and q varying. Left, the whole range: total variation is a tent with straight sides; the two KL divergences are smooth at $q = p$, differ from each other away from it, and KL(Bern(p), Bern(q)) (blue) goes to infinity as q approaches 0 or 1, where Bern(q) calls one of the two outcomes impossible. Right, near $q = p$: both KL curves hug the parabola $(p - q)^2/(2p(1 - p))$. At $q = \frac{1}{2}$: total variation 0.145, the two KL divergences 0.0428 and 0.0440, the parabola 0.0460.

the normalizing constants cancel in the ratio, and

$$\log \frac{p(x)}{q(x)} = \frac{(x-b)^2 - (x-a)^2}{2\sigma^2} = \frac{(a-b)(2x-a-b)}{2\sigma^2}.$$

This is linear in x, so its expectation under $X \sim \text{Gauss}(a, \sigma^2)$ is the same expression at $x = a$:

$$\text{KL}(\text{Gauss}(a, \sigma^2), \text{Gauss}(b, \sigma^2)) = \frac{(a-b)(2a-a-b)}{2\sigma^2} = \frac{(a-b)^2}{2\sigma^2}.$$

Half the squared distance between the means in standard-deviation units. For the left pair of Figure 10.1, means 0 and 0.01 with $\sigma = 1$ give 0.00005. The right pair, standard deviations 0.01 and 0.02, does not have a common variance and needs the general formula of Exercise 10.6: KL(Gauss$(0, 0.01^2)$, Gauss$(0, 0.02^2)$) = 0.318, and 0.807 the other way.

- **Two exponentials.** For Expo(λ) against Expo(μ), $\log(p(x)/q(x)) = \log(\lambda/\mu) - (\lambda - \mu)x$, again linear in x, with $\mathbb{E}[X] = 1/\lambda$:

$$\text{KL}(\text{Expo}(\lambda), \text{Expo}(\mu)) = \log \frac{\lambda}{\mu} - \frac{\lambda - \mu}{\lambda} = \log \frac{\lambda}{\mu} + \frac{\mu}{\lambda} - 1.$$

KL(Expo(1), Expo(2)) = 0.307 and KL(Expo(2), Expo(1)) = 0.193.

Does this make sense? The Gaussian formula

Check $(a-b)^2/(2\sigma^2)$ three ways. At $a=b$ it is 0, which a divergence between a distribution and itself must be. It is symmetric in a and b, which is right here (swapping the two Gaussians is the same as reflecting the line, and reflecting changes nothing), though KL is not symmetric in general. And it is unchanged if a, b and σ are all converted to a new unit, say minutes instead of hours: a divergence between distributions can't depend on the units of the measurement, and $(a-b)/\sigma$ doesn't.

Proposition 10.6 (Properties of the KL divergence)**.** *For distributions $\mathbb{P}$ and $\mathbb{Q}$ on E:*

1. ***nonnegativity:*** $\mathrm{KL}(\mathbb{P},\mathbb{Q}) \geq 0$*;*
2. ***definiteness:*** $\mathrm{KL}(\mathbb{P},\mathbb{Q}) = 0$ *only if* $\mathbb{P}=\mathbb{Q}$*;*
3. ***no symmetry:*** *in general* $\mathrm{KL}(\mathbb{P},\mathbb{Q}) \neq \mathrm{KL}(\mathbb{Q},\mathbb{P})$*;*
4. ***no triangle inequality:*** *in general* $\mathrm{KL}(\mathbb{P},\mathbb{Q}) \not\leq \mathrm{KL}(\mathbb{P},\mathbb{V}) + \mathrm{KL}(\mathbb{V},\mathbb{Q})$.

Properties 1 and 2 make KL behave like a distance in the way that matters most: it is zero exactly when the distributions agree. Properties 3 and 4 mean it is not a distance, which is why it is called a divergence.

Proof strategy (core). The log-ratio $\log(p/q)$ is $-\log$ of the ratio q/p. The curve $-\log y$ lies above its tangent line $1-y$ at $y=1$, touching it only there. The average of the ratio $q(X)/p(X)$ under $\mathbb{P}$ is at most 1, since it is the total mass of q on the points where p is positive, so the average of the tangent line is at least 0.

Proof. (1) If $\mathrm{KL}(\mathbb{P},\mathbb{Q}) = +\infty$ there is nothing to prove, so assume $q>0$ wherever $p>0$. Let $X\sim\mathbb{P}$ and $Y = q(X)/p(X)$, a positive random variable. Then $\log(p(X)/q(X)) = -\log Y$. The function $g(y) = -\log y$ is convex ($g''(y) = 1/y^2 > 0$), so it lies above its tangent line at $y=1$, which is $1-y$, and touches it only at $y=1$ (toolkit A.8):

$$-\log y \geq 1-y \quad \text{for all } y>0, \text{ with equality only at } y=1.$$

Taking expectations,

$$\mathrm{KL}(\mathbb{P},\mathbb{Q}) = \mathbb{E}[-\log Y] \geq 1 - \mathbb{E}[Y].$$

In the discrete case, $\mathbb{E}[Y] = \sum_{x:p(x)>0} p(x)\cdot q(x)/p(x) = \sum_{x:p(x)>0} q(x) \leq 1$; in the continuous case the same with an integral. So $1-\mathbb{E}[Y] \geq 0$. (This is the same idea as Jensen's inequality, Theorem 7.8: the tangent-line argument is its proof.)

(2) If $\mathrm{KL}(\mathbb{P},\mathbb{Q}) = 0$, both inequalities in (1) are equalities. The first, $\mathbb{E}[-\log Y] = 1-\mathbb{E}[Y]$, is an average of the nonnegative gap $-\log Y - (1-Y)$ that comes out 0, so the gap is 0 with probability 1, which forces $Y=1$: $q(x)=p(x)$ at every point where $p(x)>0$. Then q already has total mass 1 on those points, so it has none anywhere else, and $q=p$.

(3) $\mathrm{KL}(\mathrm{Bern}(\frac{1}{2}),\mathrm{Bern}(\frac{1}{10})) = 0.511$ and $\mathrm{KL}(\mathrm{Bern}(\frac{1}{10}),\mathrm{Bern}(\frac{1}{2})) = 0.368$.

(4) With the Gaussian formula,

$$\mathrm{KL}(\mathrm{Gauss}(0,1),\mathrm{Gauss}(2,1)) = 2, \quad \text{while}$$
$$\mathrm{KL}(\mathrm{Gauss}(0,1),\mathrm{Gauss}(1,1)) + \mathrm{KL}(\mathrm{Gauss}(1,1),\mathrm{Gauss}(2,1)) = 0.5 + 0.5 = 1.$$

The direct route is twice as long as the detour: KL grows like a squared distance, and squares of distances don't satisfy the triangle inequality. □

Intuition: why the order matters: who is surprised by what

Think of $\log(1/q(x))$ as how surprised you are to see x if you believe $\mathbb{Q}$. Then $\mathrm{KL}(\mathbb{P},\mathbb{Q})$ is how much more surprised, on average, a believer in $\mathbb{Q}$ is than a believer in $\mathbb{P}$, when the data really come from $\mathbb{P}$. Take $\mathbb{P} = \mathrm{Bern}(\frac{1}{2})$, a fair coin, and $\mathbb{Q} = \mathrm{Bern}(\frac{1}{10})$, a coin that rarely shows heads. With the fair coin as the truth, each head is 5 times likelier under the truth than under $\mathbb{Q}$, worth $\log 5 = 1.609$ in favor of the truth, and heads come half the time; each tail favors $\mathbb{Q}$ by $\log(0.9/0.5) = 0.588$. On average, $0.5 \times 1.609 - 0.5 \times 0.588 = 0.511$. Reverse the roles, so the $\frac{1}{10}$ coin is the truth. Now each tail favors the truth by 0.588 and comes 90% of the time, while the rare heads favor the fair-coin believer by 1.609: on average $0.9 \times 0.588 - 0.1 \times 1.609 = 0.368$. The decisive outcome, heads, is common in the first case and rare in the second. The divergence is largest when the alternative calls rare what the truth makes common. At the extreme, $\mathrm{KL}(\mathrm{Bern}(\frac{1}{2}), \mathrm{Bern}(0)) = +\infty$, since the first head refutes 'heads never happen'; but $\mathrm{KL}(\mathrm{Bern}(0), \mathrm{Bern}(\frac{1}{2})) = \log 2 = 0.693$, finite, since a stream of tails is perfectly possible for a fair coin, just less likely each time.

10.5.1 Independent observations add

The single most useful property of the KL divergence is the one total variation lacks.

Proposition 10.7 (Additivity over independent observations)**.** *If $\mathbb{P}^{\otimes n}$ and $\mathbb{Q}^{\otimes n}$ denote the distributions of n independent draws from $\mathbb{P}$ and from $\mathbb{Q}$, then*

$$\mathrm{KL}(\mathbb{P}^{\otimes n}, \mathbb{Q}^{\otimes n}) = n\,\mathrm{KL}(\mathbb{P},\mathbb{Q}).$$

Proof strategy (for completeness; skip on a first reading). The log of a product is a sum, and an expectation of a sum is a sum of expectations.

Proof. The joint PMF or density of n independent draws is the product, so the log-ratio of the joint distributions is a sum:

$$\log \frac{\prod_{i=1}^n p(x_i)}{\prod_{i=1}^n q(x_i)} = \sum_{i=1}^n \log \frac{p(x_i)}{q(x_i)}.$$

Under $\mathbb{P}^{\otimes n}$ each $X_i \sim \mathbb{P}$, and the expectation of a sum is the sum of the expectations, each equal to $\mathrm{KL}(\mathbb{P},\mathbb{Q})$. □

Evidence from independent observations accumulates, and on the log scale it accumulates by addition. For the kiss data, 124 couples supply on average $124 \times 0.0428 = 5.30$ nats of evidence for $p = 0.645$ over $p = \frac{1}{2}$, a likelihood ratio of about $e^{5.30} \approx 200$. The count of right turns is enough to compute that evidence, so the same additivity gives $\mathrm{KL}(\mathrm{Bin}(n,p), \mathrm{Bin}(n,q)) = n\,\mathrm{KL}(\mathrm{Bern}(p), \mathrm{Bern}(q))$, which Exercise 10.5 asks you to check directly. Compare total variation, which went from 0.145 for one couple to 0.900 for 124, by no rule you could state, and could never pass 1.

10.5.2 Near the truth, KL is a parabola

Figure 10.4 shows something about the shape of KL that the formulas don't make obvious. Total variation between Bernoullis is $|p - q|$, a tent with a corner at $q = p$. KL is smooth there: flat

at the bottom, curving up on both sides, the shape of a parabola. It must be flat at $q = p$, since KL has a minimum of 0 there and is differentiable. A second-order Taylor expansion (toolkit A.9) in q around p gives the curvature. With $f(q) = p\log(p/q) + (1-p)\log((1-p)/(1-q))$, $f(p) = 0$; $f'(q) = -p/q + (1-p)/(1-q)$, so $f'(p) = 0$; and $f''(q) = p/q^2 + (1-p)/(1-q)^2$, so $f''(p) = 1/p + 1/(1-p) = 1/(p(1-p))$. Hence

$$\mathrm{KL}(\mathrm{Bern}(p), \mathrm{Bern}(q)) \approx \frac{1}{2} f''(p)(q-p)^2 = \frac{(p-q)^2}{2p(1-p)} \quad \text{for } q \text{ near } p.$$

For the kiss data the approximation gives 0.0460 against the exact 0.0428.

Look at what the curvature is: $1/(p(1-p))$, one over the variance of a single Bernoulli observation. The Gaussian formula has the same shape, $(a-b)^2/(2\sigma^2)$, with curvature $1/\sigma^2$. In both, KL near the truth is half the squared parameter distance measured in units of the noise in one observation. Where an observation is noisy, nearby parameters are hard to tell apart and KL rises slowly; where it is precise, KL rises fast. Chapter 12 gives the curvature its name, the Fisher information, and shows that it sets the variance of the best estimators.

From undergraduate statistics: the G-test and Pearson's chi-squared

Some introductory courses teach a test of a proportion, or of a table of counts, called the G-test or likelihood-ratio chi-squared test: $G^2 = 2\sum O\log(O/E)$, summed over cells, with O the observed and E the expected counts. That statistic is a KL divergence. For the kiss data, the cells are right and left, $O = (80, 44)$ and, under $p = \frac{1}{2}$, $E = (62, 62)$:

$$G^2 = 2\left(80\log\tfrac{80}{62} + 44\log\tfrac{44}{62}\right) = 2 \times 124 \times \mathrm{KL}\big(\mathrm{Bern}(\hat{p}), \mathrm{Bern}(\tfrac{1}{2})\big) = 10.60,$$

twice the 5.30 nats of evidence above. Pearson's chi-squared statistic for the same table is $\sum(O-E)^2/E = 10.45$, which is also $z^2 = 3.233^2$ from the test of Section 9.6. Pearson's statistic is a parabola approximation to G^2, with the variance taken at the null value. The expansion above was in q around p; expanding instead in p around q gives the same parabola with the variance at q, $(p-q)^2/(2q(1-q))$. With $q = \frac{1}{2}$ that is $2(\hat{p} - \frac{1}{2})^2 = 0.0421$, and $248 \times 0.0421 = 10.45$, Pearson's statistic. With the variance at $\hat{p}$ instead, $248 \times 0.0460 = 11.41$, the square of the Wald statistic $\sqrt{124}(\hat{p} - \frac{1}{2})/\sqrt{\hat{p}(1-\hat{p})} = 3.38$ of Chapter 15. Both are compared with a chi-squared distribution with one degree of freedom, whose 95% quantile is $3.84 = 1.96^2$, and both reject. Chapter 16 develops the chi-squared test in general.

10.5.3 KL controls total variation

KL is not a distance, but small KL forces small total variation.

Theorem 10.8 (Pinsker's inequality)**.** *For any distributions* $\mathbb{P}$ *and* $\mathbb{Q}$*,*

$$\mathrm{TV}(\mathbb{P}, \mathbb{Q}) \le \sqrt{\tfrac{1}{2}\mathrm{KL}(\mathbb{P}, \mathbb{Q})}.$$

Nothing later in the book depends on Pinsker's inequality, so its proof is in Appendix C.1 of Appendix C. It is a page of calculus that reuses the tangent-line bound from the proof of Proposition 10.6. For the kiss data, $\sqrt{0.0428/2} = 0.146$, against TV = 0.145: nearly tight. The inequality is what lets you work with KL, which is convenient, and still conclude something about total variation, which is what you wanted in Section 10.1: if the fitted distribution is close to the truth in KL, then

every probability it reports is close to the true probability. The converse fails: total variation can be small while KL is infinite. $\mathbb{P} = \text{Pois}(1/n)$ and $\mathbb{Q}$, the distribution that puts all its mass at 0, have $\text{TV} = 1 - e^{-1/n}$, which goes to 0, while $\text{KL}(\mathbb{P}, \mathbb{Q}) = +\infty$ for every n, since $\mathbb{P}$ gives positive probability to 1 and $\mathbb{Q}$ calls it impossible (Exercise 10.8).

10.6 Estimating the KL divergence: the road to maximum likelihood

Return to the plan of Section 10.4 with KL in place of total variation. The quantity to make small is $\theta \mapsto \text{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta)$, the divergence of the candidate from the truth. The truth goes first: the average is taken under the distribution the data come from, which is the only distribution you can average under. Split the logarithm of the ratio into a difference:

$$\text{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta) = \mathop{\mathbb{E}}_{\theta^*}\big[\log p_{\theta^*}(X)\big] - \mathop{\mathbb{E}}_{\theta^*}\big[\log p_\theta(X)\big].$$

The first term involves the unknown θ^*, but it does not involve θ: as a function of θ it is a constant. The second term is an expectation under the truth of a function, $\log p_\theta$, that you can compute for any θ you like.

Recall. The Law of Large Numbers (Theorem 3.5): an average of $h(X_1), \dots, h(X_n)$ over iid data converges in probability to $\mathbb{E}_{\theta^*}[h(X)]$, provided $h(X)$ has finite variance.

Apply it to $h = \log p_\theta$, one θ at a time:

$$\widehat{\text{KL}}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta) = \text{constant} - \frac{1}{n}\sum_{i=1}^{n} \log p_\theta(X_i).$$

You don't know the constant, and you don't need to. The θ that minimizes $\widehat{\text{KL}}$ is the same whatever the constant is, and it is the θ that maximizes $\frac{1}{n}\sum_i \log p_\theta(X_i)$. Since log turns a product into a sum and is increasing (toolkit A.10),

$$\operatorname*{argmin}_\theta \widehat{\text{KL}}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta) = \operatorname*{argmax}_\theta \sum_{i=1}^{n} \log p_\theta(X_i) = \operatorname*{argmax}_\theta \prod_{i=1}^{n} p_\theta(X_i).$$

The product $\prod_i p_\theta(X_i)$ is the probability, or density, of the data you observed, computed as if θ were the truth. The θ that makes your data most probable is the **maximum likelihood estimator**, and minimizing the estimated KL divergence is the same thing as finding it. Chapter 11 takes this as its starting point.

Worked example: the kiss data

For Bernoulli data, $\log p_q(x) = x \log q + (1 - x)\log(1 - q)$, and the average over 124 couples of whom 80 turned right is

$$\frac{1}{124}\sum_{i=1}^{124} \log p_q(R_i) = \hat{p}\log q + (1 - \hat{p})\log(1 - q), \qquad \hat{p} = \tfrac{80}{124} = 0.645.$$

Its derivative in q is $\hat{p}/q - (1 - \hat{p})/(1 - q)$, which is zero at $q = \hat{p}$, positive before and negative after. The estimated KL is smallest at the sample proportion, $\hat{p} = 0.645$, the estimator you

would have used anyway.

Figure 10.5 shows the estimate at work. To draw it, the unknown constant has to be filled in, so the figure does what only a simulation can: it fixes the truth at $p^* = 0.645$ and draws samples from it. The black curve is the true $\mathrm{KL}(\mathrm{Bern}(p^*), \mathrm{Bern}(q))$. Each colored curve is the estimate from one simulated sample, $\frac{1}{n}\sum_i \log(p^*(R_i)/p_q(R_i))$, the same curve as the true one with the expectation replaced by an average. Each estimated curve has its minimum at its own sample proportion. As n grows the estimated curves close in on the true one, and their minimizers close in on p^*, which is the true curve's minimizer. That is the argument Chapter 12 turns into a proof that the maximum likelihood estimator is consistent.

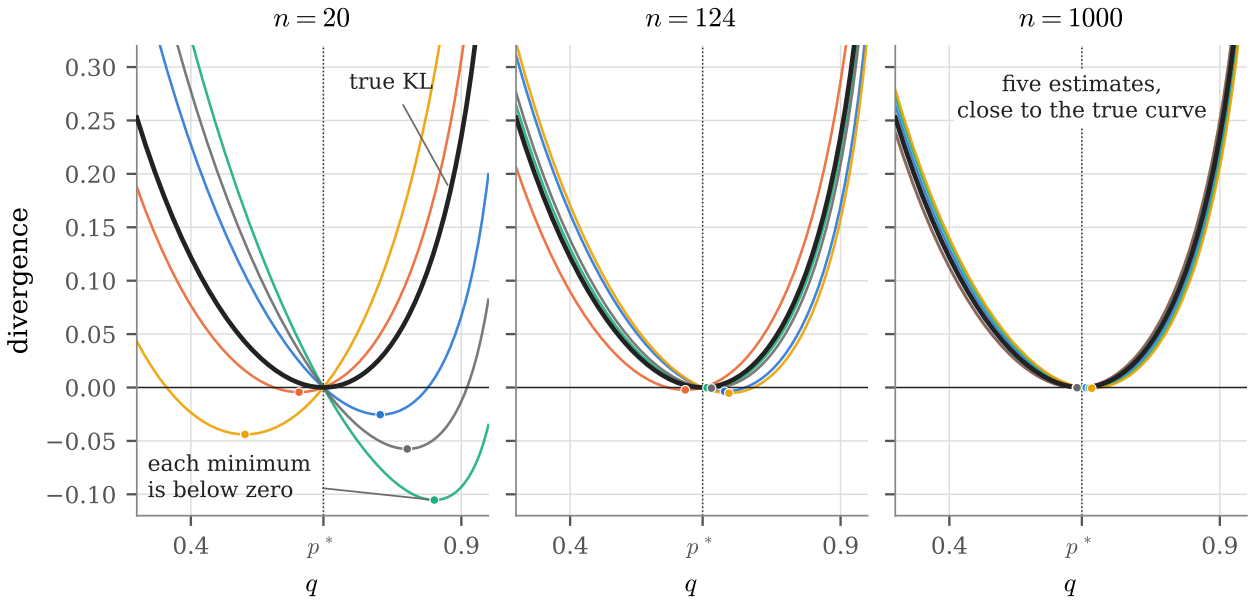


Figure 10.5: The true $\mathrm{KL}(\mathrm{Bern}(p^*), \mathrm{Bern}(q))$ for $p^* = 0.645$ (black) and five estimates of it, each from one simulated sample of n couples (colors), for $n = 20$, 124 and 1,000. Each estimated curve is lowest at its own sample proportion (dots), which is the maximum likelihood estimate. Every estimated curve passes through 0 at $q = p^*$ and dips below zero near its minimum. As n grows the estimates close in on the true curve, and their minimizers on p^*.

Does this make sense? An estimated divergence below zero

Every estimated curve in Figure 10.5 dips below zero, and a KL divergence is never negative. Is the estimate broken? No. At $q = p^*$ the estimate is exactly 0 (each term is $\log 1$), and at its own minimizer $\hat{p}$ it must be lower still. The estimate is an average, and an average of quantities whose expectation is positive can come out negative in a finite sample; the LLN promises only that it ends up near the expectation. What the dip measures is how much better the sample proportion fits this sample than the truth does, which is always at least a little, since $\hat{p}$ was chosen by looking at the sample. For $n = 20$ the dips are visible; for $n = 1{,}000$ they are invisible at this scale.

Pitfall: the constant you dropped was the distance

Minimizing $\widehat{\mathrm{KL}}$ finds the best θ, but the minimum value of $-\frac{1}{n}\sum_i \log p_\theta(X_i)$ is not an estimate of the distance from the truth. The dropped constant $\mathbb{E}_{\theta^*}[\log p_{\theta^*}(X)]$ is unknown, so the number you computed is the divergence plus an unknown shift. The maximized log-likelihood of the kiss data, $80\log 0.645 + 44\log 0.355 = -80.6$, says nothing on its own about how well the Bernoulli model fits. Only differences of log-likelihoods, between two values of θ or two models fitted to the same data, cancel the constant and mean something. That is why the likelihood ratio tests of Chapter 15 compare two maximized log-likelihoods and never interpret one alone.

10.7 When the model is wrong: the closest member of the family

Everything so far assumed the truth is $\mathbb{P}_{\theta^*}$ for some θ^* in the model. Chapter 2 warned that the assumption can fail, and promised to say what estimators estimate when it does. Here is the answer.

Let the truth be any distribution $\mathbb{P}$, in the family or not.

Recall. The Law of Large Numbers (Theorem 3.5): for iid data from $\mathbb{P}$, $\frac{1}{n}\sum_i h(X_i) \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}_{\mathbb{P}}[h(X)]$ when $h(X)$ has finite variance. Nothing in it requires $\mathbb{P}$ to belong to a model.

The estimated divergence is built the same way: $\mathrm{KL}(\mathbb{P}, \mathbb{P}_\theta) = \mathbb{E}_{\mathbb{P}}[\log p(X)] - \mathbb{E}_{\mathbb{P}}[\log p_\theta(X)]$, the first term is still a constant, and the LLN still estimates the second by $\frac{1}{n}\sum_i \log p_\theta(X_i)$, since the data still come from $\mathbb{P}$. So maximum likelihood still minimizes an estimate of $\theta \mapsto \mathrm{KL}(\mathbb{P}, \mathbb{P}_\theta)$. What has changed is the minimum: no θ makes the divergence 0, so the minimizer is the member of the family that is closest to the truth, not the truth.

Definition 10.9 (KL projection)**.** Let $\mathbb{P}$ be a distribution and $(\mathbb{P}_\theta)_{\theta\in\Theta}$ a family on the same sample space. A **KL projection** of $\mathbb{P}$ onto the family is a parameter

$$\theta^\circ = \operatorname*{argmin}_{\theta\in\Theta} \mathrm{KL}(\mathbb{P}, \mathbb{P}_\theta),$$

when the minimum exists and is attained at a single point. If $\mathbb{P} = \mathbb{P}_{\theta^*}$ is in the family and the model is identifiable, $\theta^\circ = \theta^*$.

> **In words.** Look over every distribution the model allows and pick the one the truth diverges from least. If the truth is in the model, that is the truth itself, at divergence 0. If not, it is the model's best impersonation of the truth. Maximum likelihood, applied to data from a truth outside the model, estimates θ°: under conditions that Chapter 12 states, $\hat\theta_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta^\circ$. So a wrong model fitted by maximum likelihood is not meaningless. It converges to something definite, the closest member of the family, and the question is only whether the closest member is close enough for your purpose.
>
> **For the kiss data.** The Bernoulli family contains every distribution on $\{0, 1\}$, so for the kiss experiment the model cannot be wrong in this sense: whatever the truth is, one couple's turn is some Bern(p). What can fail is the iid assumption of Chapter 1: couples influencing each other, or one observer's couples differing from another's. A Bernoulli model with dependent data is misspecified in a way that no projection onto a family of iid models repairs.

Proposition 10.10 (KL projection onto the Gaussian and Poisson families)**.** *Let $\mathbb{P}$ be the distribution of a random variable X.*

1. *If X has a density p with $\mathbb{E}|\log p(X)| < \infty$ and finite variance, the KL projection onto $(Gauss(\mu, \sigma^2))_{\mu\in\mathbb{R},\sigma^2>0}$ is $\mu^\circ = \mathbb{E}[X]$, $(\sigma^2)^\circ = \mathbb{V}ar(X)$: the Gaussian with the same mean and variance.*
2. *If X takes values in $\{0, 1, 2, \ldots\}$, with $\mathbb{E}[X] > 0$, $\mathbb{E}[\log X!] < \infty$ and $\mathbb{E}|\log p(X)| < \infty$ for its PMF p, the KL projection onto $(Pois(\lambda))_{\lambda>0}$ is $\lambda^\circ = \mathbb{E}[X]$: the Poisson with the same mean.*

The conditions on $\log p(X)$ keep the constant $\mathbb{E}[\log p(X)]$ finite. Without them the divergence can be $+\infty$ for every member of the family (a discrete X against any Gaussian, say), and what maximum likelihood targets is then the minimizer of $-\mathbb{E}[\log p_\theta(X)]$ alone, which the same calculation finds.

Proof strategy (for completeness; skip on a first reading). Only the term $-\mathbb{E}[\log p_\theta(X)]$ depends on the parameter. For both families it involves X only through $\mathbb{E}[X]$ and $\mathbb{E}[X^2]$ (plus a term free of the parameter), so it is an explicit function of the parameters. Minimize it by calculus.

Proof. (1) For the Gaussian density, $-\log p_{\mu,\sigma^2}(x) = \frac{1}{2}\log(2\pi\sigma^2) + (x-\mu)^2/(2\sigma^2)$, and

$$-\mathbb{E}[\log p_{\mu,\sigma^2}(X)] = \tfrac{1}{2}\log(2\pi\sigma^2) + \frac{\mathbb{E}[(X-\mu)^2]}{2\sigma^2} = \tfrac{1}{2}\log(2\pi\sigma^2) + \frac{\mathbb{V}\mathrm{ar}(X) + (\mathbb{E}[X]-\mu)^2}{2\sigma^2},$$

using the bias-variance split $\mathbb{E}[(X-\mu)^2] = \mathbb{V}\mathrm{ar}(X) + (\mathbb{E}[X]-\mu)^2$ from Chapter 7. For any σ^2, the μ that minimizes this is $\mu = \mathbb{E}[X]$, which kills the squared term. With that μ, the derivative in σ^2 of $\frac{1}{2}\log(2\pi\sigma^2) + \mathbb{V}\mathrm{ar}(X)/(2\sigma^2)$ is $\frac{1}{2\sigma^2} - \frac{\mathbb{V}\mathrm{ar}(X)}{2\sigma^4}$, zero at $\sigma^2 = \mathbb{V}\mathrm{ar}(X)$, negative before and positive after, so this is the minimum.

(2) For the Poisson PMF, $-\log p_\lambda(x) = \lambda - x\log\lambda + \log(x!)$, and $-\mathbb{E}[\log p_\lambda(X)] = \lambda - \mathbb{E}[X]\log\lambda + \mathbb{E}[\log X!]$. The last term doesn't depend on λ. The derivative of the rest is $1 - \mathbb{E}[X]/\lambda$, zero at $\lambda = \mathbb{E}[X]$, negative before and positive after. □

Worked example: fitting a Gaussian to exponential data

Let the truth be Expo(1), which has mean 1 and variance 1, and fit the Gaussian family. The KL projection is Gauss(1, 1). From the proof, with $\mathbb{V}\mathrm{ar}(X) = 1$ and $\mathbb{E}[X] = 1$, and with the constant term $\mathbb{E}[\log e^{-X}] = -\mathbb{E}[X] = -1$,

$$\mathrm{KL}(\mathrm{Expo}(1), \mathrm{Gauss}(\mu, \sigma^2)) = \tfrac{1}{2}\log(2\pi\sigma^2) + \frac{1 + (1-\mu)^2}{2\sigma^2} - 1.$$

At the projection the divergence is $\frac{1}{2}\log(2\pi) + \frac{1}{2} - 1 = 0.419$. The minimum is not 0, since no Gaussian is an exponential. Figure 10.6 shows the projection against two other candidates, Gauss(0.7, 0.6), which is taller and sits closer to the exponential's peak at 0, and Gauss(1.5, 1.5), which covers its long right tail better. Both diverge more, 0.572 and 0.538. The divergence weights each point by how often the truth produces it, and it penalizes a candidate that calls likely outcomes unlikely; matching the mean and the variance is the best a Gaussian can do on that criterion.

The Poisson half of Proposition 10.10 pays off the debt of Chapter 2. The siblings data of Figure 2.1 were fitted with $X - 1 \sim \mathrm{Pois}(\hat{\lambda})$ and $\hat{\lambda}$ the sample mean minus one. If the true distribution of the number of siblings is not a shifted Poisson, say with an extra bump at 5, the fit is still estimating something definite: the shifted Poisson whose mean matches the true mean, $\lambda^\circ = \mathbb{E}[X] - 1$, which is the KL projection. The sample mean minus one estimates it by the LLN, which needs no assumption

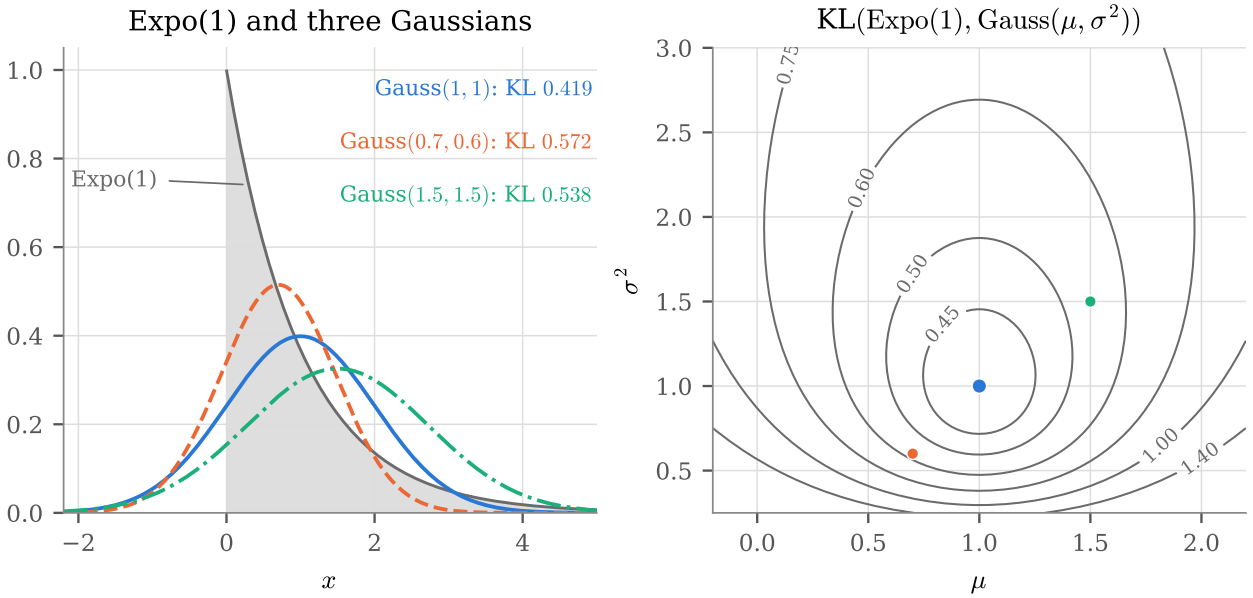


Figure 10.6: Projecting Expo(1) onto the Gaussian family. Left: the exponential density (gray) and three Gaussians with their KL divergences from it. Right: $\mathrm{KL}(\mathrm{Expo}(1), \mathrm{Gauss}(\mu, \sigma^2))$ over a grid of means and variances, with its minimum 0.419 at mean 1 and variance 1, the exponential's own mean and variance; the dots mark the three Gaussians on the left.

that the shifted Poisson is right. Whether that projection is a good description of the truth is a separate question, and it is the question the goodness-of-fit tests of Chapter 16 ask.

From undergraduate statistics: 'fit a normal with the sample mean and standard deviation'

An introductory course fits a normal distribution to data by plugging in $\bar{x}$ and s, and justifies it, if at all, by saying that the normal is determined by its mean and standard deviation. Proposition 10.10 is a better justification, and it holds whether or not the data are normal: $\mathrm{Gauss}(\mathbb{E}[X], \mathbb{V}\mathrm{ar}(X))$ is the Gaussian closest to the truth in KL, and $\bar{x}$ and s^2 estimate its two parameters. What the procedure can't tell you is whether the closest Gaussian is close. For exponential data it is 0.419 nats away, and probabilities computed from it can be badly wrong: under $\mathrm{Gauss}(1, 1)$ the probability of a negative value is $\Phi(-1) = 0.16$, under Expo(1) it is 0. A QQ plot (Chapter 17) is how you see that in data.

Pitfall: the order of the arguments is not a choice

$\mathrm{KL}(\mathbb{P}, \mathbb{P}_\theta)$ and $\mathrm{KL}(\mathbb{P}_\theta, \mathbb{P})$ are different functions of θ with different minimizers in general. Statistics minimizes the first, with the truth first, and not because it is better in some abstract sense: it is the only one that can be estimated. The first is an expectation under the truth $\mathbb{P}$, which the data average. The second is an expectation under the candidate $\mathbb{P}_\theta$, and it has the unknown p inside it in a way that does not split off as a constant. Because KL is not symmetric, the two orders are different problems, and only the truth-first order is an average under the distribution the data come from.

10.8 What comes next

This chapter ended at maximum likelihood without defining a likelihood. Chapter 11 defines it, for discrete and continuous models, and computes the maximum likelihood estimator for the families of this book: Bernoulli, Poisson, Gaussian, exponential, uniform. Chapter 12 proves what Figure 10.5 shows, that the estimator converges to the minimizer of the KL divergence, and then uses the curvature of KL at its minimum, the $1/(p(1-p))$ of Figure 10.4, to find its asymptotic variance: that curvature is the Fisher information. Distances between distributions return in Chapter 17, where the largest gap between an empirical CDF and a candidate CDF is the test statistic.

10.9 Summary

- Closeness should be measured between distributions, not parameters: a parameter difference of 0.01 is TV $= 0.004$ for two means with $\sigma = 1$ and TV $= 0.323$ for standard deviations 0.01 and 0.02. A distance between distributions also makes sense when the model is wrong.
- Total variation: $\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = \max_A |\mathbb{P}(A) - \mathbb{Q}(A)| = \frac{1}{2}\sum |p - q|$ (or $\frac{1}{2}\int$), attained at $\{p > q\}$; the area where one density is above the other; $1 - \mathrm{TV}$ is the shared mass. A true distance with values in $[0, 1]$. $\mathrm{TV}(\mathrm{Bern}(p), \mathrm{Bern}(q)) = |p - q|$; kiss data 0.145 for one couple, 0.900 for 124.
- The best test of $\mathbb{P}$ against $\mathbb{Q}$ has $\alpha + \beta = 1 - \mathrm{TV}$. For 124 couples it rejects $p = \frac{1}{2}$ from 72 right turns, with $\alpha = 0.0438$ and $\beta = 0.0567$.
- Total variation ignores the geometry of the sample space: $\mathrm{TV}(X, X + 0.001) = 1$, and a discrete distribution is at distance 1 from every continuous one, so the CLT is not convergence in total variation.
- Total variation can't be estimated by an average: written as an expectation under the truth, it has the unknown density inside an absolute value. Plugging in frequencies works for a few discrete outcomes and fails for continuous data.
- KL divergence: $\mathrm{KL}(\mathbb{P}, \mathbb{Q}) = \mathbb{E}_{\mathbb{P}}[\log(p(X)/q(X))]$, the average evidence per observation for $\mathbb{P}$ over $\mathbb{Q}$; $+\infty$ if $\mathbb{Q}$ calls impossible what $\mathbb{P}$ produces. Nonnegative and zero only for $\mathbb{P} = \mathbb{Q}$ (Jensen), but not symmetric and no triangle inequality. Kiss data: 0.0428 and, reversed, 0.0440.
- Formulas: Bernoulli $p\log(p/q) + (1-p)\log((1-p)/(1-q))$; Gaussians with a common variance $(a-b)^2/(2\sigma^2)$; exponentials $\log(\lambda/\mu) + \mu/\lambda - 1$.
- KL adds over independent observations: $\mathrm{KL}(\mathbb{P}^{\otimes n}, \mathbb{Q}^{\otimes n}) = n\mathrm{KL}(\mathbb{P}, \mathbb{Q})$; $124 \times 0.0428 = 5.30$ nats for the kiss data. Near the truth it is a parabola with curvature one over the variance of one observation, $(p-q)^2/(2p(1-p))$ for Bernoulli. $2n\mathrm{KL}$ is the G^2 statistic, and its parabola approximation is Pearson's chi-squared.
- Pinsker: $\mathrm{TV} \le \sqrt{\mathrm{KL}/2}$; the converse fails.
- $\mathrm{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta)$ is a constant minus $\mathbb{E}_{\theta^*}[\log p_\theta(X)]$, and the LLN estimates the second term by $\frac{1}{n}\sum \log p_\theta(X_i)$. Minimizing the estimate is maximizing the likelihood $\prod p_\theta(X_i)$. Kiss data: the minimizer is $\hat{p} = 0.645$.
- When the model is wrong, maximum likelihood estimates the KL projection, the member of the family closest to the truth. For the Gaussian family it has the true mean and variance; for the Poisson family, the true mean. Expo(1) projects to Gauss$(1, 1)$ at divergence 0.419.

10.10 Exercises

Exercise 10.1. Compute $\mathrm{TV}(\mathrm{Bern}(\frac{1}{2}), \mathrm{Bern}(\frac{1}{2}))$, $\mathrm{TV}(\mathrm{Bern}(\frac{1}{2}), \mathrm{Bern}(\frac{1}{3}))$ and $\mathrm{TV}(\mathrm{Bern}(\frac{1}{2}), \mathrm{Bern}(\frac{1}{10}))$ from Definition 10.1, by listing the events, and check each against the formula $|p-q|$. *(Solution on page 227.)*

Exercise 10.2. Show that $\mathrm{TV}(\mathrm{Expo}(1), \mathrm{Unif}([0,1])) = 1/e$, and identify the event $A^\star$ of Proposition 10.2. *(Solution on page 227.)*

Exercise 10.3. Let $X \sim \mathrm{Bern}(\frac{1}{2})$ and $a \in (0,1)$. Show that $\mathrm{TV}(X, X+a) = 1$, meaning the total variation distance between the distributions of X and $X+a$. Then show that $\mathrm{TV}(2\sqrt{n}(\bar{X}_n - \frac{1}{2}), Z) = 1$ for every n, where $X_1, \dots, X_n \overset{iid}{\sim} \mathrm{Bern}(\frac{1}{2})$ and $Z \sim \mathrm{Gauss}(0,1)$, and explain why this does not contradict the CLT. *(Solution on page 227.)*

Exercise 10.4. For one observation from $\mathrm{Bern}(p_0)$ or $\mathrm{Bern}(p_1)$ with $p_0 < p_1$, list all four tests (rejection regions $\emptyset$, $\{0\}$, $\{1\}$, $\{0,1\}$), compute $\alpha + \beta$ for each, and confirm that the smallest is $1 - |p_1 - p_0|$, as Proposition 10.4 says. Evaluate for $p_0 = \frac{1}{2}$, $p_1 = 0.645$. *(Solution on page 227.)*

Exercise 10.5. Compute $\mathrm{KL}(\mathrm{Bin}(n,p), \mathrm{Bin}(n,q))$ directly from the Binomial PMF and show that it equals $n\,\mathrm{KL}(\mathrm{Bern}(p), \mathrm{Bern}(q))$. Evaluate for $n = 124$, $p = 80/124$, $q = \frac{1}{2}$. *(Solution on page 227.)*

Exercise 10.6. Show that

$$\mathrm{KL}(\mathrm{Gauss}(a,\sigma^2), \mathrm{Gauss}(b,\tau^2)) = \log\frac{\tau}{\sigma} + \frac{\sigma^2 + (a-b)^2}{2\tau^2} - \frac{1}{2}.$$

Check that it reduces to $(a-b)^2/(2\sigma^2)$ when $\tau = \sigma$, and evaluate it for $\mathrm{Gauss}(0, 0.01^2)$ against $\mathrm{Gauss}(0, 0.02^2)$ and the other way round. *(Solution on page 228.)*

Exercise 10.7. Show that $\mathrm{KL}(\mathrm{Expo}(\lambda), \mathrm{Expo}(\mu)) = \log(\lambda/\mu) + \mu/\lambda - 1$, check that it is nonnegative by studying $t \mapsto t - 1 - \log t$ with $t = \mu/\lambda$, and check that it is 0 only when $\lambda = \mu$. *(Solution on page 228.)*

Exercise 10.8. Let $\mathbb{P}_n = \mathrm{Pois}(1/n)$ and let δ_0 be the distribution that puts all its mass at 0. Compute $\mathrm{TV}(\mathbb{P}_n, \delta_0)$, $\mathrm{KL}(\mathbb{P}_n, \delta_0)$ and $\mathrm{KL}(\delta_0, \mathbb{P}_n)$, and say what happens as $n \to \infty$. *(Solution on page 228.)*

Exercise 10.9. Let $X_1, \dots, X_n \overset{iid}{\sim} \mathrm{Pois}(\lambda^*)$. Write $\widehat{\mathrm{KL}}(\mathrm{Pois}(\lambda^*), \mathrm{Pois}(\lambda))$ up to its constant, and find the λ that minimizes it. *(Solution on page 228.)*

Exercise 10.10. Data $X_i \in \{0,1,2\}$ come from a distribution with $\mathbb{P}(X=0) = 0.5$, $\mathbb{P}(X=1) = 0.2$, $\mathbb{P}(X=2) = 0.3$, and are fitted with the model $X \sim \mathrm{Bin}(2,\theta)$. Find the KL projection θ° and the divergence at it. Is the model well-specified? *(Solution on page 228.)*

Chapter 11

The Likelihood and Maximum Likelihood

In this chapter

- The likelihood: the probability (or density) of the data you actually observed, read as a function of the parameter. One formula, two readings, and why the switch between them is the whole idea.
- Why a density can stand in for a probability when the data are continuous.
- The likelihoods of the standard families: Bernoulli, Poisson, Gaussian, exponential, uniform, and a mixture of Gaussians.
- The maximum likelihood estimator, and how to find a maximum: concave functions, the first-order condition, and the second-derivative test in one and in several variables.
- The maximum likelihood estimators of the standard families, computed for the kiss data, the siblings data, the subway data and the boarding data. The uniform family, where calculus fails and the answer is the largest observation.
- When the estimator doesn't exist or isn't unique, and what to do when there is no formula: the EM algorithm for a mixture, run on the waiting times of Old Faithful.

Chapter 10 ended with a principle: minimizing an estimate of the KL divergence from the truth is the same as maximizing $\prod_i p_\theta(X_i)$. This chapter names that product and learns to maximize it.

11.1 From the KL divergence to the likelihood

Recall. Chapter 10: for data $X_1, \ldots, X_n \overset{iid}{\sim} \mathbb{P}_{\theta^*}$, $\mathrm{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta) = \text{constant} - \mathbb{E}_{\theta^*}[\log p_\theta(X)]$; the Law of Large Numbers estimates the expectation by $\frac{1}{n}\sum_i \log p_\theta(X_i)$ (when $\log p_\theta(X)$ has finite variance); and the θ that minimizes the estimated divergence is the θ that maximizes $\prod_{i=1}^n p_\theta(X_i)$.

That product is the subject of this chapter. It is the probability, or the density, of the data you have, computed as if θ were the truth. Chapter 10 reached it from a question about distances. There is a second road to it, more direct and older. Ask of each candidate θ: if this were the truth, how probable would the data I actually saw be? A candidate under which your data would be a fluke is a poor candidate. A candidate under which your data are typical is a good one. Ranking candidates

by the probability they give your data, and picking the top one, is maximum likelihood. The two roads arrive at the same estimator, and Chapter 12 uses the first one to prove that it works.

11.2 The likelihood of a discrete model

Start with a model whose observations take values in a countable set E: the kiss data in $\{0,1\}$, the siblings counts in $\{1,2,3,\ldots\}$. Each $\mathbb{P}_\theta$ has a PMF $p_\theta(x) = \mathbb{P}_\theta(X = x)$.

Definition 11.1 (Likelihood, discrete model)**.** Let $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ be a statistical model for $X_1, \ldots, X_n$ i.i.d., with E countable and $\mathbb{P}_\theta$ having PMF p_θ. The **likelihood** of the model is the function

$$L_n : E^n \times \Theta \to \mathbb{R}$$

$$(x_1, \ldots, x_n; \theta) \mapsto \mathbb{P}_\theta(X_1 = x_1, \ldots, X_n = x_n) = \prod_{i=1}^{n} p_\theta(x_i).$$

The product form is independence: the joint PMF of independent observations is the product of the individual PMFs. We write L for L_n when n is clear, and $L(\theta)$ when the data are fixed.

In words. Take the data you observed and a candidate θ, and compute the probability that the model with parameter θ would produce exactly those data. That number, for every candidate at once, is the likelihood. The data are held fixed; the parameter varies.

For the kiss data. Each couple contributes p if it turned right and $1-p$ if it turned left, so for the observed 80 right turns and 44 left turns

$$L(p) = p^{80}(1-p)^{44}.$$

At $p = 0.645$ this is 9.4×10^{-36}; at $p = \frac{1}{2}$ it is 4.7×10^{-38}. Both are tiny, and they should be: there are 2^{124} possible sequences of turns, and any one sequence with 80 right turns is improbable under every p, its probability never more than 9.4×10^{-36}. What carries information is the ratio. The observed sequence is about 200 times more probable under $p = 0.645$ than under $p = \frac{1}{2}$, and $\log 200 = 5.30$, the evidence in nats that Chapter 10 found.

Intuition: one formula, two readings

You have written $\mathbb{P}_\theta(X = x)$ many times, always as a function of x with θ held fixed: the PMF. The same expression is a function of both x and θ, and the likelihood is what you get by switching which one you hold fixed. Figure 11.1 makes the switch visible for the kiss model. The left of the figure shades $\mathbb{P}_p(\text{count} = k)$ over a grid of counts k and parameters p. Read it along a horizontal line, p fixed: you get the distribution of the count, the probability question of Chapter 1, which sums to 1 over k. Read it along a vertical line, the count fixed at the 80 you observed: you get the likelihood, the statistics question, a curve over p that peaks where the data are most probable. Probability fixes the parameter and asks about the data; statistics fixes the data and asks about the parameter.

Pitfall: the likelihood is not a distribution for θ

The likelihood curve in Figure 11.1 looks like a density for p, and it isn't one. The curve drawn there (the version for the count) has area 1/125, not 1, and the sequence likelihood $p^{80}(1-p)^{44}$

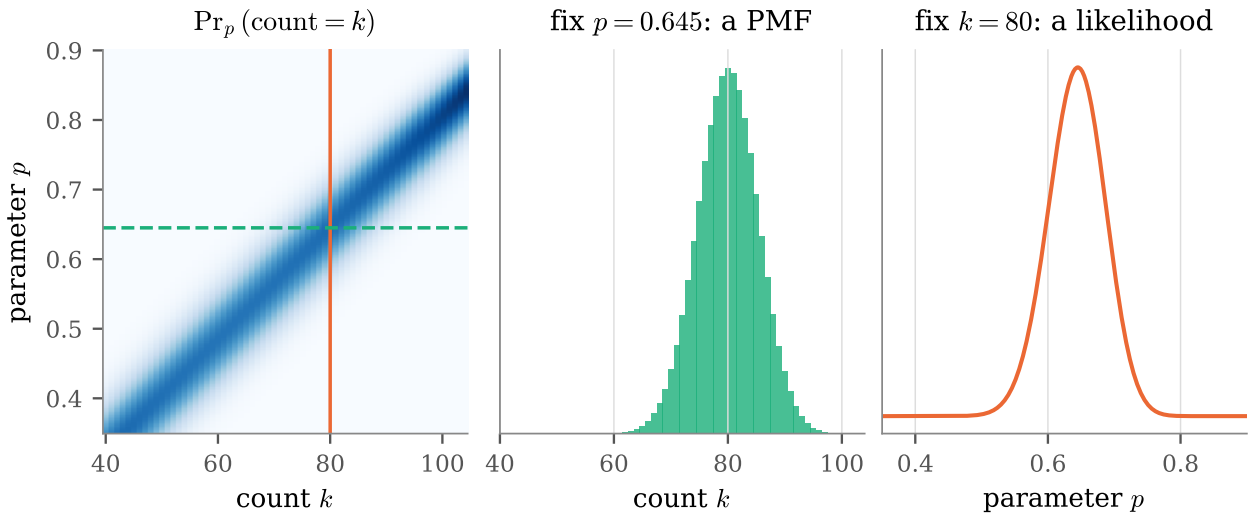


Figure 11.1: One function, two readings, for 124 couples. Left: $\mathbb{P}_p(\text{count} = k)$ over counts k (horizontal) and parameters p (vertical); darker is more probable. Middle: the slice along the dashed aqua line, $p = 0.645$, a PMF over k that sums to 1. Right: the slice along the orange line, the observed $k = 80$, which is the likelihood; it peaks at $80/124 = 0.645$. The likelihood of the full sequence of 124 turns is this curve divided by the constant $\binom{124}{80}$, which doesn't change where it peaks.

has area 1.0×10^{-36}; nothing in the construction makes the area mean anything: rescale it by any constant and every comparison it supports is unchanged. Until Chapter 18, p is a fixed unknown number, not a random variable, so 'the probability that p is between 0.6 and 0.7' has no meaning here. The likelihood is a plausibility ranking of candidate truths by how well they predicted your data. This is why the right of Figure 1.1 in Chapter 1, which is this kind of curve for the drug data scaled to peak at 1, has no vertical scale that could be read as a probability. That chapter is where the parameter gets a distribution, and the likelihood is one of its two ingredients.

Worked example: the likelihoods of the two discrete families

- **Bernoulli.** For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$, $E = \{0, 1\}$ and $\Theta = (0, 1)$. The PMF can be written in one line as $p_p(x) = p^x(1-p)^{1-x}$, which is p at $x = 1$ and $1 - p$ at $x = 0$. Then

$$L(x_1, \ldots, x_n; p) = \prod_{i=1}^{n} p^{x_i}(1-p)^{1-x_i} = p^{\sum_i x_i}(1-p)^{n-\sum_i x_i}.$$

  The data enter only through the number of ones, $\sum_i x_i$.

- **Poisson.** For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$, $E = \{0, 1, 2, \ldots\}$ and $\Theta = (0, \infty)$:

$$L(x_1, \ldots, x_n; \lambda) = \prod_{i=1}^{n} e^{-\lambda}\frac{\lambda^{x_i}}{x_i!} = e^{-n\lambda}\frac{\lambda^{\sum_i x_i}}{x_1! \cdots x_n!}.$$

  For the siblings data of Chapter 2, modeled as $X_i - 1 \sim \text{Pois}(\lambda)$, the 24 values of $x_i - 1$ add up to 24, and the likelihood is $e^{-24\lambda}\lambda^{24}$ divided by a constant that doesn't involve λ.

Three properties of a likelihood follow straight from the product.

- **Order doesn't matter.** A product can be taken in any order, so the likelihood of $(1, 2, 3)$ is the likelihood of $(2, 3, 1)$. For the Bernoulli and Poisson families it depends on the data only through their sum, up to a factor that does not involve the parameter; the sum is the sufficient statistic of Chapter 7: the likelihood never looks at anything else in the data.
- **It updates one observation at a time.** If L_n is the likelihood after n observations, then $L_{n+1}(\theta) = L_n(\theta) \cdot p_\theta(x_{n+1})$. For Poisson data $1, 2$ the likelihood is $e^{-2\lambda}\lambda^3/2$; a third observation 3 multiplies it by $e^{-\lambda}\lambda^3/6$, giving $e^{-3\lambda}\lambda^6/12$ (Exercise 11.1).
- **It is as smooth in θ as the PMF is.** If each $p_\theta(x)$ is a continuous function of θ, so is the product. A discrete model can have a perfectly smooth likelihood: the data are discrete, the parameter is not.

11.3 The likelihood of a continuous model

For continuous data the probability of observing exactly $x_1, \dots, x_n$ is 0 under every candidate, which would make every likelihood 0 and rank nothing. The fix is to use densities.

Definition 11.2 (Likelihood, continuous model)**.** Let $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ be a statistical model for $X_1, \dots, X_n$ i.i.d., where each $\mathbb{P}_\theta$ has a density f_θ (toolkit A.6). The **likelihood** of the model is

$$L_n : E^n \times \Theta \to \mathbb{R}$$
$$(x_1, \dots, x_n; \theta) \mapsto \prod_{i=1}^{n} f_\theta(x_i).$$

In words. The same recipe as for a discrete model, with each observation contributing the height of its density instead of its probability. A candidate whose density is high where your data fell scores well; one whose density is low there scores badly.

For the kiss data. The kiss data are discrete, so the discrete definition applies and this one doesn't. The two definitions are one idea, and from here on the chapter writes p_θ for either a PMF or a density.

Intuition: why the density is enough

A recorded measurement is never exact. A subway gap written down as 5 minutes was somewhere near 5, say within $h/2$ of it for a small h. The probability of that is the area of a thin strip under the density, about $f_\lambda(5) \cdot h$. Compare two candidate rates, and the h cancels:

$$\frac{\mathbb{P}_{\lambda_1}(\text{gap within } h/2 \text{ of } 5)}{\mathbb{P}_{\lambda_2}(\text{gap within } h/2 \text{ of } 5)} \approx \frac{f_{\lambda_1}(5)\,h}{f_{\lambda_2}(5)\,h} = \frac{f_{\lambda_1}(5)}{f_{\lambda_2}(5)}.$$

Figure 11.2 shows the ratio of strip probabilities settling at the ratio of densities as the strip narrows. Every use of a likelihood is a comparison between parameters, so the density is all it needs.

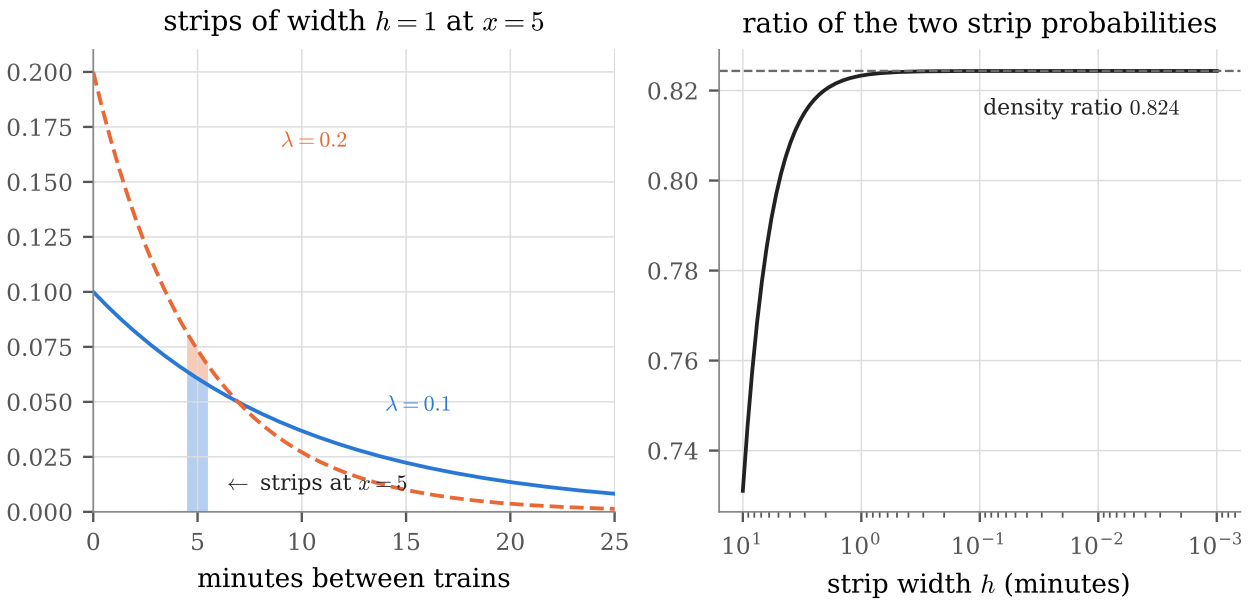


Figure 11.2: Why a density can stand in for a probability. Left: the Expo(0.1) and Expo(0.2) densities, with the probabilities of a gap within $\frac{1}{2}$ minute of 5 shown as strips; the Expo(0.1) strip is blue, and the taller Expo(0.2) strip contains it, its extra height in orange. Right: the ratio of the two strip probabilities as the width h shrinks; it settles at the ratio of the densities, $f_{0.1}(5)/f_{0.2}(5) = 0.824$.

Worked example: the likelihoods of the continuous families

- **Gaussian.** For $X_1, \dots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$, $E = \mathbb{R}$ and $\Theta = \mathbb{R} \times (0, \infty)$:

$$L(x_1, \dots, x_n; \mu, \sigma^2) = \prod_{i=1}^{n} \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$$
$$= \frac{1}{(\sigma\sqrt{2\pi})^n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^{n} (x_i - \mu)^2\right).$$

- **Exponential.** For $X_1, \dots, X_n \overset{iid}{\sim} \text{Expo}(\lambda)$, $E = (0, \infty)$ and $\Theta = (0, \infty)$. The density is $f_\lambda(x) = \lambda e^{-\lambda x} \mathbb{1}(x > 0)$, so

$$L(x_1, \dots, x_n; \lambda) = \lambda^n e^{-\lambda \sum_i x_i} \prod_{i=1}^{n} \mathbb{1}(x_i > 0) = \lambda^n e^{-\lambda \sum_i x_i} \mathbb{1}\big(\min_i x_i > 0\big).$$

  The product of indicators is 1 exactly when every x_i is positive, which is when the smallest one is (toolkit A.4). The indicator doesn't involve λ, and data from the model are always positive, so it is usually dropped: $L(\lambda) = \lambda^n e^{-\lambda \sum_i x_i}$.

- **Uniform.** For $X_1, \dots, X_n \overset{iid}{\sim} \text{Unif}([0, \theta])$, $E = [0, \infty)$ and $\Theta = (0, \infty)$. The density is $f_\theta(x) = \frac{1}{\theta} \mathbb{1}(0 \le x \le \theta)$, so

$$L(x_1, \dots, x_n; \theta) = \frac{1}{\theta^n} \prod_{i=1}^{n} \mathbb{1}(0 \le x_i \le \theta) = \frac{1}{\theta^n} \mathbb{1}\big(\min_i x_i \ge 0\big)\, \mathbb{1}\big(\max_i x_i \le \theta\big).$$

Every $x_i \le \theta$ holds exactly when the largest does. Here the indicator involving the maximum *does* depend on θ, and it cannot be dropped: it is what makes this likelihood unusual (Section 11.6). The other indicator doesn't involve θ and is dropped as before: $L(\theta) = \theta^{-n}\mathbb{1}(\theta \ge \max_i x_i)$.

- **Mixture of two Gaussians.** For the model of Section 2.7, with weight π and components $\text{Gauss}(\mu_1, \sigma_1^2)$ and $\text{Gauss}(\mu_2, \sigma_2^2)$, the density of one observation is a sum, and

$$L = \prod_{i=1}^{n} \left(\frac{\pi}{\sigma_1\sqrt{2\pi}} e^{-(x_i-\mu_1)^2/(2\sigma_1^2)} + \frac{1-\pi}{\sigma_2\sqrt{2\pi}} e^{-(x_i-\mu_2)^2/(2\sigma_2^2)} \right).$$

A product of sums does not simplify, and that is the source of the trouble in Section 11.8.

Pitfall: an indicator you drop must not involve the parameter

Dropping a factor from a likelihood is safe exactly when the factor doesn't depend on θ: it multiplies every candidate's score by the same number and changes no comparison. The exponential's $\mathbb{1}(\min_i x_i > 0)$ is such a factor. The uniform's $\mathbb{1}(\max_i x_i \le \theta)$ is not, and dropping it leaves θ^{-n}, which grows without bound as $\theta \to 0$: the best candidate would be one under which your data were impossible.

11.4 The maximum likelihood estimator

Definition 11.3 (Maximum likelihood estimator)**.** Let $X_1, \dots, X_n$ be an i.i.d. sample from a statistical model $(E, (\mathbb{P}_\theta)_{\theta\in\Theta})$ with likelihood L. The **maximum likelihood estimator** (MLE) of θ is

$$\hat{\theta}_n^{\text{MLE}} = \operatorname*{argmax}_{\theta\in\Theta} L(X_1, \dots, X_n; \theta),$$

provided the maximum exists and is attained at a single point.

In words. Among all the candidate truths the model allows, pick the one under which the data you actually got were most probable. Before the data arrive, $\hat{\theta}_n^{\text{MLE}}$ is a random variable, a function of $X_1, \dots, X_n$ like every estimator in Chapter 7; once they arrive it is a number, the maximum likelihood estimate.

For the kiss data. $L(p) = p^{80}(1-p)^{44}$ is largest at $p = 80/124 = 0.645$ (Section 11.6 does the calculus), so the maximum likelihood estimate is the sample proportion. In general $\hat{p}^{\text{MLE}} = \bar{R}_n$, the estimator Chapter 1 started from because it was the obvious one. In models where no estimator is obvious, the same principle still produces one.

In practice you maximize the logarithm of the likelihood, the **log-likelihood**

$$\ell_n(\theta) = \log L(X_1, \dots, X_n; \theta) = \sum_{i=1}^{n} \log p_\theta(X_i).$$

Because log is strictly increasing, L and ℓ_n have their maximum at the same θ (toolkit A.10). There are three reasons to prefer ℓ_n. It turns a product of n factors into a sum of n terms, and sums are easier to differentiate. It keeps the numbers in a sane range: $L(0.645) = 9.4 \times 10^{-36}$ for the

kiss data, and for a few thousand observations a likelihood is smaller than any number a computer stores, while $\ell_n(0.645) = -80.65$ is an ordinary number. And $\frac{1}{n}\ell_n(\theta)$ is an average, which is what the Law of Large Numbers and the Central Limit Theorem know how to handle; Chapter 12 uses that.

Figure 11.3 shows both curves for the kiss data. The likelihood is a narrow peak at 0.645 on a scale of 10^{-36}. The log-likelihood is a smooth hill whose top is at the same place, flat there, and 5.30 higher than its value at $p = \frac{1}{2}$.

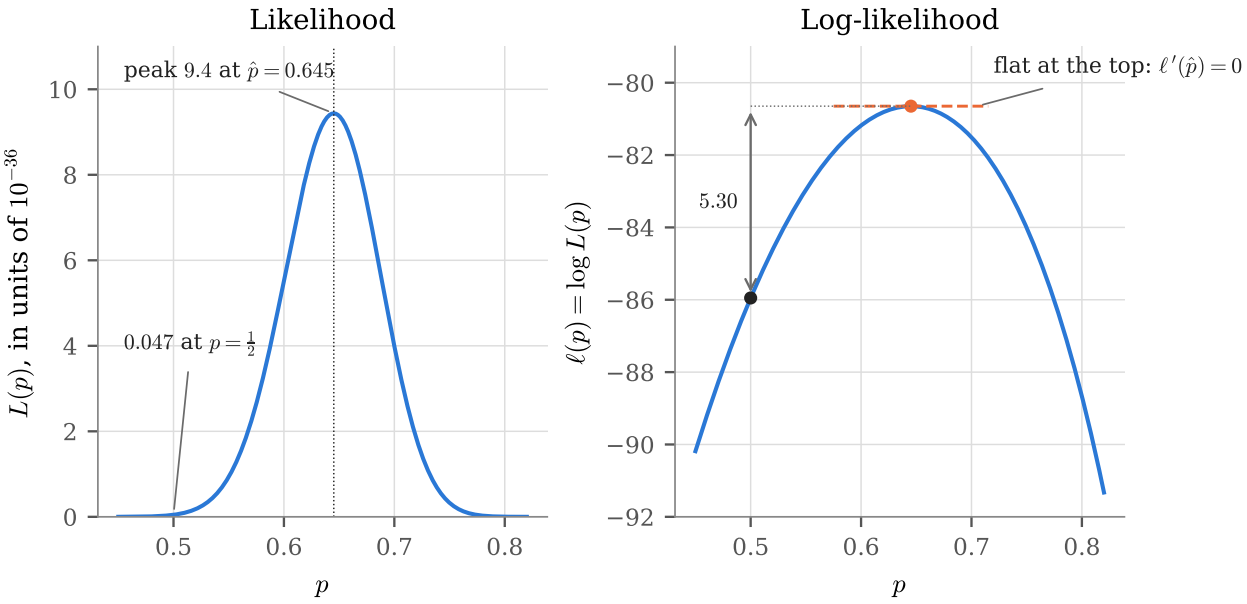


Figure 11.3: The kiss data. Left: the likelihood $L(p) = p^{80}(1-p)^{44}$ in units of 10^{-36}, with its peak 9.4 at $\hat{p} = 0.645$ and the value 0.047 at $p = \frac{1}{2}$. Right: the log-likelihood $\ell(p) = 80\log p + 44\log(1-p)$, concave, with its maximum -80.65 at 0.645 where the tangent is flat, and the value -85.95 at $\frac{1}{2}$; the gap, 5.30, is 124 times the KL divergence 0.0428 of Chapter 10.

Does this make sense? The sharpness of the peak

How sharply the log-likelihood curves down at its top should say something about how precisely the data pin down p: a flat top means many values of p fit about equally well. Check it. For the kiss data $\ell(p) = 80\log p + 44\log(1-p)$, so $\ell''(p) = -80/p^2 - 44/(1-p)^2$, and at $\hat{p}$ this is $-n/(\hat{p}(1-\hat{p})) = -541.7$. Then $1/\sqrt{541.7} = 0.0430$, which is exactly the standard error $\sqrt{\hat{p}(1-\hat{p})/n}$ of Chapter 8, since $\ell''(\hat{p}) = -n/(\hat{p}(1-\hat{p}))$. For the subway data the same computation gives $\hat{\lambda}/\sqrt{n} = 0.020$, the standard error of Chapter 5. It is not a coincidence, and Chapter 12 proves it in general: the curvature of the log-likelihood at its peak estimates n times the Fisher information, and one over its square root is the standard error of the MLE.

11.5 Finding a maximum

The definition says 'argmax', and an argmax has to be found. Maximizing an arbitrary function is hard: it can have many hills, and a point where the slope is zero can be the top of a small hill, the bottom of a valley or neither. The right of Figure 11.4 is $\theta \mapsto \prod_i(\theta - x_i)$ for five numbers x_i: four

points with zero slope, and on the whole line no maximum at all. The functions that are easy to maximize are the concave ones.

Definition 11.4 (Concave and convex functions)**.** A twice-differentiable function $h : \Theta \to \mathbb{R}$ on an interval $\Theta \subset \mathbb{R}$ is **concave** if $h''(\theta) \leq 0$ for every $\theta \in \Theta$, and **strictly concave** if $h''(\theta) < 0$ for every θ. (The 'strictly' version here is stronger than the usual chord definition, which also allows $h'' = 0$ at isolated points, as for $-\theta^4$; it is the version this book uses.) It is (strictly) **convex** if $-h$ is (strictly) concave, that is, $h'' \geq 0$ ($h'' > 0$).

> **In words.** A concave function's slope never increases: it is a hill, or a straight line, never a valley. Every chord lies below the graph (toolkit A.8). On their natural domains, $-\theta^2$, $\sqrt{\theta}$, $\log\theta$ and $\sin\theta$ on $[0, \pi]$ are concave; $2\theta - 3$ is both concave and convex; θ^2 and e^θ are convex.
>
> **For the kiss data.** $\ell''(p) = -80/p^2 - 44/(1-p)^2$ is negative for every $p \in (0, 1)$, so the kiss log-likelihood is strictly concave: one hill, no valleys, as the right of Figure 11.3 shows.

Proposition 11.5 (Maximizing a strictly concave function)**.** *Let h be strictly concave and differentiable on an open interval Θ. If $h'(\hat\theta) = 0$ at some $\hat\theta \in \Theta$, then $\hat\theta$ is the unique maximizer of h on Θ. If h' never vanishes, h has no maximum on Θ.*

Proof strategy (for completeness; skip on a first reading). Strict concavity makes the slope strictly decreasing, so it can cross zero at most once, and on each side of the crossing the function moves away from the top.

Proof. Since $h'' < 0$, h' is strictly decreasing. If $h'(\hat\theta) = 0$, then $h' > 0$ to the left of $\hat\theta$ and $h' < 0$ to the right: h increases up to $\hat\theta$ and decreases after, so $h(\theta) < h(\hat\theta)$ for every $\theta \neq \hat\theta$. At an interior maximizer the derivative must be 0, since otherwise moving a little in the direction of the slope would increase h; so if h' never vanishes there is no interior maximizer, and on an open interval there is no other kind. □

So the recipe for a one-parameter model is: write ℓ_n, check that it is strictly concave, solve $\ell_n'(\theta) = 0$. The Bernoulli, Poisson and exponential log-likelihoods of Section 11.6 pass the check; the Gaussian needs a two-step argument instead; the uniform fails it entirely.

For a parameter with several coordinates, $\theta = (\theta_1, \ldots, \theta_d)$, the derivative becomes the **gradient** $\nabla\ell_n(\theta)$, the vector of partial derivatives, and the second derivative becomes the **Hessian** $\nabla^2\ell_n(\theta)$, the $d \times d$ matrix of second partial derivatives (toolkit A.12). The first-order condition is $\nabla\ell_n(\hat\theta) = 0$, a system of d equations. A sufficient condition for strict concavity (on a Θ that contains the segment between any two of its points) is a negative definite Hessian at every θ, which for $d = 2$ means its top-left entry is negative and its determinant positive. Under it, the solution of $\nabla\ell_n = 0$, if there is one, is the unique maximizer, as in one dimension. When there is no formula for the solution, numerical methods find it; for a concave ℓ_n they are reliable, and the subject of convex optimization is how to do it fast.

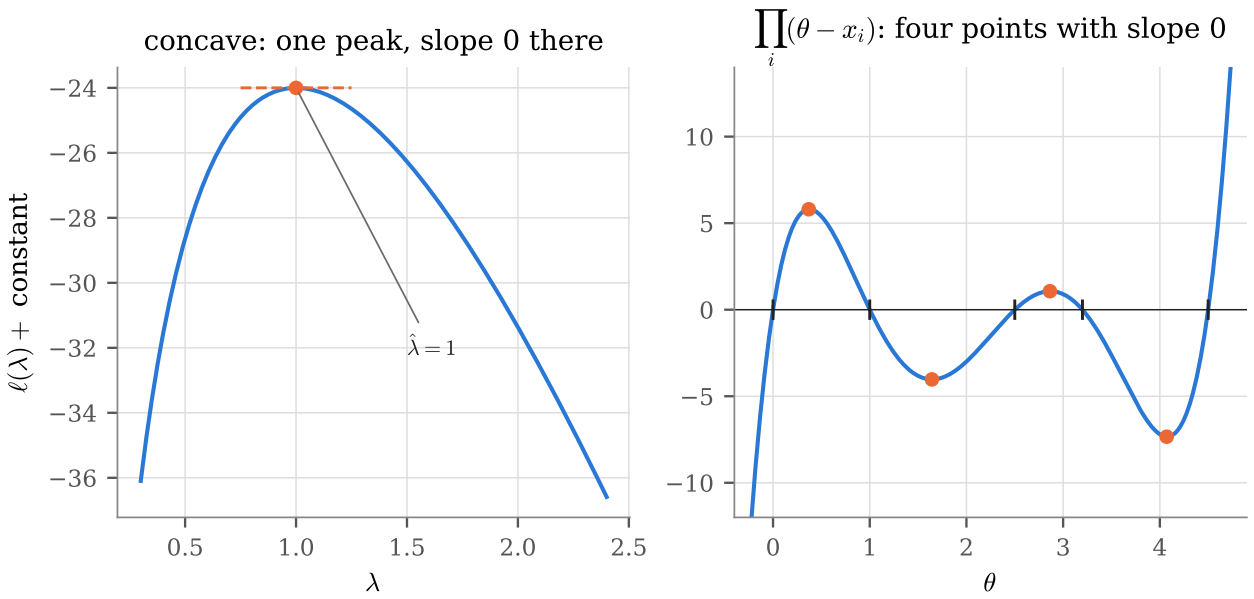


Figure 11.4: Easy and hard functions to maximize. Left: the Poisson log-likelihood of the siblings data, $-24\lambda + 24\log\lambda$ up to a constant, strictly concave with its single maximum at $\hat{\lambda} = 1$, where the slope is 0. Right: $\prod_i(\theta - x_i)$ for the five values marked on the axis; four points have slope 0 (dots), and setting the derivative to 0 finds all four without saying which, if any, is a maximum.

11.6 Maximum likelihood estimators of the standard families

Worked example: Bernoulli: the kiss data

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Bern}(p)$ with $S = \sum_i X_i$ ones,

$$\ell_n(p) = S\log p + (n - S)\log(1 - p), \quad \ell_n'(p) = \frac{S}{p} - \frac{n - S}{1 - p}, \quad \ell_n''(p) = -\frac{S}{p^2} - \frac{n - S}{(1 - p)^2} < 0.$$

Setting $\ell_n'(p) = 0$ gives $S(1 - p) = (n - S)p$, so $p = S/n$:

$$\hat{p}^{\text{MLE}} = \bar{X}_n.$$

(If $S = 0$ or $S = n$ there is no solution in $(0, 1)$; Section 11.7 deals with that.) For the kiss data, $\hat{p} = 80/124 = 0.645$, the estimate of every earlier chapter.

Worked example: Poisson: the siblings data

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$,

$$\ell_n(\lambda) = -n\lambda + \Big(\sum_i X_i\Big)\log\lambda - \sum_i \log(X_i!), \quad \ell_n'(\lambda) = -n + \frac{\sum_i X_i}{\lambda}, \quad \ell_n''(\lambda) = -\frac{\sum_i X_i}{\lambda^2} < 0$$

(when at least one X_i is positive). The first-order condition gives

$$\hat{\lambda}^{\text{MLE}} = \bar{X}_n.$$

For the siblings data, with $X_i - 1 \sim \text{Pois}(\lambda)$, the sample mean of the $x_i - 1$ is $24/24 = 1$, so $\hat{\lambda} = 1.0$: the fit drawn in Figure 2.1, and the left of Figure 11.4.

Worked example: exponential: the subway data

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Expo}(\lambda)$,

$$\ell_n(\lambda) = n \log \lambda - \lambda \sum_i X_i, \quad \ell_n'(\lambda) = \frac{n}{\lambda} - \sum_i X_i, \quad \ell_n''(\lambda) = -\frac{n}{\lambda^2} < 0,$$

and

$$\hat{\lambda}^{\text{MLE}} = \frac{n}{\sum_i X_i} = \frac{1}{\bar{X}_n}.$$

For the 64 subway gaps averaging 6.23 minutes, $\hat{\lambda} = 0.1605$ per minute, the estimator of Chapter 5. Chapter 7 showed it is biased upward, $\mathbb{E}[1/\bar{T}_n] = \frac{n}{n-1}\lambda$: a maximum likelihood estimator need not be unbiased.

Worked example: Gaussian: the boarding data

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$ the parameter has two coordinates, and

$$\ell_n(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^{n} (X_i - \mu)^2.$$

Maximize in μ first, for a fixed σ^2. The only term involving μ is $-\frac{1}{2\sigma^2} \sum_i (X_i - \mu)^2$, and the split

$$\sum_{i=1}^{n} (X_i - \mu)^2 = \sum_{i=1}^{n} (X_i - \bar{X}_n)^2 + n(\bar{X}_n - \mu)^2$$

(write $X_i - \mu = (X_i - \bar{X}_n) + (\bar{X}_n - \mu)$, square, and note the cross terms add to 0) shows it is largest at $\mu = \bar{X}_n$, whatever σ^2 is. Now put $\mu = \bar{X}_n$ and maximize in σ^2. With $D = \sum_i (X_i - \bar{X}_n)^2$, the function $-\frac{n}{2} \log \sigma^2 - D/(2\sigma^2)$ has derivative $-\frac{n}{2\sigma^2} + \frac{D}{2\sigma^4}$, positive for $\sigma^2 < D/n$ and negative after. Chaining the two steps, for every (μ, σ^2),

$$\ell_n(\mu, \sigma^2) \le \ell_n(\bar{X}_n, \sigma^2) \le \ell_n(\bar{X}_n, D/n),$$

with equality only at $(\bar{X}_n, D/n)$, so

$$\hat{\mu}^{\text{MLE}} = \bar{X}_n, \qquad \widehat{\sigma^2}^{\text{MLE}} = \frac{1}{n} \sum_{i=1}^{n} (X_i - \bar{X}_n)^2.$$

At the maximum the gradient is 0 and the Hessian is the diagonal matrix with entries $-n/\widehat{\sigma^2}$ and $-n/(2\widehat{\sigma^2}^2)$, negative definite. (The log-likelihood is not concave in σ^2 everywhere, which is why the argument went one coordinate at a time instead of checking the Hessian at every point. It also needs $D > 0$: if every observation is the same number, the likelihood grows without bound as $\sigma^2 \to 0$.)

For the 72 flights boarded rear to front, with mean 24.2 minutes and standard deviation $s = 2.1$ (taken to be computed with $n - 1$, the usual convention), $\hat{\mu} = 24.2$ and $\widehat{\sigma^2} = \frac{71}{72} \times 2.1^2 = 4.35$. Figure 11.5 draws the log-likelihood surface.

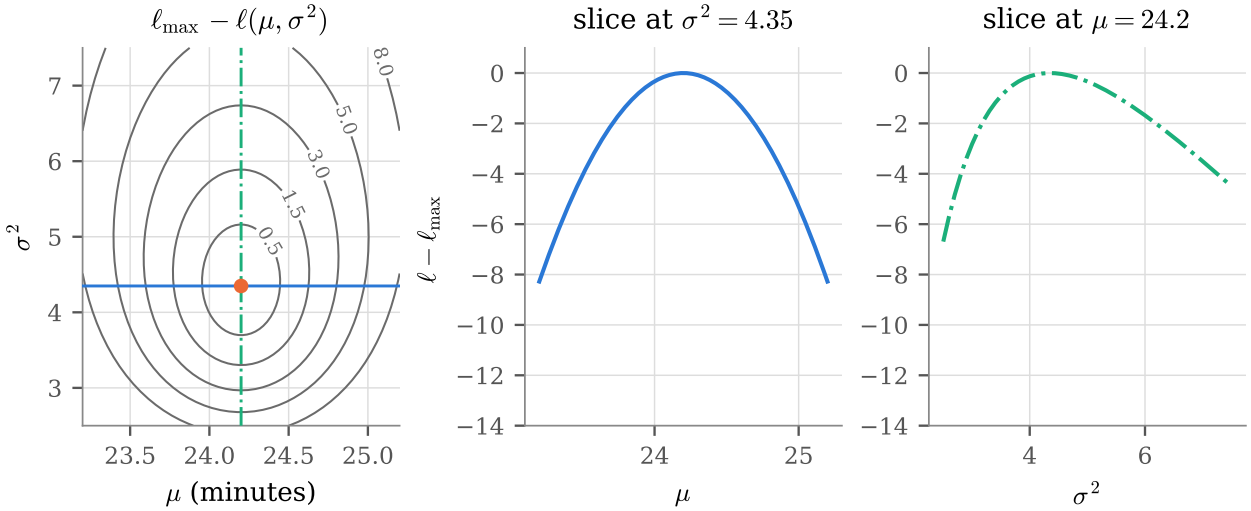


Figure 11.5: The Gaussian log-likelihood of the boarding data ($n = 72$, mean 24.2, standard deviation 2.1), measured from its maximum. Left: contours over (μ, σ^2), labeled by how far each lies below the maximum, with the maximum at $(24.2, 4.35)$ and the two slices drawn as lines (blue solid, aqua dash-dot). Middle: the slice at $\sigma^2 = 4.35$, symmetric in μ. Right: the slice at $\mu = 24.2$, lopsided: variances above the maximum are penalized less than variances below it.

From undergraduate statistics: n or $n-1$?

The maximum likelihood estimator of the variance divides by n. The sample variance you were taught divides by $n-1$. They differ by the factor $(n-1)/n$, which is 0.986 for the boarding data: 4.35 against 4.41. Chapter 7 showed that dividing by n is biased low by exactly that factor and that dividing by $n-1$ removes the bias; it also showed that removing the bias is not the same as minimizing the risk. Maximum likelihood doesn't aim at unbiasedness, so it lands on n. For large n the difference is negligible, and every large-sample result in Chapter 12 holds for either.

From undergraduate statistics: least squares is maximum likelihood

The Gaussian log-likelihood contains $-\frac{1}{2\sigma^2}\sum_i (X_i - \mu)^2$, so maximizing it over μ is minimizing a sum of squared deviations. The same is true when μ is replaced by a regression line $a + bx_i$: if the errors around the line are Gaussian with a common variance, the maximum likelihood estimates of a and b are the least squares estimates of an introductory regression course. Chapter 20 starts from that observation.

Worked example: uniform: where calculus fails

For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, \theta])$ the likelihood is $L(\theta) = \theta^{-n} \mathbb{1}(\theta \geq \max_i X_i)$. Below the largest observation the likelihood is 0 (a θ smaller than an observation makes that observation impossible). From the largest observation on it is θ^{-n}, which decreases, with derivative $-n\theta^{-n-1}$, never 0. The maximum is at the jump:

$$\hat{\theta}^{\text{MLE}} = \max_i X_i.$$

Figure 11.6 draws it for five observations whose largest is 7.6. Setting the derivative to 0 would have found only the flat stretch below $\max_i X_i$, where the likelihood is 0; the recipe

of Section 11.5 assumes a smooth function with its maximum inside Θ, and this likelihood is neither smooth at its maximum nor concave. The estimator is also biased: the largest of n draws always falls short of θ, and $\mathbb{E}[\max_i X_i] = \frac{n}{n+1}\theta$ (Exercise 11.4). Multiplying by $\frac{n+1}{n}$ removes the bias.

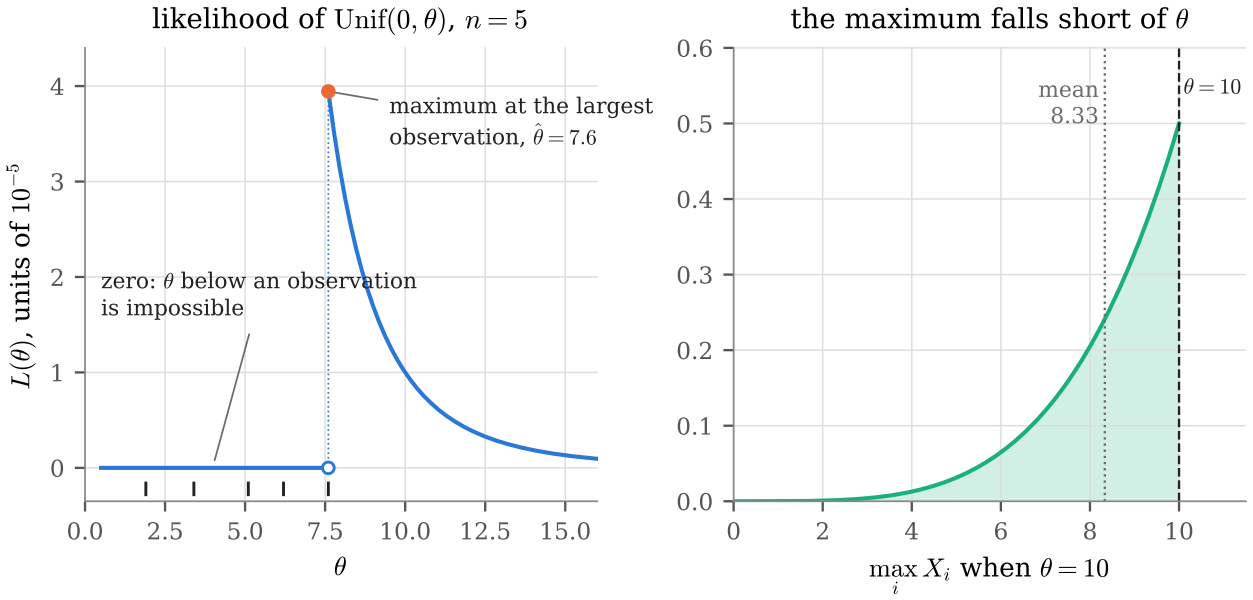


Figure 11.6: Left: the likelihood of Unif$([0,\theta])$ for five observations (ticks below the axis), zero until the largest observation, 7.6, then θ^{-5}, decreasing. The maximum is at the jump, where no derivative is zero. Right: the density $5m^4/10^5$ of the maximum of five draws when $\theta = 10$; it never exceeds 10, and its mean is 8.33.

Intuition: the German tank problem

The uniform model once estimated an enemy's tank production. During the Second World War, Allied statisticians estimated German tank production from the serial numbers on captured tanks. If serial numbers run $1, 2, \ldots, N$ and the captured ones are a random sample, N plays the role of θ, the largest serial number seen is the maximum likelihood estimate, and correcting it for the shortfall of the maximum gives a better one. Their serial-number estimates, which refined this idea and used markings on several components (chassis, gearboxes, wheels), for three months, beside the figures from conventional intelligence and the German records found after the war:[a]

Month	Statistical estimate	Intelligence estimate	German records
June 1940	169	1,000	122
June 1941	244	1,550	271
August 1942	327	1,550	342

The model is crude (serial numbers are not quite a random sample), and its estimates were within 50 tanks of the records while the intelligence estimates were off by a factor of about four to eight.

[a] R. Ruggles and H. Brodie, 'An empirical approach to economic intelligence in World War II', *Journal of the American Statistical Association* 42 (1947), 72–91.

One more property makes maximum likelihood convenient: it doesn't care how the parameter is written.

Proposition 11.6 (Invariance)**.** *Let $g : \Theta \to g(\Theta)$ be one-to-one, and reparameterize the model by $\eta = g(\theta)$. If $\hat{\theta}$ is the MLE of θ, then $g(\hat{\theta})$ is the MLE of η.*

Proof strategy (for completeness; skip on a first reading). Reparameterizing relabels the candidates without changing any candidate's likelihood, so the winner is the same distribution under its new label.

Proof. The candidate labeled η in the new parameterization is the distribution $\mathbb{P}_{g^{-1}(\eta)}$, so its likelihood is $L(g^{-1}(\eta))$. This is largest when $g^{-1}(\eta) = \hat{\theta}$, that is, at $\eta = g(\hat{\theta})$. □

For the subway data, the mean gap $\beta = 1/\lambda$ has MLE $1/\hat{\lambda} = \bar{X}_n = 6.23$ minutes. For the boarding data, the standard deviation has MLE $\sqrt{4.35} = 2.09$ minutes. Invariance does not carry unbiasedness along: $\bar{X}_n$ is unbiased for β and $1/\bar{X}_n$ is biased for λ, though both are maximum likelihood estimators.

11.7 When the maximum likelihood estimator fails to exist or isn't unique

Definition 11.3 said 'provided the maximum exists and is attained at a single point'. Both can fail, in ordinary models.

Pitfall: three ways the definition can fail

- **No maximum: the data sit on the boundary.** If every observation in a Bernoulli sample is 0 ($S = 0$), then $L(p) = (1-p)^n$ increases as p decreases and has no maximum on $\Theta = (0, 1)$: the supremum is at $p = 0$, which the model excludes. For the kiss data this can't happen in practice (under $p = 0.645$ its probability is 0.355^{124}), but for $n = 5$ observations of an outcome with $p = 0.1$ it happens 59% of the time. The usual repair is to allow $p \in [0, 1]$, and then $\hat{p} = 0$.

- **More than one maximum.** For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([\theta, \theta + 1])$, $L(\theta) = \mathbb{1}(\theta \leq \min_i X_i)\,\mathbb{1}(\theta + 1 \geq \max_i X_i)$, which equals 1 on the whole interval $[\max_i X_i - 1, \min_i X_i]$ and 0 elsewhere. Every point of that interval is a maximizer (Exercise 11.5).

- **An unbounded likelihood.** In the mixture of two Gaussians with unknown variances, put $\mu_1 = x_1$ and let $\sigma_1 \to 0$. The first factor of the likelihood contains $\pi/(\sigma_1\sqrt{2\pi})$, which goes to infinity, while every other factor stays above $(1 - \pi)$ times the second component's density at its point, which is positive. The likelihood is unbounded: a component of zero width sitting on one data point 'explains' that point infinitely well. The maximum likelihood estimator of a mixture with free variances does not exist, and what is computed in practice is the best local maximum away from these spikes.

11.8 No formula: mixtures and the EM algorithm

The left of Figure 11.7 is a histogram of the 272 waiting times, in minutes, between eruptions of the Old Faithful geyser in Yellowstone National Park.[1] It has two humps, one near 55 minutes and one near 80, and a single Gaussian (orange) fits it badly. A mixture of two Gaussians is the natural model: each wait comes from a short-wait regime or a long-wait regime, and you don't get to see which.

The log-likelihood of the mixture,

$$\ell_n = \sum_{i=1}^{n} \log\Big(\pi\, f_{\mu_1,\sigma_1^2}(X_i) + (1-\pi)\, f_{\mu_2,\sigma_2^2}(X_i)\Big),$$

has the log of a sum in every term, and setting its gradient to 0 gives five equations with no closed-form solution. The **EM algorithm** (expectation-maximization, given its general form and its name in 1977 and still in everyday use) gets around this with one idea: pretend you know which regime each observation came from.

The complete data

Recall the generative story of Section 2.7: draw $Z_i \sim \text{Bern}(\pi)$; if $Z_i = 1$, draw X_i from component 1, otherwise from component 2. If you observed the labels Z_i along with the X_i, the log-likelihood of the *complete data* (X_i, Z_i) would split:

$$\ell_n^{\text{comp}} = \sum_{i=1}^{n} \Big[Z_i\big(\log\pi + \log f_{\mu_1,\sigma_1^2}(X_i)\big) + (1-Z_i)\big(\log(1-\pi) + \log f_{\mu_2,\sigma_2^2}(X_i)\big)\Big].$$

There is no log of a sum: each observation contributes to exactly one component. Maximizing it is two separate Gaussian fits and one Bernoulli fit, each done in Section 11.6: $\hat\pi$ is the share of observations labeled 1, and $\hat\mu_1$, $\hat\sigma_1^2$ are the mean and divide-by-n variance of the observations labeled 1 (and likewise for component 2).

The E step: guess the labels

You don't observe the Z_i. The best guess for Z_i given X_i, under the current parameter values, is its conditional expectation, which is the probability that X_i came from component 1. By Bayes' formula,

$$w_i = \mathbb{E}[Z_i \mid X_i] = \mathbb{P}(Z_i = 1 \mid X_i) = \frac{\pi\, f_{\mu_1,\sigma_1^2}(X_i)}{\pi\, f_{\mu_1,\sigma_1^2}(X_i) + (1-\pi)\, f_{\mu_2,\sigma_2^2}(X_i)} \in (0,1).$$

A wait of 50 minutes gets w_i close to 1; a wait of 85 gets w_i close to 0; a wait of 67, between the humps, gets something in the middle. Guessing $Z_i \approx \mathbb{E}[Z_i] = \pi$, the same for every observation, would ignore the data entirely; conditioning on X_i is what makes the guess useful.

[1] The version of the data distributed with W. Härdle, *Smoothing Techniques* (Springer, 1991); A. Azzalini and A. W. Bowman, 'A look at some data on the Old Faithful geyser', *Applied Statistics* 39 (1990), 357–365, analyze a longer record.

The M step: fit as if the guesses were labels

Replace each Z_i in ℓ_n^{comp} by w_i and maximize. This is legitimate because ℓ_n^{comp} is linear in each Z_i, so its expectation given the data, under the current parameter values, is ℓ_n^{comp} with each Z_i replaced by w_i. The same calculus as before, with weights, gives

$$\hat{\pi} = \frac{1}{n}\sum_i w_i, \quad \hat{\mu}_1 = \frac{\sum_i w_i X_i}{\sum_i w_i}, \quad \hat{\sigma}_1^2 = \frac{\sum_i w_i (X_i - \hat{\mu}_1)^2}{\sum_i w_i},$$

and the same for component 2 with weights $1 - w_i$ (Exercise 11.9). Each observation is split between the components in proportion to how likely it is to belong to each.

Worked example: the EM algorithm

1. Start from some values of $\pi, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2$.
2. **E step:** compute w_i for every observation from the current values.
3. **M step:** update all five parameters by the weighted formulas.
4. Repeat steps 2 and 3 until the parameters stop changing.

For Old Faithful, starting from means 50 and 70, both standard deviations 10 and $\pi = \frac{1}{2}$, the algorithm settles on weight 0.36 for Gauss$(54.6, 5.87^2)$ and 0.64 for Gauss$(80.1, 5.87^2)$: about a third of the waits are short, averaging 55 minutes, and the rest are long, averaging 80. The fitted mixture is the blue curve in Figure 11.7. Its log-likelihood is -1034.0, against -1095.3 for the best single Gaussian.

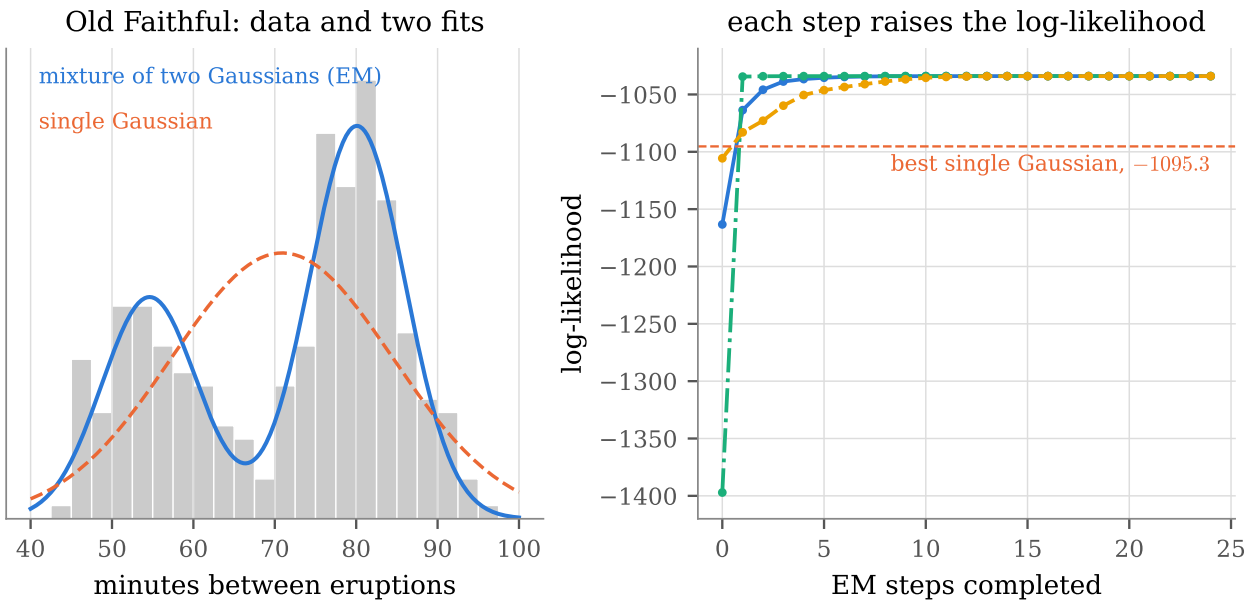


Figure 11.7: The EM algorithm on the 272 waiting times between eruptions of Old Faithful. Left: the histogram, the best single Gaussian (orange, dashed) and the two-component mixture fitted by EM (blue): weight 0.36 on Gauss$(54.6, 5.87^2)$ and 0.64 on Gauss$(80.1, 5.87^2)$. Right: the log-likelihood at the start and after each EM step, from three starting points: means 50 and 70 (blue, solid), 40 and 95 (aqua, dash-dot), 65 and 75 (yellow, dashed). It never decreases, and all three runs reach -1034.0, well above the single Gaussian's -1095.3 (orange line).

The right of Figure 11.7 shows the property that makes EM trustworthy as far as it goes: every step

raises the log-likelihood or leaves it unchanged. That fact is proved in Appendix C.2 of Appendix C, with an argument that uses the nonnegativity of the KL divergence; nothing later in the book depends on it. What it does not guarantee is that EM reaches the highest point. It climbs whatever hill it starts near, which may be a local maximum or one of the spikes of Section 11.7. In practice EM is run from several starting points and the run with the largest log-likelihood is kept. EM is guaranteed to climb, not to reach the top: the results that guarantee it reaches the maximum likelihood estimate assume a starting point close enough to it, or well-separated components.

Pitfall: which component is which

EM's labels 'component 1' and 'component 2' mean nothing. Start it with the means the other way round and it returns the same fit with the labels swapped: weight 0.64 on Gauss$(80.1, 5.87^2)$ and 0.36 on Gauss$(54.6, 5.87^2)$. Both have the same likelihood, since they are the same distribution. This is the label switching of Section 2.7: the mixture's parameter is identifiable only up to relabeling, so every peak of its likelihood has a twin with the labels swapped.

11.9 What comes next

This chapter defined the maximum likelihood estimator and computed it. Chapter 12 asks how good it is, and answers with three results that hold for every well-behaved model at once. It is consistent, by the argument Figure 10.5 drew: the average log-likelihood converges to a function maximized at the truth. It is asymptotically normal, by the Central Limit Theorem applied to the derivative of the log-likelihood. And its asymptotic variance is one over the Fisher information, the curvature that the sanity check of Section 11.4 found matching the standard errors of earlier chapters. Chapter 13 generalizes in two directions: the method of moments, an older and simpler recipe that often gives the same answer, and M-estimation, which maximizes or minimizes averages other than the log-likelihood.

11.10 Summary

- The likelihood $L(\theta) = \prod_i p_\theta(x_i)$ (PMFs) or $\prod_i f_\theta(x_i)$ (densities) is the probability or density of the observed data under each candidate θ: the joint PMF read with the data fixed and the parameter varying. It is not a distribution for θ; only ratios of likelihoods mean something. Kiss data: $L(p) = p^{80}(1-p)^{44}$, 200 times larger at 0.645 than at $\frac{1}{2}$.
- A density works in place of a probability because a recorded value stands for a small interval, whose width cancels in every comparison between parameters.
- The likelihood ignores the order of the data, updates one observation at a time, and is as smooth in θ as the PMF or density. Factors that don't involve θ (the exponential's indicator) can be dropped; factors that do (the uniform's $\mathbb{1}(\theta \geq \max_i x_i)$) cannot.
- MLE: $\hat{\theta}^{\text{MLE}} = \operatorname{argmax}_\theta L = \operatorname{argmax}_\theta \ell_n$, with $\ell_n = \sum_i \log p_\theta(X_i)$. If ℓ_n is strictly concave, the solution of $\ell_n' = 0$ (or $\nabla \ell_n = 0$) is the unique maximizer; in several dimensions, a negative definite Hessian everywhere is enough for strict concavity.
- Bernoulli $\bar{X}_n$ (0.645 for the kiss data); Poisson $\bar{X}_n$ (1.0 for the siblings); exponential $1/\bar{X}_n$ (0.1605 per minute for the subway); Gaussian $\bar{X}_n$ and $\frac{1}{n}\sum(X_i - \bar{X}_n)^2$ (24.2 and 4.35 for boarding); uniform $\max_i X_i$, found at a jump rather than where a derivative vanishes, and biased low by the factor $n/(n+1)$.

- The curvature of the log-likelihood at its peak gives the standard error: $1/\sqrt{541.7} = 0.043$ for the kiss data.
- Invariance: the MLE of $g(\theta)$ is $g(\hat{\theta})$ for one-to-one g. Unbiasedness is not invariant.
- The MLE can fail to exist (all-zero Bernoulli data on $(0,1)$; mixtures with free variances, whose likelihood is unbounded) or be non-unique ($\text{Unif}([\theta, \theta+1])$).
- EM for mixtures: E step $w_i = \mathbb{P}(Z_i = 1 \mid X_i)$ by Bayes' formula; M step, weighted proportions, means and variances. No step lowers the log-likelihood; it can stop at a local maximum, so run it from several starts. Old Faithful: weight 0.36 on mean 54.6, 0.64 on mean 80.1, log-likelihood -1034.0 against -1095.3 for one Gaussian.

11.11 Exercises

Exercise 11.1. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$. You observe $X_1 = 1$, $X_2 = 2$. Compute $L_2(1, 2; \lambda)$. A third observation $X_3 = 3$ arrives: compute $L_3(1, 2, 3; \lambda)$ in two ways, from the definition and by updating L_2. Then compute $L_3(2, 3, 1; \lambda)$ and explain why it is the same. *(Solution on page 228.)*

Exercise 11.2. Let $(\mathbb{Z}, (\mathbb{P}_\theta)_{\theta \in \mathbb{R}})$ be a discrete model whose PMF $p_\theta(x)$ is a continuous function of θ for each fixed x. True or false: for fixed data $x_1, \ldots, x_n$, the likelihood $\theta \mapsto L_n(x_1, \ldots, x_n; \theta)$ is continuous. Give a discrete model with a smooth likelihood. *(Solution on page 228.)*

Exercise 11.3. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(0, \sigma^2)$ with $\sigma^2 > 0$ unknown. Find the MLE of σ^2, check that the log-likelihood has a unique maximum, and use invariance to find the MLE of σ. *(Solution on page 229.)*

Exercise 11.4. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, \theta])$ and $M = \max_i X_i$. Show that $\mathbb{P}(M \le m) = (m/\theta)^n$ for $0 \le m \le \theta$, find the density of M, and show that $\mathbb{E}[M] = \frac{n}{n+1}\theta$. Which multiple of M is unbiased? *(Solution on page 229.)*

Exercise 11.5. Let $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([\theta, \theta+1])$ with $\theta \in \mathbb{R}$. Write the likelihood and show that every θ in $[\max_i X_i - 1, \min_i X_i]$ maximizes it. Why is that interval never empty? *(Solution on page 229.)*

Exercise 11.6. A basketball player takes shots until the first one goes in; the number of shots is $X \sim \text{Geom}(p)$, with $\mathbb{P}(X = k) = (1-p)^{k-1}p$ for $k = 1, 2, \ldots$. Over n independent rounds the numbers of shots are $X_1, \ldots, X_n$. Find the MLE of p, and evaluate it for rounds of 3, 1, 2, 5 and 1 shots. *(Solution on page 229.)*

Exercise 11.7. Let $X_1, \ldots, X_n$ be i.i.d. with density $f_a(x) = e^{-(x-a)}\mathbb{1}(x \ge a)$, an exponential shifted to start at an unknown $a \in \mathbb{R}$. Find the MLE of a. *(Solution on page 229.)*

Exercise 11.8. Let $X_1, \ldots, X_n$ be i.i.d. with the Laplace density $f_m(x) = \frac{1}{2}e^{-|x-m|}$, $m \in \mathbb{R}$, and n odd. Show that the MLE of m is the sample median. *(Solution on page 229.)*

Exercise 11.9. Derive the M-step formulas for π and μ_1: maximize $\sum_i [w_i(\log \pi + \log f_{\mu_1, \sigma_1^2}(X_i)) + (1 - w_i)(\log(1-\pi) + \log f_{\mu_2, \sigma_2^2}(X_i))]$ over π and over μ_1. *(Solution on page 229.)*

Exercise 11.10. In the Bernoulli model with $\Theta = (0, 1)$, show that the MLE fails to exist when all n observations are 0. If $p = 0.1$, how likely is that for $n = 5$ and for $n = 50$? *(Solution on page 230.)*

Chapter 12

Properties of the MLE: Consistency, Fisher Information and Asymptotic Normality

In this chapter

- Three facts that hold for the maximum likelihood estimator of every well-behaved model at once: it is consistent, it is asymptotically normal, and its asymptotic variance is one over a number you can compute from the model alone, the Fisher information.
- Why it is consistent: the average log-likelihood is an average, the Law of Large Numbers sends it to a fixed curve, and that curve peaks at the truth. The picture Chapter 10 drew, made into an argument.
- The score: the slope of one observation's log-likelihood. It averages to zero at the truth, and its variance is the Fisher information, which is also the curvature of the log-likelihood at its peak. Both definitions, and why they agree.
- The Fisher information of the Bernoulli, Poisson, exponential and Gaussian models, and what each says about how precisely its parameter can be learned.
- Why the MLE is asymptotically normal: the Central Limit Theorem applied to the score, then one step of the delta method's reasoning. The result is the interval $\hat{\theta} \pm 1.96/\sqrt{nI(\hat{\theta})}$, for any model.
- The Cramér–Rao bound: no unbiased estimator has smaller variance than $1/(nI(\theta))$, so the MLE's precision is the best available.
- What the theorem needs, and the models it does not cover: the uniform model, whose error shrinks like $1/n$, not $1/\sqrt{n}$.

Chapter 11 computed the maximum likelihood estimator. This chapter says how good it is, and makes good on two promises: the consistency argument sketched under Figure 10.5, and the curvature-equals-standard-error coincidence of the sanity check in Section 11.4.

12.1 Three questions about an estimator

Chapter 7 set out what to ask of an estimator: does it get to the truth as the sample grows (consistency), how fast and with what shape of error (asymptotic normality and the asymptotic variance), and how does it compare with the alternatives. For the sample proportion $\bar{R}_n$ those questions were answered by the Law of Large Numbers and the Central Limit Theorem, applied directly to an average. The maximum likelihood estimator of an arbitrary model is not an average. It is the argmax of a random function, and no theorem in Part II says anything about argmaxes.

This chapter's whole trick is to turn the argmax back into an average. It happens twice. For consistency, the thing averaged is the log-likelihood of one observation: $\frac{1}{n}\ell_n(\theta)$ is an average, and the LLN handles it. For asymptotic normality, the thing averaged is the slope of the log-likelihood of one observation, the score: $\frac{1}{n}\ell_n'(\theta)$ is an average, and the CLT handles it. Everything else is bookkeeping, with one Taylor expansion.

12.2 Consistency

Recall. The Law of Large Numbers (Theorem 3.5): if $X_1, \dots, X_n$ are i.i.d. and $h(X)$ has finite variance, then $\frac{1}{n}\sum_{i=1}^{n} h(X_i) \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}[h(X)]$.

Fix a candidate θ and look at the average log-likelihood,

$$\frac{1}{n}\ell_n(\theta) = \frac{1}{n}\sum_{i=1}^{n} \log p_\theta(X_i).$$

With $h = \log p_\theta$, the LLN says this converges in probability, under the truth $\mathbb{P}_{\theta^*}$, to

$$M(\theta) = \mathop{\mathbb{E}}_{\theta^*}\big[\log p_\theta(X)\big] = \mathop{\mathbb{E}}_{\theta^*}\big[\log p_{\theta^*}(X)\big] - \mathrm{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta),$$

the split Chapter 10 made. The first term does not depend on θ. The second is a KL divergence, which is 0 at $\theta = \theta^*$ and positive everywhere else if the model is identifiable (Proposition 10.6). So M is a fixed curve, not random, and it has a unique peak at the truth.

Figure 12.1 draws what happens. The random curve $\frac{1}{n}\ell_n$ settles onto the fixed curve M as n grows, and its peak, the MLE, slides toward the peak of M, the truth. At $n = 10$ the sample curve peaks at 0.80, far from $p^* = 0.645$; at $n = 1{,}000$ the two curves are indistinguishable and the peaks nearly coincide.

Intuition: the peaks follow the curves

The LLN works one θ at a time: for each fixed candidate, the average log-likelihood at that candidate converges. The theorem below needs more, that the whole random curve stays close to the whole fixed curve at once, so that the location of its peak converges too. Here is why that is plausible, and where it can fail. If two curves are within ϵ of each other everywhere, the peak of one is within ϵ of the peak of the other in height. If the fixed curve M drops off decisively away from its peak, as it does in Figure 12.1, a point far from θ^* has M well below $M(\theta^*)$, and the random curve there is also well below its value near θ^*, so the random peak cannot be far away. What can go wrong is a parameter space where M never drops off decisively (an infinite Θ with M flattening toward its edge), or a model where the random curve is close to M at

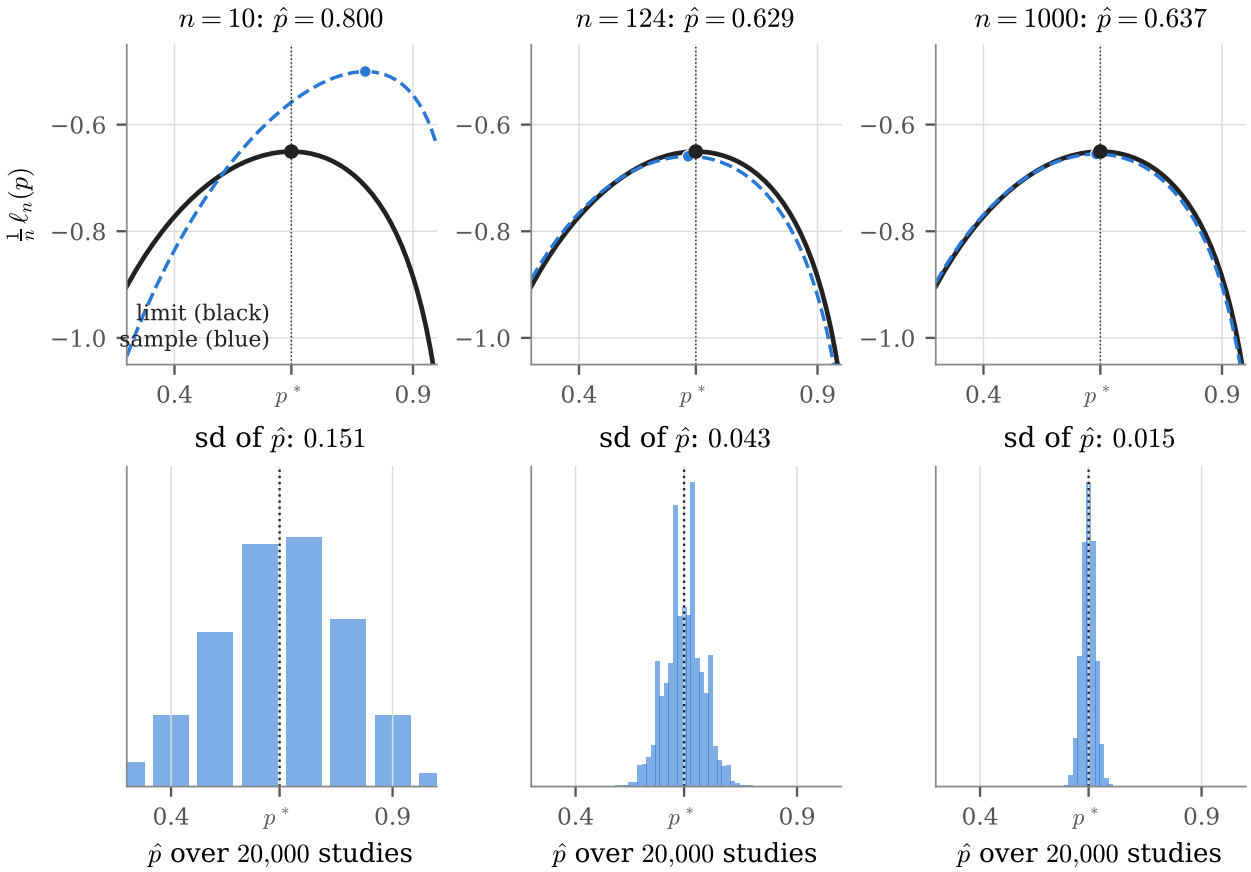


Figure 12.1: Consistency of the MLE in the kiss model, with the truth at $p^* = 0.645$. Top: the average log-likelihood of one simulated sample of n couples (blue, dashed), peaking at its own $\hat{p}$ (blue dot), against its limit $M(p) = p^* \log p + (1 - p^*) \log(1 - p)$ (black), which peaks at p^* (black dot). Bottom: the sampling distribution of $\hat{p}$ at the same n, from 20,000 simulated studies; its standard deviation $\sqrt{p^*(1 - p^*)/n}$ is 0.151, 0.043 and 0.015.

every single θ but not uniformly, with thin spikes that move around. The 'technical conditions' in the theorem rule those out, and once the parameter space is restricted to a bounded closed set around the truth, as the paragraph after the theorem explains, they hold for every model in this book except the ones Section 12.6 names.

Theorem 12.1 (Consistency of the MLE)**.** *Let* $X_1, \dots, X_n \overset{iid}{\sim} \mathbb{P}_{\theta^*}$ *in an identifiable model* $(E, (\mathbb{P}_\theta)_{\theta \in \Theta})$ *with* $\Theta \subset \mathbb{R}^d$*, and let* $\hat{\theta}_n$ *be the maximum likelihood estimator. Suppose that* Θ *is bounded and closed, that* $\log p_\theta(x)$ *is continuous in* θ *for every* x*, and that* $\mathbb{E}_{\theta^*}[\sup_{\theta \in \Theta} |\log p_\theta(X)|] < \infty$*. Then*

$$\hat{\theta}_n \xrightarrow[n \to \infty]{\mathbb{P}} \theta^* \quad \text{under } \mathbb{P}_{\theta^*}.$$

If the model is misspecified, so that the data come from a $\mathbb{P}$ *outside the family, the same conclusion holds with* θ^* *replaced by the KL projection* θ° *of Definition 10.9, provided it is unique.*

Proof strategy (core). Under the hypotheses, a uniform law of large numbers (proved in Appendix C.3 of Appendix C) upgrades the pointwise LLN to $\sup_\theta |\frac{1}{n}\ell_n(\theta) - M(\theta)| \xrightarrow[n \to \infty]{\mathbb{P}} 0$. Then the argument of the

intuition box applies: outside any neighborhood of θ^*, M is at least some $\delta > 0$ below $M(\theta^*)$, and once the random curve is uniformly within $\delta/2$ of M, its maximizer must be inside the neighborhood.

Proof. Write $A_n = \sup_{\theta\in\Theta}|\frac{1}{n}\ell_n(\theta) - M(\theta)|$. The uniform law of large numbers (Theorem C.3) says $A_n \xrightarrow[n\to\infty]{\mathbb{P}} 0$ under the stated hypotheses; the one-θ-at-a-time version is the LLN with $h = \log p_\theta$, and the uniform version is where the continuity, boundedness and the integrable envelope $\sup_\theta|\log p_\theta(X)|$ are used. It is the one step of this chapter proved outside it.

Fix $\epsilon > 0$ and let $N = \{\theta : \|\theta - \theta^*\| < \epsilon\}$. Its complement in Θ, $\Theta \setminus N$, is bounded and closed, and M is continuous on it (continuity of $\log p_\theta$ in θ plus the envelope condition), so M attains a maximum on $\Theta \setminus N$ at some θ_1. By identifiability and Proposition 10.6, $M(\theta_1) < M(\theta^*)$; write $\delta = M(\theta^*) - M(\theta_1) > 0$.

Now suppose $A_n < \delta/2$. For every $\theta \in \Theta \setminus N$,

$$\tfrac{1}{n}\ell_n(\theta) \le M(\theta) + A_n < M(\theta_1) + \delta/2 = M(\theta^*) - \delta/2 < M(\theta^*) - A_n \le \tfrac{1}{n}\ell_n(\theta^*).$$

So every θ outside N has a strictly smaller average log-likelihood than θ^* does, and the maximizer $\hat\theta_n$ must lie in N. Hence $\mathbb{P}_{\theta^*}(\|\hat\theta_n - \theta^*\| \ge \epsilon) \le \mathbb{P}_{\theta^*}(A_n \ge \delta/2) \to 0$.

For the misspecified case replace θ^* by θ° and $\mathrm{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta)$ by $\mathrm{KL}(\mathbb{P}, \mathbb{P}_\theta)$ throughout; uniqueness of θ° is what gives $\delta > 0$. □

None of the families in this book has a bounded closed parameter space: the Bernoulli's is $(0, 1)$, the Poisson's $(0, \infty)$, and on the closed interval $[0, 1]$ the Bernoulli envelope $\sup_p|\log p_p(X)|$ is infinite. The theorem is applied to them by restricting Θ to a bounded closed set Θ' that contains θ^* in its interior, for instance $[a, 1-a]$ for the Bernoulli with $0 < a < \min(p^*, 1-p^*)$. The conclusion is unchanged as long as the MLE over all of Θ lands in Θ' with probability tending to 1, and for every family in this chapter it does, since ℓ_n falls to $-\infty$ at the edge of Θ: for the kiss data, $80\log p + 44\log(1-p) \to -\infty$ at both ends of $(0, 1)$, so $\bar R_n \in [a, 1-a]$ once $\bar R_n$ is within $\min(p^*, 1-p^*) - a$ of p^*, which happens with probability tending to 1.

> **In words.** The average log-likelihood is an average, so it converges to a fixed curve. That curve is a constant minus the KL divergence from the truth, so it peaks at the truth. If the random curve hugs the fixed one closely enough everywhere, its peak is near the truth. The same argument, with the truth outside the model, lands on the member of the family closest to the truth.
>
> **For the kiss data.** $M(p) = 0.645\log p + 0.355\log(1-p)$, the black curve in Figure 12.1, peaks at 0.645. The theorem says $\hat p = \bar R_n \xrightarrow[n\to\infty]{\mathbb{P}} p^*$, which Chapter 7 already knew from the LLN directly. The theorem's value is that it needs no formula for the estimator: for the mixture of Section 11.8, once the labels are fixed by requiring $\mu_1 < \mu_2$ so that the model is identifiable, it says the same thing, though the MLE has no formula.

12.3 The score and the Fisher information

Consistency says the MLE gets there. The next question is how fast, and the answer is in the shape of the log-likelihood near its peak. Chapter 11 ended with a sanity check: for the kiss data, the curvature $-\ell_n''(\hat p) = 541.7$ and $1/\sqrt{541.7} = 0.043$, the standard error. This section defines the quantity behind that.

Everything is built from one observation. Write

$$\ell(\theta) = \log p_\theta(X)$$

for the log-likelihood of a single $X \sim \mathbb{P}_\theta$, a random function of θ, and assume it is twice differentiable in θ. Its derivative $\ell'(\theta)$ is called the **score**. For n observations, $\ell_n(\theta) = \sum_i \log p_\theta(X_i)$ is a sum of n independent copies of ℓ, so ℓ_n' is a sum of n independent scores and ℓ_n'' a sum of n independent second derivatives.

Intuition: what the score is

Fix the truth θ^* and draw one observation. The score $\ell'(\theta^*)$ is the slope, at the truth, of that observation's log-likelihood: how much the observation pushes the estimate of θ up or down from θ^*. For the kiss model (Figure 12.4, left) a couple turning right has log-likelihood $\log p$, with slope $1/p^* = 1.55$ at $p^* = 0.645$: it pushes the estimate up. A couple turning left has log-likelihood $\log(1-p)$, with slope $-1/(1-p^*) = -2.82$: it pushes down, harder, since left turns are the rarer event. Weighted by how often each happens, $0.645 \times 1.55 - 0.355 \times 2.82 = 0$, to rounding. At the truth the pushes cancel on average, which is the first fact below. How much they vary around zero, their variance, measures how much one observation can move the estimate: that is the Fisher information.

Proposition 12.2 (The score has mean zero)**.** *Under regularity conditions that allow differentiating under the integral sign (stated in Section 12.6), for every $\theta \in \Theta$,*

$$\mathop{\mathbb{E}}_{\theta}[\ell'(\theta)] = 0.$$

Proof strategy (for completeness; skip on a first reading). The density integrates to 1 for every θ, so the derivative of that integral in θ is 0. Moving the derivative inside turns it into the expectation of the score.

Proof. Take the continuous case; the discrete case replaces the integral by a sum. For every θ, $\int p_\theta(x)\,\mathrm{d}x = 1$. Differentiate both sides in θ, moving the derivative inside the integral:

$$0 = \frac{\mathrm{d}}{\mathrm{d}\theta}\int p_\theta(x)\,\mathrm{d}x = \int \frac{\partial p_\theta(x)}{\partial \theta}\,\mathrm{d}x. \tag{$*$}$$

Now $\ell'(\theta) = \partial_\theta \log p_\theta(X) = \partial_\theta p_\theta(X)/p_\theta(X)$, so

$$\mathop{\mathbb{E}}_{\theta}[\ell'(\theta)] = \int \frac{\partial_\theta p_\theta(x)}{p_\theta(x)}\,p_\theta(x)\,\mathrm{d}x = \int \partial_\theta p_\theta(x)\,\mathrm{d}x = 0$$

by $(*)$. □

Definition 12.3 (Fisher information)**.** For a one-parameter model satisfying the regularity conditions of Section 12.6, the **Fisher information** at θ is

$$I(\theta) = \mathop{\mathbb{V}\mathrm{ar}}_{\theta}\left(\ell'(\theta)\right) = \mathop{\mathbb{E}}_{\theta}\left[\ell'(\theta)^2\right],$$

the variance of the score of one observation. For a parameter $\theta \in \mathbb{R}^d$ the score is the gradient $\nabla\ell(\theta)$ and the **Fisher information matrix** is its $d \times d$ covariance matrix, $I(\theta) = \mathrm{Cov}_\theta(\nabla\ell(\theta))$.

In words. Draw one observation from $\mathbb{P}_\theta$ and ask how steeply its log-likelihood rises or falls at θ. That slope is zero on average. Its spread, the Fisher information, measures how much one observation typically says about θ: a large spread means observations push hard in one direction or the other, and each one carries a lot of information. n independent observations carry n times as much, since the variance of a sum of independent scores is the sum of their variances.

For the kiss data. $\ell'(p) = R/p - (1-R)/(1-p)$. The score takes the value $1/p$ with probability p and $-1/(1-p)$ with probability $1-p$, and its variance is $I(p) = 1/(p(1-p))$ (Section 12.3.1). At $p = 0.645$, $I = 4.37$, and 124 couples carry $nI(\hat{p}) = 541.7$ units of information, the curvature the sanity check of Chapter 11 found.

There is a second way to compute the same number, and it is the one that explains the sanity check.

Theorem 12.4 (Fisher information as curvature)**.** *Under the regularity conditions of Section 12.6,*

$$I(\theta) = -\mathop{\mathbb{E}}_{\theta}\big[\ell''(\theta)\big].$$

In several dimensions, $I(\theta) = -\mathbb{E}_\theta[\nabla^2 \ell(\theta)]$, the negative of the expected Hessian.

Proof strategy (core). Differentiate the identity $(*)$ once more, to get $\int \partial_\theta^2 p_\theta = 0$. Write ℓ'' by the quotient rule as $\partial_\theta^2 p_\theta / p_\theta$ minus the squared score. The expectation of the first piece is 0 by the new identity, and the expectation of the second is $I(\theta)$.

Proof. Differentiating $(*)$ in θ once more, again moving the derivative inside,

$$\int \frac{\partial^2 p_\theta(x)}{\partial \theta^2}\,\mathrm{d}x = 0. \tag{$**$}$$

By the quotient rule applied to $\ell'(\theta) = \partial_\theta p_\theta(X)/p_\theta(X)$,

$$\ell''(\theta) = \frac{\partial_\theta^2 p_\theta(X)\cdot p_\theta(X) - \big(\partial_\theta p_\theta(X)\big)^2}{p_\theta(X)^2} = \frac{\partial_\theta^2 p_\theta(X)}{p_\theta(X)} - \left(\frac{\partial_\theta p_\theta(X)}{p_\theta(X)}\right)^2 = \frac{\partial_\theta^2 p_\theta(X)}{p_\theta(X)} - \ell'(\theta)^2.$$

Take expectations under $\mathbb{P}_\theta$. The first term gives $\int \partial_\theta^2 p_\theta(x)\,\mathrm{d}x = 0$ by $(**)$. The second gives $-\mathbb{E}_\theta[\ell'(\theta)^2] = -I(\theta)$ by Definition 12.3. So $\mathbb{E}_\theta[\ell''(\theta)] = -I(\theta)$. The matrix version is the same computation entry by entry. □

Intuition: information is curvature

$\ell''(\theta)$ is how sharply one observation's log-likelihood curves at θ, and the theorem says the Fisher information is the average of that curvature, with the sign flipped so that it comes out positive. Figure 12.2 draws it for the kiss model. The two possible log-likelihoods, $\log p$ and $\log(1-p)$, are each curved; their average under p^*, the black curve, peaks at p^* and curves down there with second derivative $-I(p^*) = -4.37$. The aqua parabola has exactly that curvature. For n observations the log-likelihood is a sum, its curvature near θ^* is about $nI(\theta^*)$, and a sharply curved peak is one whose location is well determined: move a little away from $\hat{\theta}$ and the log-likelihood drops a lot, so few candidates fit nearly as well.

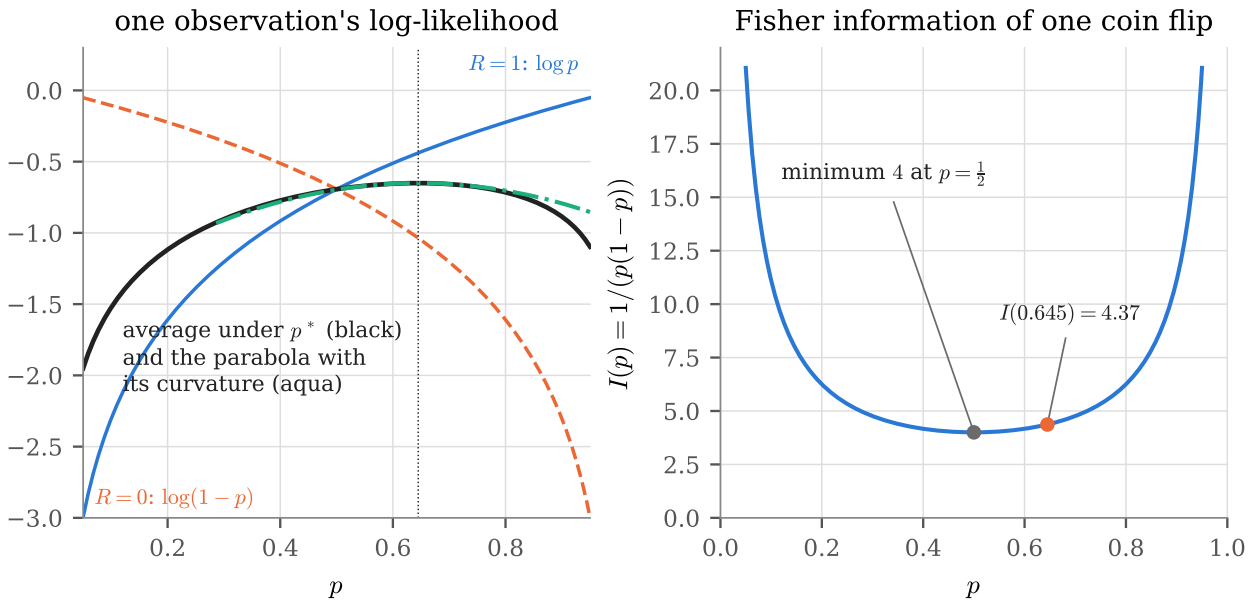


Figure 12.2: Fisher information is curvature. Left: the log-likelihood of one couple's turn, $\log p$ for a right turn (blue) and $\log(1-p)$ for a left turn (orange, dashed), and their average under $p^* = 0.645$ (black), which peaks at p^*; the aqua parabola through the peak has second derivative $-I(p^*) = -4.37$. Right: $I(p) = 1/(p(1-p))$ for the Bernoulli model, smallest at $p = \frac{1}{2}$ and growing without bound at either end.

Does this make sense? Which coin is most informative?

The right of Figure 12.2 says a single flip of a fair coin carries the least information about p, $I(\frac{1}{2}) = 4$, and a flip of a heavily biased coin carries the most, $I(0.9) = 11.1$. Is that backwards? A fair coin's outcome is the most unpredictable, so it feels the most informative. But the question is not how surprising the flip is; it is how much the flip changes your estimate of p. The variance of $\hat{p}$ is $p(1-p)/n$, which is largest at $p = \frac{1}{2}$: a fair coin's p is the hardest to pin down in absolute terms, since its flips vary the most. For $p = 0.9$ nearly every flip is heads, the sample proportion barely moves from sample to sample, and p is learned to within a narrow margin. The Fisher information is the reciprocal of that variance per observation, $1/(p(1-p))$, so it says exactly this.

12.3.1 Fisher information of the standard families

Worked example: four Fisher informations

In each case the curvature formula of Theorem 12.4 is the quicker route; the variance of the score is computed once as a check.

- **Bernoulli.** $\ell(p) = R\log p + (1-R)\log(1-p)$, so $\ell'(p) = R/p - (1-R)/(1-p)$ and $\ell''(p) = -R/p^2 - (1-R)/(1-p)^2$. With $\mathbb{E}_p[R] = p$,

$$I(p) = -\mathop{\mathbb{E}}_p[\ell''(p)] = \frac{p}{p^2} + \frac{1-p}{(1-p)^2} = \frac{1}{p} + \frac{1}{1-p} = \frac{1}{p(1-p)}.$$

  As a check, $\ell'(p) = R\left(\frac{1}{p} + \frac{1}{1-p}\right) - \frac{1}{1-p}$ is a linear function of R, so $\mathbb{V}\mathrm{ar}_p(\ell'(p)) =$

$\left(\frac{1}{p(1-p)}\right)^2 \mathbb{V}\mathrm{ar}(R) = \frac{p(1-p)}{p^2(1-p)^2} = \frac{1}{p(1-p)}$. The same.

- **Poisson.** $\ell(\lambda) = -\lambda + X\log\lambda - \log(X!)$, $\ell''(\lambda) = -X/\lambda^2$, and $\mathbb{E}_\lambda[X] = \lambda$ gives $I(\lambda) = 1/\lambda$.
- **Exponential.** $\ell(\lambda) = \log\lambda - \lambda X$, $\ell''(\lambda) = -1/\lambda^2$ with no X in it, so $I(\lambda) = 1/\lambda^2$. For the subway rate $\hat\lambda = 0.1605$ per minute, $I = 38.8$.
- **Gaussian, both parameters unknown.** With $\theta = (\mu, \sigma^2)$ and $v = \sigma^2$, $\ell = -\frac{1}{2}\log(2\pi v) - (X-\mu)^2/(2v)$. The second partial derivatives are

$$\frac{\partial^2\ell}{\partial\mu^2} = -\frac{1}{v}, \qquad \frac{\partial^2\ell}{\partial\mu\,\partial v} = -\frac{X-\mu}{v^2}, \qquad \frac{\partial^2\ell}{\partial v^2} = \frac{1}{2v^2} - \frac{(X-\mu)^2}{v^3}.$$

Under $\mathbb{P}_\theta$, $\mathbb{E}[X-\mu] = 0$ and $\mathbb{E}[(X-\mu)^2] = v$, so

$$I(\mu, \sigma^2) = \begin{pmatrix} 1/\sigma^2 & 0 \\ 0 & 1/(2\sigma^4) \end{pmatrix}.$$

The off-diagonal zero says that, for a Gaussian, learning about μ and learning about σ^2 don't interfere with each other.

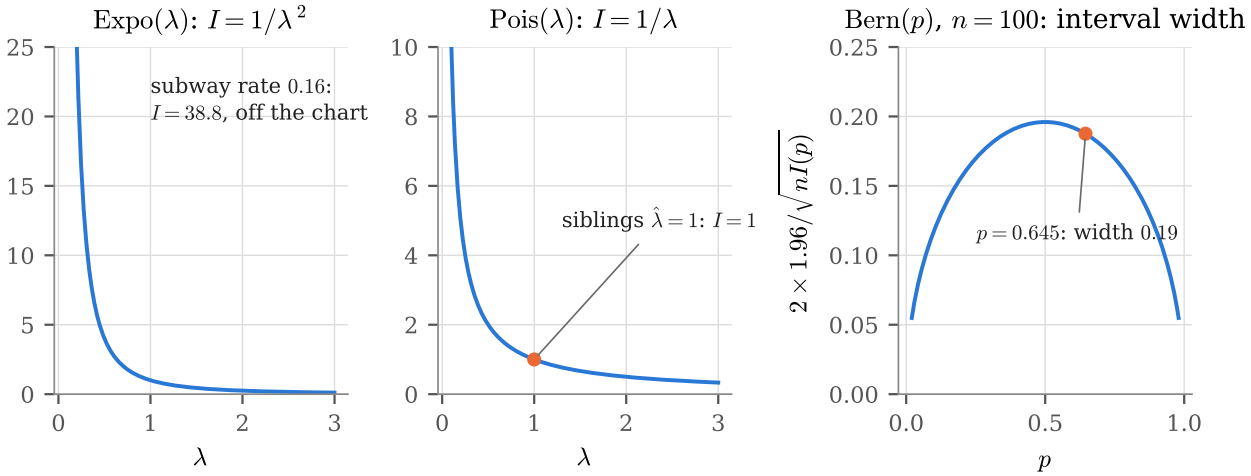


Figure 12.3: Fisher information in three one-parameter families. Left: $I(\lambda) = 1/\lambda^2$ for Expo(λ); the subway rate 0.16 has $I = 38.8$. Middle: $I(\lambda) = 1/\lambda$ for Pois(λ); the siblings fit $\hat\lambda = 1$ has $I = 1$. Right: what the information buys, the width $2 \times 1.96/\sqrt{nI(p)}$ of the 95% interval for a proportion at $n = 100$, widest at $p = \frac{1}{2}$ and 0.19 at $p = 0.645$.

Pitfall: Fisher information depends on how the parameter is written

The exponential model has $I(\lambda) = 1/\lambda^2$ in the rate λ. Written in the mean gap $\beta = 1/\lambda$ it has $I(\beta) = 1/\beta^2$, a different function (Exercise 12.2). The information is a property of the model *and* the parameterization: it measures information per unit of θ, and a unit of λ is not a unit of β. The general rule is $I(\beta) = I(\lambda)\,(\mathrm{d}\lambda/\mathrm{d}\beta)^2$, which is the delta method's $g'(\theta)^2$ seen from the other side. What does not depend on the parameterization is the conclusion drawn from it: the 95% interval for β obtained from $I(\beta)$ is the image under $\beta = 1/\lambda$ of the interval for λ, to first

order.

12.4 Asymptotic normality

Recall. The Central Limit Theorem (Theorem 4.8): if $Y_1, Y_2, \dots$ are i.i.d. with mean m and variance $s^2 < \infty$, then $\sqrt{n}(\bar{Y}_n - m) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, s^2)$.

Apply it to the scores $Y_i = \ell_i'(\theta^*)$, where $\ell_i(\theta) = \log p_\theta(X_i)$. They are i.i.d., with mean 0 (Proposition 12.2) and variance $I(\theta^*)$ (Definition 12.3), so

$$\sqrt{n} \cdot \frac{1}{n}\ell_n'(\theta^*) = \frac{1}{\sqrt{n}} \sum_{i=1}^{n} \ell_i'(\theta^*) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}\big(0, I(\theta^*)\big).$$

The middle of Figure 12.4 shows it for the kiss model at $n = 124$: the average score at p^* over 20,000 simulated studies, with the Gaussian the CLT predicts.

That is a statement about the slope of the log-likelihood at the truth. The MLE is where the slope is zero. To connect them, expand the slope around the truth.

Recall. The mean value theorem (toolkit A.9): for a differentiable g, $g(b) - g(a) = g'(c)(b-a)$ for some c between a and b. Slutsky's theorem (Theorem 5.1): if $X_n \xrightarrow[n\to\infty]{\text{(d)}} X$ and $Y_n \xrightarrow[n\to\infty]{\mathbb{P}} y$, a constant, then $X_n Y_n \xrightarrow[n\to\infty]{\text{(d)}} yX$ and $X_n + Y_n \xrightarrow[n\to\infty]{\text{(d)}} X + y$. The continuous mapping theorem (Theorem 5.2): if $Y_n \xrightarrow[n\to\infty]{\mathbb{P}} y$ and f is continuous at y, then $f(Y_n) \xrightarrow[n\to\infty]{\mathbb{P}} f(y)$.

Theorem 12.5 (Asymptotic normality of the MLE)**.** *Let $X_1, \dots, X_n \overset{iid}{\sim} \mathbb{P}_{\theta^*}$ in a model satisfying the regularity conditions of Section 12.6, with θ^* in the interior of Θ and $I(\theta^*) > 0$, and let $\hat{\theta}_n$ be the MLE. Then*

$$\sqrt{n}\,(\hat{\theta}_n - \theta^*) \xrightarrow[n\to\infty]{(d)} \textit{Gauss}\Big(0, \frac{1}{I(\theta^*)}\Big) \quad \textit{under } \mathbb{P}_{\theta^*}.$$

For $\theta \in \mathbb{R}^d$, the limit is the d-dimensional Gaussian with mean 0 and covariance matrix $I(\theta^*)^{-1}$.

Proof strategy (core). The slope of ℓ_n is 0 at $\hat{\theta}_n$ and, by the CLT, approximately $\text{Gauss}(0, nI(\theta^*))$ at θ^*. The mean value theorem writes the difference between the two slopes as (curvature somewhere in between) $\times$ $(\hat{\theta}_n - \theta^*)$. The curvature in between is an average of second derivatives, which the LLN sends to $-I(\theta^*)$ once consistency pulls the in-between point to θ^*. Dividing the Gaussian by $-I(\theta^*)$ gives the result, with Slutsky doing the division.

Proof. Since θ^* is interior and $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta^*$ (Theorem 12.1), with probability tending to 1 the MLE is interior too, and then the first-order condition holds: $\ell_n'(\hat{\theta}_n) = 0$. Apply the mean value theorem to ℓ_n' between θ^* and $\hat{\theta}_n$: there is a $\tilde{\theta}_n$ between them with

$$0 = \ell_n'(\hat{\theta}_n) = \ell_n'(\theta^*) + \ell_n''(\tilde{\theta}_n)\,(\hat{\theta}_n - \theta^*).$$

Rearrange and scale by $\sqrt{n}$:

$$\sqrt{n}\,(\hat{\theta}_n - \theta^*) = \frac{\frac{1}{\sqrt{n}}\ell_n'(\theta^*)}{-\frac{1}{n}\ell_n''(\tilde{\theta}_n)}.$$

The numerator converges in distribution to $\text{Gauss}(0, I(\theta^*))$ by the CLT applied to the scores, as above. For the denominator, $-\frac{1}{n}\ell_n''(\theta^*) = -\frac{1}{n}\sum_i \ell_i''(\theta^*) \xrightarrow[n\to\infty]{\mathbb{P}} -\mathbb{E}_{\theta^*}[\ell''(\theta^*)] = I(\theta^*)$ by the LLN and Theorem 12.4; and since $\tilde{\theta}_n$ is squeezed between θ^* and $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta^*$, the bound $|\ell_i'''(\theta)| \le K(X_i)$ near θ^* from the regularity conditions gives, by the mean value theorem once more,

$$\left|\tfrac{1}{n}\ell_n''(\tilde{\theta}_n) - \tfrac{1}{n}\ell_n''(\theta^*)\right| \le \left(\tfrac{1}{n}\sum_{i=1}^{n} K(X_i)\right)|\hat{\theta}_n - \theta^*| \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}_{\theta^*}[K(X)]\cdot 0 = 0,$$

by the LLN for the average of the $K(X_i)$ and Slutsky's product rule; so the denominator converges in probability to $I(\theta^*)$ as well. Since $I(\theta^*) > 0$, the function $x \mapsto 1/x$ is continuous there, and the continuous mapping theorem gives $1/(-\frac{1}{n}\ell_n''(\tilde{\theta}_n)) \xrightarrow[n\to\infty]{\mathbb{P}} 1/I(\theta^*)$. Slutsky's theorem multiplies the two:

$$\sqrt{n}\,(\hat{\theta}_n - \theta^*) \xrightarrow[n\to\infty]{\text{(d)}} \frac{1}{I(\theta^*)}\cdot \text{Gauss}(0, I(\theta^*)) = \text{Gauss}\Big(0, \frac{I(\theta^*)}{I(\theta^*)^2}\Big) = \text{Gauss}\Big(0, \frac{1}{I(\theta^*)}\Big).$$

In d dimensions (toolkit A.12) each coordinate of the gradient gets its own in-between point $\tilde{\theta}_{n,j}$, all squeezed to θ^*; the denominator becomes the matrix whose row j is $-\frac{1}{n}$ times row j of the Hessian evaluated at $\tilde{\theta}_{n,j}$, which converges in probability to $I(\theta^*)$ entry by entry; the numerator converges to a Gaussian vector with covariance $I(\theta^*)$ by the multivariate CLT (Section 5.5); and the inverse matrix multiplies it, giving covariance $I^{-1}II^{-1} = I(\theta^*)^{-1}$. □

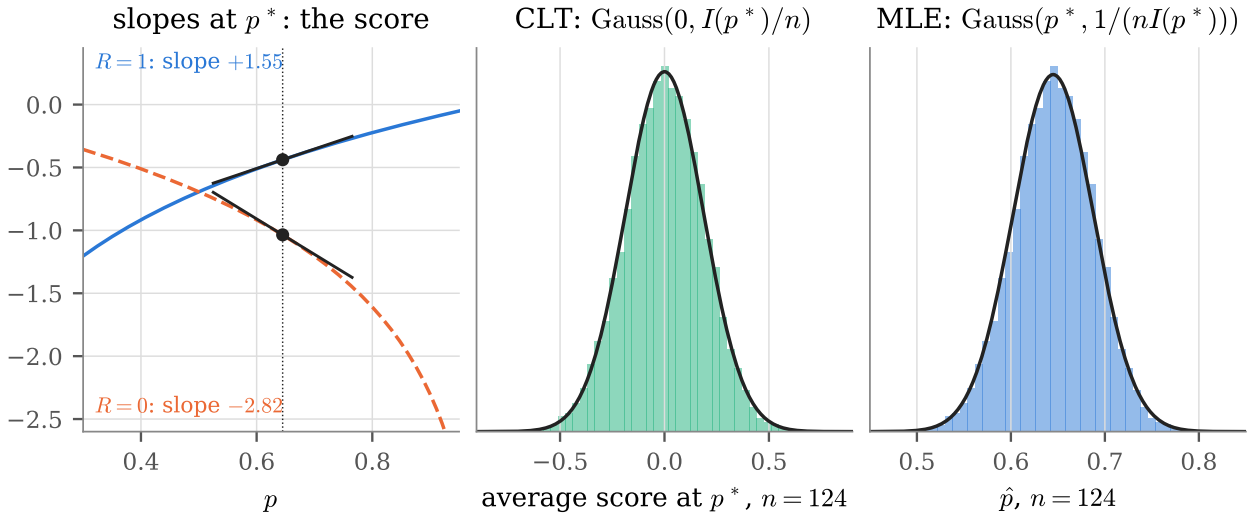


Figure 12.4: Why the MLE is asymptotically normal, for the kiss model with $p^* = 0.645$ and $n = 124$. Left: one couple's log-likelihood, $\log p$ (blue) or $\log(1-p)$ (orange, dashed), with the tangent at p^* drawn on each; the slopes, $+1.55$ and -2.82, are the two possible values of the score, which averages to 0 under p^*. Middle: the average score at p^* over 20,000 simulated studies, with the $\text{Gauss}(0, I(p^*)/n)$ the CLT predicts (standard deviation 0.188). Right: $\hat{p}$ over the same studies, with the $\text{Gauss}(p^*, 1/(nI(p^*)))$ the theorem predicts (standard deviation 0.043). The two histograms are the same picture: $\hat{p} - p^*$ is the average score divided by $I(p^*)$, to first order.

Intuition: why the variance is one over the information

Two forces set the size of the MLE's error. The slope of the log-likelihood at the truth is random, with standard deviation $\sqrt{nI}$ (the CLT): that is how far the data push. The curvature near the peak is about nI (the LLN): that is how hard the log-likelihood resists. Displacement is push divided by resistance, $\sqrt{nI}/(nI) = 1/\sqrt{nI}$, which is the standard deviation of $\hat{\theta}_n$. More information means both a bigger push and a bigger resistance, and the resistance wins, since it grows like n while the push grows like $\sqrt{n}$.

In words. For a large sample, the maximum likelihood estimator is approximately Gaussian, centered at the truth, with standard deviation $1/\sqrt{nI(\theta^*)}$. Everything about its precision is in one number computed from the model, and the recipe estimate $\pm 1.96\,/\sqrt{nI(\hat{\theta}_n)}$ gives a 95% confidence interval for any regular model, with no need to work out the estimator's distribution by hand.

For the kiss data. $I(p) = 1/(p(1-p))$, so $1/(nI(p^*)) = p^*(1-p^*)/n$: the theorem reproduces the CLT for $\bar{R}_n$ exactly, with asymptotic variance 0.229 and standard error 0.043 at $n = 124$. The interval $0.645 \pm 1.96 \times 0.043 = [0.561, 0.729]$ is the plug-in interval of Chapter 8. In natural frequencies: of 1,000 studies of 124 couples with $p^* = 0.645$, about 683 land within 0.043 of the truth (the exact Binomial count is 698).

From undergraduate statistics: one recipe for every standard error

In an introductory course each estimator came with its own standard error formula: $\sqrt{\hat{p}(1-\hat{p})/n}$ for a proportion, $s/\sqrt{n}$ for a mean, something from software for a regression coefficient. Theorem 12.5 is where those formulas come from, and it is one formula: the standard error of a maximum likelihood estimator is $1/\sqrt{nI(\hat{\theta})}$, or in practice $1/\sqrt{-\ell_n''(\hat{\theta})}$, the reciprocal square root of the curvature of the log-likelihood at its peak. For a proportion that is $\sqrt{\hat{p}(1-\hat{p})/n}$, exactly. For a Gaussian mean it is $\sqrt{\widehat{\sigma^2}/n}$, which is $s/\sqrt{n}$ up to the factor $\sqrt{(n-1)/n}$. For a regression coefficient, a logistic regression coefficient, a Poisson rate, the same recipe, which is what software computes when it reports a standard error next to a maximum likelihood estimate.

Worked example: standard errors from the Fisher information

- **Subway rate.** $I(\lambda) = 1/\lambda^2$, so the standard error of $\hat{\lambda}$ is $\hat{\lambda}/\sqrt{n} = 0.1605/8 = 0.020$ per minute, the number Chapter 5 obtained by the delta method. A 95% interval: 0.1605 ± 0.039, or $[0.121, 0.200]$.
- **Siblings.** $I(\lambda) = 1/\lambda$, so the standard error of $\hat{\lambda} = 1.0$ from 24 students is $\sqrt{1.0/24} = 0.204$: with so few observations, the Poisson rate is known only to about ± 0.4.
- **Boarding times.** The inverse information matrix is $\mathrm{diag}(\sigma^2, 2\sigma^4)$, so $\hat{\mu}$ has standard error $\sqrt{4.35/72} = 0.25$ minutes and $\widehat{\sigma^2}$ has standard error $\sqrt{2 \times 4.35^2/72} = 0.73$. The mean boarding time 24.2 is known to a quarter of a minute; the variance 4.35 only to within about 1.4.

12.5 The Cramér–Rao bound

Is $1/(nI(\theta))$ a good variance, or merely the MLE's variance? It is the best possible, in the following sense.

Theorem 12.6 (Cramér–Rao bound)**.** *Under the regularity conditions of Section 12.6, let $T = T(X_1, \dots, X_n)$ be any unbiased estimator of θ with finite variance. Then for every θ,*

$$\operatorname*{\mathbb{V}ar}_{\theta}(T) \geq \frac{1}{nI(\theta)}.$$

Proof strategy (for completeness; skip on a first reading). Unbiasedness, differentiated in θ, says the covariance of T with the total score $\ell_n'(\theta)$ is exactly 1. The Cauchy–Schwarz inequality bounds a covariance by the product of standard deviations, and the total score has variance $nI(\theta)$.

Proof. Unbiasedness says $\int T(x)\, p_\theta^{(n)}(x)\, \mathrm{d}x = \theta$ for every θ, where $p_\theta^{(n)} = \prod_i p_\theta(x_i)$ is the joint density. Differentiate in θ, moving the derivative inside (which the finite variance of T and the regularity conditions permit), and use $\partial_\theta p_\theta^{(n)} = p_\theta^{(n)}\, \ell_n'(\theta)$:

$$1 = \int T(x)\, \ell_n'(\theta)\, p_\theta^{(n)}(x)\, \mathrm{d}x = \operatorname*{\mathbb{E}}_{\theta}\left[T\, \ell_n'(\theta)\right] = \operatorname{Cov}_\theta\big(T, \ell_n'(\theta)\big),$$

the last step because $\mathbb{E}_\theta[\ell_n'(\theta)] = 0$. By the Cauchy–Schwarz inequality for covariances, $\operatorname{Cov}(T,S)^2 \leq \operatorname{\mathbb{V}ar}(T)\operatorname{\mathbb{V}ar}(S)$ (take $c = \operatorname{Cov}(T,S)/\operatorname{\mathbb{V}ar}(S)$ in $0 \leq \operatorname{\mathbb{V}ar}(T - cS) = \operatorname{\mathbb{V}ar}(T) - 2c\operatorname{Cov}(T,S) + c^2 \operatorname{\mathbb{V}ar}(S)$), with $S = \ell_n'(\theta)$ and $\operatorname{\mathbb{V}ar}_\theta(S) = nI(\theta)$ (a sum of n independent scores):

$$1 \leq \operatorname*{\mathbb{V}ar}_{\theta}(T) \cdot nI(\theta).$$

□

An unbiased estimator whose variance equals the bound is called **efficient**. The sample proportion is efficient for p: $\operatorname{\mathbb{V}ar}(\bar{R}_n) = p(1-p)/n = 1/(nI(p))$ exactly, at every n. Figure 12.5 draws the bound for the kiss model and two other estimators from Chapter 7. The MLE in general is not unbiased at finite n (the subway $1/\bar{T}_n$ is not), so the bound does not apply to it directly; what Theorem 12.5 says is that its asymptotic variance equals the bound, so that in large samples nothing unbiased can do better.

Pitfall: a bound for unbiased estimators only

The shrinkage estimator in Figure 12.5 has variance below the Cramér–Rao bound for every p. That does not contradict the theorem: the estimator is biased, and the theorem is about unbiased ones. Bias can buy variance, as Chapter 7 showed, and the honest comparison for a biased estimator is its mean squared error, which for the shrinkage estimator is below $\bar{R}_n$'s only near $p = \frac{1}{2}$. A general version of the bound, with the bias's derivative in the numerator, covers biased estimators too; it is not needed here.

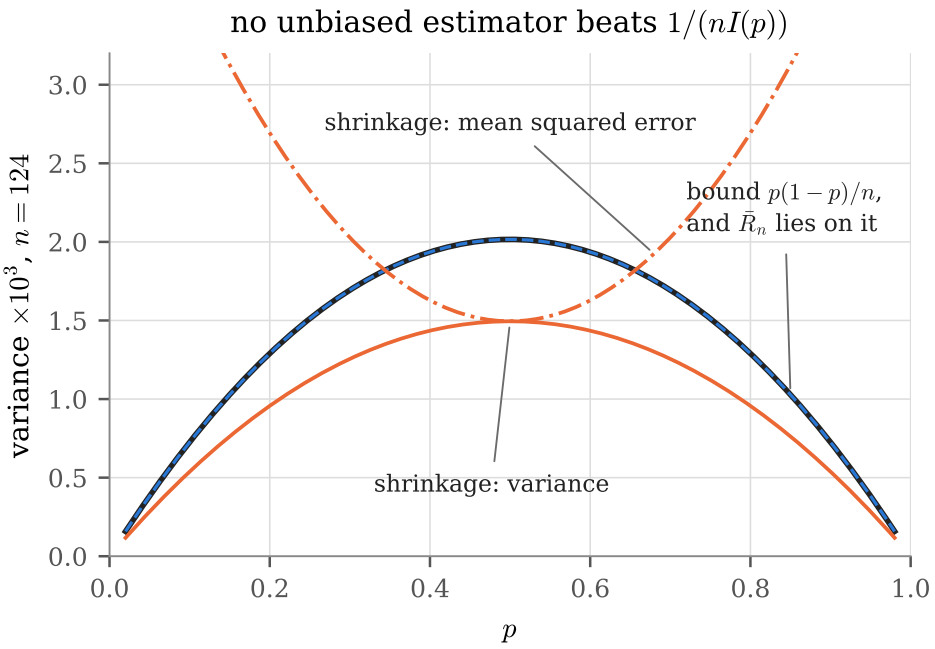


Figure 12.5: The Cramér–Rao bound $p(1-p)/n$ for the kiss model at $n = 124$ (black), with the variance of $\bar{R}_n$ lying exactly on it (blue, dashed). The shrinkage estimator $(n\bar{R}_n + 10)/(n + 20)$ of Chapter 7 has a smaller variance (orange, solid), which the bound permits because the estimator is biased; its mean squared error (orange, dash-dot) is what to compare, and it rises above the bound once p is more than about 0.16 from $\frac{1}{2}$.

12.6 What the theorems need, and the models they do not cover

The results above hold for **regular** models, which means:

1. the parameter is identifiable (Section 2.6);
2. the support of $\mathbb{P}_\theta$, the set where $p_\theta > 0$, does not depend on θ;
3. the true θ^* is not on the boundary of Θ;
4. $p_\theta(x)$ is three times differentiable in θ, the derivatives can be moved inside the integral (which gave $(*)$ and $(**)$, and the corresponding step in the Cramér–Rao proof), $\ell''(\theta^*)$ has finite variance, and $|\ell'''(\theta)| \le K(X)$ near θ^* for a K with finite variance;
5. $I(\theta) > 0$ near θ^*, and in d dimensions $I(\theta)$ is invertible.

Every family in this book except the uniform satisfies all five. Here is what each one guards against.

- **Identifiability** is what makes the limit curve M peak at one point. In the Gaussian mixture of Section 2.7, two parameter values give the same distribution, so M has two peaks and the MLE converges to one of them, with the labels decided by the starting point.
- **Fixed support** is what allows differentiating under the integral sign. For $\text{Unif}([0, \theta])$ the support is $[0, \theta]$, which moves with θ; the identity $(*)$ fails, the score is not defined at the observations' maximum, and the theorems do not apply. And they are false there: the MLE $\max_i X_i$ has error of exact order $1/n$, not $1/\sqrt{n}$, as Exercise 11.4 found. The limit of $n(\theta - \max_i X_i)$ is exponential, not Gaussian (Exercise 12.5). The uniform MLE is better than regular, not worse: its error shrinks faster than any regular estimator's.

- **Interior truth** is what allows $\ell_n'(\hat\theta_n) = 0$. If θ^* sits on the boundary, say $\Theta = [0, \infty)$ for a Gaussian mean with $\mu^* = 0$, the MLE is $\max(\bar X_n, 0)$: equal to 0 half the time and to the positive half of a Gaussian otherwise, not Gaussian.
- **Smoothness** is used in the mean value step and the LLN for ℓ''. The Laplace model of Exercise 11.8, whose log-likelihood has a corner at each observation, needs a separate argument (and its MLE, the median, is asymptotically normal all the same, with variance 1).
- **Positive information** is what makes the division in the proof legitimate. $I(\theta^*) = 0$ means the log-likelihood is flat at the truth to second order: the data say nothing about θ to first order, and the error shrinks more slowly than $1/\sqrt{n}$.

Does this make sense? Faster than $1/\sqrt{n}$

For Unif$([0, \theta])$ with $\theta = 10$, Exercise 11.4 gave $\mathbb{E}[\max_i X_i] = \frac{n}{n+1}\theta$, so the bias of the MLE is $-\theta/(n+1)$: about -0.1 at $n = 100$ and -0.01 at $n = 1{,}000$, shrinking tenfold when n grows tenfold. A regular estimator's error shrinks by $\sqrt{10} = 3.2$ when n grows tenfold. So the uniform MLE is converging faster than any regular model allows, and the Cramér–Rao bound, which would put a floor of order $1/n$ under its variance, does not hold for it either. No contradiction: the bound needs $(*)$, which needs a fixed support.

12.7 What comes next

This chapter finished the theory of maximum likelihood: consistent, asymptotically normal, with asymptotic variance $1/I(\theta^*)$ and the Cramér–Rao bound saying that is as good as unbiased estimation gets. Chapter 13 puts two other recipes beside it. The method of moments matches sample moments to population moments, and the delta method gives its asymptotic variance, which is at least $1/I(\theta^*)$ and usually larger. M-estimation replaces $\log p_\theta$ by any function ρ whose expectation peaks at the quantity of interest; the argument of this chapter, an average, a limit curve, a score and a curvature, goes through word for word, with $I(\theta)$ replaced by a sandwich of two matrices. Chapter 15 builds tests from this chapter's theorem: Wald's test standardizes $\hat\theta_n - \theta_0$ by $1/\sqrt{nI(\hat\theta_n)}$, and the likelihood ratio test compares $\ell_n(\hat\theta_n)$ with $\ell_n(\theta_0)$, whose difference is approximately $\frac{1}{2}nI(\theta^*)(\hat\theta_n - \theta_0)^2$ by the same Taylor expansion, and so approximately half a chi-squared.

12.8 Summary

- Consistency: $\frac{1}{n}\ell_n(\theta) \xrightarrow[n\to\infty]{\mathbb{P}} M(\theta) = \text{constant} - \mathrm{KL}(\mathbb{P}_{\theta^*}, \mathbb{P}_\theta)$ by the LLN; M peaks at θ^* (or at the KL projection θ° if the model is wrong); under a uniform version of the LLN the peak of $\frac{1}{n}\ell_n$ follows the peak of M, so $\hat\theta_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta^*$.
- The score $\ell'(\theta)$ is the slope of one observation's log-likelihood. $\mathbb{E}_\theta[\ell'(\theta)] = 0$, from differentiating $\int p_\theta = 1$.
- Fisher information: $I(\theta) = \mathbb{V}\mathrm{ar}_\theta(\ell'(\theta)) = -\mathbb{E}_\theta[\ell''(\theta)]$, the variance of the score and the expected curvature of the log-likelihood. n observations carry $nI(\theta)$. Bernoulli $1/(p(1-p))$ (kiss: 4.37, 124 couples: 541.7); Poisson $1/\lambda$; exponential $1/\lambda^2$; Gaussian $\mathrm{diag}(1/\sigma^2, 1/(2\sigma^4))$. It changes with the parameterization: $I(\beta) = I(\lambda)(\mathrm{d}\lambda/\mathrm{d}\beta)^2$.

- Asymptotic normality: $\sqrt{n}(\hat{\theta}_n - \theta^*) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1/I(\theta^*))$, by the CLT on the scores, the mean value theorem on ℓ_n', the LLN on ℓ_n'' and Slutsky. Push $\sqrt{nI}$ over resistance nI. The 95% interval is $\hat{\theta}_n \pm 1.96/\sqrt{nI(\hat{\theta}_n)}$ for any regular model: kiss $[0.561, 0.729]$, subway $[0.121, 0.200]$.
- Cramér–Rao: every unbiased estimator has $\mathbb{V}\text{ar} \geq 1/(nI(\theta))$. $\bar{R}_n$ attains it. Biased estimators can go below it (shrinkage), and are judged by mean squared error.
- Regular models: identifiable, fixed support, interior truth, three derivatives that pass under the integral, positive information. The uniform model fails fixed support, and its MLE converges at rate $1/n$ with a non-Gaussian limit; mixtures fail identifiability.

12.9 Exercises

Exercise 12.1. For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Pois}(\lambda)$, compute $I(\lambda)$ both as $\mathbb{V}\text{ar}(\ell'(\lambda))$ and as $-\mathbb{E}[\ell''(\lambda)]$, state the asymptotic distribution of $\hat{\lambda} = \bar{X}_n$, and check it against the CLT applied directly to $\bar{X}_n$. *(Solution on page 230.)*

Exercise 12.2. For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Expo}(\lambda)$, let $\beta = 1/\lambda$ be the mean gap. Write the log-likelihood in β, compute $I(\beta)$, and verify $I(\beta) = I(\lambda)(\mathrm{d}\lambda/\mathrm{d}\beta)^2$. State the asymptotic distribution of $\hat{\beta} = \bar{X}_n$ and check it against the CLT. *(Solution on page 230.)*

Exercise 12.3. For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Gauss}(\mu, \sigma^2)$ with σ^2 known, compute $I(\mu)$ and show that $\bar{X}_n$ is efficient: its variance equals the Cramér–Rao bound at every n. *(Solution on page 230.)*

Exercise 12.4. For the kiss data, compute the total score $\ell_n'(p)$ at $p = \frac{1}{2}$, 0.6, 0.645 and 0.7, and the standardized score $\ell_n'(p_0)/\sqrt{nI(p_0)}$ at $p_0 = \frac{1}{2}$. Compare the latter with the test statistic of Section 9.6. *(Solution on page 230.)*

Exercise 12.5. For $X_1, \ldots, X_n \overset{iid}{\sim} \text{Unif}([0, \theta])$ and $M = \max_i X_i$, show that $n(\theta - M)/\theta \xrightarrow[n\to\infty]{\text{(d)}} \text{Expo}(1)$. Conclude that the MLE's error is of order $1/n$ and its limit is not Gaussian. *(Solution on page 230.)*

Exercise 12.6. For the geometric model of Exercise 11.6, $\mathbb{P}(X = k) = (1-p)^{k-1}p$, compute $I(p)$ and the standard error of $\hat{p} = 1/\bar{X}_n$ for the five rounds of $3, 1, 2, 5, 1$ shots. *(Solution on page 230.)*

Exercise 12.7. Show directly that the Fisher information of n i.i.d. observations, defined as the variance of the total score $\ell_n'(\theta)$, is $nI(\theta)$, and that it is also $-\mathbb{E}_\theta[\ell_n''(\theta)]$. *(Solution on page 230.)*

Exercise 12.8. Data come from Expo(1) but are fitted with the Gaussian family. Proposition 10.10 says the KL projection is Gauss(1, 1). What does Theorem 12.1 say the Gaussian MLEs $(\bar{X}_n, \frac{1}{n}\sum(X_i - \bar{X}_n)^2)$ converge to, and does that agree with the LLN applied directly? *(Solution on page 231.)*

Chapter 13

Method of Moments and M-Estimation

Part V

Parametric Hypothesis Tests

The tests of Chapter 9 were built on one standardized average. This part supplies the distributions the remaining parametric tests use, chi-squared, Student's t and Cochran's theorem (Chapter 14), and the two general-purpose tests built on maximum likelihood, Wald's and the likelihood ratio (Chapter 15).

Chapter 14

Chi-Squared, Student's t, Cochran's Theorem and the t-Test

Chapter 15

Wald and Likelihood Ratio Tests

Part VI

Nonparametric Hypothesis Tests

Every method so far assumed the model was well-specified. This part tests that assumption: goodness of fit for discrete distributions by the chi-squared test (Chapter 16), and for continuous ones by comparing the empirical CDF with the candidate, the Kolmogorov–Smirnov and Lilliefors tests and the QQ plot (Chapter 17).

Chapter 16

Goodness of Fit for Discrete Distributions

Chapter 17

Kolmogorov-Smirnov, Lilliefors and QQ Plots

Part VII

Bayesian Statistics

The one part of the book where the parameter gets a distribution. Priors, posteriors and conjugacy (Chapter 18) set up the machinery; Jeffreys priors and Bayesian confidence regions (Chapter 19) finish it, and say how the result relates to the frequentist intervals of Chapter 8.

Chapter 18

Priors, Posteriors and Conjugacy

Chapter 19

Jeffreys Priors and Bayesian Intervals

Part VIII

Linear Regression

Least squares as a statistical model. One predictor first (Chapter 20), then the matrix form with several, and the inference that comes with Gaussian errors (Chapter 21), where the identifiability of Chapter 2 returns as the rank of the design matrix.

Chapter 20

Univariate Linear Regression

Chapter 21

Multivariate Linear Regression and Inference

Part IX

Generalized Linear Models

Regression for responses that are not Gaussian. Exponential families (Chapter 22) are the distributions the method works for; link functions and the fitting algorithm (Chapter 23) connect them to a linear predictor, with maximum likelihood doing the estimation as in Part IV.

Chapter 22

Exponential Families

Chapter 23

Generalized Linear Models: Link Functions and Fitting

Part X

Principal Component Analysis

The book's one chapter of unsupervised statistics: finding the directions in which multivariate data vary most, through the eigenvectors of the sample covariance matrix, and what the Law of Large Numbers says about them (Chapter 24).

Chapter 24

Principal Component Analysis

Appendices

Appendix A

Mathematical Toolkit

This appendix collects the mathematical habits the book uses: not statistics, but the handful of ideas from calculus, sets and probability that the proofs lean on. Each entry says what the idea is in plain words, draws it, and names the first place in the book that needs it. When a chapter marks a step with '(toolkit A.9)', this is where to come. Nothing here is needed to state a result; everything here is needed to follow a proof. Later chapters will add entries of their own (vectors and matrices for regression, the chi-squared family for tests) at the end.

A.1 Limits of sequences of numbers

A sequence of numbers $x_1, x_2, x_3, \ldots$ **converges** to x if the terms eventually stay as close to x as you like. The definition makes 'as close as you like' and 'eventually' precise:

$$\text{for every } \epsilon > 0 \text{ there is an index } N \text{ such that } |x_n - x| < \epsilon \text{ for all } n \geq N.$$

Read it as a game with an opponent. The opponent names a tolerance ϵ, any positive number, however small. You must answer with an index N past which every term of the sequence is within ϵ of x. If you can always answer, the sequence converges. The order of the two moves is the whole content: ϵ is chosen first and N is allowed to depend on it, which is why a smaller ϵ is usually met with a larger N (Figure A.1). Writing the moves in the other order ('there is an N that works for every ϵ') would say the sequence is eventually *equal* to x, a much stronger and usually false statement.

Where the book uses it. Section 3.1 starts from this definition and explains why it cannot be applied directly to random variables; Definition 3.1 keeps its shape but applies it to the *probabilities* $\mathbb{P}(|X_n - X| \geq \epsilon)$, which are ordinary numbers. Every '$\to 0$' in a proof of convergence in probability is this definition applied to a sequence of numbers.

Limits inside an expectation. One further fact is used once, in the uniform law of large numbers of Appendix C.3: if random variables W_r converge to 0 as $r \to 0$ for every outcome, and all of them are bounded by one fixed random variable G with $\mathbb{E}[G] < \infty$, then $\mathbb{E}[W_r] \to 0$ as well. This is the dominated convergence theorem. The bound is what makes it true: without it, a W_r that is huge on a set of tiny probability can have an expectation that refuses to go to 0 even though W_r itself goes to 0 at every outcome.

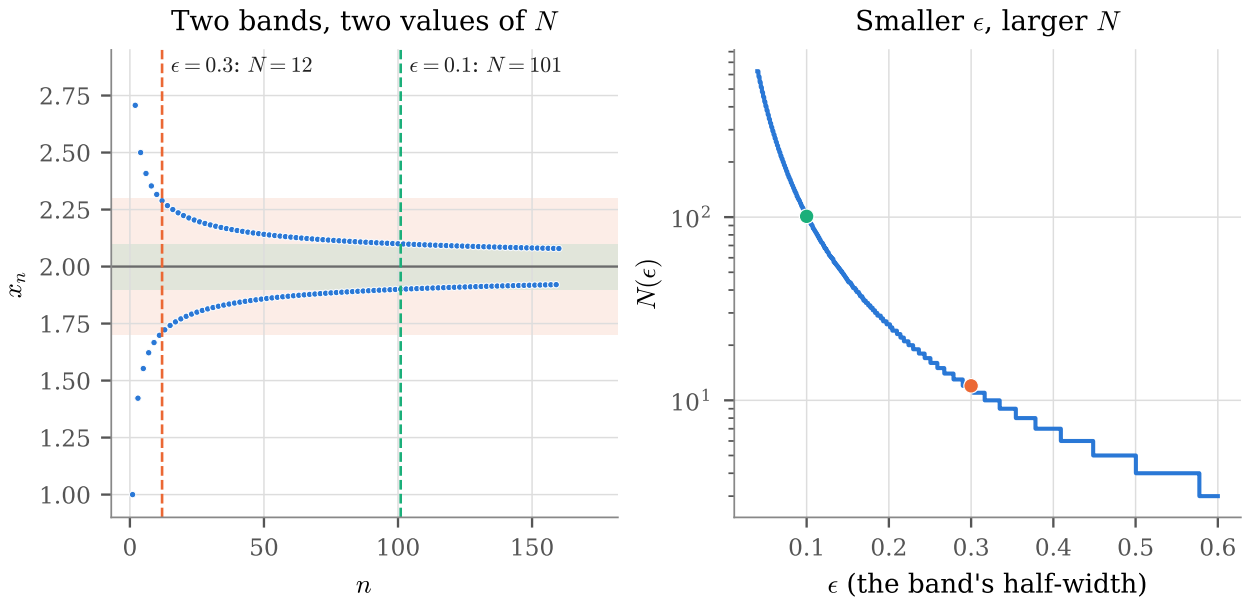


Figure A.1: The ϵ-N definition as a game. Left: the sequence $x_n = 2 + (-1)^n/\sqrt{n}$ with two tolerances; for $\epsilon = 0.3$ every term past $N = 12$ is inside the band, for $\epsilon = 0.1$ you need $N = 101$ (at $n = 100$ the term sits exactly on the band's edge, and the definition asks for strictly inside). Right: the index that works, as a function of the tolerance, climbs as ϵ shrinks. Convergence means the staircase is defined for every $\epsilon > 0$, however small.

A.2 Supremum and infimum

The **maximum** of a set of numbers is its largest member. Some sets have none: $\{1 - \frac{1}{n} : n = 1, 2, \ldots\}$ climbs toward 1 and never reaches it. The **supremum**, written sup, is the least upper bound: the smallest number that no member exceeds. For that set the supremum is 1, which is not a member. When a maximum exists it equals the supremum; the supremum exists more often, which is why proofs prefer it. The **infimum**, inf, is the greatest lower bound, and the same remarks apply with 'minimum'. Figure A.2 shows one set whose supremum is not attained and one function on a closed interval whose supremum and infimum both are (a continuous function on a closed, bounded interval always attains both).

Where the book uses it. The conservative confidence interval of Section 8.4 replaces $\sqrt{p(1-p)}$ by its supremum over $p \in (0, 1)$, which is $\frac{1}{2}$ and is attained; the subway version has $\sup_\lambda \lambda = \infty$, which is why no conservative interval exists there. Chapter 7 mentions that a statistic may be defined through a sup or inf; the maximum of a sample, $\max_i X_i$, is a maximum of finitely many numbers and causes no trouble.

A.3 Sets, maps and one-to-one

A **map** (or function) f from a set A to a set B assigns to each element of A exactly one element of B. Its **image** is the set of elements of B that are actually hit. The map is **injective** (one-to-one) if different inputs always give different outputs: $a_1 \neq a_2$ implies $f(a_1) \neq f(a_2)$, or, turned around, $f(a_1) = f(a_2)$ implies $a_1 = a_2$. An injective map can be undone: from the output you can recover

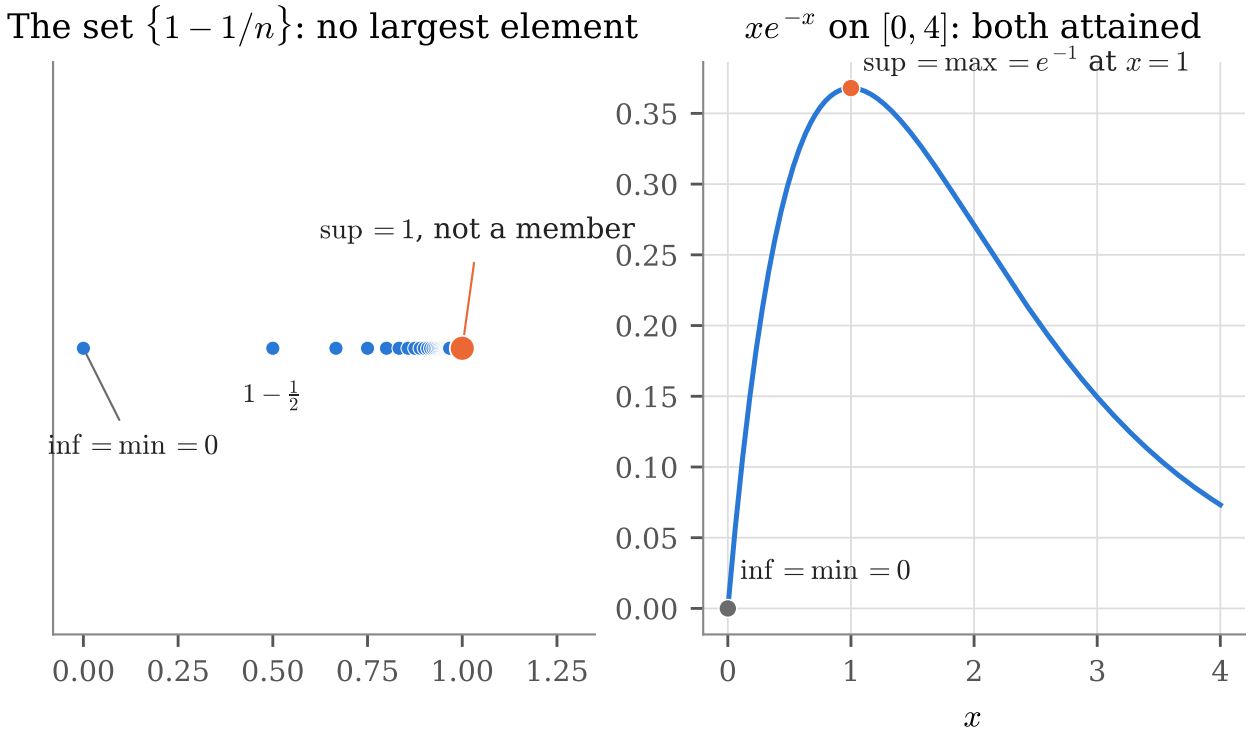


Figure A.2: Left: the set $\{1 - 1/n\}$ has no largest element, but 1 is its least upper bound. Right: xe^{-x} on $[0, 4]$ attains its supremum at $x = 1$ and its infimum at $x = 0$, so here sup = max and inf = min.

the input. For a function of one real variable the test is visual, the horizontal-line test of Figure A.3: if some horizontal line crosses the graph twice, two inputs share an output and the map is not injective. A strictly increasing or strictly decreasing function always passes the test.

A **partition** of a set is a way of cutting it into pieces that do not overlap and together cover everything; any map partitions its domain into the sets of inputs that share an output.

Where the book uses it. Identifiability (Definition 2.3) is exactly the statement that the map $\theta \mapsto \mathbb{P}_\theta$ is injective; every identifiability argument in Chapter 2 is a horizontal-line test on that map. A statistic partitions the sample space (Section 7.2).

A.4 Indicator functions

The **indicator** of an event or a condition is the function that is 1 when the condition holds and 0 when it does not: $\mathbb{1}(x \geq 5)$ is 0 for $x < 5$ and 1 from 5 on. It is the simplest random variable there is, a coin that comes up 1 with probability $\mathbb{P}(\text{condition})$, and its expectation is that probability:

$$\mathbb{E}[\mathbb{1}(X \geq 5)] = 1 \cdot \mathbb{P}(X \geq 5) + 0 \cdot \mathbb{P}(X < 5) = \mathbb{P}(X \geq 5).$$

That one line, *the expectation of an indicator is a probability*, is how probabilities enter computations that are written with expectations, and how a sum of indicators becomes a count (the number of couples turning right is $\sum_i \mathbb{1}(R_i = 1)$).

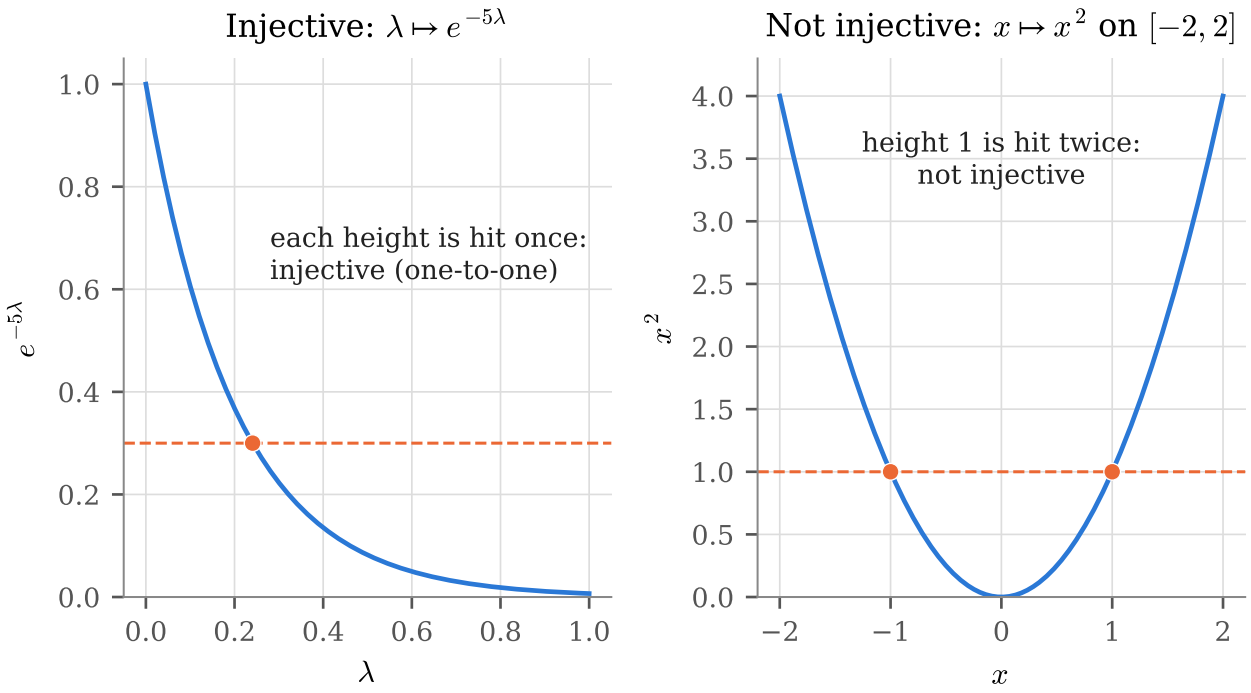


Figure A.3: The horizontal-line test. Left: $\lambda \mapsto e^{-5\lambda}$ is strictly decreasing, so every height is hit once and λ can be read off the output; this is the censored-exponential example of Example 2.5. Right: $x \mapsto x^2$ on $[-2, 2]$ hits the height 1 twice, so the input cannot be recovered from the output.

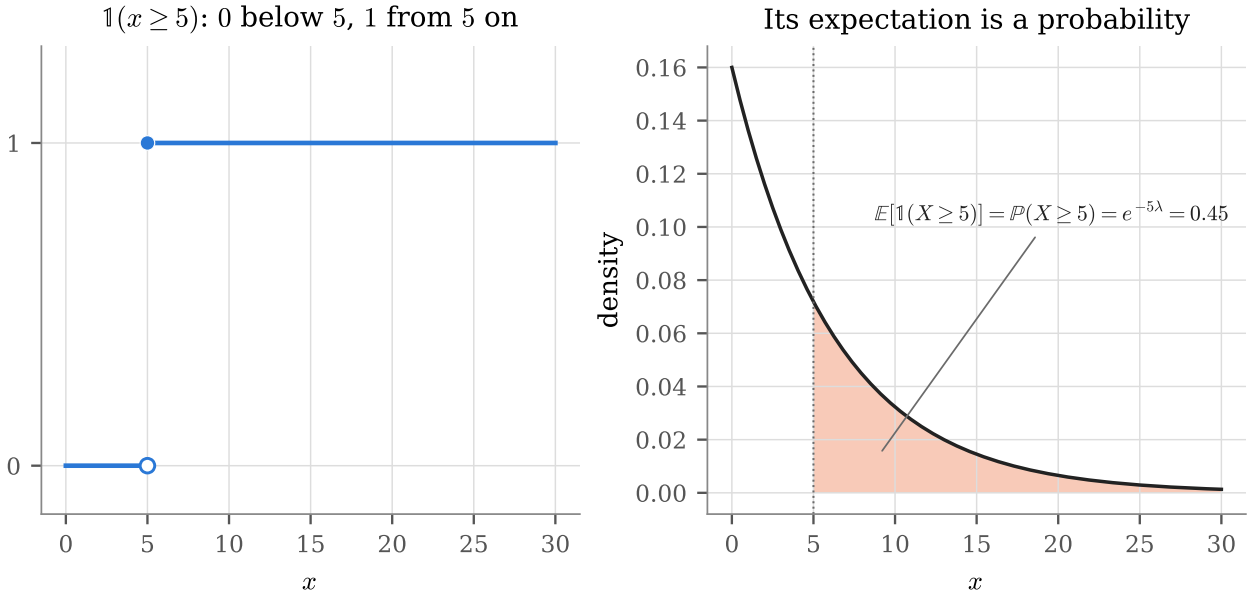


Figure A.4: Left: the indicator $\mathbb{1}(x \geq 5)$ is a step from 0 to 1. Right: for $X \sim \text{Expo}(0.16)$ its expectation is the shaded probability $\mathbb{P}(X \geq 5) = e^{-5\times 0.16} = 0.45$.

Where the book uses it. The CDF of a constant is a step function (Example 4.2); the censored observations of Chapter 2 are indicators of the unobserved variable; the Strong Law in Appendix B counts failures with a sum of indicators; Markov's inequality below is proved with one.

A.5 Random variables as functions, and 'almost surely'

A **random variable** is a function from the sample space (the set of possible outcomes of an experiment) to the real line. Figure A.5 draws the simplest case: the 36 outcomes of two dice, and the function $X =$ 'the sum'. The probability $\mathbb{P}(X = 7)$ is the probability of the set of outcomes that X sends to 7. Two random variables on the same sample space, say the sum X and the larger die Y, are two functions of the *same* outcome, which is what gives meaning to $X - Y$ and to 'X and Y are close': you compare them outcome by outcome. This is the viewpoint Section 3.1 needs for a sequence of random variables, which is a sequence of functions on one sample space.

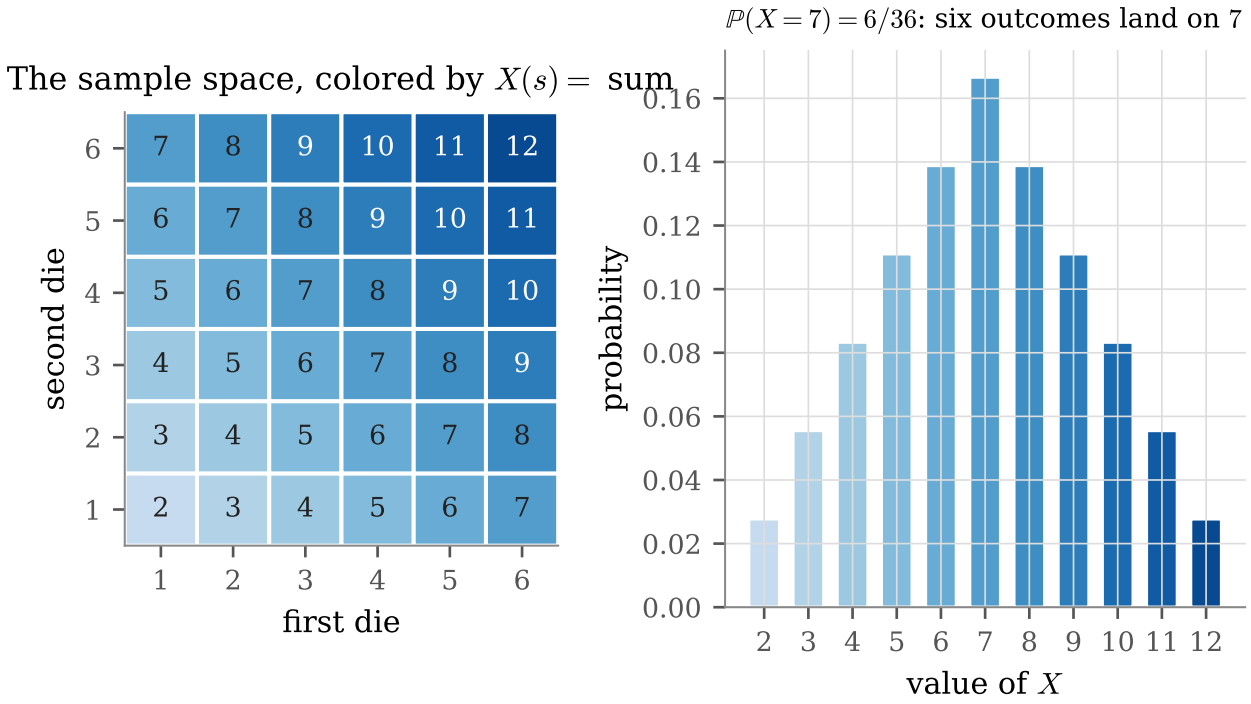


Figure A.5: A random variable is a function on the sample space. Left: the 36 outcomes of two dice, each colored by the value the function $X =$ 'sum' assigns it. Right: the distribution of X on the real line; $\mathbb{P}(X = 7) = 6/36$ because six outcomes are sent to 7.

Three phrases from probability theory appear in the book and mean less than they sound.

- **Almost surely** means 'with probability one'. A continuous random variable equals any particular number with probability zero, so '$X \neq 5$ almost surely' is true even though the outcome $X = 5$ is not impossible; the set of outcomes where a statement fails is allowed to be nonempty as long as it has probability zero. The book uses it in Hoeffding's inequality ('bounded almost surely': the probability of landing outside $[a, b]$ is zero) and in Appendix B.
- **Measurable** (a measurable function, a measurable sample space) is the technical condition that makes probabilities of the relevant events well defined. Every function you can write

down with formulas, limits, sums and integrals is measurable; the usual rule of thumb, that anything you can compute from the data is a statistic, is correct in practice. The book never needs more.

- A **probability measure** (or distribution) assigns a probability to every event. A PMF or PDF is a recipe for one: sum or integrate it over the event. The two determine each other, and the book writes down families of distributions through their PMFs and PDFs (Definition 2.1).

A.6 Densities, CDFs and quantiles

For a continuous random variable, **probability is area**: $\mathbb{P}(a \le X \le b)$ is the area under the density f between a and b. The **CDF** $F(x) = \mathbb{P}(X \le x)$ is the area accumulated up to x, so $\mathbb{P}(a \le X \le b) = F(b) - F(a)$, and F is the antiderivative of f: $f = F'$. The **quantile** reads the CDF backwards: given a probability u, the quantile is the x with $F(x) = u$; the book's q_α is the case $u = 1 - \alpha$, indexed by the probability *above* the cutoff. Figure A.6 draws the three on one example. The same three ideas work for a discrete random variable with sums in place of integrals: the PMF is the list of probabilities, the CDF the running total, and the CDF jumps at each possible value instead of rising smoothly.

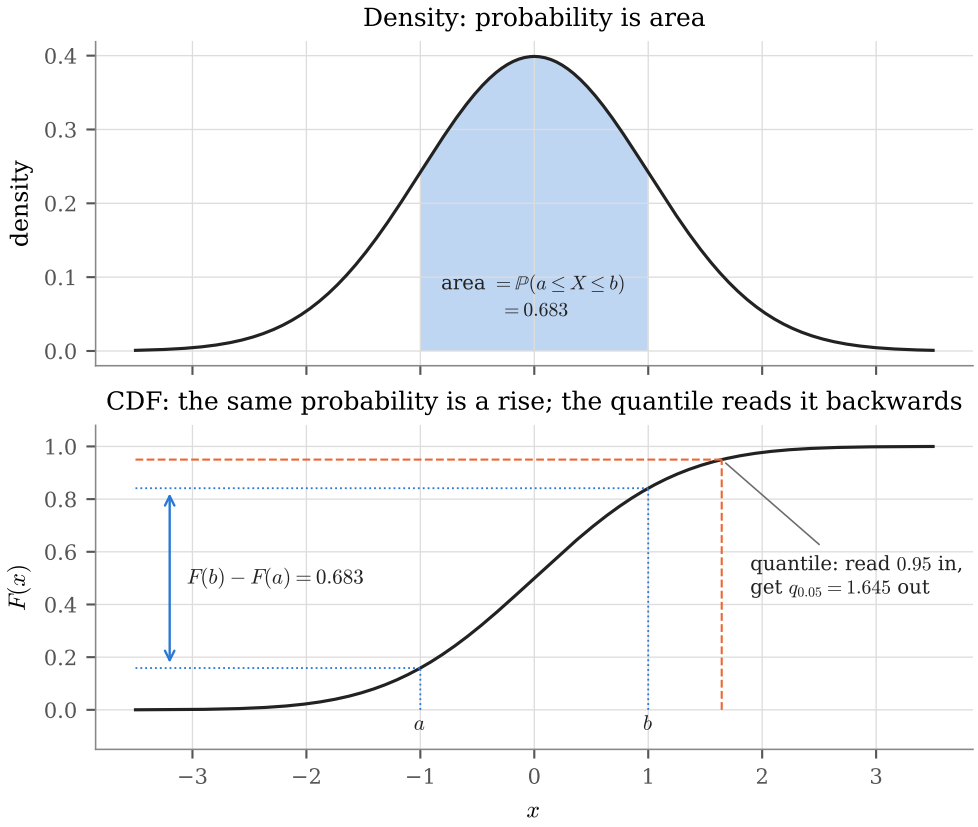


Figure A.6: Density, CDF and quantile as one picture. Top: the shaded area under the Gauss$(0, 1)$ density between $a = -1$ and $b = 1$ is $\mathbb{P}(a \le X \le b) = 0.683$. Bottom: the same probability is the rise of the CDF from $F(a)$ to $F(b)$; reading the CDF sideways at 0.95 gives the quantile $q_{0.05} = 1.645$ of Section 6.5.

Where the book uses it. Convergence in distribution (Definition 4.1) compares CDFs, and the

jumps of a discrete CDF are why only continuity points count. Every probability in Chapter 6 is a CDF lookup. The density of a transformed variable, such as $aX + b$ or X^2, is found by writing its CDF and differentiating (Section 6.3, Chapter 2).

A.7 Three inequalities

The triangle inequality: $|x + y| \le |x| + |y|$, and so $|x - z| \le |x - y| + |y - z|$: a detour through y is never shorter than the direct distance. Proofs use it to split an error into two pieces, for instance $|T_n - \mu| \le |T_n - \mathbb{E}[T_n]| + |\mathbb{E}[T_n] - \mu|$ in Proposition 3.3: if the total is at least ϵ, one of the two pieces is at least $\epsilon/2$.

Markov's inequality: if $Y \ge 0$ then $\mathbb{P}(Y \ge a) \le \mathbb{E}[Y]/a$ for every $a > 0$. The picture (Figure A.7, left) is the proof: the step function $a \cdot \mathbb{1}(Y \ge a)$ never exceeds Y, so its expectation, $a\mathbb{P}(Y \ge a)$, is at most $\mathbb{E}[Y]$. A non-negative variable with a small mean is rarely large.

Chebyshev's inequality: for any random variable with finite variance, $\mathbb{P}(|T - \mathbb{E}[T]| \ge \epsilon) \le \mathbb{V}\mathrm{ar}(T)/\epsilon^2$. It is Markov's inequality applied to $Y = (T - \mathbb{E}[T])^2$ with $a = \epsilon^2$. In standard-deviation units it says $\mathbb{P}(|T - \mathbb{E}[T]| \ge k\sigma) \le 1/k^2$ for every distribution whatsoever, which is both its strength (universal) and its weakness (loose: for a Gaussian the true two-standard-deviation tail is 0.046, the bound says 0.25).

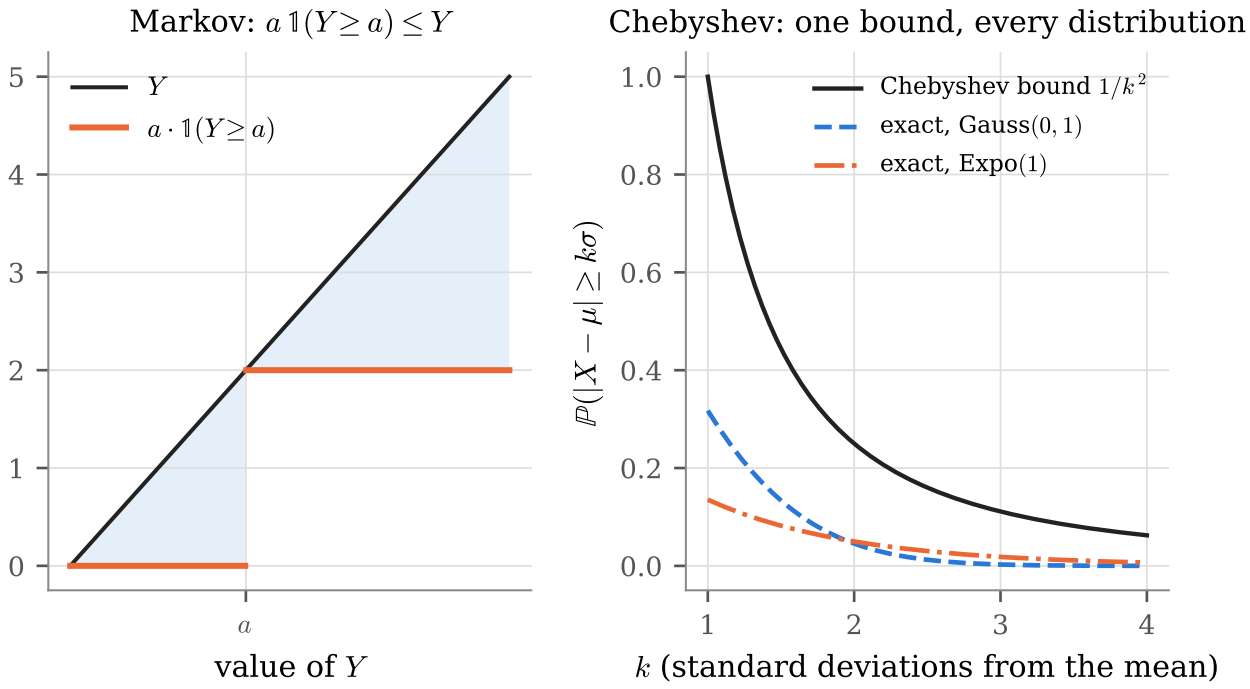


Figure A.7: Left: Markov's inequality is the picture $a \cdot \mathbb{1}(Y \ge a) \le Y$; take expectations. Right: Chebyshev's bound $1/k^2$ against the exact tail probabilities of a Gaussian and of an Expo(1) variable. The bound holds for every distribution and is far from tight for these two.

Where the book uses it. The Law of Large Numbers (Theorem 3.5) is Chebyshev's inequality with $n \to \infty$; the shrinking-variance criterion is the same with the triangle inequality in front; convergence in mean implies convergence in probability by Markov (Appendix B); Hoeffding's inequality (Theorem 4.9) is the sharper cousin for bounded variables.

A.8 Convex functions

A function is **convex** if the chord joining any two points of its graph lies on or above the graph, and **concave** if every chord lies on or below (Figure A.8). For a function with a second derivative, convex means $g'' \geq 0$ (the slope keeps increasing) and concave means $g'' \leq 0$. x^2, e^x and $1/x$ on $(0, \infty)$ are convex; $\log x$ and $\sqrt{x}$ are concave; $\sin x$ is neither. Averaging interacts with convexity in one direction only: for convex g, the average of g over two points is at least g of their average, which is the chord sitting above the curve at its midpoint.

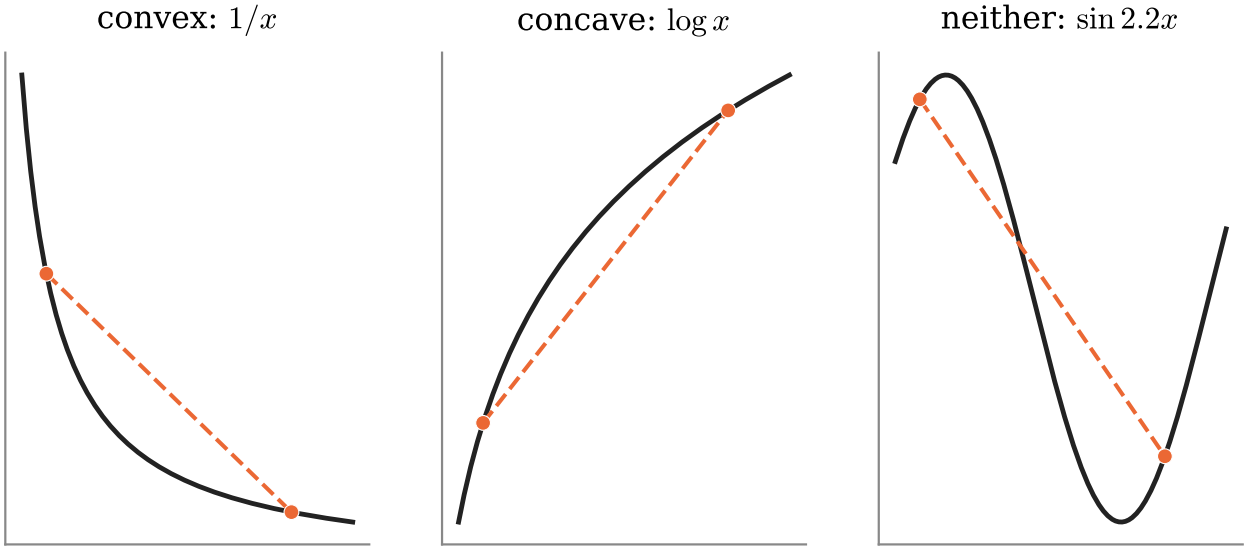


Figure A.8: Convex: chords above the graph. Concave: chords below. Neither: a chord can cross the graph.

Where the book uses it. Jensen's inequality (Theorem 7.8) is the chord statement with the two-point average replaced by an expectation; it explains why $1/\bar{T}_n$ overestimates λ and why $\sqrt{\cdot}$ of an unbiased estimator is biased (Chapter 7). The quadratic in the solve interval of Chapter 8 is convex, which is why the inequality holds between its roots. Likelihoods in Part IV are maximized by finding where a concave function's slope is zero.

A.9 Tangent lines, the mean value theorem and Taylor's expansion

The derivative $g'(\theta)$ is the slope of the tangent line to the graph of g at θ, and the tangent line $g(\theta) + g'(\theta)(x - \theta)$ is the best straight-line stand-in for g near θ. Two facts make this precise.

The mean value theorem. For a smooth g and any $a < b$ there is a point c between them where the tangent is parallel to the chord:

$$g(b) - g(a) = g'(c)\,(b - a) \quad \text{for some } c \in (a, b).$$

The change in g over an interval equals the slope somewhere inside the interval times the interval's length (Figure A.9, left). You never learn where c is, and you never need to: in the delta method it is squeezed between two points that converge to θ, so $g'(c)$ converges to $g'(\theta)$.

The first-order Taylor expansion. Near θ,

$$g(x) = g(\theta) + g'(\theta)(x - \theta) + (\text{an error of order } (x - \theta)^2),$$

so on a window of width h around θ the tangent line is wrong by about h^2, which is negligible compared with the h-sized changes along the line (Figure A.9, right). Adding the next term, $\frac{1}{2}g''(\theta)(x - \theta)^2$, gives the second-order expansion, which Part IV uses for the likelihood ratio test.

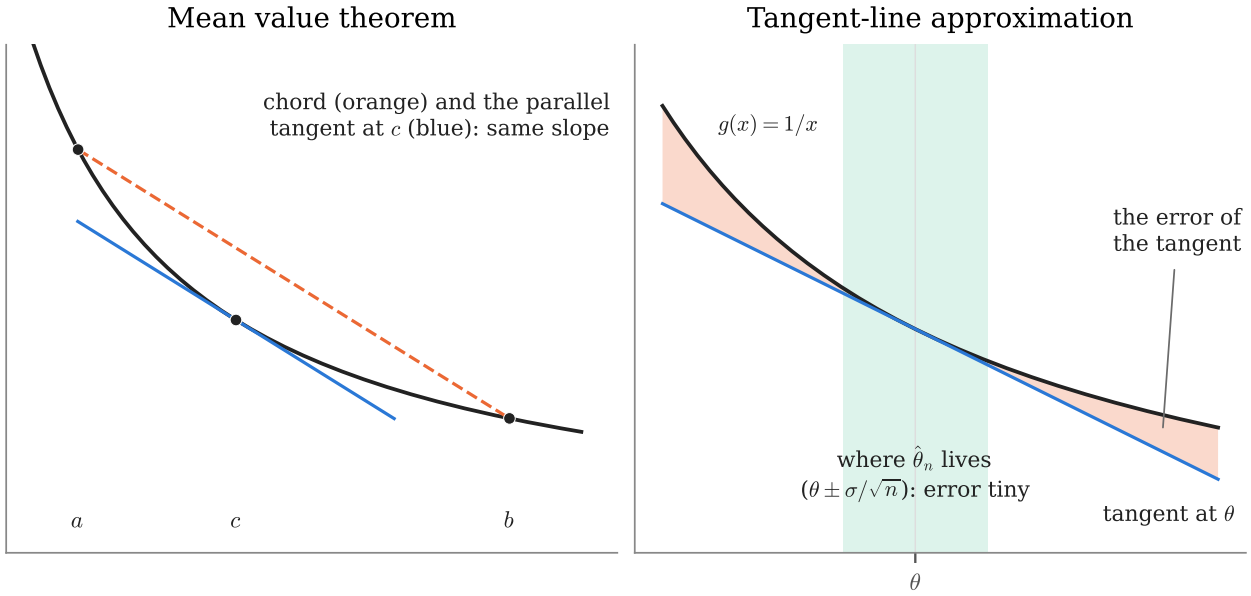


Figure A.9: Left: the mean value theorem. Between a and b there is a c where the tangent (blue) is parallel to the chord (orange). Right: the tangent line at θ approximates $g(x) = 1/x$; the error (shaded) is tiny on the window $\theta \pm \sigma/\sqrt{n}$ where an estimator lives, which is why the delta method can treat g as linear there.

Where the book uses it. The delta method (Theorem 5.4) is the mean value theorem followed by Slutsky's theorem; its intuition box is the tangent-line picture. The second-order expansion is behind the remark on $g'(\theta) = 0$ in Section 5.4 and behind Wilks' theorem in Chapter 15.

A.10 Exponentials and logarithms

Two facts about e recur. First, $(1 - x/n)^n \to e^{-x}$ as $n \to \infty$ (Figure A.10, left); more generally $(1 + a/n)^n \to e^a$. It is the step that turns the finite-n formula in the maximum-of-uniforms proof (Example 4.6) into an exponential limit, and it reappears whenever a product of n nearly-equal factors is pushed to the limit. Second, the logarithm turns multiplication into addition, $\log(ab) = \log a + \log b$ (Figure A.10, right). A likelihood is a product of n densities, a number so small it underflows a computer for moderate n; its logarithm is a sum of n ordinary terms, and a sum of n terms is what the Law of Large Numbers and the Central Limit Theorem know how to handle. That is the entire reason Part IV works with log-likelihoods. Because log is strictly increasing, maximizing the likelihood and maximizing the log-likelihood give the same answer.

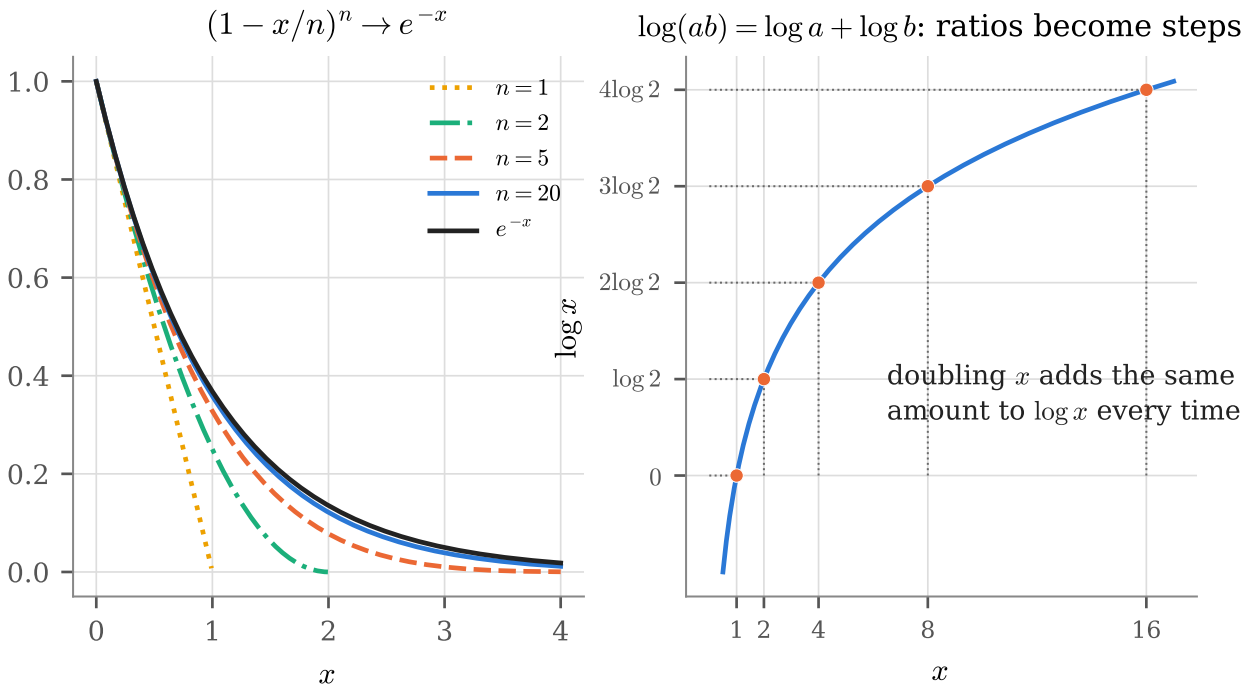


Figure A.10: Left: $(1 - x/n)^n$ for $n = 1, 2, 5, 20$ climbing toward e^{-x}; each curve is drawn up to $x = n$, where it reaches 0. Right: the logarithm turns equal ratios into equal steps, so a product becomes a sum.

A.11 Orders of magnitude: how fast is 'fast'

The book keeps saying that an error is 'of order $1/\sqrt{n}$', that a bound 'goes to zero like $1/n$', or 'exponentially fast'. These are statements about how a quantity scales with n, ignoring constants. Figure A.11 draws the three rates that matter here on ordinary axes and on log-log axes, where a power law n^{-c} is a straight line of slope $-c$ and an exponential bends away below every power law. Practical translation of each rate:

- $1/\sqrt{n}$: four times the data halves the error; a hundred times the data buys one decimal place. This is the rate of every estimator built from an average and of every confidence interval in this book.
- $1/n$: twice the data halves the error. Chebyshev's bound and the maximum-of-uniforms example have this rate; so does the bias of $1/\bar{T}_n$ in Chapter 7, which is why the bias is swamped by the $1/\sqrt{n}$ standard error.
- e^{-cn}: each fixed addition to n multiplies the error by the same factor below 1. Hoeffding's bound has this rate, which is why it overtakes Chebyshev's for large n (Figure 4.5).

A.12 Functions of several variables: gradient and Hessian

A parameter with several coordinates, like (μ, σ^2) for a Gaussian, turns a log-likelihood into a function of several variables, and maximizing it needs the several-variable versions of the derivative and the second derivative.

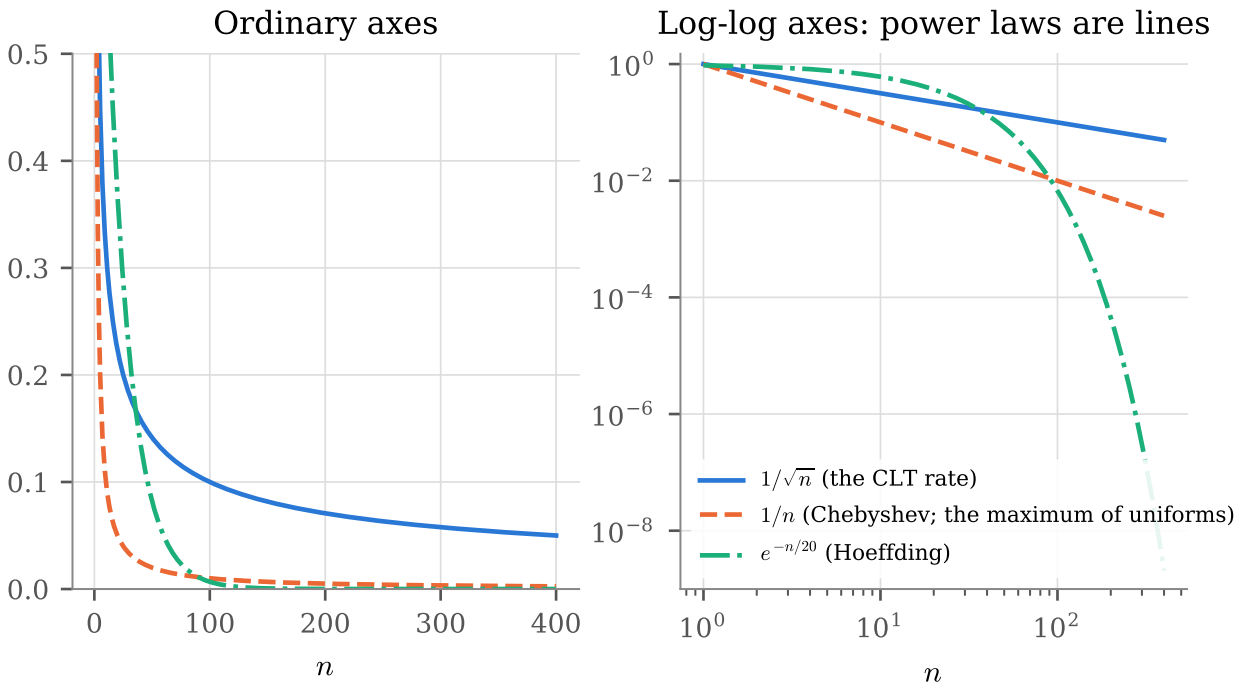


Figure A.11: Three rates of decay. Left: on ordinary axes they all look alike at first. Right: on log-log axes $1/\sqrt{n}$ and $1/n$ are straight lines whose slopes are the exponents, and the exponential drops away below both.

Partial derivatives and the gradient. For $h(a, b)$, the partial derivative $\partial h/\partial a$ is the ordinary derivative in a with b held fixed: the slope of the surface walking in the a direction. The **gradient** collects them in a vector, $\nabla h = (\partial h/\partial a, \partial h/\partial b)$. It points in the direction in which h climbs fastest, it is perpendicular to the contour line through its point, and its length is the steepness of the climb (Figure A.12, left). At the top of a hill the surface is flat in every direction, so the gradient is the zero vector: $\nabla h = 0$ is the first-order condition, one equation per coordinate.

The Hessian and the second-derivative test. The **Hessian** is the matrix of second partial derivatives,

$$\nabla^2 h = \begin{pmatrix} \partial^2 h/\partial a^2 & \partial^2 h/\partial a\,\partial b \\ \partial^2 h/\partial b\,\partial a & \partial^2 h/\partial b^2 \end{pmatrix},$$

which is symmetric for the smooth functions of this book. A zero gradient alone does not certify a maximum: the saddle $h(a, b) = a^2 - b^2$ has $\nabla h = 0$ at the origin, which is a minimum along a and a maximum along b (Figure A.12, right). The second-derivative test settles it. A symmetric matrix H is **negative definite** if $u^\top H u < 0$ for every nonzero vector u: the function curves down in every direction. If at a point where $\nabla h = 0$ the Hessian is negative definite, the point is a local maximum. For a 2×2 Hessian, negative definite means the top-left entry is negative and the determinant is positive. For the saddle the Hessian is $\begin{pmatrix} 2 & 0 \\ 0 & -2 \end{pmatrix}$, whose determinant is -4: not negative definite. For $h(a, b) = -(a^2 + ab + b^2)$ it is $\begin{pmatrix} -2 & -1 \\ -1 & -2 \end{pmatrix}$, top-left $-2 < 0$, determinant $3 > 0$: negative definite, and the origin is the maximum.

Concave in several variables. Suppose the Hessian is negative definite at every point of a domain

that contains the segment between any two of its points. Along the line $\theta + tu$, the second derivative in t is $u^\top \nabla^2 h\, u < 0$, so every slice through the function, in every direction, is strictly concave in one variable. If $\nabla h(\hat{\theta}) = 0$, the slice from $\hat{\theta}$ toward any other point θ has zero slope at $\hat{\theta}$, so by the one-variable result (Proposition 11.5) it is highest there: $h(\theta) < h(\hat{\theta})$. A point with zero gradient, if there is one, is the unique maximizer, exactly as in one variable.

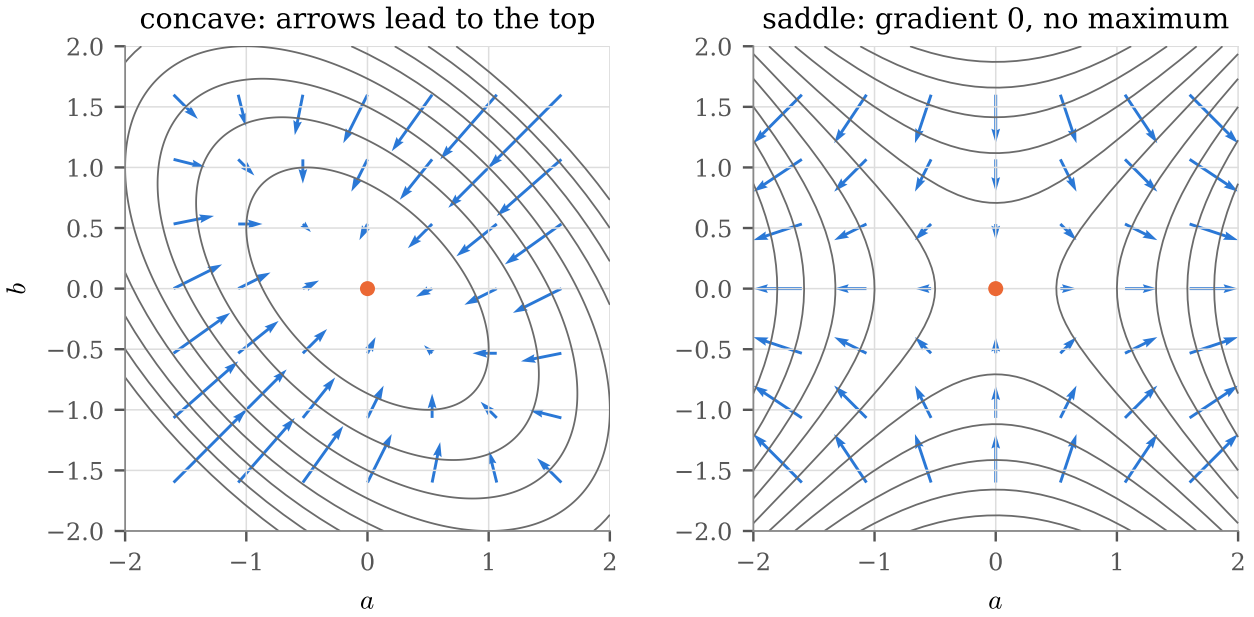


Figure A.12: Left: contours of the concave function $-(a^2 + ab + b^2)$ with its gradient drawn at a grid of points; each arrow points uphill, perpendicular to the contour, and the arrows shrink to nothing at the top, the one point where the gradient is 0. Right: the saddle $a^2 - b^2$; its gradient is also 0 at the origin, which is not a maximum.

Where the book uses it. The Gaussian maximum likelihood estimator (Section 11.6) is found by a two-step argument, and its Hessian at the maximum is negative definite; the EM algorithm (Section 11.8) maximizes in five coordinates; Chapter 12 turns the Hessian of the log-likelihood into the Fisher information matrix.

A.13 To be added with later chapters

Vectors, matrices, the rank of a matrix and projections (for Chapter 21); the chi-squared, Student's t and F families and how they arise from Gaussians (for Chapter 14); eigenvalues and eigenvectors (for Chapter 24).

Appendix B

The Modes of Convergence in Full

Chapters 3 and 4 used two modes of convergence, in probability and in distribution: the Law of Large Numbers lives in the first and the Central Limit Theorem in the second. There is a third, almost sure convergence, and a fourth, convergence in mean; all four can be ranked by strength. Nothing later in the book is proved with the two extra modes, which is why they're here rather than in the main sequence; where a later chapter mentions almost sure convergence, as Chapters 5 and 7 do, the in-probability version is the one in use. This appendix is for orientation, and for the reader who wants to know what the Strong Law of Large Numbers actually says.

B.1 Intuition for the three main modes

1. **Convergence in probability:** Imagine you have a sequence of r.v.s (like a sequence of coin flips or dice rolls). If this sequence converges in probability to a specific value, it means that as you progress through the sequence, the probability of the r.v.s being 'close' to that specific value tends to 1. In other words, with enough trials you can make it as likely as you please that the result is close to the value the sequence is converging to.

2. **Almost sure convergence (convergence almost surely):** Almost sure convergence is a stronger notion than convergence in probability. If a sequence of r.v.s converges almost surely to a value, it means that the probability of the sequence eventually staying close to that value is 1. More concretely, if you were to run the experiment and watch the whole infinite sequence unfold, the sequence you see converges, in the plain calculus sense, with probability 1. Almost sure convergence implies convergence in probability, but not conversely. Given that a sequence of r.v.s converging almost surely means that it converges to the value with probability 1, this mode of convergence is the closest analog to convergence of sequences of numbers as studied in calculus; it's a deterministic understanding of convergence applied one outcome at a time.

3. **Convergence in distribution:** Convergence in distribution is a weaker notion compared to the previous two. It doesn't focus on the actual values of the random variables but rather on their distribution. If a sequence of r.v.s converges in distribution to an r.v., it means that the shape of the distribution of the sequence becomes increasingly similar to the distribution of the limiting r.v. In simpler terms, the probabilities of the sequence falling within certain ranges will approximate the probabilities of the limiting variable falling within the same ranges.

The modes come in the order the book uses them. Convergence in probability first since the Law of Large Numbers lives there and it's where estimators are shown to work; convergence in distribution second since the Central Limit Theorem lives there and it's where confidence intervals and tests come from; and almost sure convergence last since the Strong Law is worth stating on its own account and is never needed again.

B.2 Convergence almost surely

Like convergence in probability, convergence almost surely involves a limiting process between random variables, say X_n and X, that live on the same sample space. Unlike convergence in probability, though, the limiting process for almost sure convergence is located *within* the probability measure. Compare the two definitions side by side:

$$\begin{aligned} &\text{in probability:} && \lim_{n\to\infty} \mathbb{P}\left(|X_n - X| \geq \epsilon\right) = 0 \quad \text{for every } \epsilon > 0, \\ &\text{almost surely:} && \mathbb{P}\left(\lim_{n\to\infty} X_n = X\right) = 1. \end{aligned}$$

In the first, you compute a probability for each n and then take a limit of numbers. In the second, you fix an outcome s of the experiment, which fixes the entire sequence of numbers $X_1(s), X_2(s), \ldots$, ask whether that sequence of numbers converges to the number $X(s)$ in the calculus sense, and then ask for the probability of the set of outcomes for which the answer is yes.

Definition B.1 (Convergence almost surely)**.** Let $(X_n)_{n\geq 1}$ be a sequence of r.v.s defined on a sample space S. The sequence **converges almost surely** to an r.v. X if:

$$\mathbb{P}\left(\left\{s \in S : \lim_{n\to\infty} X_n(s) = X(s)\right\}\right) = 1.$$

If $(X_n)_{n\geq 1}$ converges almost surely to the r.v. X, this is denoted by $X_n \xrightarrow[n\to\infty]{\text{a.s.}} X$.

> **In words.** Run the experiment once and watch the whole infinite sequence of numbers it produces; with probability one, that sequence converges in the plain calculus sense (toolkit A.1). The set of runs where it fails is allowed to be nonempty, as long as it has probability zero (toolkit A.5).
>
> **For the kiss data.** Keep observing couples forever and keep updating the running proportion; the Strong Law says the running proportion settles on p and stays there, for all but a probability-zero set of infinite runs.

'Almost surely' is probability's phrase for 'with probability 1'. The set of outcomes on which the sequence fails to converge is allowed to be nonempty, as long as it has probability 0.

Example B.2. Let $U \sim \text{Unif}([0,1])$ and define $X_n = U + U^n$. For every outcome with $U \in [0,1)$, $U^n \to 0$ as $n \to \infty$, so $X_n \to U$. For the outcome $U = 1$ the sequence is constant at 2 and does not converge to $U = 1$. But $\mathbb{P}(U = 1) = 0$, so $\mathbb{P}(X_n \to U) = \mathbb{P}(U \in [0,1)) = 1$ and $X_n \xrightarrow[n\to\infty]{\text{a.s.}} U$. This is the sense in which 'almost' is used: convergence can fail on a set of outcomes, provided that set has probability zero.

B.2.1 The Strong Law of Large Numbers (SLLN)

Theorem B.3 (Strong Law of Large Numbers)**.** *Let $X_1, X_2, \ldots$ be i.i.d. r.v.s with $\mathbb{E}|X_i| < \infty$ and $\mathbb{E}[X_i] = \mu$. Then:*

$$\bar{X}_n \xrightarrow[n\to\infty]{a.s.} \mu.$$

Two things are notable about the statement. First, it requires only a finite mean, not a finite variance; the Weak Law was proved in Chapter 3 using Chebyshev's inequality, which needs the variance, but the conclusion of both laws holds under the weaker assumption. Second, the conclusion is about sample paths: run the experiment once, forever, and the sequence of running averages you record converges to μ, with probability one. Figure B.1 shows six such paths for a fair coin. The two laws are often written together as

$$\bar{X}_n \xrightarrow[n\to\infty]{\text{a.s./}\mathbb{P}} \mu,$$

allowing whichever one is convenient to be used; for everything in this book the Weak Law of Chapter 3 is enough.

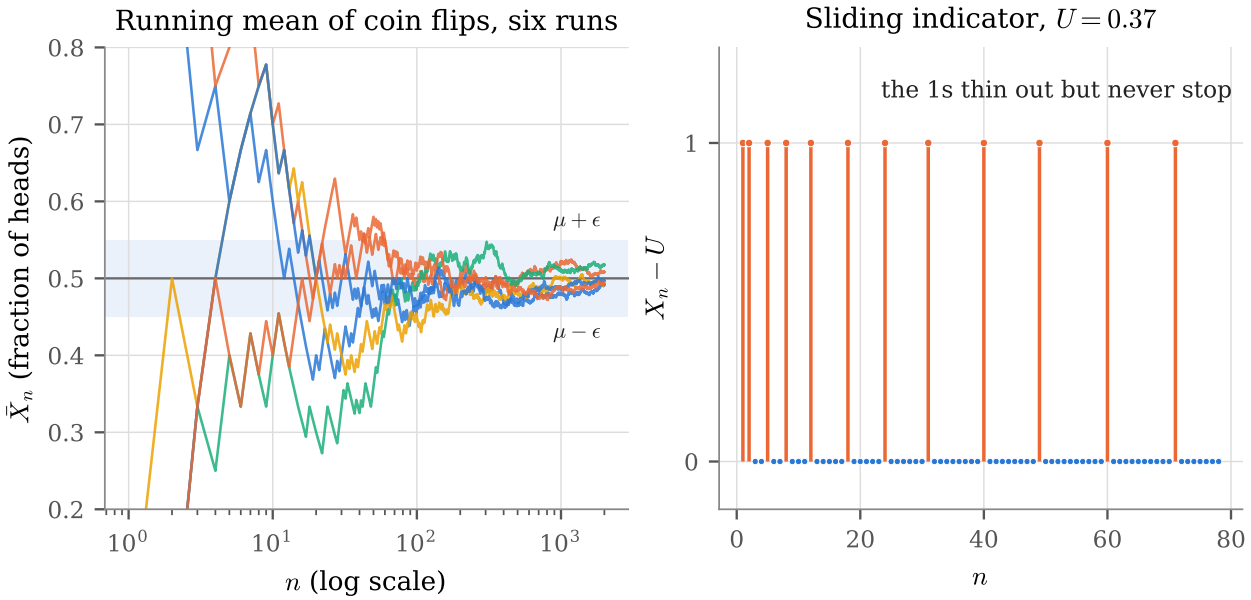


Figure B.1: Left: six realizations of the running proportion of heads in a fair coin, on a logarithmic n axis. Each path wanders early, enters the band $\mu \pm 0.05$ and, as the Strong Law says, stays there; with probability one every path eventually does this. Right: one outcome of the sliding-indicator sequence of Example B.4, which converges in probability but not almost surely. For this outcome ($U = 0.37$) the value 1 keeps recurring, at ever longer intervals, forever.

B.2.2 Convergence almost surely versus convergence in probability

While the Weak Law of Large Numbers illustrates convergence in probability and the Strong Law of Large Numbers shows convergence almost surely, what's the difference between the two? The following provides some intuition on the distinction between these two modes of convergence.

Suppose you have some device that improves over time. So every time you use the device the probability of it failing is less than before.

Convergence in probability says that the chance of failure goes to zero as the number of usages goes to infinity. So after using the device a large number of times, you can be very confident of it working correctly; it still might fail, but that's very unlikely.

On the other hand, convergence almost surely is a bit stronger. It says that the total number of failures is *finite*. That is, if you count the number of failures as the number of usages goes to infinity, you will get a finite number. The impact of this is as follows: As you use the device more and more, you will, after some finite number of usages, exhaust all failures. From then on the device will work perfectly. You don't actually know when you have exhausted all failures, so from a purely practical point of view, there isn't much difference between the two modes of convergence.

It's a good thing that the Strong Law of Large Numbers exists, as opposed to just the Weak Law. For example, a scientific experiment is justified in taking averages: after obtaining enough data, you can get arbitrarily close to the true average, and there will be only finitely many failures (however improbable) in the averaging process, after which the average stays close.

Here is what 'failures (however improbable) in the averaging process' means. Choose some $\delta > 0$ arbitrarily small. You obtain n estimates $X_1, \ldots, X_n$ of some other quantity that has some 'true' value, say μ. You compute the average:

$$S_n = \frac{1}{n} \sum_{i=1}^{n} X_i.$$

As you obtain more data (n increases) you can compute S_n for $n = 1, 2, \ldots$. The Weak Law says (under some assumptions about the X_n) that the probability $\mathbb{P}(|S_n - \mu| > \delta) \to 0$, as n goes to ∞. The Strong Law says that the number of times that $|S_n - \mu|$ is larger than δ is finite (with probability 1). That is, if we define the indicator function $\mathbb{1}(|S_n - \mu| > \delta)$ that returns one when $|S_n - \mu| > \delta$ and zero otherwise, then $\sum_{n=1}^{\infty} \mathbb{1}(|S_n - \mu| > \delta)$ is finite. This gives you considerable confidence in the value of S_n because it guarantees, with probability 1, the existence of some finite n_0 such that $|S_n - \mu| \leq \delta$ for all $n > n_0$. That is, the average never fails for $n > n_0$. Note that the Weak Law gives no such guarantee: it says failures become rare, not that they stop.

The following example, from Casella and Berger, is the standard illustration of a sequence for which failures become rare but never stop.

Example B.4 (The sliding indicator)**.** Let $U \sim \text{Unif}([0, 1])$ and define

$$\begin{aligned}
&X_1 = U + \mathbb{1}(U \in [0, 1]), \\
&X_2 = U + \mathbb{1}(U \in [0, \tfrac{1}{2}]), \qquad X_3 = U + \mathbb{1}(U \in [\tfrac{1}{2}, 1]), \\
&X_4 = U + \mathbb{1}(U \in [0, \tfrac{1}{3}]), \qquad X_5 = U + \mathbb{1}(U \in [\tfrac{1}{3}, \tfrac{2}{3}]), \qquad X_6 = U + \mathbb{1}(U \in [\tfrac{2}{3}, 1]), \\
&\quad \vdots
\end{aligned}$$

so the indicator's interval slides across $[0, 1]$ in blocks, and the intervals in the m-th block have length $1/m$. Each X_n equals either U or $U + 1$. *Convergence in probability holds:* for any $0 < \epsilon < 1$, $\mathbb{P}(|X_n - U| \geq \epsilon) = \mathbb{P}(U \in [a_n, b_n]) = b_n - a_n$, the length of the n-th interval, which goes to zero. So $X_n \xrightarrow[n\to\infty]{\mathbb{P}} U$. *Almost sure convergence fails:* fix any outcome $U = u$. In every block, exactly one interval (or two, at a shared endpoint) contains u, so $X_n(u) - u$ equals 1 at least once in every block, forever. The sequence of numbers $X_n(u)$ therefore takes the value $u + 1$ infinitely often and

cannot converge to u. This happens for every outcome, so $\mathbb{P}(X_n \to U) = 0$, the opposite of what almost sure convergence requires. The right panel of Figure B.1 shows one outcome.

In the device analogy, the sliding indicator is a device whose failures get rarer and rarer but never run out.

B.3 Convergence in mean

There's a fourth mode, which this book needs only in passing but which you've already used: the shrinking-variance criterion of Proposition 3.3 is a statement about it.

Definition B.5 (Convergence in r-th mean)**.** For $r \geq 1$, a sequence $(X_n)_{n\geq 1}$ **converges in r-th mean** (or in L^r) to X, written $X_n \xrightarrow[n\to\infty]{L^r} X$, if

$$\lim_{n\to\infty} \mathbb{E}\left[|X_n - X|^r\right] = 0.$$

The case $r = 2$ is called convergence in mean square (or quadratic mean).

> **In words.** The average size of the error, measured by its r-th power, goes to zero. For $r = 2$ this is 'the quadratic risk goes to zero'.
>
> **For the kiss data.** $\mathbb{E}[(\bar{R}_n - p)^2] = p(1-p)/n \to 0$, so $\bar{R}_n \to p$ in mean square; by Markov's inequality (toolkit A.7) that implies convergence in probability, which is the shrinking-variance criterion of Chapter 3 again.

Convergence in mean square to a constant μ is exactly the condition $\mathbb{E}[(X_n - \mu)^2] \to 0$, and since $\mathbb{E}[(X_n - \mu)^2] = \mathbb{V}\mathrm{ar}(X_n) + (\mathbb{E}[X_n] - \mu)^2$ (variance plus squared bias, the identity behind the quadratic risk of Chapter 7), it holds if and only if $\mathbb{E}[X_n] \to \mu$ and $\mathbb{V}\mathrm{ar}(X_n) \to 0$. Convergence in r-th mean implies convergence in probability, by Markov's inequality applied to $|X_n - X|^r$:

$$\mathbb{P}(|X_n - X| \geq \epsilon) = \mathbb{P}(|X_n - X|^r \geq \epsilon^r) \leq \frac{\mathbb{E}|X_n - X|^r}{\epsilon^r} \to 0.$$

Neither convergence in mean nor almost sure convergence implies the other.

B.4 The hierarchy of convergence types

A sequence of r.v.s might converge for one type of convergence, but not for another. Some convergence types are 'stronger' than others by which it's meant that if Type A convergence is stronger than Type B convergence, then Type A convergence implies Type B convergence. In Figure B.2, the stronger types of convergence are given at the top and as you move to the bottom the type of convergence becomes weaker. By way of example, if a sequence of r.v.s converges in probability, then it converges in distribution as well.

Why does this hierarchy exist? Let's focus on the 3 main modes of convergence. Consider the following points:

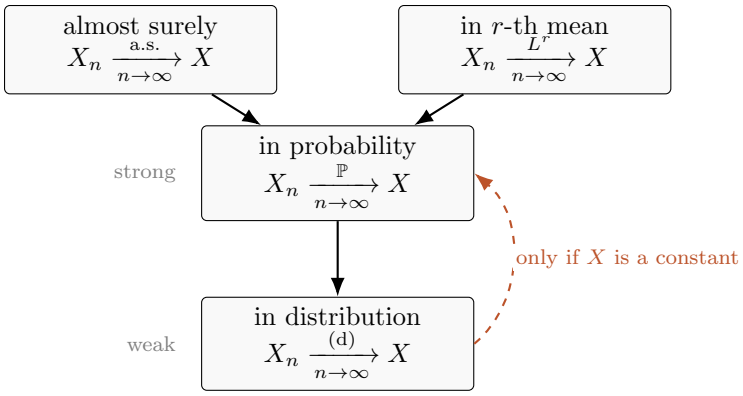


Figure B.2: The hierarchy. Solid arrows are implications that always hold; no solid arrow reverses. The dashed arrow is the one partial converse that matters: convergence in distribution to a *constant* is the same as convergence in probability to it. Almost sure convergence and convergence in mean don't imply each other.

- **Almost sure convergence:** If a sequence of r.v.s converges almost surely to an r.v., it means that the probability of the sequence eventually staying close to the limiting r.v. is 1. This is the strongest form of convergence because it tells us that the individual values of the r.v.s in the sequence will eventually be close to the limiting r.v. with probability 1. Since the sequence eventually stays within ϵ of the limit, the probability of being more than ϵ away at time n must go to zero, which is convergence in probability.

- **Convergence in probability:** If a sequence of r.v.s converges in probability to an r.v., it means that for any given $\epsilon > 0$, the probability of the difference between the random variables in the sequence and the limiting r.v. being greater than ϵ goes to 0 as the sequence progresses. This is weaker than almost sure convergence because it only tells us about the probability of the r.v.s being close to the limiting r.v. at each n separately, not that they will eventually be close with probability 1. The sliding indicator of Example B.4 is the example of the gap.

- **Convergence in distribution:** If a sequence of r.v.s converges in distribution to an r.v., it means that the cumulative distribution functions (CDFs) of the r.v.s in the sequence converge pointwise to the CDF of the limiting r.v. This is the weakest form of convergence because it only tells us about the behavior of the distribution of the r.v.s, not about the individual values or their probabilities of being close to the limiting r.v. If X_n is close to X with high probability then their CDFs must be close, which is why convergence in probability implies convergence in distribution; the $X_n = -Z$ example in the pitfall box of Section 4.1 shows the converse fails.

The dashed arrow in Figure B.2 is Proposition 4.7: convergence in distribution to a constant is convergence in probability to it.

B.5 Summary

Mode	Definition	Where it's used in this book
almost surely, $X_n \xrightarrow[n\to\infty]{\text{a.s.}} X$	$\mathbb{P}(X_n \to X) = 1$	Strong Law of Large Numbers
in probability, $X_n \xrightarrow[n\to\infty]{\mathbb{P}} X$	$\mathbb{P}(\lvert X_n - X\rvert \geq \epsilon) \to 0$ for every $\epsilon > 0$	Weak LLN; consistency of estimators; the constant in Slutsky's theorem
in distribution, $X_n \xrightarrow[n\to\infty]{\text{(d)}} X$	$F_{X_n}(x) \to F_X(x)$ at every continuity point of F_X	CLT; confidence intervals; test statistics
in r-th mean, $X_n \xrightarrow[n\to\infty]{L^r} X$	$\mathbb{E}\lvert X_n - X\rvert^r \to 0$	shrinking-variance criterion ($r = 2$); quadratic risk

B.6 Exercise

Exercise B.1. Is the following statement correct? 'If $(T_n)_{n\geq 1}$ converges in probability, then it also converges almost surely.' *(Solution on page 231.)*

Appendix C

Supplementary Proofs

Some results are stated in a chapter because they explain something there, but nothing later in the book depends on them. Their proofs are collected here so the chapters can keep moving. Each section says where the result is stated and confirms that no later argument uses it.

C.1 Pinsker's inequality

Theorem 10.8 in Chapter 10 states that for any distributions $\mathbb{P}$ and $\mathbb{Q}$ on the same sample space,

$$\mathrm{TV}(\mathbb{P}, \mathbb{Q}) \leq \sqrt{\tfrac{1}{2}\mathrm{KL}(\mathbb{P}, \mathbb{Q})}.$$

It is there to say that a fit that is good in KL is good for every probability computed from it. **Nothing later in the book depends on it**: no theorem, proof or exercise after Chapter 10 uses Pinsker's inequality, and you can skip this section without losing anything.

Proof strategy (for completeness; skip on a first reading). Two steps. First prove the inequality for two Bernoulli distributions, by calculus in one variable. Then show that collapsing $\mathbb{P}$ and $\mathbb{Q}$ to 'in the event $A^\star$ or not' turns them into two Bernoullis with the same total variation and no larger KL divergence. The second step uses the same tangent-line bound $-\log y \geq 1 - y$ as the proof that KL is nonnegative (Proposition 10.6).

Step 1: two Bernoulli distributions

Lemma C.1. *For $p, q \in (0, 1)$,*

$$\mathrm{KL}(\mathit{Bern}(p), \mathit{Bern}(q)) = p \log \frac{p}{q} + (1 - p) \log \frac{1 - p}{1 - q} \geq 2(p - q)^2.$$

The inequality also holds for $p \in [0, 1]$ and $q \in (0, 1)$, with the convention $0 \log 0 = 0$.

Proof strategy (for completeness; skip on a first reading). Fix p, treat the gap between the two sides as a function of q, and show by its derivative that it is smallest, and equal to 0, at $q = p$.

Proof. Fix p and look at the gap as a function of q:

$$g(q) = p \log \frac{p}{q} + (1 - p) \log \frac{1 - p}{1 - q} - 2(p - q)^2.$$

At $q = p$ both terms are 0, so $g(p) = 0$. Differentiate in q:

$$g'(q) = -\frac{p}{q} + \frac{1-p}{1-q} + 4(p-q).$$

The first two terms combine over the common denominator $q(1-q)$: $-p(1-q) + (1-p)q = q - p$. So

$$g'(q) = \frac{q-p}{q(1-q)} - 4(q-p) = (q-p)\left(\frac{1}{q(1-q)} - 4\right).$$

The product $q(1-q)$ is at most $\frac{1}{4}$, since $\frac{1}{4} - q(1-q) = (q - \frac{1}{2})^2 \geq 0$ (it is the variance of one Bernoulli observation, largest for a fair coin). So $1/(q(1-q)) \geq 4$ and the bracket is never negative. The sign of $g'(q)$ is therefore the sign of $q - p$: g decreases until $q = p$ and increases after. Its smallest value is $g(p) = 0$, so $g(q) \geq 0$ for every q. □

Figure C.1 draws the lemma. The divergence curve sits above the parabola $2(p-q)^2$ and touches it only at $q = p$. Near $q = p$ the two are close: the divergence is approximately $(p-q)^2/(2p(1-p))$ (Section 10.5), and $1/(2p(1-p)) \geq 2$, with equality only for $p = \frac{1}{2}$. The parabola $2(p-q)^2$ is the divergence's own parabola for the hardest case, the fair coin, and every other Bernoulli curves up faster.

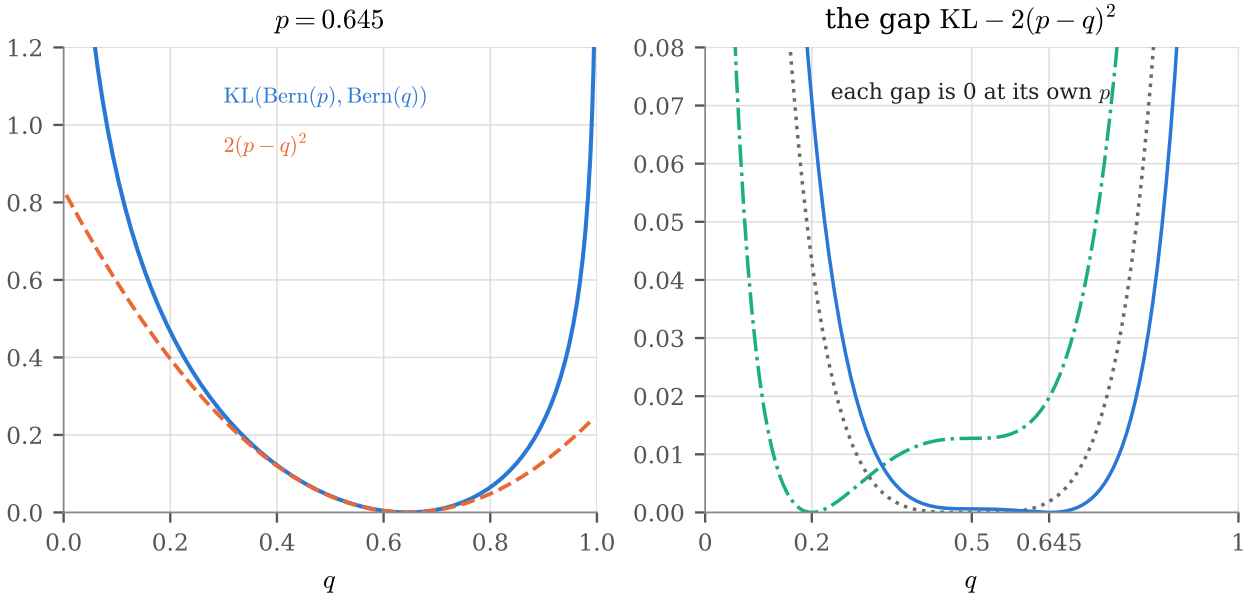


Figure C.1: The Bernoulli case of Pinsker's inequality. Left: with $p = 0.645$, $\mathrm{KL}(\mathrm{Bern}(p), \mathrm{Bern}(q))$ (blue) lies above $2(p-q)^2$ (orange, dashed) and touches it only at $q = p$. Right: the gap between the two for $p = 0.2$ (aqua, dash-dot), 0.5 (gray, dotted) and 0.645 (blue, solid); each gap is 0 at $q = p$ and positive elsewhere.

For the kiss data, $p = 80/124$ and $q = \frac{1}{2}$: KL $= 0.0428$ and $2(p-q)^2 = 2 \times 0.1452^2 = 0.0422$. Taking square roots, $|p-q| = 0.1452 \leq \sqrt{0.0428/2} = 0.1462$. Since $\mathrm{TV}(\mathrm{Bern}(p), \mathrm{Bern}(q)) = |p-q|$, the lemma is Pinsker's inequality for Bernoulli distributions.

Step 2: reducing any pair of distributions to two Bernoullis

Lemma C.2 (The log-sum inequality)**.** *For nonnegative numbers $a_1, \dots, a_k$ and nonnegative numbers $b_1, \dots, b_k$ with $b_i > 0$ whenever $a_i > 0$, and with $a = \sum_i a_i > 0$ and $b = \sum_i b_i$,*

$$\sum_{i=1}^{k} a_i \log \frac{a_i}{b_i} \geq a \log \frac{a}{b},$$

with the convention $0 \log 0 = 0$. The same holds with the sums replaced by integrals over a region.

Proof strategy (for completeness; skip on a first reading). Rewrite the difference of the two sides as a weighted average of $-\log$ of ratios and apply the tangent-line bound $-\log y \geq 1 - y$ to each.

Proof. Only the terms with $a_i > 0$ contribute on the left. For those, write

$$\sum_{i:a_i>0} a_i \log \frac{a_i}{b_i} - a \log \frac{a}{b} = \sum_{i:a_i>0} a_i \log \left(\frac{a_i}{b_i} \cdot \frac{b}{a} \right) = \sum_{i:a_i>0} a_i \left[-\log \left(\frac{b_i}{a_i} \cdot \frac{a}{b} \right) \right],$$

using $\sum_{i:a_i>0} a_i = a$ to bring $a \log(a/b)$ inside the sum. Apply $-\log y \geq 1 - y$ to each bracket:

$$\geq \sum_{i:a_i>0} a_i \left(1 - \frac{b_i}{a_i} \cdot \frac{a}{b} \right) = a - \frac{a}{b} \sum_{i:a_i>0} b_i \geq a - \frac{a}{b} \cdot b = 0.$$

The last step uses that the b_i with $a_i > 0$ add up to at most b. For integrals the same lines go through with $\int$ in place of $\sum$. □

In words: lumping several outcomes into one can only lower the divergence. If you record only whether an outcome fell in a set, and not which outcome it was, the evidence it carries cannot go up.

Proof strategy (for completeness; skip on a first reading). Collapse both distributions to the indicator of $A^\star$: total variation is unchanged, KL can only fall (Lemma C.2), and for the two resulting Bernoullis Lemma C.1 applies.

Proof of Theorem 10.8. If $\mathrm{KL}(\mathbb{P}, \mathbb{Q}) = +\infty$ there is nothing to prove, so assume it is finite; then $q(x) > 0$ wherever $p(x) > 0$. Let $A^\star = \{x : p(x) > q(x)\}$, and write $a = \mathbb{P}(A^\star)$, $b = \mathbb{Q}(A^\star)$. By Proposition 10.2, $\mathrm{TV}(\mathbb{P}, \mathbb{Q}) = a - b$. If $a = b$, TV $= 0$ and there is nothing to prove, so assume $a > b$; then $0 < a$ and $b < 1$.

Split the sum (or integral) defining $\mathrm{KL}(\mathbb{P}, \mathbb{Q})$ into the part over $A^\star$ and the part over its complement, and apply Lemma C.2 to each part, with $a_i = p(x)$ and $b_i = q(x)$:

$$\mathrm{KL}(\mathbb{P}, \mathbb{Q}) = \sum_{x \in A^\star} p(x) \log \frac{p(x)}{q(x)} + \sum_{x \notin A^\star} p(x) \log \frac{p(x)}{q(x)} \geq a \log \frac{a}{b} + (1 - a) \log \frac{1 - a}{1 - b}.$$

(If a part has total $\mathbb{P}$-probability 0, it contributes 0 on both sides.) The right-hand side is $\mathrm{KL}(\mathrm{Bern}(a), \mathrm{Bern}(b))$, the divergence between the two distributions of the indicator $\mathbb{1}\{X \in A^\star\}$. By Lemma C.1 (in the extended form for $a \in [0, 1]$; if $b = 0$ the divergence on the right is $+\infty$, which contradicts the finite left side, so $b \in (0, 1)$),

$$\mathrm{KL}(\mathbb{P}, \mathbb{Q}) \geq \mathrm{KL}(\mathrm{Bern}(a), \mathrm{Bern}(b)) \geq 2(a - b)^2 = 2\,\mathrm{TV}(\mathbb{P}, \mathbb{Q})^2.$$

Divide by 2 and take square roots. □

Intuition: why the proof works

Total variation is decided by one event, $A^\star$, where $\mathbb{P}$ beats $\mathbb{Q}$. Throw away everything about an outcome except whether it fell in $A^\star$. That leaves a coin, with probability a under $\mathbb{P}$ and b under $\mathbb{Q}$, and the same total variation $a - b$. Lumping can only lower KL, so the coin's KL is at most the original one. And for coins, the lemma says KL is at least twice the squared total variation.

C.2 Each step of the EM algorithm raises the likelihood

Section 11.8 in Chapter 11 states that one E step followed by one M step never lowers the log-likelihood of a mixture. It explains the climbing curves of Figure 11.7. **Nothing later in the book depends on it.**

Proof strategy (for completeness; skip on a first reading). Split the log-likelihood of the observed data into two pieces: the quantity the M step maximizes, minus a term that measures how well the parameter predicts the hidden labels. The M step raises the first piece. The second piece can only fall away from its value at the current parameter, by an amount that is a KL divergence and so is never negative.

Proof. Write θ for all five parameters, θ_t for the current values and θ_{t+1} for the values after one E step and one M step. For observation i, with hidden label $z \in \{0, 1\}$, let $p_\theta(x_i, z)$ be the joint probability (density in x_i) of the observation and the label, $p_\theta(x_i) = \sum_z p_\theta(x_i, z)$ the mixture density, and $p_\theta(z \mid x_i) = p_\theta(x_i, z)/p_\theta(x_i)$ the conditional probability of the label. Write $q_i(z) = p_{\theta_t}(z \mid x_i)$ for the current weights: $q_i(1) = w_i$ and $q_i(0) = 1 - w_i$.

For any θ and each z, $\log p_\theta(x_i) = \log p_\theta(x_i, z) - \log p_\theta(z \mid x_i)$. Average this over z with the weights $q_i(z)$, which add up to 1, and add over i:

$$\ell_n(\theta) = \underbrace{\sum_{i=1}^{n} \sum_{z} q_i(z) \log p_\theta(x_i, z)}_{Q(\theta)} \; - \; \underbrace{\sum_{i=1}^{n} \sum_{z} q_i(z) \log p_\theta(z \mid x_i)}_{H(\theta)} .$$

$Q(\theta)$ is the complete-data log-likelihood with each label replaced by its weight, which is exactly what the M step maximizes over all five parameters (Exercise 11.9), so $Q(\theta_{t+1}) \geq Q(\theta_t)$. This holds for the sum over i, not observation by observation: a step can lower the density of some observations while raising the total. For the second term,

$$H(\theta_t) - H(\theta) = \sum_{i=1}^{n} \sum_{z} q_i(z) \log \frac{p_{\theta_t}(z \mid x_i)}{p_\theta(z \mid x_i)} = \sum_{i=1}^{n} \mathrm{KL}\big(p_{\theta_t}(\cdot \mid x_i),\, p_\theta(\cdot \mid x_i)\big) \geq 0$$

by Proposition 10.6, term by term, since $q_i = p_{\theta_t}(\cdot \mid x_i)$. So $H(\theta_{t+1}) \leq H(\theta_t)$, and

$$\ell_n(\theta_{t+1}) - \ell_n(\theta_t) = \big(Q(\theta_{t+1}) - Q(\theta_t)\big) + \big(H(\theta_t) - H(\theta_{t+1})\big) \geq 0.$$

□

Ascent says nothing about where the climb ends: EM stops where the gradient of ℓ_n is 0, which can be a local maximum or a saddle point.

C.3 The uniform law of large numbers

The consistency proof of Theorem 12.1 in Chapter 12 needs the Law of Large Numbers to hold for all candidates θ at once, not one at a time. That is the one step of that chapter not proved there. **Nothing later in the book depends on this section** beyond Theorem 12.1 itself; the asymptotic normality theorem uses consistency, not the uniform law directly.

Recall. The Law of Large Numbers (Theorem 3.5): for i.i.d. $X_1, X_2, \dots$ and a function h with $\mathbb{V}\mathrm{ar}(h(X)) < \infty$, $\frac{1}{n}\sum_{i=1}^{n} h(X_i) \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}[h(X)]$.

Theorem C.3 (Uniform law of large numbers)**.** *Let $X_1, \dots, X_n$ be i.i.d. and let $g(x, \theta)$ be a function of an observation and a parameter $\theta \in \Theta$, with $\Theta \subset \mathbb{R}^d$ bounded and closed. Suppose that $g(x, \theta)$ is continuous in θ for every x, and that there is an envelope $G(x) \geq \sup_\theta |g(x, \theta)|$ with $\mathbb{E}[G(X)^2] < \infty$. Then, with $M(\theta) = \mathbb{E}[g(X, \theta)]$,*

$$\sup_{\theta\in\Theta} \left| \frac{1}{n}\sum_{i=1}^{n} g(X_i, \theta) - M(\theta) \right| \xrightarrow[n\to\infty]{\mathbb{P}} 0.$$

Proof strategy (for completeness; skip on a first reading). Cover Θ by finitely many small balls. On each ball the function g varies little, so the average over the ball's center controls the average everywhere in the ball; the ordinary LLN handles the finitely many centers at once, and continuity plus the envelope make the within-ball variation as small as you like.

Proof. Fix $\epsilon > 0$. For $\theta \in \Theta$ and $r > 0$ let $B(\theta, r)$ be the ball of radius r around θ, and define the local oscillation

$$w_r(x, \theta) = \sup_{\theta' \in B(\theta, r) \cap \Theta} |g(x, \theta') - g(x, \theta)|.$$

For each fixed x and θ, continuity of g in θ gives $w_r(x, \theta) \to 0$ as $r \to 0$, and $0 \leq w_r(x, \theta) \leq 2G(x)$. So by dominated convergence (toolkit A.1) (the dominating function is $2G$, integrable), $\mathbb{E}[w_r(X, \theta)] \to 0$ as $r \to 0$, for each θ.

A bounded closed set has the following property: for any way of attaching a ball $B(\theta, r_\theta)$ to each of its points, finitely many of those balls already cover it. Choose r_θ for each θ so small that $\mathbb{E}[w_{r_\theta}(X, \theta)] < \epsilon$, and pick a finite cover $B(\theta_1, r_1), \dots, B(\theta_K, r_K)$.

Now take any $\theta \in \Theta$, say in the ball around θ_k. Split the error at θ into three pieces:

$$\left| \frac{1}{n}\sum_i g(X_i, \theta) - M(\theta) \right| \leq \underbrace{\frac{1}{n}\sum_i |g(X_i, \theta) - g(X_i, \theta_k)|}_{\leq \frac{1}{n}\sum_i w_{r_k}(X_i, \theta_k)} + \underbrace{\left| \frac{1}{n}\sum_i g(X_i, \theta_k) - M(\theta_k) \right|}_{\text{LLN at } \theta_k} + \underbrace{|M(\theta_k) - M(\theta)|}_{\leq \mathbb{E}[w_{r_k}(X, \theta_k)] < \epsilon}.$$

Taking the supremum over θ amounts to taking the maximum over the finitely many k, so

$$\sup_\theta \left| \frac{1}{n}\sum_i g(X_i, \theta) - M(\theta) \right| \leq \max_{k\leq K} \frac{1}{n}\sum_i w_{r_k}(X_i, \theta_k) + \max_{k\leq K} \left| \frac{1}{n}\sum_i g(X_i, \theta_k) - M(\theta_k) \right| + \epsilon.$$

Each of the $2K$ averages on the right is an ordinary average of i.i.d. terms with finite variance (bounded by $2G$ or G, with $\mathbb{E}[G^2] < \infty$), so by the LLN the first maximum converges in probability

to $\max_k \mathbb{E}[w_{r_k}(X, \theta_k)] < \epsilon$ and the second to 0. A maximum of finitely many sequences converging in probability converges in probability to the maximum of the limits. Hence with probability tending to 1 the right side is below 3ϵ, and since ϵ was arbitrary, the left side converges in probability to 0. $\square$

The theorem is applied in Theorem 12.1 with $g(x, \theta) = \log p_\theta(x)$, where the envelope $G = \sup_\theta |\log p_\theta(X)|$ is the hypothesis stated there (with a finite second moment, which the book's version of the LLN needs).

Solutions to the Exercises

Chapter 1. The Kiss Experiment and What Statistics Is For

Exercise 1.1. Call 7: six of the 36 equally likely pairs sum to 7, more than any other total, so the probability is $1/6$ and the expected winning is $100/6 \approx \$16.67$. The mechanism (fair, independent dice) is known in full, so there is nothing to learn from data; the answer is deduced.

Exercise 1.2. The count is $\text{Bin}(100, 0.8)$ with mean $np = 80$ and standard deviation $\sqrt{np(1-p)} = \sqrt{16} = 4$. The value 65 is $15/4 = 3.75$ standard deviations below the mean; a Gaussian approximation gives $\mathbb{P}(\text{at least } 65) \approx \Phi(3.75) \approx 0.9999$, and the exact Binomial computation gives 0.99985.

Exercise 1.3. $\mathbb{P}(\text{at least } 2) = 3p^2(1-p) + p^3$, which is 0.282 at $p = 0.35$ and 0.718 at $p = 0.65$. The data are only about 2.5 times more probable under the larger p, so three couples cannot distinguish the two truths; with 124 couples the corresponding ratio is astronomically large (Figure 1.2).

Exercise 1.4. (a) Independence: the 124 records are really 62 independent observations, so $\hat{p}$ is as variable as a sample of 62, not 124, though its expectation is still p. (b) Identical distribution and independence: those couples share a p_i near 1 that is not the population's p, and they all followed the same instruction, so their outcomes are not independent; $\hat{p}$ is pushed upward. (c) The sample is not drawn from the population of interest at all: if being visible is related to the direction of the turn (it plausibly is, since from one side the observer sees the faces of couples who turn one way and the backs of the heads of couples who turn the other), the R_i are Bernoulli with a parameter that differs from p, and no amount of data fixes it. The third is the most dangerous because it is invisible in the data.

Chapter 2. Statistical Models and Identifiability

Exercise 2.1. (2). The first has a sample space that excludes the outcome 0, which every $\text{Bern}(p)$ puts mass on. The second is a legitimate model with a restricted parameter space; whether it's well-specified is a separate question.

Exercise 2.2. The intended answer is (2), and (1) is the one that is definitely wrong: the sample space can't depend on the unknown a. Both $\mathbb{R}_+$ and $\mathbb{R}$ contain every observation for every a, so by Definition 2.1 (3) is a legitimate model too; the convention is to take $E = \mathbb{R}_+$, the smallest natural set that works, and that convention is what the question is testing.

Exercise 2.3. B, C, D, E, F. In (A) a member of the family needs infinitely many numbers (a probability for every non-negative integer). (B), (C), (F) need one number, (E) two, (D) three (the fourth p_i is determined by the other three). In (G) a member is a continuous function on $[0, 1]$, which no finite list of numbers specifies.

Exercise 2.4. $\mathbb{P}(Y_i \geq c) = 1 - \Phi(c - \mu) = \Phi(\mu - c)$, strictly increasing in μ, so μ is identifiable: $\mu = c + \Phi^{-1}(p)$. With $\mu = 0$ known and σ unknown, $\mathbb{P}(Y_i \geq c) = \Phi(-c/\sigma)$: for $c \neq 0$ this is strictly monotone in σ and σ is identifiable; for $c = 0$ it equals $1/2$ for every σ, and σ is not. The threshold matters: thresholding at the known mean throws away all information about the spread, as in the uniform example.

Exercise 2.5. Ordering the means removes the label-switching ambiguity, and it's a theorem (not proved here) that finite Gaussian mixtures are then identifiable. The endpoint $\pi = 1$ is excluded from Θ for a good reason: at $\pi = 1$ the density is f_1 whatever (μ_2, σ_2^2) are, so every value of the second component's parameters gives the same distribution. A component with zero weight is invisible, which is why the weight is required to be in the open interval $(0, 1)$.

Chapter 3. The Law of Large Numbers

Exercise 3.1. (a) $\mathbb{P}(|X_n| > \epsilon) = \mathbb{P}(X_n = 1) = 1/n$. (b) Yes, to 0, since $1/n \to 0$. (c) Yes, to 0: convergence in probability implies convergence in distribution (Proposition 4.4 in the next chapter), or directly, the CDF of X_n is $1 - 1/n$ on $[0, 1)$ and 1 from 1 on, which converges to the CDF of the constant 0 at every x (at $x = 0$ too, here). (d) $\mathbb{E}[\cos X_n] = (1 - \frac{1}{n}) \cos 0 + \frac{1}{n} \cos 1 \to 1$; the next chapter's Proposition 4.3 says this is no accident, since cos is continuous and bounded.

Exercise 3.2. $\mathbb{E}[\bar{X}_n] = 1/\lambda$ for every n and $\mathbb{V}\mathrm{ar}(\bar{X}_n) = 1/(n\lambda^2) \to 0$, so Proposition 3.3 applies. Chebyshev with $\epsilon = 0.1/\lambda$ gives $\mathbb{P}(|\bar{X}_n - 1/\lambda| \geq 0.1/\lambda) \leq \frac{1/(n\lambda^2)}{0.01/\lambda^2} = \frac{100}{n}$, which is at most 0.05 once $n \geq 2000$. The CLT of the next chapter, used as in its exercise 4, gives $n \approx 1.96^2/0.01 \approx 385$, which is the usual gap between a bound and the truth.

Exercise 3.3. $\mathbb{P}(|\bar{R}_n - p| \geq 0.02) \leq p(1-p)/(n \times 0.0004) \leq 0.05$ requires $n \geq p(1-p)/(0.00002) = 0.229/0.00002 \approx 11{,}450$. With the worst case $p(1-p) = 1/4$ the requirement is 12,500; the dependence on p is mild because $p(1-p)$ is flat near $1/2$.

Chapter 4. The Central Limit Theorem

Exercise 4.1. (a) $|T_n - 5| \leq \frac{1}{2n}$ always, so the probability is 0 as soon as $\frac{1}{2n} \leq \delta$, i.e., for $n \geq \frac{1}{2\delta}$; any $N \geq \frac{1}{2\delta}$ works. (b) $T_n \xrightarrow[n\to\infty]{\mathbb{P}} 5$ (and so also in distribution). (c) For all $t \neq 5$: $F_n(5) = \frac{1}{2}$ for every n while $F(5) = 1$. See Example 4.2 and Figure 4.1.

Exercise 4.2. (a) Since $M_n \leq 1$, $\mathbb{P}(|M_n - 1| > \delta) = \mathbb{P}(M_n < 1 - \delta) = \mathbb{P}(Y_1 < 1 - \delta)^n = (1 - \delta)^n$. (b) Yes: $(1 - \delta)^n \to 0$, so $M_n \xrightarrow[n\to\infty]{\mathbb{P}} 1$. (c) $F_{M_n}(x) = \mathbb{P}(Y_1 \leq x, \ldots, Y_n \leq x) = x^n$. As $n \to \infty$ this tends to 0 for $x < 1$ and equals 1 at $x = 1$: the CDF of the constant 1, at its continuity points.

Exercise 4.3. Chebyshev with $\epsilon = c\sigma/\sqrt{n}$ gives $\frac{\sigma^2/n}{c^2\sigma^2/n} = 1/c^2$: 0.25 for $c = 2$ and 0.028 for $c = 6$. Hoeffding with $(b-a)^2 = b^2$ gives $2\exp(-2n\epsilon^2/b^2) = 2\exp(-2c^2\sigma^2/b^2) = 2e^{-c^2/6}$: $2e^{-2/3} = 1.03$ (useless, since a probability is at most 1) for $c = 2$, and $2e^{-6} = 0.005$ for $c = 6$. Chebyshev wins for small c; Hoeffding's bound decays like $e^{-c^2/6}$ against Chebyshev's $1/c^2$, so Hoeffding wins once c is large, exactly as in Figure 4.5 with n in place of c.

Exercise 4.4. The CLT gives $\mathbb{P}(|\bar{R}_n - p| \geq \epsilon) \approx 2(1 - \Phi(\epsilon\sqrt{n}/\sqrt{p(1-p)}))$, which equals 0.05 when $\epsilon\sqrt{n}/\sqrt{p(1-p)} = 1.96$, i.e., $n = 1.96^2 \times 0.229/0.02^2 \approx 2{,}200$. Chebyshev demanded about 11,450, five times as many, for the same guarantee.

Chapter 5. Working with Limits

Exercise 5.1. With $g(p) = p(1-p)$, $g'(p) = 1 - 2p$, and the CLT gives asymptotic variance $p(1-p)$ for $\bar{X}_n$, so the limit is $\text{Gauss}(0, (1-2p)^2 p(1-p))$. At $p = \frac{1}{2}$ the derivative is 0 and the limit is degenerate: the sample variance of a fair coin is estimated with error of order $1/n$ rather than $1/\sqrt{n}$, and the second-order remark of Section 5.4 applies.

Exercise 5.2. The CLT gives $\sqrt{n}(\bar{X}_n - 1/\lambda) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, 1/\lambda^2)$. With $g(x) = \log x$, $g'(1/\lambda) = \lambda$, so the asymptotic variance of $\log \bar{X}_n$ is $\lambda^2 \cdot \lambda^{-2} = 1$, free of λ. A limit whose variance doesn't depend on the unknown parameter needs no plug-in: the interval $\log \bar{X}_n \pm 1.96/\sqrt{n}$ for $\log(1/\lambda)$ has asymptotic level 95% with no further work, and can be exponentiated afterwards. This is called a variance-stabilizing transformation.

Exercise 5.3. By the CLT, $\sqrt{n}(\bar{X}_n - p)/\sqrt{p(1-p)} \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$. By the LLN, $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} p$, and $x \mapsto \sqrt{p(1-p)/(x(1-x))}$ is continuous at p, so by the continuous mapping theorem $\sqrt{p(1-p)/(\bar{X}_n(1-\bar{X}_n))} \xrightarrow[n\to\infty]{\mathbb{P}} 1$. Slutsky's multiplicative property with $y = 1$ gives the product's limit, $1 \cdot \text{Gauss}(0,1)$.

Chapter 6. The Gaussian Distribution and Quantiles

Exercise 6.1. From the table, $\Phi(1) = 0.8413$. By symmetry $\mathbb{P}(Z \geq -1) = \mathbb{P}(Z \leq 1) = 0.8413$, and $\mathbb{P}(|Z| > 1) = 2(1 - \Phi(1)) = 0.3173$. So about a third of standard Gaussian draws fall more than one standard deviation from the mean, the complement of the 68% in Figure 6.1.

Exercise 6.2. By the affine property $Y \sim \text{Gauss}(2\mu, 4\sigma^2)$, so $f_Y(y) = \frac{1}{2\sigma\sqrt{2\pi}} \exp\left(-\frac{(y-2\mu)^2}{2\cdot 4\sigma^2}\right)$. By change of variables with $g(x) = 2x$, $g^{-1}(y) = y/2$ and $g' = 2$: $f_Y(y) = \frac{1}{2} f_X(y/2) = \frac{1}{2\sigma\sqrt{2\pi}} \exp\left(-\frac{(y/2-\mu)^2}{2\sigma^2}\right)$, and $(y/2-\mu)^2/(2\sigma^2) = (y-2\mu)^2/(8\sigma^2)$, the same thing.

Exercise 6.3. Standardize: $(X-2)/\sqrt{2} \sim \text{Gauss}(0,1)$, so $\mathbb{P}(X \geq -x) = \mathbb{P}\left(Z \geq \frac{-x-2}{\sqrt{2}}\right) = 1 - \Phi\left(-\frac{x+2}{\sqrt{2}}\right) = \Phi\left(\frac{x+2}{\sqrt{2}}\right)$ by symmetry. For $x = 1$ this is $\Phi(2.12) = 0.983$.

Exercise 6.4. $2\Phi(1)-1 = 0.6827$, $2\Phi(2)-1 = 0.9545$, $2\Phi(3)-1 = 0.9973$. The complement of the second is $\mathbb{P}(|X-\mu| > 2\sigma) = \mathbb{P}(|Z| > 2) = 0.0455$. The event $|X-\mu| > 2\sigma$ is the event $|Z| > 2$ for the standardized variable, whose distribution is the same whatever μ and σ are.

Exercise 6.5. $c = q_{0.10} = 1.28$ and $c' = q_{0.05} = 1.645$. A confidence interval is a two-sided statement (the truth may be on either side of the estimate), so it needs $c' = 1.645$: $\mathbb{P}(|\bar{X}_n - \mu| \le 1.645\,\sigma/\sqrt{n}) \approx 0.90$. Using 1.28 would give an 80% interval.

Chapter 7. Estimators: Bias, Variance and Risk

Exercise 7.1. $\mathbb{E}[\bar{X}_n] = a+\frac{1}{2}$, bias $\frac{1}{2}$. $\mathbb{V}\mathrm{ar}(X_1) = \int_a^{a+1} x^2\,\mathrm{d}x - (a+\frac{1}{2})^2 = a^2+a+\frac{1}{3}-a^2-a-\frac{1}{4} = \frac{1}{12}$, so $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \frac{1}{12n}$. The estimator $\bar{X}_n - \frac{1}{2}$ is unbiased with the same variance, so its risk is $\frac{1}{12n}$, which tends to 0.

Exercise 7.2. $\hat{p} \sim \mathrm{Bern}(p^2)$, so bias $= p^2-p$, variance $= p^2(1-p^2)$, and risk $= p^2(1-p^2)+(p^2-p)^2 = p^2(1-p^2)+p^2(1-p)^2$. None of these depend on n, so the risk does not go to 0. The estimator is not consistent either, for the direct reason that its distribution never changes with n: $\mathbb{P}(|\hat{p}-p| > \epsilon)$ is the same positive number for every n. It never uses $X_3, \ldots, X_n$.

Exercise 7.3. $\mathbb{E}[\tilde{p}_n] = (np+10)/(n+20)$ and $\mathbb{V}\mathrm{ar}(\tilde{p}_n) = np(1-p)/(n+20)^2$. At $p = 0.645$: mean 0.6249, bias -0.0201, variance 0.001369, risk 0.001775, against $\bar{R}_n$'s risk 0.001847. At $p = 0.9$: mean 0.8444, bias -0.0556, variance 0.000538, risk 0.003625, against $\bar{R}_n$'s 0.000726. Shrinking toward $\frac{1}{2}$ pays when p is near $\frac{1}{2}$ and costs five times over when it isn't; Figure 7.5 shows the crossover at $p \approx 0.34$ and 0.66.

Exercise 7.4. $g(x) = 1/x$ is convex on $(0,\infty)$, so $\mathbb{E}[1/\bar{X}_n] \ge 1/\,\mathbb{E}[\bar{X}_n] = \lambda$: the bias is non-negative. Exactly, the bias is $\frac{n}{n-1}\lambda - \lambda = \frac{\lambda}{n-1}$, positive and of order $1/n$. The estimator $\frac{n-1}{n}\cdot\frac{1}{\bar{X}_n} = (n-1)/\sum X_i$ is unbiased.

Exercise 7.5. Bias and variance tending to 0 means $R(\hat{\theta}_n) = \mathbb{E}[(\hat{\theta}_n - \theta)^2] \to 0$, which implies convergence in probability by Chebyshev's inequality (Proposition 3.3; in the language of Appendix B, convergence in quadratic mean implies convergence in probability). The converse fails: convergence in probability does not require any moments to exist. For an example, take $\hat{\theta}_n = \theta$ with probability $1-1/n$ and $\hat{\theta}_n = n$ with probability $1/n$; then $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$, but $\mathbb{E}[\hat{\theta}_n] = \theta(1-1/n)+1 \to \theta+1$, so the bias tends to 1, not 0.

Chapter 8. Confidence Intervals

Exercise 8.1. $q_{0.01} = 2.326$. Conservative: the half-width is $2.326/(2\sqrt{124}) = 0.1045$, so $0.6452 \pm 0.1045 = [0.541, 0.750]$. Plug-in: the half-width is $2.326 \times 0.0430 = 0.1000$, so $0.6452 \pm 0.1000 = [0.545, 0.745]$. Both contain their 95% versions, as they must for intervals of the same type. A 95% interval is a 90% interval (level is a floor) but is not in general a 98% interval, so the answers are No to the 98% question and Yes to the 90% one.

Exercise 8.2. (1) is the realization of $[0, \bar{R}_n]$, and $\mathbb{P}_p([0, \bar{R}_n] \ni p) = \mathbb{P}_p(\bar{R}_n \geq p) \to \mathbb{P}(Z \geq 0) = \frac{1}{2}$ for every p, by the CLT: a 50% asymptotic interval, a strange one. (2) is $[\bar{R}_n, 1]$, the same argument. (3) is the realization at $n = 150$ of the rule $\bar{R}_n \pm c/\sqrt{n}$ with $c = 0.0335\sqrt{150} = 0.41$ (a fixed half-width would have limiting coverage 1, so the rule has to shrink with n), and the conservative bound gives $\lim_{n\to\infty} \mathbb{P}_p(|\bar{R}_n - p| \leq c/\sqrt{n}) \geq \mathbb{P}(|Z| \leq 2c) = \mathbb{P}(|Z| \leq 0.82) = 0.59 \geq 0.5$ for every p: yes. (4) is the rule $\bar{R}_n \pm c/\sqrt{n}$ with $c = 0.02\sqrt{150} = 0.245$, and at $p = \frac{1}{2}$ its limiting coverage is $\mathbb{P}(|Z| \leq 2c) = \mathbb{P}(|Z| \leq 0.49) = 0.38 < 0.5$. The definition asks for coverage at every p, and $p = \frac{1}{2}$ is a p, so (4) is not a 50% interval even though it is one for p far from $\frac{1}{2}$. The realizations of 50% asymptotic intervals are therefore (1), (2) and (3).

Exercise 8.3. $\hat{p}(1-\hat{p}) \leq \frac{1}{4}$ for every $\hat{p} \in [0,1]$, so the plug-in half-width $q_{\alpha/2}\sqrt{\hat{p}(1-\hat{p})/n}$ is at most the conservative one $q_{\alpha/2}/(2\sqrt{n})$, with equality exactly when $\hat{p} = \frac{1}{2}$. From Equation (8.2), the center of $\mathcal{I}_{\text{solve}}$ is $(\bar{R}_n + q^2_{\alpha/2}/(2n))/(1 + q^2_{\alpha/2}/n)$, which equals $\bar{R}_n$, the center of $\mathcal{I}_{\text{plug-in}}$, if and only if $\bar{R}_n + q^2/(2n) = \bar{R}_n + \bar{R}_n q^2/n$, that is, $\bar{R}_n = \frac{1}{2}$. Otherwise the solve interval's center lies strictly between $\bar{R}_n$ and $\frac{1}{2}$.

Exercise 8.4. (a) $\bar{X}_n \sim \text{Gauss}(\theta, \theta/n)$ exactly, so $\sqrt{n/\theta}\,(\bar{X}_n - \theta) \sim \text{Gauss}(0,1)$ for every n and θ. (b) With $t = q_{0.05} = 1.645$: $\mathcal{I}_\theta = [\bar{X}_n - t\sqrt{\theta/n},\ \bar{X}_n + t\sqrt{\theta/n}]$, exact coverage 0.90, not computable. (c) $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta$, so Slutsky gives $\sqrt{n/\bar{X}_n}\,(\bar{X}_n - \theta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$ and $\mathcal{I}_{\text{plug-in}} = \bar{X}_n \pm t\sqrt{\bar{X}_n/n}$, asymptotic level 90%. (d) $\frac{n}{\theta}(\bar{X}_n - \theta)^2 \leq t^2$ is $\theta^2 - (2\bar{X}_n + t^2/n)\theta + \bar{X}_n^2 \leq 0$, so

$$\mathcal{I}_{\text{solve}} = \left[\bar{X}_n + \frac{t^2}{2n} - \sqrt{\Delta},\ \bar{X}_n + \frac{t^2}{2n} + \sqrt{\Delta}\right], \qquad \Delta = \frac{t^4}{4n^2} + \frac{t^2\bar{X}_n}{n}.$$

Because the pivot in (a) is exact, so is this interval's level: it is a confidence interval of level 90% for every n, not only asymptotically. Solving adds no approximation beyond the pivot's own; plugging in always adds one.

Exercise 8.5. (a) $\mathbb{E}[X_1] = a + 1$ and $\mathbb{V}\text{ar}(X_1) = 1$ (a shift doesn't change the variance of Expo(1)), so $\hat{a}_1 \xrightarrow[n\to\infty]{\mathbb{P}} a$ by the LLN and $\sqrt{n}(\hat{a}_1 - a) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0,1)$ by the CLT. The asymptotic variance is known, so there is nothing to plug in: $\mathcal{I}_1 = \hat{a}_1 \pm q_{\alpha/2}/\sqrt{n}$. (b) For $t \geq a$, $\mathbb{P}(\hat{a}_2 > t) = \mathbb{P}(X_i > t\ \forall i) = (e^{-(t-a)})^n = e^{-n(t-a)}$, so $\mathbb{P}(n(\hat{a}_2 - a) > s) = e^{-s}$: $n(\hat{a}_2 - a) \sim \text{Expo}(1)$, for every n. Since $\hat{a}_2 \geq a$ always, the natural interval is $[\hat{a}_2 - s, \hat{a}_2]$, and $\mathbb{P}(a \in [\hat{a}_2 - s, \hat{a}_2]) = \mathbb{P}(n(\hat{a}_2 - a) \leq ns) = 1 - e^{-ns}$. Setting this to $1-\alpha$ gives $s = \log(1/\alpha)/n$ and $\mathcal{I}_2 = [\hat{a}_2 - \log(1/\alpha)/n,\ \hat{a}_2]$, exact level $1-\alpha$. (c) $\mathcal{I}_1$ has width $2q_{\alpha/2}/\sqrt{n}$ and $\mathcal{I}_2$ has width $\log(1/\alpha)/n$: at $n = 100$ and $\alpha = 0.05$, 0.39 against 0.03. The minimum is tailored to the model, uses its exact distribution, and needs no asymptotics. It is also fragile: one observation below a, impossible under the model but possible under a slightly wrong one, breaks it completely, while the sample mean barely notices. The generic $1/\sqrt{n}$ interval is the price of robustness.

Exercise 8.6. One-sided, so $q_{0.05} = 1.645$. Plug-in: $[0.645 - 1.645 \times 0.0430,\ 1] = [0.574, 1]$. Conservative: $[0.645 - 1.645/(2\sqrt{124}),\ 1] = [0.571, 1]$. With 95% confidence, at least 57% of couples turn right. The conservative lower bound exceeds $\frac{1}{2}$ when $\bar{R}_n - 1.645/(2\sqrt{n}) > \frac{1}{2}$, that is, $\bar{R}_n > 0.5 + 0.0739 = 0.5739$, that is, $n\bar{R}_n > 71.2$, that is, $n\bar{R}_n \geq 72$. A confidence interval excluding the value $\frac{1}{2}$ and a test rejecting $p = \frac{1}{2}$ are the same decision looked at from two sides, and the conservative scale $\frac{1}{2}$ is exactly the scale the test uses under its null hypothesis; Chapter 9 makes this duality general.

Exercise 8.7. The pivot is $G = \lambda n\bar{T}_n$, whose distribution doesn't involve λ. Trap it: $\mathbb{P}(49.29 \le \lambda n\bar{T}_n \le 80.60) = 0.95$ for every λ and this n. Solve for λ: $\lambda \in [49.29/(n\bar{T}_n),\ 80.60/(n\bar{T}_n)]$, an interval of exact level 95% with computable endpoints. With $n\bar{T}_n = 398.7$: $[0.124, 0.202]$, between the solve interval $[0.129, 0.213]$ and the plug-in interval $[0.121, 0.200]$ of Section 8.6, and like the solve interval not centered at $\hat{\lambda} = 0.1605$. This is the exact interval Section 4.4 mentioned; it exists because the parent distribution is known in full, which is exactly what the CLT lets you do without.

Chapter 9. Hypothesis Tests

Exercise 9.1. $H_0 : p = \frac{1}{2}$ against $H_1 : p \neq \frac{1}{2}$ in the model $(\{0,1\}, (\text{Bern}(p))_{p\in(0,1)})$. $T_n = 2\sqrt{124}(0.6452 - 0.5) = 3.23$. CLT p-value: $2(1 - \Phi(3.23)) = 0.0012$. Exact: $\mathbb{P}_{1/2}(n\bar{R}_n \ge 80) = 0.0008$, doubled for two sides, 0.0016. Both are far below 0.05 (and below 0.01): reject H_0 at any conventional level; the data are not compatible with equal appeal. One-sided, $H_1 : p > \frac{1}{2}$: p-values $1 - \Phi(3.23) = 0.0006$ (CLT) and 0.0008 (exact), half the two-sided ones.

Exercise 9.2. Let μ be the mean number of chips. The claim is the null, $H_0 : \mu \ge 20$, and the student's doubt is the alternative, $H_1 : \mu < 20$: a one-sided test with the threshold computed at the boundary $\mu_0 = 20$. With σ replaced by the sample standard deviation, $T_n = \sqrt{32}\,(14.77 - 20)/4.37 = -6.77$. By the CLT and Slutsky the p-value is approximately $\Phi(-6.77) \approx 6 \times 10^{-12}$: reject at every conventional level; the difference is not chance. The approximation is the Gaussian limit at $n = 32$ with an estimated σ; Chapter 14 replaces it by Student's t with 31 degrees of freedom, for which -6.77 is just as decisive.

Exercise 9.3. The claim needs evidence, so it is the alternative: $H_0 : p \ge 0.95$ against $H_1 : p < 0.95$, one-sided, threshold at the boundary $p_0 = 0.95$. $\hat{p} = 206/297 = 0.6936$ and $T_n = \sqrt{297}\,(0.6936 - 0.95)/\sqrt{0.95 \times 0.05} = -20.3$. The p-value $\Phi(-20.3)$ is zero to every decimal place anyone prints: reject H_0; the procedure's false positive rate is below 95%. (Whether it is below by enough to matter is the confidence interval's question: $0.694 \pm 1.96\sqrt{0.694 \times 0.306/297} = [0.64, 0.75]$.)

Exercise 9.4. $T_n = -0.73$, CLT p-value $2\Phi(-0.73) = 0.47$. Exactly, $\mathbb{P}_{1/2}(\text{Bin}(30, \frac{1}{2}) \le 13) = 0.292$, doubled 0.58. Neither is close to rejecting. The gap between 0.47 and 0.58 is the discreteness of Figure 9.4 at $n = 30$: the Gaussian tail beyond 13 undercounts the mass of the bar at 13 itself, which the continuity correction (evaluating at 13.5) largely repairs: $2\Phi(\sqrt{30}(13.5/30 - 0.5)/0.5) = 2\Phi(-0.55) = 0.58$.

Exercise 9.5. $\Theta_0 = [1, 3]$. Reject when $\hat{\lambda}$ is far outside it on either side, with a margin that shrinks like $1/\sqrt{n}$ and uses the standard deviation $\sqrt{\lambda}$ at its own endpoint: $\psi = \mathbb{1}\{\hat{\lambda} > 3 + s\sqrt{3/n}\} + \mathbb{1}\{\hat{\lambda} < 1 - s\sqrt{1/n}\}$ for a number s to be chosen (the two events are disjoint, since $\hat{\lambda}$ cannot be both above 3 and below 1). The type 1 error at $\lambda \in [1, 3]$ is largest at the endpoints: at $\lambda = 3$ the first indicator is the event $\sqrt{n}(\hat{\lambda} - 3)/\sqrt{3} > s$, whose probability tends to $\mathbb{P}(Z > s)$ by the CLT, and the second indicator's probability tends to 0 by the LLN; at $\lambda = 1$ the reverse, with limit $\mathbb{P}(Z < -s)$. At any interior λ both probabilities tend to 0. Choosing $s = 1.645$ makes each endpoint's limiting type 1 error 0.05, so the test has asymptotic level 5% (not 2.5% per side: at any one λ only one side can fire in the limit). Each side uses the standard deviation $\sqrt{\lambda}$ at its own endpoint, which is the 'parameter known under the null' idea applied endpoint by endpoint.

Exercise 9.6. For $u \in (0,1)$, $\mathbb{P}_{H_0}(P \le u) = \mathbb{P}_{H_0}(1 - F_0(T) \le u) = \mathbb{P}_{H_0}(F_0(T) \ge 1-u) = \mathbb{P}_{H_0}(T \ge F_0^{-1}(1-u)) = 1 - F_0(F_0^{-1}(1-u)) = u$, using that F_0 is continuous and strictly increasing so that F_0^{-1} exists and $\mathbb{P}(T = t) = 0$ (toolkit A.6). A CDF equal to u on $(0,1)$ is the uniform CDF. Then $\mathbb{P}_{H_0}(P \le \alpha) = \alpha$, which is the golden rule's 'reject when $P \le \alpha$' having type 1 error exactly α. For a discrete T the equality becomes $\le \alpha$, the staircase of Section 9.6.

Exercise 9.7. Not rejecting p_0 means $|\bar{R}_n - p_0| \le q_{\alpha/2}\sqrt{p_0(1-p_0)/n}$, that is, $(\bar{R}_n - p_0)^2 \le q_{\alpha/2}^2\, p_0(1-p_0)/n$, the quadratic inequality in p_0 whose solution set is the solve interval (Equation (8.2)). With $\sqrt{\bar{R}_n(1-\bar{R}_n)/n}$ in place of $\sqrt{p_0(1-p_0)/n}$ the inequality is $|\bar{R}_n - p_0| \le q_{\alpha/2}\sqrt{\bar{R}_n(1-\bar{R}_n)/n}$, whose solution set is $\bar{R}_n \pm q_{\alpha/2}\sqrt{\bar{R}_n(1-\bar{R}_n)/n}$, the plug-in interval. The conservative interval inverts the test that uses the worst-case standard error $1/(2\sqrt{n})$ for every p_0.

Chapter 10. Distances Between Distributions

Exercise 10.1. On $E = \{0,1\}$ the events are $\emptyset$, $\{0\}$, $\{1\}$ and E, and the first and last have the same probability under any two distributions. For Bern($\frac{1}{2}$) against Bern(q), the event $\{1\}$ has probabilities $\frac{1}{2}$ and q, and $\{0\}$ has $\frac{1}{2}$ and $1-q$; both differences are $|\frac{1}{2} - q|$. So the three distances are 0, $\frac{1}{6} = 0.167$ and 0.4, matching $|p - q|$. The first is 0 whether or not two Bernoulli variables with these distributions are independent: total variation compares distributions, not random variables, and their joint behavior doesn't enter.

Exercise 10.2. The computation is the third item of the worked example in Section 10.2. The event on which Expo(1) puts more mass is $A^\star = \{x > 1\}$, where the uniform density is 0; its probabilities are e^{-1} and 0, so $|\mathbb{P}(A^\star) - \mathbb{Q}(A^\star)| = e^{-1} = 0.368$, the same answer as the integral. The complementary event $[0,1]$ favors the uniform by the same amount, $1 - (1 - e^{-1}) = e^{-1}$.

Exercise 10.3. X takes values in $\{0,1\}$ and $X + a$ in $\{a, 1+a\}$, disjoint from it since $0 < a < 1$. The event $A = \{0,1\}$ has $\mathbb{P}(X \in A) = 1$ and $\mathbb{P}(X + a \in A) = 0$, so TV $= 1$, the largest possible value. For the second, $2\sqrt{n}(\bar{X}_n - \frac{1}{2})$ takes values in the finite set $S_n = \{2\sqrt{n}(k/n - \frac{1}{2}) : k = 0, \dots, n\}$, so $\mathbb{P}(2\sqrt{n}(\bar{X}_n - \frac{1}{2}) \in S_n) = 1$, while $\mathbb{P}(Z \in S_n) = 0$ because a continuous random variable has probability 0 of landing in any finite set. Hence TV $= 1$. The CLT says $\mathbb{P}(2\sqrt{n}(\bar{X}_n - \frac{1}{2}) \le t) \to \Phi(t)$ for every t: convergence of the probabilities of the intervals $(-\infty, t]$. Total variation asks for agreement on all events at once, including the finite set S_n, which changes with n. Convergence in distribution is the weaker requirement.

Exercise 10.4. $\alpha = \mathbb{P}_{p_0}(R)$, $\beta = \mathbb{P}_{p_1}(R^c)$. $R = \emptyset$: $\alpha = 0$, $\beta = 1$, sum 1. $R = \{0,1\}$: $\alpha = 1$, $\beta = 0$, sum 1. $R = \{1\}$: $\alpha = p_0$, $\beta = 1 - p_1$, sum $1 - (p_1 - p_0)$. $R = \{0\}$: $\alpha = 1 - p_0$, $\beta = p_1$, sum $1 + (p_1 - p_0)$. The smallest is $1 - (p_1 - p_0) = 1 - \text{TV}$, from the test that rejects when $X = 1$, the outcome Bern(p_1) favors. For the kiss values, 0.855: one couple barely distinguishes the two hypotheses.

Exercise 10.5. The binomial coefficients cancel in the ratio: $\log\left(p_{\text{Bin}(n,p)}(k)/p_{\text{Bin}(n,q)}(k)\right) = k\log(p/q) + (n-k)\log((1-p)/(1-q))$. This is linear in k, and under Bin(n,p), $\mathbb{E}[K] = np$, so the expectation is $np\log(p/q) + n(1-p)\log((1-p)/(1-q)) = n\,\text{KL}(\text{Bern}(p), \text{Bern}(q))$. For the kiss data, $124 \times 0.0428 = 5.30$. The count of right turns carries the same evidence as the full sequence of turns, which is why the distributions of the sequence and of the count give the same divergence.

Exercise 10.6. The log-ratio of the densities is

$$\log \frac{p(x)}{q(x)} = \log \frac{\tau}{\sigma} - \frac{(x-a)^2}{2\sigma^2} + \frac{(x-b)^2}{2\tau^2}.$$

Under $X \sim \text{Gauss}(a, \sigma^2)$, $\mathbb{E}[(X-a)^2] = \sigma^2$ and $\mathbb{E}[(X-b)^2] = \sigma^2 + (a-b)^2$, so the expectation is $\log(\tau/\sigma) - \frac{1}{2} + (\sigma^2 + (a-b)^2)/(2\tau^2)$. With $\tau = \sigma$ the first term is 0 and the rest is $\frac{1}{2} + (a-b)^2/(2\sigma^2) - \frac{1}{2}$. For $\sigma = 0.01$, $\tau = 0.02$, $a = b = 0$: $\log 2 + \frac{1}{8} - \frac{1}{2} = 0.318$. The other way round, $\sigma = 0.02$, $\tau = 0.01$: $\log \frac{1}{2} + 2 - \frac{1}{2} = 0.807$. Believing the narrow distribution when the wide one is true is costlier: the wide one regularly produces values the narrow one calls very unlikely.

Exercise 10.7. The derivation is the third item of the worked example in Section 10.5. With $t = \mu/\lambda > 0$ the divergence is $-\log t + t - 1 = g(t)$. Then $g(1) = 0$, and $g'(t) = 1 - 1/t$ is negative for $t < 1$ and positive for $t > 1$, so g has its minimum 0 at $t = 1$ and is positive elsewhere. This is Proposition 10.6 for this family, by direct calculation instead of Jensen.

Exercise 10.8. $\mathbb{P}_n$ puts mass $e^{-1/n}$ at 0 and $1 - e^{-1/n}$ on $\{1, 2, \ldots\}$. By Proposition 10.2, TV $= \frac{1}{2}\big(|e^{-1/n} - 1| + (1 - e^{-1/n})\big) = 1 - e^{-1/n} \to 0$ (0.095 at $n = 10$). $\text{KL}(\mathbb{P}_n, \delta_0) = +\infty$ for every n, since $\mathbb{P}_n$ gives positive probability to 1, which δ_0 calls impossible. $\text{KL}(\delta_0, \mathbb{P}_n)$ has a single term, at $x = 0$: $1 \cdot \log\big(1/e^{-1/n}\big) = 1/n \to 0$. So $\mathbb{P}_n$ approaches δ_0 in total variation and in one direction of KL, and stays infinitely far in the other.

Exercise 10.9. $\log p_\lambda(x) = -\lambda + x \log \lambda - \log(x!)$, so $\widehat{\text{KL}} = \text{constant} - \frac{1}{n}\sum_i \log p_\lambda(X_i) = \text{constant}' + \lambda - \bar{X}_n \log \lambda$, where $\frac{1}{n}\sum_i \log(X_i!)$ has been absorbed into the constant since it doesn't depend on λ. The derivative $1 - \bar{X}_n/\lambda$ is zero at $\lambda = \bar{X}_n$, negative before and positive after. The minimizer, which is the maximum likelihood estimator, is the sample mean. It is the sample version of Proposition 10.10(2), with the true mean replaced by the sample mean.

Exercise 10.10. $-\mathbb{E}[\log p_\theta(X)] = -\mathbb{E}[\log \binom{2}{X}] - \mathbb{E}[X] \log \theta - (2 - \mathbb{E}[X]) \log(1-\theta)$, and only the last two terms depend on θ. Setting the derivative $-\mathbb{E}[X]/\theta + (2 - \mathbb{E}[X])/(1-\theta)$ to zero gives $\theta^\circ = \mathbb{E}[X]/2 = (0.2 + 0.6)/2 = 0.4$: the Binomial with the same mean. Bin$(2, 0.4)$ has PMF $(0.36, 0.48, 0.16)$, and $\text{KL} = 0.5 \log(0.5/0.36) + 0.2 \log(0.2/0.48) + 0.3 \log(0.3/0.16) = 0.164 - 0.175 + 0.189 = 0.178$. The model is not well-specified: every Bin$(2, \theta)$ satisfies $\mathbb{P}(X = 1)^2 = 4\mathbb{P}(X = 0)\mathbb{P}(X = 2)$, and here $0.04 \neq 0.6$. The fitted model matches the mean but puts far too much mass on 1 and too little on 0 and 2.

Chapter 11. The Likelihood and Maximum Likelihood

Exercise 11.1. $L_2(1, 2; \lambda) = e^{-\lambda}\frac{\lambda^1}{1!} \cdot e^{-\lambda}\frac{\lambda^2}{2!} = e^{-2\lambda}\frac{\lambda^3}{2}$. From the definition, $L_3(1, 2, 3; \lambda) = e^{-3\lambda}\frac{\lambda^{1+2+3}}{1!\,2!\,3!} = e^{-3\lambda}\frac{\lambda^6}{12}$. By updating, $L_3 = L_2 \cdot e^{-\lambda}\frac{\lambda^3}{3!} = e^{-2\lambda}\frac{\lambda^3}{2} \cdot e^{-\lambda}\frac{\lambda^3}{6} = e^{-3\lambda}\frac{\lambda^6}{12}$. $L_3(2, 3, 1; \lambda)$ is the same product with the factors in a different order, so it is also $e^{-3\lambda}\lambda^6/12$: the likelihood of i.i.d. data does not depend on the order of the observations. The MLE from the three observations is $\bar{X}_3 = 2$.

Exercise 11.2. True. $L_n(\theta) = \prod_{i=1}^n p_\theta(x_i)$ is a product of n continuous functions of θ, and a product of continuous functions is continuous. The Bernoulli likelihood $p^S(1-p)^{n-S}$ is a polynomial in p, as smooth as a function can be, although the data take only the values 0 and 1. The Poisson likelihood $e^{-n\lambda}\lambda^{\sum x_i}/\prod x_i!$ is another.

Exercise 11.3. $\ell_n(\sigma^2) = -\frac{n}{2}\log(2\pi\sigma^2) - \frac{1}{2\sigma^2}\sum_i X_i^2$. With $v = \sigma^2$ and $Q = \sum_i X_i^2$, $\ell_n'(v) = -\frac{n}{2v} + \frac{Q}{2v^2} = \frac{Q-nv}{2v^2}$, positive for $v < Q/n$ and negative for $v > Q/n$, so the unique maximizer is $\widehat{\sigma^2} = \frac{1}{n}\sum_i X_i^2$. (The log-likelihood is not concave in v for all v, which is why the sign of the derivative is used instead of the second derivative.) It is unbiased here, since the mean is known: $\mathbb{E}[X_i^2] = \sigma^2$. By invariance, $\hat\sigma = \sqrt{\frac{1}{n}\sum_i X_i^2}$.

Exercise 11.4. $M \le m$ exactly when every $X_i \le m$, and by independence $\mathbb{P}(M \le m) = \prod_i \mathbb{P}(X_i \le m) = (m/\theta)^n$. Differentiating, M has density nm^{n-1}/θ^n on $[0,\theta]$. Then $\mathbb{E}[M] = \int_0^\theta m \cdot \frac{nm^{n-1}}{\theta^n}\,\mathrm{d}m = \frac{n}{\theta^n}\cdot\frac{\theta^{n+1}}{n+1} = \frac{n}{n+1}\theta$. So $\frac{n+1}{n}M$ is unbiased. For the five observations of Figure 11.6, $\frac{6}{5}\times 7.6 = 9.12$. The bias of M is $-\theta/(n+1)$, of order $1/n$, and its standard deviation is also of order $1/n$, not $1/\sqrt{n}$: this model is not one of the regular models whose large-sample theory Chapter 12 develops.

Exercise 11.5. Each density is $\mathbb{1}(\theta \le x_i \le \theta+1)$, so $L(\theta) = \prod_i \mathbb{1}(\theta \le x_i)\mathbb{1}(x_i \le \theta+1) = \mathbb{1}(\theta \le \min_i x_i)\,\mathbb{1}(\theta \ge \max_i x_i - 1)$. It is 1 on $[\max_i x_i - 1, \min_i x_i]$ and 0 elsewhere, so every point of that interval maximizes it and the MLE is not unique. The interval is nonempty because all the data lie in an interval of length 1, $[\theta^*, \theta^*+1]$, so $\max_i x_i - \min_i x_i \le 1$. Its midpoint is a sensible choice of estimate, but the likelihood alone doesn't pick it.

Exercise 11.6. $\ell_n(p) = \sum_i \big[(X_i-1)\log(1-p) + \log p\big] = (S-n)\log(1-p) + n\log p$ with $S = \sum_i X_i$. $\ell_n'(p) = -\frac{S-n}{1-p} + \frac{n}{p} = 0$ gives $n(1-p) = (S-n)p$, so $p = n/S$: $\hat p = 1/\bar X_n$. $\ell_n''(p) = -\frac{S-n}{(1-p)^2} - \frac{n}{p^2} < 0$, so this is the unique maximum. For the data, $S = 12$ and $n = 5$, so $\hat p = 5/12 = 0.417$: five makes in twelve shots.

Exercise 11.7. $L(a) = \prod_i e^{-(x_i-a)}\mathbb{1}(x_i \ge a) = e^{na - \sum_i x_i}\,\mathbb{1}(a \le \min_i x_i)$. For $a \le \min_i x_i$ this is increasing in a; beyond, it is 0. So $\hat a = \min_i X_i$, found at the jump, as for the uniform. Again no derivative vanishes where $L > 0$, and again the estimator is biased: $\min_i X_i - a \sim \mathrm{Expo}(n)$, so $\mathbb{E}[\hat a] = a + 1/n$.

Exercise 11.8. $\ell_n(m) = -n\log 2 - \sum_i |X_i - m|$, so maximizing it is minimizing $g(m) = \sum_i |X_i - m|$. Away from the data points, $g'(m) = (\text{number of } X_i < m) - (\text{number of } X_i > m)$, which is negative when fewer than half the observations lie below m and positive when more than half do. With n odd, g decreases until the middle observation and increases after it, so the minimizer is the sample median. The Gaussian's squared deviations give the mean; the Laplace's absolute deviations give the median. Chapter 13 takes this up as M-estimation.

Exercise 11.9. The terms involving π are $\big(\sum_i w_i\big)\log\pi + \big(\sum_i (1-w_i)\big)\log(1-\pi)$: the Bernoulli log-likelihood with $\sum_i w_i$ 'successes' out of n, maximized at $\hat\pi = \frac{1}{n}\sum_i w_i$. The terms involving μ_1 are $-\sum_i w_i(X_i-\mu_1)^2/(2\sigma_1^2)$, a concave quadratic in μ_1 with derivative $\sum_i w_i(X_i-\mu_1)/\sigma_1^2$, zero at $\hat\mu_1 = \sum_i w_i X_i / \sum_i w_i$. For the variance, write $W = \sum_i w_i$ and $D_1 = \sum_i w_i(X_i - \hat\mu_1)^2$. The terms involving (μ_1, σ_1^2) are $-\frac{W}{2}\log(2\pi\sigma_1^2) - \sum_i w_i(X_i-\mu_1)^2/(2\sigma_1^2)$. For each σ_1^2 this is largest at $\mu_1 = \hat\mu_1$, as just shown; with $\mu_1 = \hat\mu_1$ it is $-\frac{W}{2}\log\sigma_1^2 - D_1/(2\sigma_1^2)$ plus a constant, whose derivative $-\frac{W}{2\sigma_1^2} + \frac{D_1}{2\sigma_1^4}$ is positive below D_1/W and negative above. Chaining the two steps as in the unweighted Gaussian case gives the joint maximum $\hat\sigma_1^2 = D_1/W$. The terms for component 2 are maximized separately in the same way, with weights $1 - w_i$.

Exercise 11.10. With $S = 0$, $L(p) = (1-p)^n$, strictly decreasing on $(0, 1)$, so for every p in Θ there is a smaller p with a larger likelihood: no maximum is attained. Allowing $p \in [0, 1]$ gives $\hat{p} = 0$. The event has probability $0.9^5 = 0.59$ for $n = 5$ and $0.9^{50} = 0.005$ for $n = 50$.

Chapter 12. Properties of the MLE

Exercise 12.1. $\ell(\lambda) = -\lambda + X \log \lambda - \log X!$, $\ell'(\lambda) = -1 + X/\lambda$, $\ell''(\lambda) = -X/\lambda^2$. $\mathbb{V}\mathrm{ar}(\ell') = \mathbb{V}\mathrm{ar}(X)/\lambda^2 = \lambda/\lambda^2 = 1/\lambda$; $-\mathbb{E}[\ell''] = \mathbb{E}[X]/\lambda^2 = 1/\lambda$. So $\sqrt{n}(\bar{X}_n - \lambda) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, \lambda)$. The CLT directly: $\bar{X}_n$ has mean λ and variance λ/n, so the same. For the siblings, standard error $\sqrt{1.0/24} = 0.20$.

Exercise 12.2. In β the density is $\beta^{-1}e^{-x/\beta}$, so $\ell(\beta) = -\log \beta - X/\beta$, $\ell'(\beta) = -1/\beta + X/\beta^2$, $\ell''(\beta) = 1/\beta^2 - 2X/\beta^3$, and with $\mathbb{E}[X] = \beta$, $I(\beta) = -\mathbb{E}[\ell''] = -1/\beta^2 + 2/\beta^2 = 1/\beta^2$. Check: $I(\lambda) = 1/\lambda^2 = \beta^2$ and $\mathrm{d}\lambda/\mathrm{d}\beta = -1/\beta^2$, so $I(\lambda)(\mathrm{d}\lambda/\mathrm{d}\beta)^2 = \beta^2/\beta^4 = 1/\beta^2$. The theorem gives $\sqrt{n}(\bar{X}_n - \beta) \xrightarrow[n\to\infty]{\text{(d)}} \text{Gauss}(0, \beta^2)$, and the CLT gives the same since $\mathbb{V}\mathrm{ar}(X) = \beta^2$. For the subway, $\hat{\beta} = 6.23$ with standard error $6.23/8 = 0.78$ minutes.

Exercise 12.3. $\ell(\mu) = -\frac{1}{2}\log(2\pi\sigma^2) - (X-\mu)^2/(2\sigma^2)$, $\ell''(\mu) = -1/\sigma^2$, so $I(\mu) = 1/\sigma^2$ and the bound is σ^2/n. $\mathbb{V}\mathrm{ar}(\bar{X}_n) = \sigma^2/n$ exactly, and $\bar{X}_n$ is unbiased, so it is efficient. (The Gaussian and Bernoulli sample means are efficient at every n; most MLEs reach the bound only in the limit.)

Exercise 12.4. $\ell_n'(p) = 80/p - 44/(1-p)$: at $\frac{1}{2}$, $160 - 88 = 72$; at 0.6, $133.3 - 110 = 23.3$; at 0.645, 0; at 0.7, $114.3 - 146.7 = -32.4$. The score is positive below $\hat{p}$ (the data push the estimate up), zero at $\hat{p}$, negative above. At $p_0 = \frac{1}{2}$, $nI(p_0) = 124 \times 4 = 496$ and $72/\sqrt{496} = 3.23$, which is the $T_n = 2\sqrt{124}(0.645 - 0.5) = 3.23$ of Chapter 9: the standardized score at the null value is the score test statistic, and under H_0 it is approximately Gauss$(0, 1)$ by the CLT for scores. Chapter 15 develops this.

Exercise 12.5. From Exercise 11.4, $\mathbb{P}(M \le m) = (m/\theta)^n$. For $t \ge 0$, $\mathbb{P}\big(n(\theta - M)/\theta > t\big) = \mathbb{P}\big(M < \theta(1 - t/n)\big) = (1 - t/n)^n \to e^{-t}$ (toolkit A.10). So the CDF of $n(\theta - M)/\theta$ converges to $1 - e^{-t}$, the Expo(1) CDF. The error $\theta - M$ is therefore about θ/n times an exponential variable: order $1/n$, always positive, with a skewed limit. Scaling by $\sqrt{n}$ instead would send $\sqrt{n}(\theta - M) \xrightarrow[n\to\infty]{\mathbb{P}} 0$, a degenerate limit, which is what asymptotic normality at rate $1/\sqrt{n}$ would look like if it held with asymptotic variance 0; the model is not regular and the theorem does not apply.

Exercise 12.6. $\ell(p) = (X-1)\log(1-p) + \log p$, $\ell''(p) = -(X-1)/(1-p)^2 - 1/p^2$, and $\mathbb{E}[X] = 1/p$ gives $\mathbb{E}[X - 1] = (1-p)/p$, so $I(p) = \frac{1}{p(1-p)} + \frac{1}{p^2} = \frac{p+(1-p)}{p^2(1-p)} = \frac{1}{p^2(1-p)}$. At $\hat{p} = 5/12$ with $n = 5$: $nI = 5/((5/12)^2 \times 7/12) = 49.4$ and the standard error is $1/\sqrt{49.4} = 0.14$. Five rounds say p is somewhere around 0.42 ± 0.28.

Exercise 12.7. $\ell_n'(\theta) = \sum_{i=1}^n \ell_i'(\theta)$ is a sum of n independent random variables each with variance $I(\theta)$, so its variance is $nI(\theta)$. And $-\mathbb{E}[\ell_n''(\theta)] = -\sum_i \mathbb{E}[\ell_i''(\theta)] = nI(\theta)$ by Theorem 12.4 applied to each term. Information adds over independent observations, the same way KL divergence did in Proposition 10.7, and for the same reason: logs turn products into sums.

Exercise 12.8. The theorem (misspecified case) says $\hat{\theta}_n \xrightarrow[n\to\infty]{\mathbb{P}} \theta^\circ = (1,1)$. Directly: $\bar{X}_n \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}[X] = 1$ by the LLN, and $\frac{1}{n}\sum(X_i - \bar{X}_n)^2 = \frac{1}{n}\sum X_i^2 - \bar{X}_n^2 \xrightarrow[n\to\infty]{\mathbb{P}} \mathbb{E}[X^2] - 1 = 2 - 1 = 1$ by the LLN and the continuous mapping theorem (Theorem 5.2, since $x \mapsto x^2$ is continuous). The two agree. The fitted Gaussian is the closest one to the truth in KL, at divergence 0.419, and no amount of data moves it closer.

Appendix B. The Modes of Convergence in Full

Exercise B.1. No. The sliding indicator of Example B.4 converges in probability to U and almost surely to nothing.